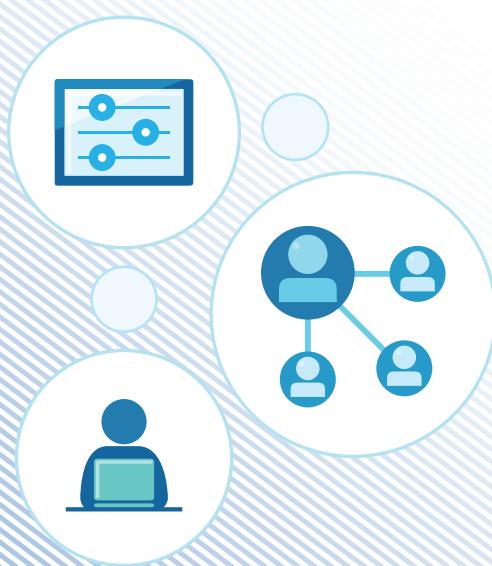


Course Guide

# **IBM Cloud Garage - Cloud Native Developer Bootcamp for IBM Cloud Private**

Course code CK105 ERC 1.0





IBM Training

IBM

## CK105 IBM Cloud Private Cloud Native Developer Bootcamp

### Course Introduction

© Copyright IBM Corporation 2018  
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

IBM Training

IBM

## Welcome!

- Instructor Introduction
- In a few minutes, you will be asked to quickly introduce yourself
  - What's your name?
  - What's your role?
  - Experience with IBM Cloud or Cloud Development?
  - Why are you here?

IBM Cloud Private Cloud Native Developer Bootcamp

2

© Copyright IBM Corporation 2018

## Course Overview

This course:

- Teaches the key features and architecture of cloud-native applications.
- Introduces concepts of Kubernetes and IBM Cloud Private
- Identifies the elements of twelve-factor applications and the characteristics of a resilient, scalable cloud application design.
- Explains the role of DevOps in the cloud-native application lifecycle.

## Course Prerequisites

- Experience with enterprise application development in Java™ or other languages
- Familiarity with agile development practices

## Course Objectives

- Understand the key features, practices, and architecture of cloud-native application development.
- Learn how to use IBM Cloud Private to develop and deploy cloud-native application.
- Explore and practice running DevOps pipeline for deploying Microservices application in IBM Cloud Private.
- Evaluate the use of Istio service mesh for managing a Microservices application.

## Agenda

- Day 1 – Introduction to IBM Cloud
  - Introduction to Cloud Native Development
  - The twelve factor application
  - IBM Cloud Private architecture
  - IBM Cloud Full Stack Development
  - Lab 1 – Environment Preparation
- Day 2 – Container and Kubernetes
  - Introduction to Docker Container
  - Lab 2 – Docker introduction
  - Introduction to Kubernetes
  - Introduction to Helm
  - Building and Deploying application in Kubernetes
  - Lab 3 - Deploy Java application
- Day 3 – DevOps
  - DevOps Concepts
  - Jenkins in IBM Cloud Private
  - ICP Monitoring
  - Lab 4 – Setting up Jenkins
- Day 4 – Microservices architecture
  - Microservices architecture
  - Cloud Data Services
  - Microservices Development
  - Service Mesh
  - Lab 5 – Microservices Design
  - Lab 6 – Deploying BlueCompute
- Day 5 – Managing Microservices with Istio
  - Managing Traffic and Telemetry
  - Microservices Network Security
  - Lab 7 – Working with Istio

IBM Training 

## Introductions

- Name
- Company
- Where you live
- Your job role
- Your current experience with the products and technologies in this course
- Do you meet the course prerequisites?
- What you expect from this class

IBM Cloud Private Cloud Native Developer Bootcamp 7 © Copyright IBM Corporation 2018

IBM Training 

## Class Logistics

- Course environment
  - You will use a browser to access a virtual desktop image running on IBM Remote Lab Platform (Skytap)
  - This image has a number of software packages pre-installed, typical of a developer's desktop
  - Once you log into the virtual desktop image, you will log into IBM Cloud from the virtual desktop image
  - The virtual desktop image provides a consistent starting point and avoids problems in installing software directly on student machines
- Start and end times
- Lab exercise procedures
- Materials in your student packet
- Topics not on the agenda
- Evaluations
- Breaks and lunch
- Outside business
- Food
- Restrooms
- Fire exits
- Local amenities

IBM Cloud Private Cloud Native Developer Bootcamp 8 © Copyright IBM Corporation 2018

IBM Training

IBM

## Introduction to Cloud-Native Application Development

© Copyright IBM Corporation 2016-18  
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

IBM Training

IBM

### Unit objectives

- Differentiate IBM Cloud developer roles
- Explain cloud native application characteristics

Introduction to Cloud-Native Application Development

2

© Copyright IBM Corporation 2016-18

IBM Training 

## Lesson 1

### Developer roles in IBM Cloud

Introduction to Cloud-Native Application Development

3

© Copyright IBM Corporation 2016-18

IBM Training 

### IBM Cloud developer personas

IBM Cloud developer personas can be thought of in these categories:

- Public cloud developer
- Enterprise cloud developer
- Probably others

Introduction to Cloud-Native Application Development

4

© Copyright IBM Corporation 2016-18

IBM Training 

## Public cloud developer

- Simple applications
  - Greenfield development, with no existing apps to migrate
  - Runtime mainly mediates between GUIs and services
  - Little business logic
- Hosted entirely in public
  - No private hosting, neither dedicated or local
- Uses only public resources
  - Twitter feeds, Watson services, and others
  - No existing on-premises systems of record (SoRs) to integrate with
- Values variety of run times and services
  - Node.js, Go, Swift, and others
  - Data and Analytics, Cognitive (Watson services), Mobile services, Internet of Things services
  - Twilio, Payeezy, and others

Introduction to Cloud-Native Application Development 5 © Copyright IBM Corporation 2016-18

IBM Training 

## Enterprise cloud developer

- Complex applications
  - Often existing traditional IT applications migrated to cloud
  - Lots of business logic
  - Most often Java and Java EE
- Hybrid cloud
  - Multiple public clouds for disaster recovery
  - Private clouds (dedicated and local) for guaranteed access, geographic location, and privacy
- Hybrid applications
  - Application components hosted across multiple public, dedicated, and local clouds
- Hybrid IT
  - Application components integrate with existing on-premises resources and systems of record (SoRs)
- Limited use of cloud services
  - The existing on-premises resources are the primary services these apps need
- Specialized skills
  - Examples: IBM Cloud integration developer (and architect), IBM Cloud migration specialist, and others

Introduction to Cloud-Native Application Development 6 © Copyright IBM Corporation 2016-18

IBM Training 

## Public and enterprise developer together

- After the enterprise developers have done their jobs, future development is mostly of the public cloud variety
  - Even if the new development is being hosted privately (dedicated and local)
- New development doesn't have to worry about enterprise concerns as much
  - Useful on-premises applications have been migrated
  - Access to useful on-premises resources has been provided securely
- New development can focus on developing new apps
  - Innovative systems of engagement
  - Rapid development
  - Using IBM Cloud's variety of services

Introduction to Cloud-Native Application Development 7 © Copyright IBM Corporation 2016-18

IBM Training 

## Lesson 2

### Cloud-native application characteristics

Introduction to Cloud-Native Application Development 8 © Copyright IBM Corporation 2016-18

IBM Training 

## Range of application's cloud readiness

- Traditional IT
  - Stateful, transactional, vertical scaling, reliable hardware, SQL databases
  - Traditional middleware, stack products like Portal and BPM
- Cloud ready
  - Minimally sufficient to run on the cloud
  - IaaS is more flexible than PaaS, more like traditional IT servers
- Cloud native
  - Stateless, eventual consistency, horizontal scaling, commodity hardware, NoSQL databases
  - Middleware function provided by services, separate from runtimes

Introduction to Cloud-Native Application Development 9 © Copyright IBM Corporation 2016-18

IBM Training 

## Common characteristics of an enterprise application

- Enterprise applications were usually not built with Cloud compatibility in mind, but they might have been built with future-proofing in mind
  - Trying to catch up to the latest versions of libraries, and others
- Might be built on WebSphere stack products
  - Such as portals with WebSphere Portal Server
  - Purpose-built (departmental) ESBs built with DataPower and IIB
- Deep ties into internal IT
  - Especially in cases where you not only read from, but update systems of record
  - May read from and update multiple systems not specified as APIs
  - Deep ties into corporate standards and security systems, such as LDAP, ISAM, SiteMinder
- Often have high and well-specified QoS attributes
  - Business customers have expectations on how often applications are available and how they perform

Introduction to Cloud-Native Application Development 10 © Copyright IBM Corporation 2016-18

**IBM Training**

**Characteristics of a cloud-native application**

- Built with the cloud in mind
  - Use cloud services
    - Such as storage, queuing, and caching
  - Have rapid and repeatable deployments to maximize agility
  - Have automated setup to minimize time and cost for new developers
  - Have clean contract with underlying OS to ensure maximum portability
- Loose ties into corporate IT
  - Security often specified using open standards
    - Such as OpenID and OAuth
  - Data might be local to the application itself
- QoS attributes are those of the cloud
  - It's always expected to be there, but sometimes sites and mobile apps become unavailable, although they must be restored quickly
  - Applications must scale elastically without significant changes to tooling, architecture, or development practice
  - Application must be resilient to inevitable failures in the infrastructure and application



Introduction to Cloud-Native Application Development

11

© Copyright IBM Corporation 2016-18

**IBM Training**

**Differing assumptions**

- Enterprise or cloud-ready assumptions
  - The infrastructure is stable
  - The components of my application are co-located
  - My ops team controls the production servers
  - If a disaster happens, it's someone else's responsibility to fix it
- Cloud-native assumptions
  - The infrastructure is constantly changing; it is elastic
  - My application components might be globally distributed
  - As a DevOps team member, I control the production servers
    - "You build it, you run it"
  - If a disaster happens, it's my responsibility to make sure my app continues to run




**Choosing one or the other has an effect on your team composition and roles**

Introduction to Cloud-Native Application Development

12

© Copyright IBM Corporation 2016-18

**IBM Training**

**Cloud native**

An **application architecture** designed to use the *strengths* and accommodate the *challenges* of a standardized cloud environment, including:

- Elastic scaling
- Immutable deployment
- Disposable instances
- Less predictable infrastructure



Introduction to Cloud-Native Application Development

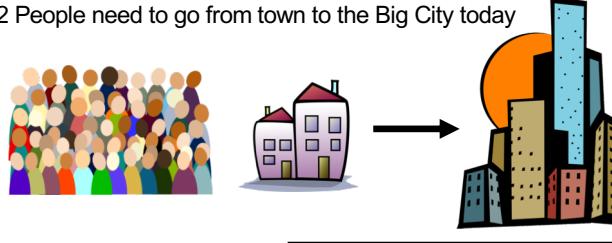
13

© Copyright IBM Corporation 2016-18

**IBM Training**

**An important Cloud-native concept: Managing Capacity**

32 People need to go from town to the Big City today



BUT

Your bus only holds 20 people



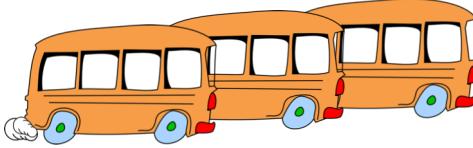
**Vertical Scaling (1 bigger instance)**

Double-decker bus holds 40 people



**Horizontal Scaling (multiple instances)**

More buses - 20 people each



**Which solution is better?**

Introduction to Cloud-Native Application Development

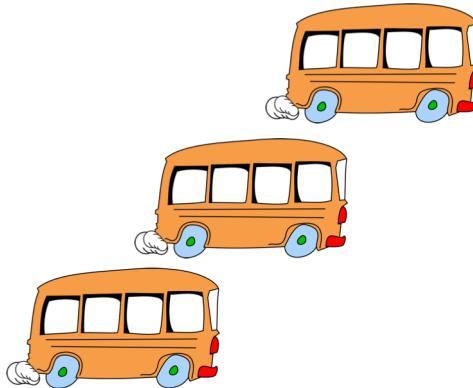
14

© Copyright IBM Corporation 2016-18

**IBM Training**

**For Cloud-native - multiple instances are preferred – Why?**

- No single point of failure
  - If a bus breaks down, you still have other buses
  - Can activate another bus to carry the stranded passengers
- Elastic scalability
  - You can add or remove buses from service to meet the demand
- Load distribution
  - Passengers don't care which bus they go on, as long as they reach their destination
- Flexibility
  - Some buses can use different roads if there is a problem on the default route
- Availability
  - You can take down buses for maintenance, and still provide your service with other buses



Introduction to Cloud-Native Application Development

15

© Copyright IBM Corporation 2016-18

**IBM Training**

**Cloud-Native Application Goals**

- Horizontal scaling
  - Application runs in multiple runtimes spread across multiple hosts (Twelve-Factor VIII)
- Immutable deployment
  - A runtime is not patched, it's replaced (Twelve-Factor IX)
  - A runtime is stateless (Twelve-Factor VI)
  - Shared functionality in backing services (Twelve-Factor IV)
- Elasticity
  - Automatic scale-out and scale-in to maintain performance
  - Achieved via containerization
- Pay-as-you-go charging model
  - Pay for what you use



Introduction to Cloud-Native Application Development

16

© Copyright IBM Corporation 2016-18

IBM Training 

### Recap: What it means to be cloud native

- Clean contract with underlying OS to ensure maximum portability
- Scale elastically without significant changes to tooling, architecture, or development practices
- Resilient to inevitable failures in the infrastructure and application
- Instrumented to provide both technical and business insight
- Use cloud services such as storage, queuing, and caching
- Rapid and repeatable deployments to maximize agility
- Automated setup to minimize time and cost for new developers



Introduction to Cloud-Native Application Development 17 © Copyright IBM Corporation 2016-18

IBM Training 

### Unit summary

- Differentiate IBM Cloud developer roles
- Explain cloud-native application characteristics

Introduction to Cloud-Native Application Development 18 © Copyright IBM Corporation 2016-18

IBM Training

IBM

## The Twelve-Factor Application

© Copyright IBM Corporation 2017-2018  
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

IBM Training

IBM

### Unit objectives

- Describe the twelve-factor app best practices for cloud applications
- Define and use IBM Cloud facilities that incorporate these best practices
- Describe how IBM Cloud incorporates these practices

The Twelve-Factor Application

2

© Copyright IBM Corporation 2017-2018

IBM Training 

## Lesson 1 The twelve factors

The Twelve-Factor Application  © Copyright IBM Corporation 2017-2018

IBM Training 

### The twelve-factor app



- A set of best practices for creating applications
  - Implementing, deploying, monitoring, and managing
- Typical modern applications
  - Deployed in the cloud
  - Accessible as web applications that deliver software-as-a-service (SaaS)
- Can be applied to any application
  - Implemented in any programming language
  - Using any backing services such as database, messaging, and caching
- Addresses common problems
  - The dynamics of the growth of an app over time
  - The dynamics of collaboration between developers
  - Avoiding the cost of software erosion
  - Systemic problems in modern application development
- Provides a shared vocabulary for addressing these problems

The Twelve-Factor Application  © Copyright IBM Corporation 2017-2018

IBM Training 

## The twelve factors

- I. **Codebase:** One codebase that is tracked in revision control, with many deployments
- II. **Dependencies:** Explicitly declare and isolate dependencies
- III. **Configuration:** Store Configuration in the environment
- IV. **Backing services:** Treat backing services as attached resources
- V. **Build, release, run:** Strictly separate build and run stages
- VI. **Processes:** Execute the app as one or more stateless processes
- VII. **Port binding:** Export services with port binding
- VIII. **Concurrency:** Scale out using the process model
- IX. **Disposability:** Maximize robustness with fast startup and efficient shutdown
- X. **Development and production parity:** Keep development, staging, and production as similar as possible
- XI. **Logs:** Treat logs as event streams
- XII. **Admin processes:** Run administrative and management tasks as one-off processes

The Twelve-Factor Application 5 © Copyright IBM Corporation 2017-2018

IBM Training 

## Lesson 2 Twelve-factor details

The Twelve-Factor Application 6 © Copyright IBM Corporation 2017-2018

IBM Training IBM

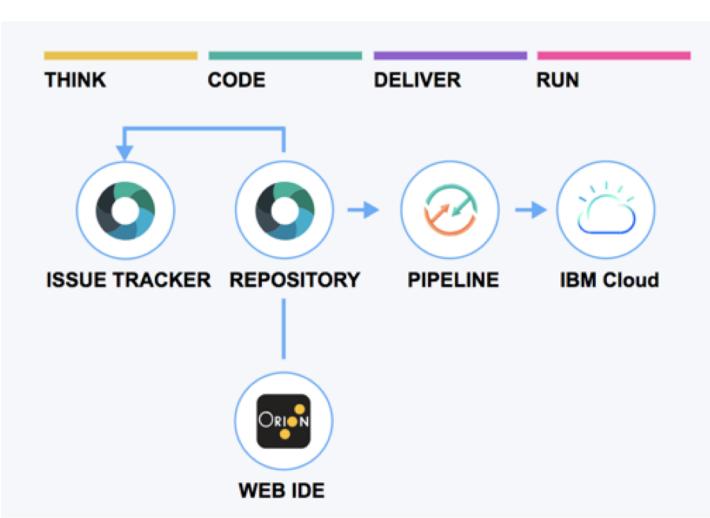
## Factor 1: Codebase

<b>I. Codebase</b> II. Dependencies III. Configuration IV. Backing services V. Build, release, run VI. Processes VII. Port binding VIII. Concurrency IX. Disposability X. Dev/prod parity XI. Logs XII. Admin processes	<ul style="list-style-type: none"> <li>• One codebase tracked in source code management (SCM) with versioning               <ul style="list-style-type: none"> <li>• git &amp; github most commonly used</li> </ul> </li> <li>• Multiple deployments from the same codebase</li>   <li>• IBM Cloud: use IBM Continuous Delivery toolchains or external automation with Cloud Foundry development tools such as these:               <ul style="list-style-type: none"> <li>• Urban Code Deploy</li> <li>• Gradle</li> <li>• Jenkins</li> </ul> </li> </ul>
--	--

The Twelve-Factor Application 7 © Copyright IBM Corporation 2017-2018

IBM Training IBM

## Codebase: IBM Cloud continuous delivery toolchains



The diagram shows a horizontal flow of four stages: THINK, CODE, DELIVER, and RUN. Below each stage is a circular icon representing a tool: ISSUE TRACKER (a target icon), REPOSITORY (a gear icon), PIPELINE (a checkmark icon), and IBM Cloud (a cloud icon). Arrows connect the icons from left to right. A blue line connects the ISSUE TRACKER and REPOSITORY icons. A blue line also connects the REPOSITORY icon to the PIPELINE icon. A blue arrow points from the PIPELINE icon to the IBM Cloud icon.

A *toolchain* is a set of tool integrations that support development, deployment, and operations tasks.

Tool integrations with the source code repository and delivery pipelines can drive multiple deployments from a single repository.

The Twelve-Factor Application 8 © Copyright IBM Corporation 2017-2018

IBM Training IBM

## Factor 2: Dependencies

- I. Codebase
- II. Dependencies**
- III. Configuration
- IV. Backing services
- V. Build, release, run
- VI. Processes
- VII. Port binding
- VIII. Concurrency
- IX. Disposability
- X. Dev/prod parity
- XI. Logs
- XII. Admin processes

- Explicitly declare and isolate dependencies
- Typically language-specific
  - Node.js: Node Package Manager (NPM)
  - Liberty: Feature manager
  - Ruby: Bundler
  - Java EE: Application resources
- Never rely on system-wide dependencies
  
- IBM Cloud: Buildpack adds external dependencies

The Twelve-Factor Application      9      © Copyright IBM Corporation 2017-2018

IBM Training IBM

### Dependencies: npm package file

- Application configuration is part of the application
- You manage it in source control



```

1  {
2    "name": "MachineTranslationNodejs",
3    "version": "0.0.1",
4    "description": "A sample nodejs app for Bluemix that use the machine translation service",
5    "dependencies": {
6      "express": "3.4.7",
7      "jade": "1.1.4",
8      "cors": "2.4.2" ←
9    },
10   "engines": {
11     "node": "0.10.26"
12   },
13   "repository": {}
14 }
15
  
```

Added cors

The Twelve-Factor Application      10      © Copyright IBM Corporation 2017-2018

IBM Training

## Dependencies: Liberty feature manager

- Server configuration is part of deploying the application
- You manage it in source control

The screenshot shows a code editor window titled "server.xml" with the following XML content:

```
<server description="new server">
    <!-- Enable features -->
    <featureManager>
        <feature>javaee-7.0</feature>
        <feature>localConnector-1.0</feature>
    </featureManager>
```

The editor has tabs for "Design" and "Source", with "Source" selected.

The bottom of the slide includes the footer: "The Twelve-Factor Application" (page 11), "© Copyright IBM Corporation 2017-2018".

IBM Training

## Dependencies: IBM Cloud application manifest

- Server configuration is part of deploying the application
- You manage it in source control

The screenshot shows a code editor window titled "manifest.yml" with the following YAML content:

```
applications:
- name: redbooklibrary
  memory: 512M
  path: RedbookLibrary.war
  host: redbook-library
  buildpack: liberty-for-java
services:
- library_db
```

The editor has tabs for "Ln 9, Col 1", "UTF-8", "LF", "YAML", and a smiley face icon.

The bottom of the slide includes the footer: "The Twelve-Factor Application" (page 12), "© Copyright IBM Corporation 2017-2018".

**IBM Training**

## Factor 3: Configuration

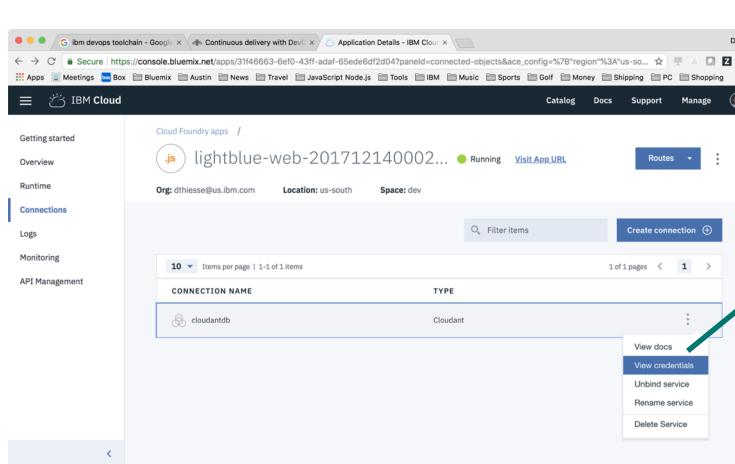
- I. Codebase
- II. Dependencies
- III. Configuration**
- IV. Backing services
- V. Build, release, run
- VI. Processes
- VII. Port binding
- VIII. Concurrency
- IX. Disposability
- X. Dev/prod parity
- XI. Logs
- XII. Admin processes

- Store configuration in the environment
- Separate configuration from source
- Enables the same code to be deployed to different environments
  
- IBM Cloud: You assign parameters to applications with environment variables, both system-provided and custom

The Twelve-Factor Application      13      © Copyright IBM Corporation 2017-2018

**IBM Training**

## Configuration: IBM Cloud VCAP\_SERVICES



```
{
  "cloudantNoSQLDB": [
    {
      "credentials": {
        "username": "4e94679f-767a-4327-926f-cafe516bee0-blue",
        "password": "e81181d22ae71b25106c4a8659778297b014d9",
        "host": "4e94679f-767a-4327-926f-cafe516bee0-bluemix.c",
        "port": 443,
        "url": "https://4e94679f-767a-4327-926f-cafe516bee0-blue
59778297b014d90cbf629aeb8dc80d2a672522d6@4e94679f-76
mix.cloudant.com"
      },
      "syslog_drain_url": null,
      "label": "cloudantNoSQLDB",
      "provider": null,
      "plan": "Lite",
      "name": "Your Application-cloudantNoSQLDB",
      "tags": []
    }
  ]
}
```

The Twelve-Factor Application      14      © Copyright IBM Corporation 2017-2018

IBM Training 

## Kubernetes configuration

- Kubernetes uses ConfigMap and Secret resources

```
kubectl create secret generic apikey --from-literal=API_KEY=123-456
```

```
kubectl create configmap language --from-literal=LANGUAGE=English
```

The Twelve-Factor Application      15      © Copyright IBM Corporation 2017-2018

IBM Training 

## Factor 4: Backing services

<ul style="list-style-type: none"><li>I. Codebase</li><li>II. Dependencies</li><li>III. Configuration</li><li><b>IV. Backing services</b></li><li>V. Build, release, run</li><li>VI. Processes</li><li>VII. Port binding</li><li>VIII. Concurrency</li><li>IX. Disposability</li><li>X. Dev/prod parity</li><li>XI. Logs</li><li>XII. Admin processes</li></ul>	<ul style="list-style-type: none"><li>• Treat backing services as attached resources:<ul style="list-style-type: none"><li>• Databases</li><li>• Messaging systems</li><li>• LDAP servers</li><li>• Others</li></ul></li><li>• Local and remote resources should be treated identically<ul style="list-style-type: none"><li>Possible locations for run time and resources:<ul style="list-style-type: none"><li>• In the same process</li><li>• On the same host</li><li>• On different hosts in the same data center</li><li>• In different data centers</li></ul></li></ul></li><li>• <b>IBM Cloud: All services are bound to the run time as though they are remote, including custom user-provided services</b></li></ul>
---	--

The Twelve-Factor Application      16      © Copyright IBM Corporation 2017-2018

IBM Training IBM

## Backing services: Cloud Foundry service bindings

```
cf bind-service appname service_instance
```

```
cf push traderback
cf bind-service traderback stsql
cf bind-service traderback stmessaging
```

- When deploying, do not manually bind services to an application
- Specify the services that an application needs in its manifest file

The Twelve-Factor Application      17      © Copyright IBM Corporation 2017-2018

IBM Training IBM

## Backing services: Kubernetes bindings

- Bind a container in Kubernetes to an IBM Cloud service
  - Add a volume to the pod definition that stores the binding in a secret
  - By mounting the Kubernetes secret as a volume to your deployment, you make the IBM Cloud service credentials available to the container that is running in your pod.
- Example: Bind to an instance of Watson Tone Analyzer named mytoneanalyzer

```
volumes:
  - name: service-bind-volume
    secret:
      defaultMode: 420
      secretName: binding-mytoneanalyzer
```

The Twelve-Factor Application      18      © Copyright IBM Corporation 2017-2018

IBM Training IBM

### Factor 5: Build, release, run

- I. Codebase
- II. Dependencies
- III. Configuration
- IV. Backing services
- V. Build, release, run**
- VI. Processes
- VII. Port binding
- VIII. Concurrency
- IX. Disposability
- X. Dev/prod parity
- XI. Logs
- XII. Admin processes

- Strictly separate build and run stages
- IBM Cloud: The output of build and release is an immutable container
  - In the IBM Cloud Kubernetes Service, the Docker build creates a container image that is stored to the private image registry.
  - You do not have to build the container again just to deploy another one

The Twelve-Factor Application      19      © Copyright IBM Corporation 2017-2018

IBM Training IBM

### Build, release, run

- Build stage
  - converts a code repo into an executable bundle known as a build
  - fetches dependencies and compiles binaries and assets
- Release stage
  - takes the build produced by the build stage and combines it with the deploy's current config
  - ready for immediate execution in the execution environment
- Run stage
  - runs the app in the execution environment

Source: <https://12factor.net>

The Twelve-Factor Application      20      © Copyright IBM Corporation 2017-2018

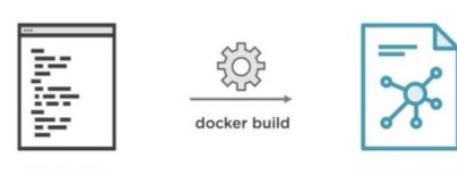
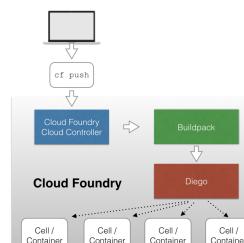
**IBM Training**

**Build & release produce an Immutable image**

- An immutable image does not get changed – only used for deploying instances
- If it needs to change, you delete it and create a new one

**Examples**

- Docker containers
  - A container image is created from a Dockerfile
  - After that, it is only deployed as a container, not changed
  - If you need to make changes, make a new container image by creating a new build and running the Dockerfile again
- Cloud Foundry
  - When an app is pushed to Cloud Foundry, the buildpack creates a droplet – from that CF creates Garden container instances that run the app
  - If you need to make changes, create a new build and push it

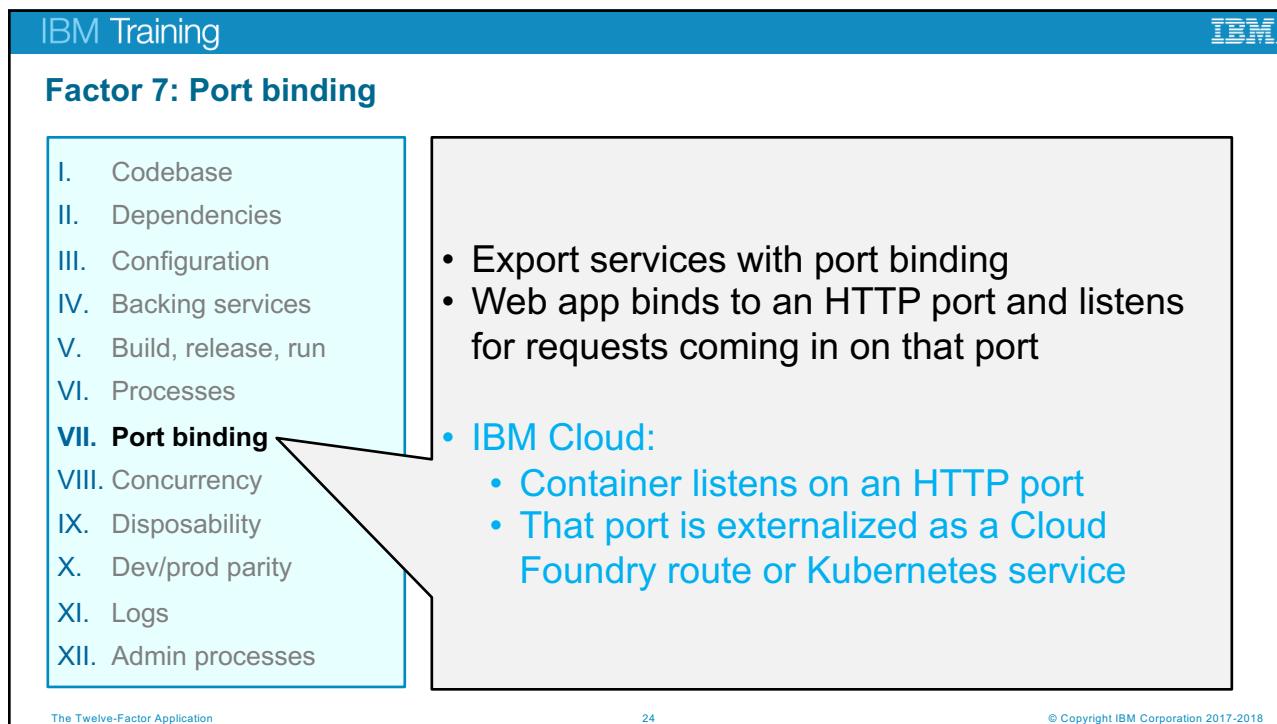
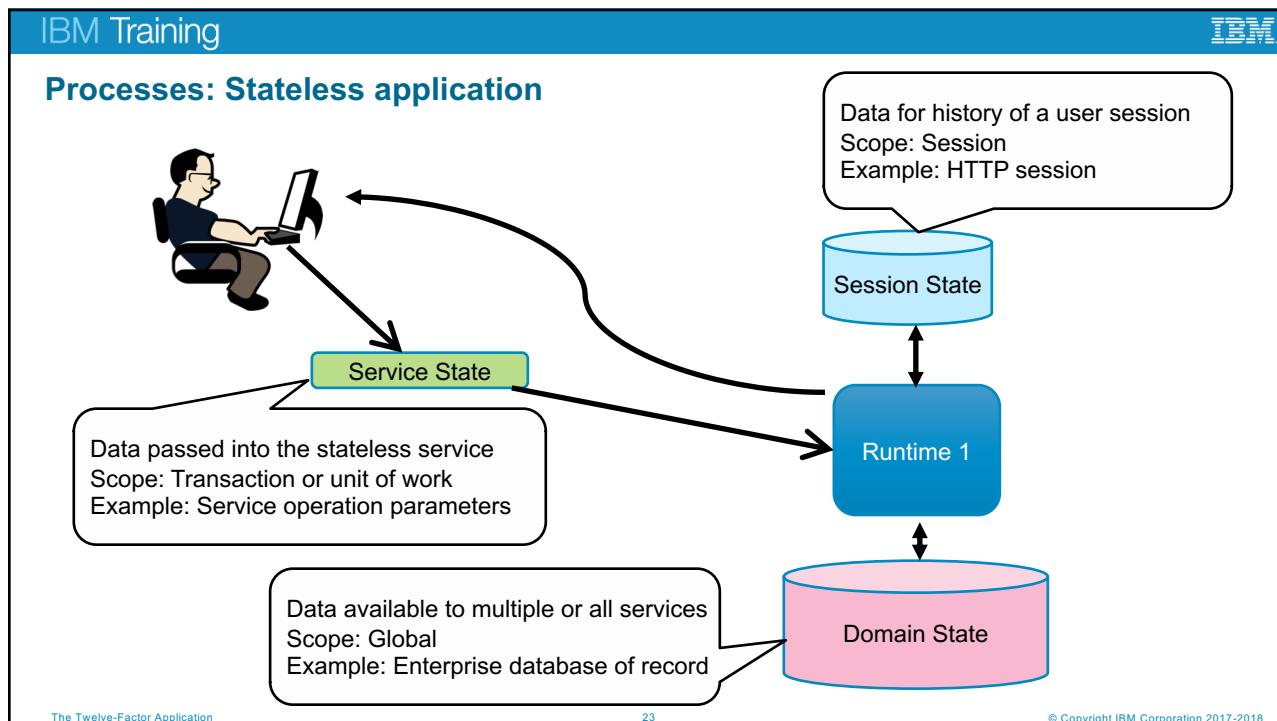
The Twelve-Factor Application      21      © Copyright IBM Corporation 2017-2018

**IBM Training**

**Factor 6: Processes**

I. Codebase II. Dependencies III. Configuration IV. Backing services V. Build, release, run <b>VI. Processes</b> VII. Port binding VIII. Concurrency IX. Disposability X. Dev/prod parity XI. Logs XII. Admin processes	<ul style="list-style-type: none"> <li>• Run the application as one or more stateless processes</li> <li>• Do not rely on session affinity, also called sticky sessions</li> <li>• <b>IBM Cloud: Application instances are stateless, with the state held by service instances</b></li> </ul>
--	---

The Twelve-Factor Application      22      © Copyright IBM Corporation 2017-2018



**IBM Training**

### Port binding: IBM Cloud application routes

The screenshot shows two parts of the IBM Cloud interface. On the left, a detailed view of an application named "order-toolchain-14890075..." shows its URL as "Orders-api-microservice.mybluemix.net". On the right, a list of "Cloud Foundry Apps" shows five entries, each with a red box around its URL column:

NAME	ROUTE
active-deploy-palistra-912_1	<a href="#">active-deploy-palistra-912.mybluemix.net</a>
BlueComputeSimple	<a href="#">bluecomputesimple.mybluemix.net</a>
dev-catalog-api-toolchain-de	<a href="#">dev-catalog-api-toolchain-demo-1488398990;</a>
dev-orders-api-toolchain-de	<a href="#">dev-orders-api-toolchain-demo-14883989903</a>
dev-ui-toolchain-demo-1488	<a href="#">dev-ui-toolchain-demo-1488398990350.myblu</a>

The bottom of the screen includes copyright information: "The Twelve-Factor Application" at the bottom left, "25" at the bottom center, and "© Copyright IBM Corporation 2017-2018" at the bottom right.

**IBM Training**

### Port binding: Kubernetes service

- Expose a set of pod replicas as a service
 

```
kubectl expose deployment/app-a
--type=LoadBalancer
--port=8080
--name=app-a-service
--target-port=8080
```

The diagram illustrates the Kubernetes architecture. It features a central "Master Node" connected to three "Worker Node" boxes. Each worker node contains multiple "Pod" boxes, each with an "App A" or "App B" component. Two "Service" boxes, "App A Service" and "App B Service", are shown. Blue lines connect the "App A Client" and "App B Client" to the "App A Service" and "App B Service" respectively, indicating how external traffic is directed through the Kubernetes components to the appropriate pods.

The bottom of the screen includes copyright information: "The Twelve-Factor Application" at the bottom left, "26" at the bottom center, and "© Copyright IBM Corporation 2017-2018" at the bottom right.

IBM Training IBM

## Factor 8: Concurrency

- I. Codebase
- II. Dependencies
- III. Configuration
- IV. Backing services
- V. Build, release, run
- VI. Processes
- VII. Port binding
- VIII. Concurrency**
- IX. Disposability
- X. Dev/prod parity
- XI. Logs
- XII. Admin processes

- Scale out using the process model  
To add capacity, run more instances
- There are limits to how far an individual process can scale
- Stateless applications make scaling simple
- IBM Cloud:
  - cf scale
  - Auto-Scaling service

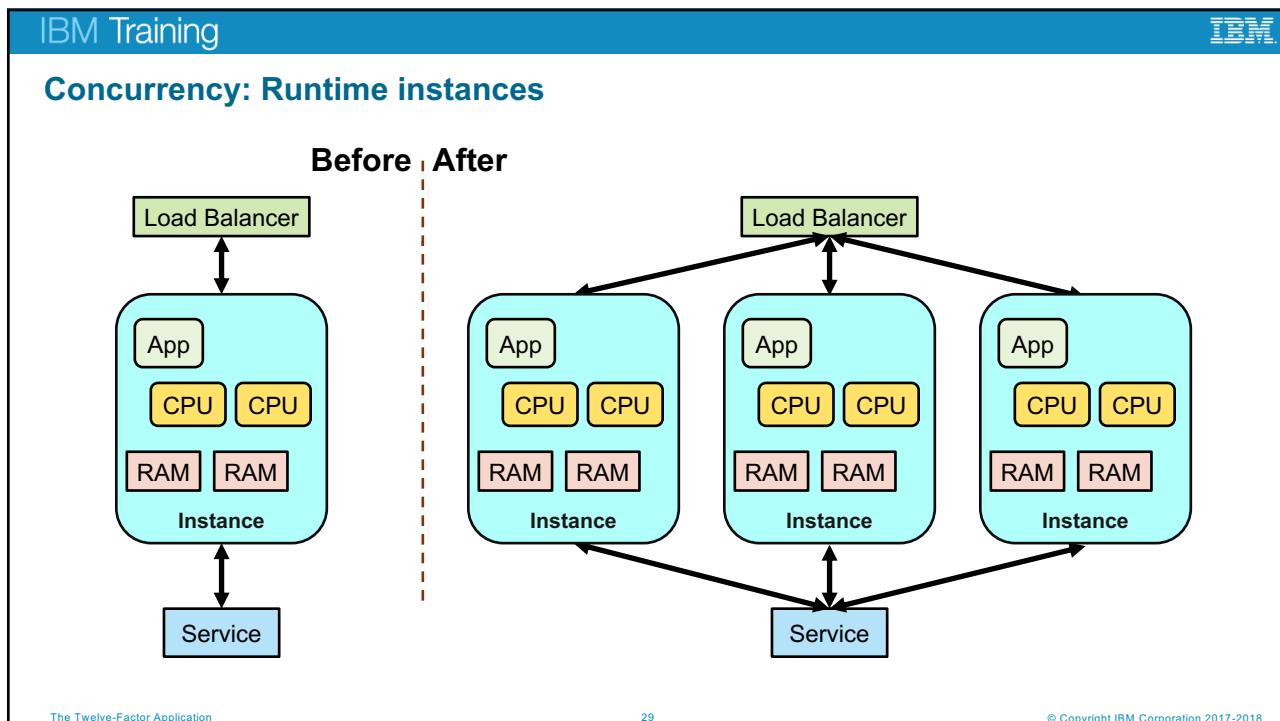
The Twelve-Factor Application      27      © Copyright IBM Corporation 2017-2018

IBM Training IBM

## Concurrency: Runtime instances

Instances				
STATUS	INSTANCE	CPU	MEMORY	DISK
Running	0	0.1%	67.4 MB / 96 MB	111 MB / 1 GB
Running	1	0.1%	66.4 MB / 96 MB	111 MB / 1 GB
Running	2	0.1%	66.2 MB / 96 MB	111 MB / 1 GB

The Twelve-Factor Application      28      © Copyright IBM Corporation 2017-2018



**cf scale**

```
cf scale appname -i number_of_instances
```

- Horizontal scaling in Cloud Foundry
  - Runs the specified number of instances of the application
  - Same as the Instances setting in the console
- Instances run independently
  - Stateless
  - Unrelated to each other
- For clustering
  - A router must distribute load
    - IBM Cloud automatically distributes across instances with the same route
  - If the app requires session state, it must be shared
    - For example, the instances can all share an instance of Redis or WebSphere eXtreme scale for session state
- The administrator manually adjusts the cluster size

The 'IBM Training' logo is at the top left, and the 'IBM' logo is at the top right. The bottom of the slide includes copyright information: 'The Twelve-Factor Application' on the left, '30' in the center, and '© Copyright IBM Corporation 2017-2018' on the right.

**IBM Training**

**Auto-Scaling service**



The Auto-Scaling service icon is a hexagon containing three stylized 3D cubes. One cube has a plus sign (+) on its front face, suggesting growth or addition. Another cube has a minus sign (-) on its front face, suggesting reduction or removal. The third cube is plain.

- Dynamically increases and decreases application capacity
  - Makes the application *elastic*
  - Number of instances grows and shrinks based on load
- Requires no manual intervention by the administrator
  - Manual: `cf scale`
  - Automated: Auto-Scaling service
- Policy-driven: *Auto-Scaling Policy*
  - Policy declaratively specifies the proper range for these items, among others:
    - JVM heap
    - Memory
    - Throughput
    - Response time
  - Automatically keeps capacity within range

The Twelve-Factor Application

31

© Copyright IBM Corporation 2017-2018

**IBM Training**

**Factor 9: Disposability**

I. Codebase  
II. Dependencies  
III. Configuration  
IV. Backing services  
V. Build, release, run  
VI. Processes  
VII. Port binding  
VIII. Concurrency  
**IX. Disposability**  
X. Dev/prod parity  
XI. Logs  
XII. Admin processes

Maximize robustness with fast startup and efficient shutdown  
Application instances are disposable  
Application should handle shutdown signal or hardware failure with crash-only design

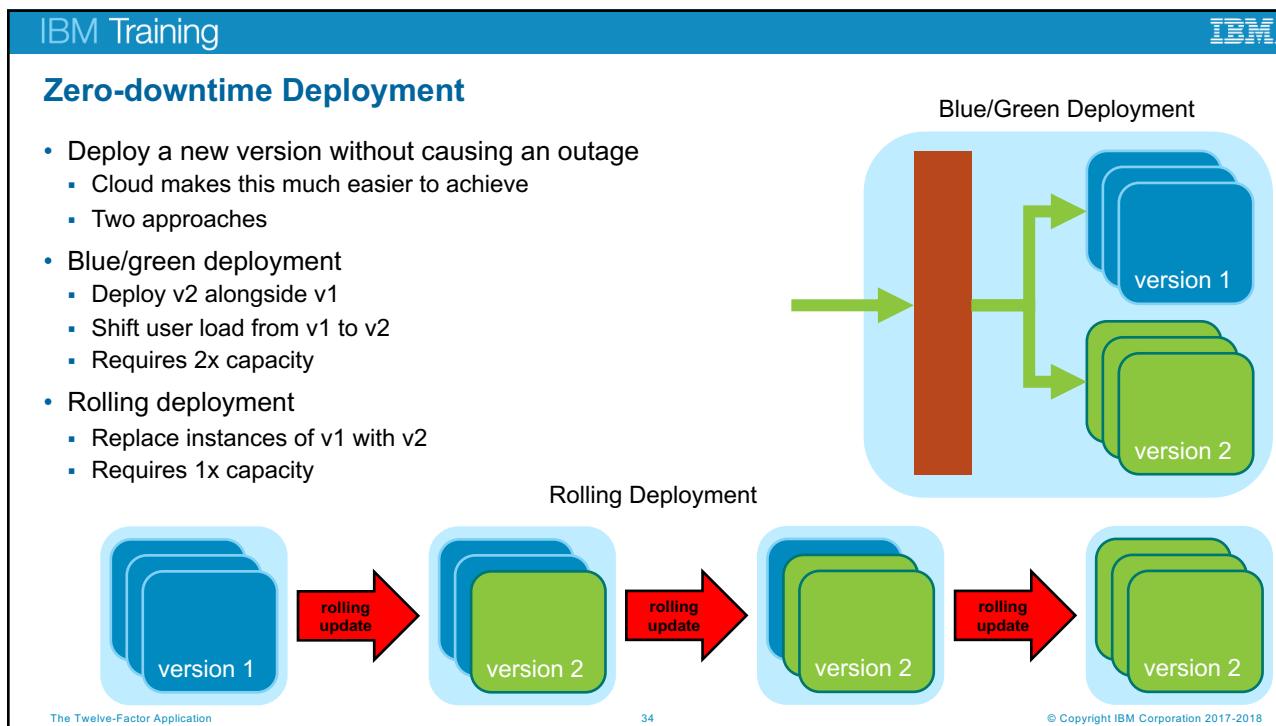
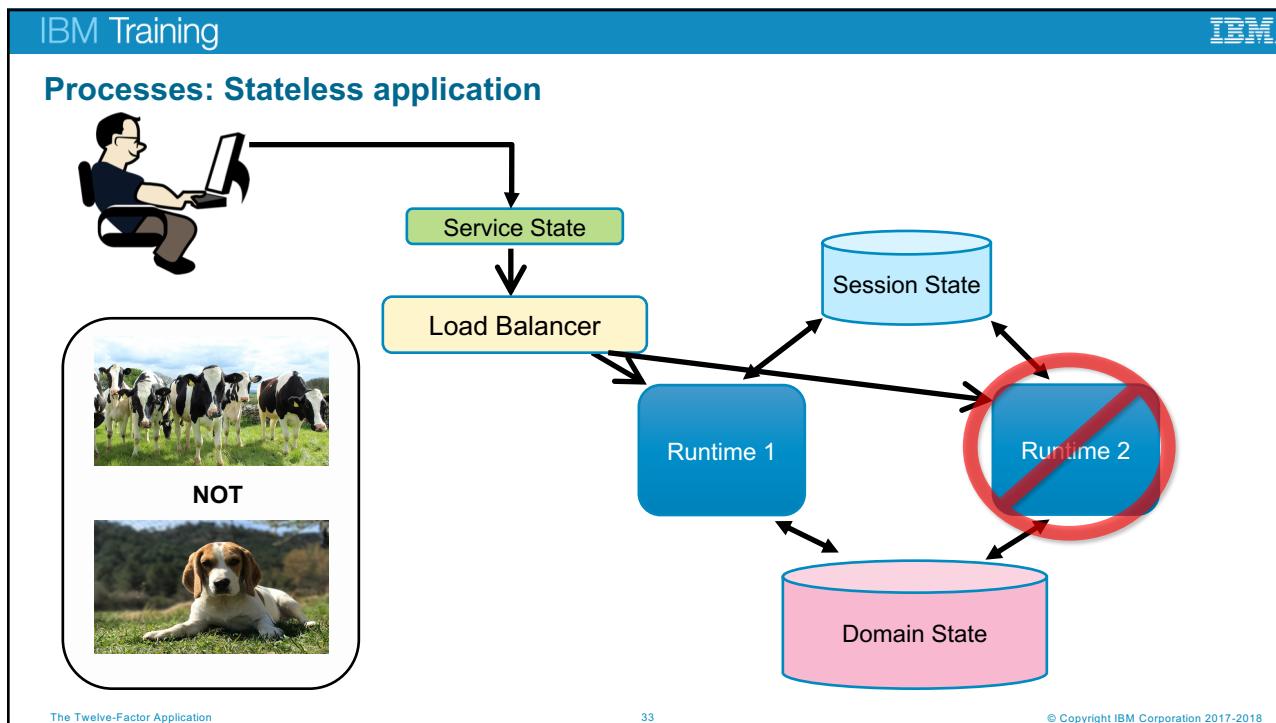
**IBM Cloud:**

- Architecture can rapidly start and stop instances
- Applications need to ensure they can respond

The Twelve-Factor Application

32

© Copyright IBM Corporation 2017-2018



**IBM Training**

**Cloud Foundry Blue-green deployment**

- A zero-downtime deployment technique
  - Also known as red-black Deployment
- Two nearly identical production environments, called Blue and Green
  1. Deploy v1 (blue environment), which users access
  2. Deploy v2 (green environment)
  3. Shift client load from Blue to Green
  4. Delete Blue

**cf rename**

- Rename Blue (v1) app to Green so that you can deploy Blue (v2)
- Both share Blue's route until you delete Green, which is v1

**cf map-route**

- Push Blue, which is v1, and Green, which is v2
- Map Blue's route to both Blue and Green
- Both share Blue's route until you delete Blue, which is v1

The Twelve-Factor Application      35      © Copyright IBM Corporation 2017-2018

**IBM Training**

**Factor 10: Development and production parity**

<ol style="list-style-type: none"> <li>I. Codebase</li> <li>II. Dependencies</li> <li>III. Configuration</li> <li>IV. Backing services</li> <li>V. Build, release, run</li> <li>VI. Processes</li> <li>VII. Port binding</li> <li>VIII. Concurrency</li> <li>IX. Disposability</li> <li>X. <b>Dev/prod parity</b></li> <li>XI. Logs</li> <li>XII. Admin processes</li> </ol>	<ul style="list-style-type: none"> <li>• Keep development, staging, and production environments as similar as possible</li> <li>• Use the same backing services in each environment</li>   <li>• <b>IBM Cloud:</b> <ul style="list-style-type: none"> <li>▪ Cloud Foundry - Use separate spaces for separate environments: (ie. Development, Test, Production)           <ul style="list-style-type: none"> <li>▪ Use different instances of the same services in each space</li> <li>▪ Deploy the same application to each space</li> </ul> </li> <li>▪ Kubernetes – Use different namespaces or other methods to maintain separation between environments in the same cluster</li> <li>▪ Resource Groups – Can be used to segregate services</li> </ul> </li> </ul>
--	---

The Twelve-Factor Application      36      © Copyright IBM Corporation 2017-2018

**IBM Training**

## Regions, organizations, spaces, users in Cloud Foundry

The diagram illustrates the hierarchical structure of Cloud Foundry entities:

- Region:** The top level contains two regions: "US South" and "United Kingdom".
- Organization:** Each region contains one organization. The "US South" region's organization is associated with the email "john@acme.com". The "United Kingdom" region's organization is also associated with the same email "john@acme.com".
- Space:** Each organization contains multiple spaces. The "john@acme.com" organization has four spaces: "dev", "test", "stage", and "prod".
- Users:** Each space is associated with a user. The "dev" and "test" spaces are associated with the user "john@acme.com". The "stage" and "prod" spaces are associated with the user "admin@acme.com".

Legend: Region (blue), Organization (green), Space (orange)

- **Region:** A *region* is a specific installation of IBM Cloud. IBM Cloud public has regions like US-South, UK-South, and AP-South.
- **Organization:** An *organization* has a quota of resources available for deploying and running applications. It can span regions. It can have one or more users.
- **Space:** A *space* is a group of runtime artifacts, like applications or services, hosted in a region and belonging to an organization. Can be used in a variety of ways to group artifacts (dev, test, production).
- **User:** A *user* will have assigned roles and permissions within an organization or space.

The Twelve-Factor Application      37      © Copyright IBM Corporation 2017-2018

**IBM Training**

## A Cloud Foundry Space can only contain Cloud Foundry Apps & Services

The diagram shows a single Cloud Foundry Space named "dev" within the "US South" region. This space contains the following components:

- CF App 1 (represented by a gear icon)
- CF Srv 1 (represented by a server icon)
- CF Srv 2 (represented by a database icon)
- CF App 2 (represented by a gear icon)

Contained in a single region

The Twelve-Factor Application      38      © Copyright IBM Corporation 2017-2018

**IBM Training**

**IBM**

### A Resource Group can contain any kind of resource

The diagram illustrates a Resource Group named "dev-rg" which spans two regions: "US South" and "United Kingdom". The "dev-rg" is represented by an orange rounded rectangle containing several resources:

- US South:** Contains "CF App 1" (represented by a blue icon), "Virtual Server" (represented by a white box with a cloud icon and an upward arrow), "Kubernetes Cluster" (represented by a blue cube icon), and "CF Srv 1" (represented by a white box with a blue server icon).
- United Kingdom:** Contains "Kubernetes Cluster" (represented by a blue cube icon), "CF App 2" (represented by a blue icon), "Virtual Server" (represented by a white box with a cloud icon and an upward arrow), and "CF Srv 2" (represented by a white box with a blue server icon).

**• Can Span regions**  
**• Users provided different levels of access**

The Twelve-Factor Application      39      © Copyright IBM Corporation 2017-2018

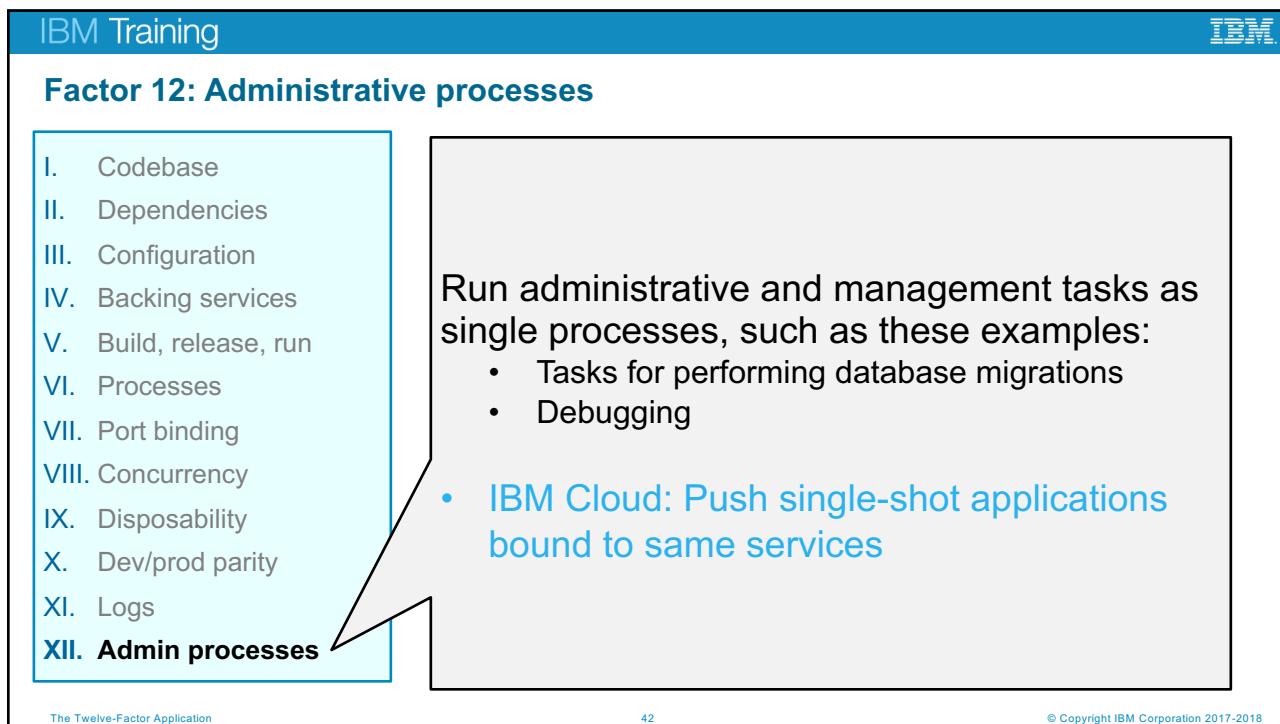
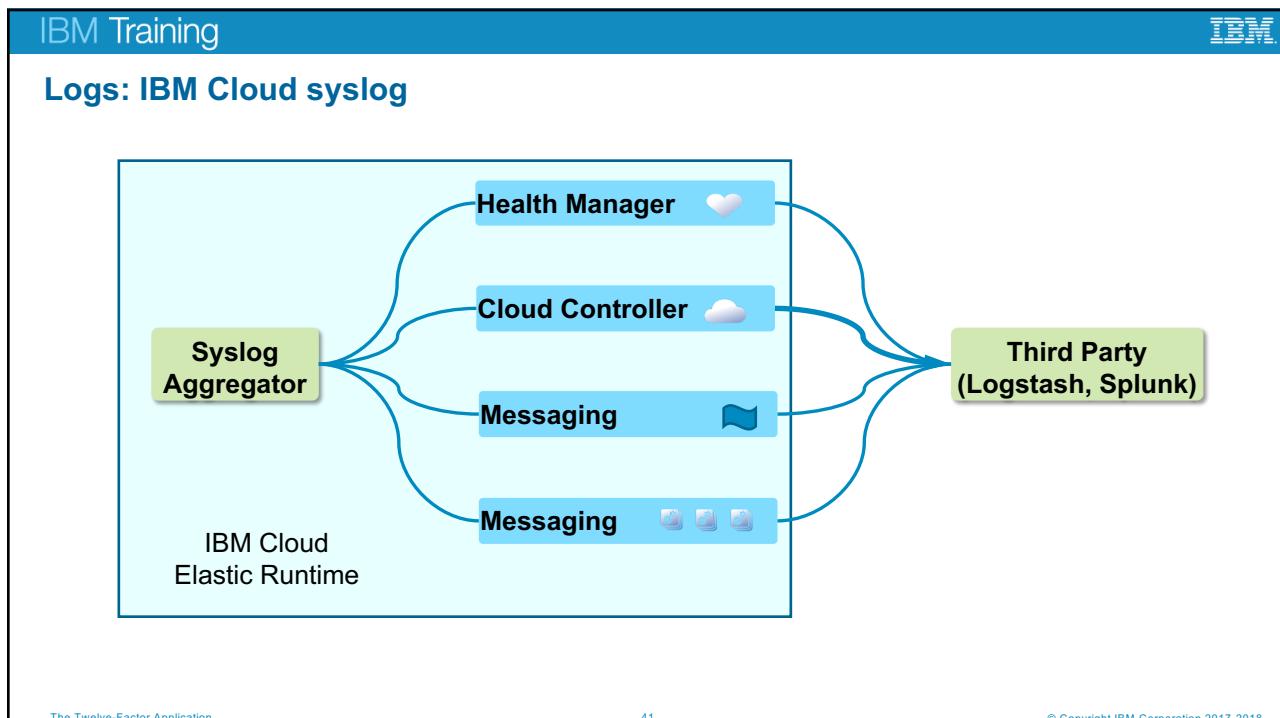
**IBM Training**

**IBM**

### Factor 11: Logs

<ol style="list-style-type: none"> <li>I. Codebase</li> <li>II. Dependencies</li> <li>III. Configuration</li> <li>IV. Backing services</li> <li>V. Build, release, run</li> <li>VI. Processes</li> <li>VII. Port binding</li> <li>VIII. Concurrency</li> <li>IX. Disposability</li> <li>X. Dev/prod parity</li> <li><b>XI. Logs</b></li> <li>XII. Admin processes</li> </ol>	<ul style="list-style-type: none"> <li>• Treat logs as event streams</li> <li>• Each process writes to stdout           <ul style="list-style-type: none"> <li>• Application should not write to specialized log files</li> <li>• Environment decides how to gather, aggregate, and persist stdout output</li> </ul> </li> <li>• <b>IBM Cloud</b> <ul style="list-style-type: none"> <li>• Syslog provides event streams for applications</li> <li>• Can be sent to third-party log management system</li> </ul> </li> </ul>
--	--

The Twelve-Factor Application      40      © Copyright IBM Corporation 2017-2018



IBM Training 

## Administrative processes: Scripting

- Create an SQL database
  - DDL script that creates the schema
  - SQL script that populates initial data
  - Script that runs these scripts as part of creating the database
- Migrate data to a new schema
- Software-defined data center (SDDC)
- Deployment scripts as part of a CI/CD pipeline
  - Creates service instances, deploys runtimes, and binds them
- Scripts are repeatable
  - Test scripts in stage environment
  - Scripts run the same in production
- Store scripts in software configuration management (SCM) system

The Twelve-Factor Application      43      © Copyright IBM Corporation 2017-2018

IBM Training 

## Unit summary

- Describe the twelve-factor app best practices for cloud applications
- Define and use IBM Cloud facilities that incorporate these best practices
- Describe how IBM Cloud incorporates these practices

The Twelve-Factor Application      44      © Copyright IBM Corporation 2017-2018

IBM Training

IBM

## Backup Charts

The Twelve-Factor Application

45

© Copyright IBM Corporation 2017-2018

IBM Training

IBM

## Factor 6 – Processes

### Restaurant Example

The Twelve-Factor Application

46

© Copyright IBM Corporation 2017-2018

IBM Training IBM

### Here's an example from real life – Your Favorite Restaurant

Notice how the person who takes your order doesn't always bring your food? Well...

The Twelve-Factor Application      47      © Copyright IBM Corporation 2017-2018

IBM Training IBM

### The Restaurant Model

uc Use Cases

System Boundary

Waiter: receive order, place order, confirm order

Client: eat food, pay

Chef: cook food, serve wine, drink wine

Cashier: accept payment

Order Food <<extend>> Order Wine  
Order Wine {if wine was ordered}

Serve Food <<extend>> Serve Wine  
Serve Wine {if wine was served}

Eat Food <<extend>> Drink Wine  
Drink Wine {if wine was consumed}

Pay for Food <<extend>> Pay for Wine  
Pay for Wine {if wine was consumed}

Think of your waiter as a Runtime

Each request you make is Service State

The food & drink you've ordered but haven't yet received is Session State

Your total order that will result in your bill is Domain State

The Twelve-Factor Application      48      © Copyright IBM Corporation 2017-2018

IBM Training 

## Stateless applications can be delicious



What if your waiter just tried to remember your whole order?

- What if he forgets?
- What if he gets too busy?
- What if he has to leave?



Guest Check  
0001  
Please call again  
Guest Receipt  
0001  
Please call again  
Remarks  
Thank You

Maintaining state externally is a better idea

- Paper check, electronic
- Available to chefs, other waiters
- Builds domain state

Now others can know what needs to be done and can help deliver the service



The Twelve-Factor Application 49 © Copyright IBM Corporation 2017-2018

IBM Training

IBM

## IBM Cloud Private architecture

© Copyright IBM Corporation 2018  
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

1

IBM Training

IBM

### Objectives

- Describe differences between IBM Cloud Private editions
- List IBM Cloud Private major components
- Differentiate kubernetes and cloud foundry environment for IBM Cloud Private

IBM Cloud Private architecture

2

© Copyright IBM Corporation 2018

IBM Training

## IBM Cloud Private – editions and packaging

IBM Cloud Private architecture

3

© Copyright IBM Corporation 2018

IBM Training

### Private cloud is critical to the IBM cloud strategy

**Choice with consistency**

**Hybrid integration**

**DevOps productivity**

**Powerful, accessible data and analytics**

**Cognitive solutions**

**PUBLIC**  
Maximize on cloud agility and economics

**DEDICATED**  
Public cloud benefits, with dedicated infrastructure

**PRIVATE**  
Behind the firewall for the most demanding workloads

**Seamless Experience**  
Regardless of which combination you choose, you get a single, seamless experience.

IBM Cloud Private architecture

4

© Copyright IBM Corporation 2018

**IBM Training**

## IBM Cloud Private – bringing cloud native to the enterprise

Rapid Innovation	Hybrid Integration	Investment Leverage	Management & Compliance
 <p><b>Rapid Innovation</b></p> <ul style="list-style-type: none"> <li>• Open Kubernetes-based container platform</li> <li>• Cloud Foundry for app dev and deployment</li> <li>• DevOps toolchain integration</li> </ul>	 <p><b>Hybrid Integration</b></p> <ul style="list-style-type: none"> <li>• Integration capabilities to unlock and connect</li> <li>• Secure access to public cloud services (AI, Blockchain)</li> <li>• Consistent experience across private/public</li> </ul>	 <p><b>Investment Leverage</b></p> <ul style="list-style-type: none"> <li>• Containerized versions of IBM Middleware</li> <li>• Prescriptive guidance to optimize workloads</li> <li>• Work with existing apps, data, skills, infrastructure</li> </ul>	 <p><b>Management &amp; Compliance</b></p> <ul style="list-style-type: none"> <li>• Core operational services including logging, monitoring, security</li> <li>• Flexibility to integrate with existing tools and processes</li> </ul>

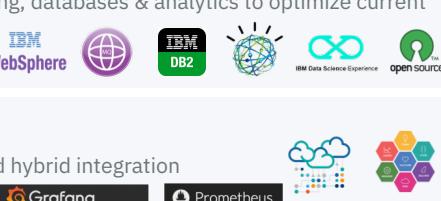
IBM Cloud Private architecture

5

© Copyright IBM Corporation 2018

**IBM Training**

## Enterprise grade content with open source technologies

 <p><b>IBM Middleware &amp; Open Source – e.g. Data, Analytics and Developer Services</b> Cloud-enabled middleware, application runtimes, messaging, databases &amp; analytics to optimize current investments and rapidly innovate</p>	 <p><b>Core Operational Services</b> To simplify Operations Management, Security, DevOps, and hybrid integration</p>	 <p><b>Kubernetes-based Container Platform</b> Industry leading container orchestration platform across private, dedicated &amp; public clouds</p>	 <p><b>Cloud Foundry</b> For prescribed application development &amp; deployment</p>	 <p><b>Terraform (CAM)</b> Infrastructure as Code for provisioning on public and on-prem cloud</p>
--	--	---	---	---

**Runs on existing IaaS (On Premise – e.g. Vmware, Openstack OR off Premise – e.g. SL, AWS) or hardware from IBM POWER and Z, Dell, Cisco, NetApp, Lenovo, ...**

IBM Cloud Private architecture

6

© Copyright IBM Corporation 2018

**IBM Training**

**Built with open standards, preventing vendor lock-in**

 <b>Containers</b> Executable package of software that includes everything needed to run it	 <b>Orchestration</b> Automate deployment, scaling, and management of containerized applications	 <b>Management</b> Define, install, and upgrade Kubernetes applications	 <b>Provisioning</b> Infrastructure as code to provision public cloud and on-premises environments
---	--	--	--

IBM Cloud | © 2018 IBM Corporation. All rights reserved. 7 © Copyright IBM Corporation 2018

**IBM Training**

**Solution Overview**

	<b>Enterprise Content Catalog</b> Open Source and IBM Middleware, Data, Analytics, and AI Software		
	<b>Core Operational Services</b> Log Management, Monitoring, Security, Alerting		<b>Strategic Values:</b>
	<b>Kubernetes Container Orchestration Platform</b>		Self-service catalog

Choose your infrastructure:

- Power Systems
- openstack
- IBM Spectrum
- intel
- vmware
- IBM Z

**Strategic Values:**

- Agility, scalability, and elasticity
- Self-healing
- Enterprise security
- No vendor lock-in

IBM Cloud Private architecture 8 © Copyright IBM Corporation 2018

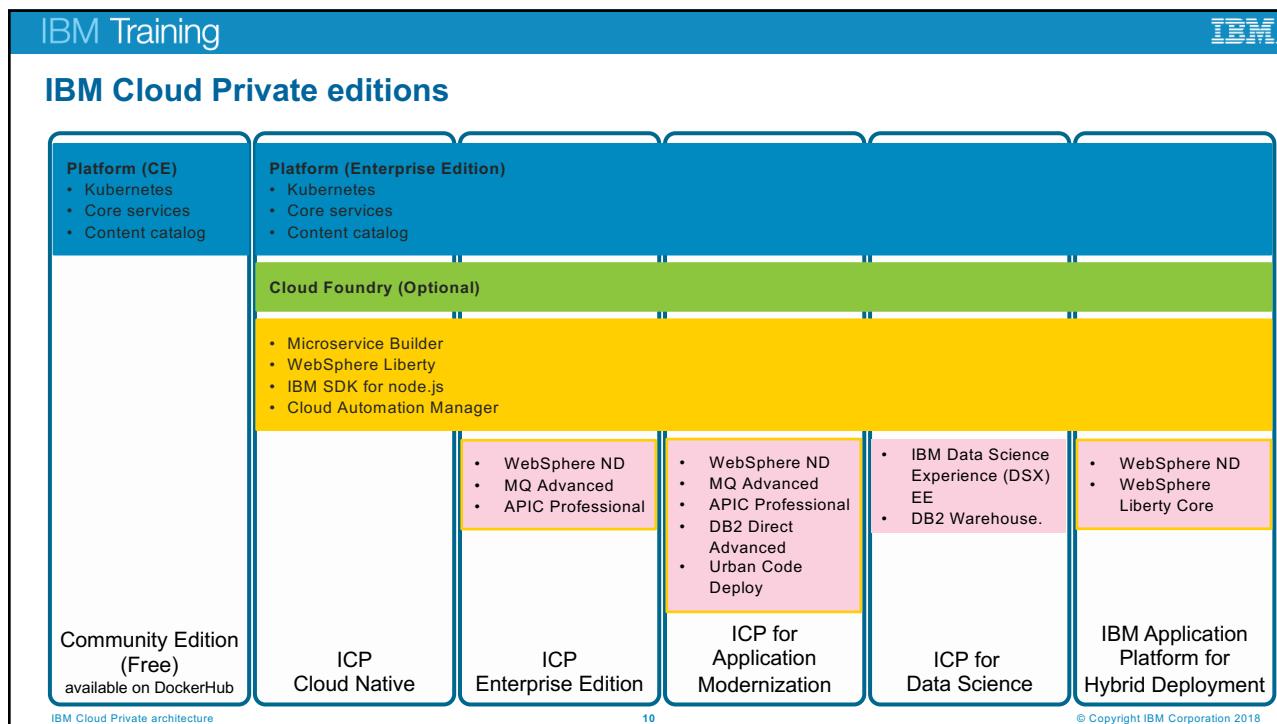
**IBM Training**

**Enterprise Content Catalog**

<b>Toolchain &amp; Runtimes</b> Microclimate Beta Microservice Builder Jenkins ( <a href="#">open source</a> ) IBM WebSphere Liberty Open Liberty ( <a href="#">open source</a> ) IBM SDK for Node.js Swift runtime ( <a href="#">open source</a> )	<b>Data Services</b> IBM Db2 Dev-C IBM Data Server Manager (for Db2 Dev-C) IBM Db2 Direct Advanced Edition / AESE with Data Server Manager IBM Db2 Warehouse Dev-C IBM Db2 Warehouse Enterprise IBM Cloudant Developer Edition MongoDB ( <a href="#">open source</a> ) PostreSQL ( <a href="#">open source</a> ) MariaDB ( <a href="#">open source</a> ) Galera with MariaDB ( <a href="#">open source</a> )	<b>Data Science and Business Analytics</b> IBM Data Science Experience Developer Edition IBM Data Science Experience Local
<b>Logging &amp; Monitoring Services</b> ELK ( <a href="#">open source</a> ) & Prometheus ( <a href="#">open source</a> )	<b>Messaging</b> IBM MQ Advanced for Developers IBM MQ Advanced Rabbit MQ ( <a href="#">open source</a> )	<b>Data Governance and Integration</b> IBM InfoSphere Information Server for Evaluation
<b>App Modernization Tooling</b> IBM Transformation Advisor	<b>Integration</b> IBM Integration Bus for Developers IBM Integration Bus IBM DataPower Gateway for Developers IBM DataPower Gateway Virtual Edition	<b>Cognitive Connectivity</b> IBM Voice Gateway Developer Trial
<b>Multi-cloud Management</b> IBM Cloud Automation Manager		<b>Tooling</b> Web Terminal ( <a href="#">open source</a> ) Skydive – network analyzer ( <a href="#">open source</a> )
<b>Mobile</b> IBM Mobile Foundation		<b>HPC / HPDA</b> IBM Spectrum LSF Community Edition IBM Spectrum Symphony Community Edition
<b>Digital Business Automation</b> IBM Operational Decision Manager for Developers		<b>Bring your own</b> Add a Helm repo or your own charts

*Catalog content is not distributed with IBM Cloud Private. Content is distributed separately, licensed under separate terms and conditions.*

IBM Cloud Private architecture © Copyright IBM Corporation 2018



**IBM Training**

**Cloud Foundry and Kubernetes**

The diagram illustrates the differences between Cloud Foundry and Kubernetes deployment models. It shows two side-by-side boxes: one for Cloud Foundry (left) and one for Docker (right). Both boxes have a dashed border labeled 'AND' between them.

**Cloud Foundry:**

- Developer brings app
- Platform provides standard Buildpack for runtime\*
- Platform provides fixed OS container image
- Platform provides fixed host OS Kernel

\*Node.js, Tomcat, Pearl, Ruby on Rails ...

**Docker:**

- Developer brings app
- Developer brings runtime Docker image
- Developer brings Docker OS image
- Platform provides fixed host OS Kernel

**kubernetes**

IBM Cloud Private architecture

- **CF and Kubernetes (K8s) are operating on different levels of abstraction**
- **Choices to make**
  - K8s: flexibility to control all of the underlying technology, and freedom to deploy whatever - wherever at large scale – Optimizing for Performance & Scalability.
  - CF: the platform takes care of all of the "plumbing" to get the code you wrote into a running application, monitor its health, and scale at the expense of control and flexibility – Optimizing for Speed & Simplicity
- **IBM's recommendation – Use Both platforms at their best**
  - Do rapid prototyping on Cloud Foundry PaaS
  - Harden prototypes & deploy on an Enterprise Grade K8 platform

**IBM Training**

**IBM Cloud Private – Kubernetes cluster**

IBM Cloud Private architecture

12

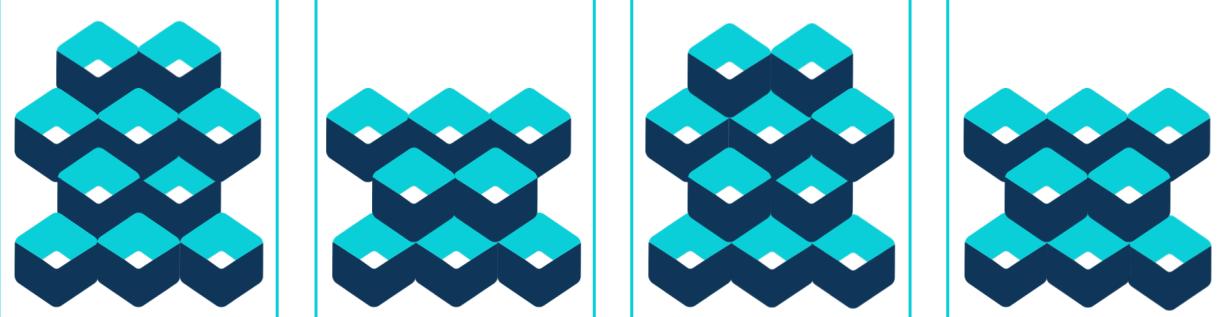
© Copyright IBM Corporation 2018

IBM Training IBM

## Kubernetes – (Κυβερνήτης - Captain in Greek)

Regain control with Containers and Kubernetes

- Organize and Govern the Container Chaos



IBM Cloud Private architecture

IBM Training IBM

## What does Kubernetes offer?

Intelligent Scheduling	 Automatically places containers based on their resource requirements and other constraints, while not sacrificing availability. Mix critical and best-effort workloads in order to drive up utilization and save even more resources.
Self Healing	 Restarts containers that fail, replaces and reschedules containers when nodes die, kills containers that don't respond to your user-defined health check, and doesn't advertise them to clients until they are ready to serve.
Horizontal Scaling	 Scale your application up and down with a simple command, with a UI, or automatically based on CPU usage.
Service Discovery and Load Balancing	 No need to modify your application to use an unfamiliar service discovery mechanism. Kubernetes gives containers their own IP addresses and a single DNS name for a set of containers, and can load-balance across them.
Automated rollout and rollback	 Kubernetes progressively rolls out changes to your application, while monitoring application health to ensure it doesn't kill all your instances at the same time. If something goes wrong, Kubernetes will rollback the change for you. Take advantage of a growing ecosystem of deployment solutions.
Secret and configuration management	 Deploy and update secrets and application configuration without rebuilding your image and without exposing secrets in your stack configuration.

IBM Cloud Private architecture

**IBM Training**

## And HELM is ...



**The package manager for Kubernetes**

**Helm is the best way to find, share, and use software built for Kubernetes.**

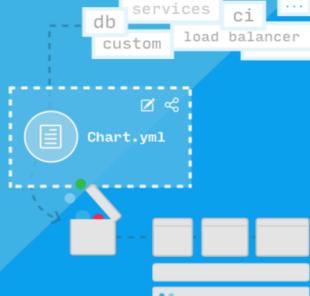
**Tells Kubernetes all it needs to know about an application its parameters and dependencies**

**What is Helm?**

Helm helps you manage Kubernetes applications — Helm Charts helps you define, install, and upgrade even the most complex Kubernetes application.

Charts are easy to create, version, share, and publish — so start using Helm and stop the copy-and-paste madness.

The latest version of Helm is maintained by the CNCF - in collaboration with Microsoft, Google, Bitnami and the Helm contributor community.

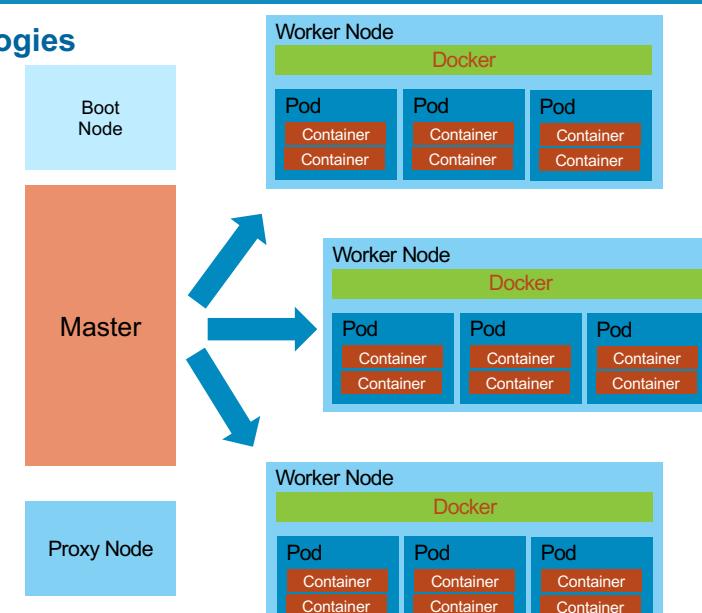


IBM Cloud Private architecture

**IBM Training**

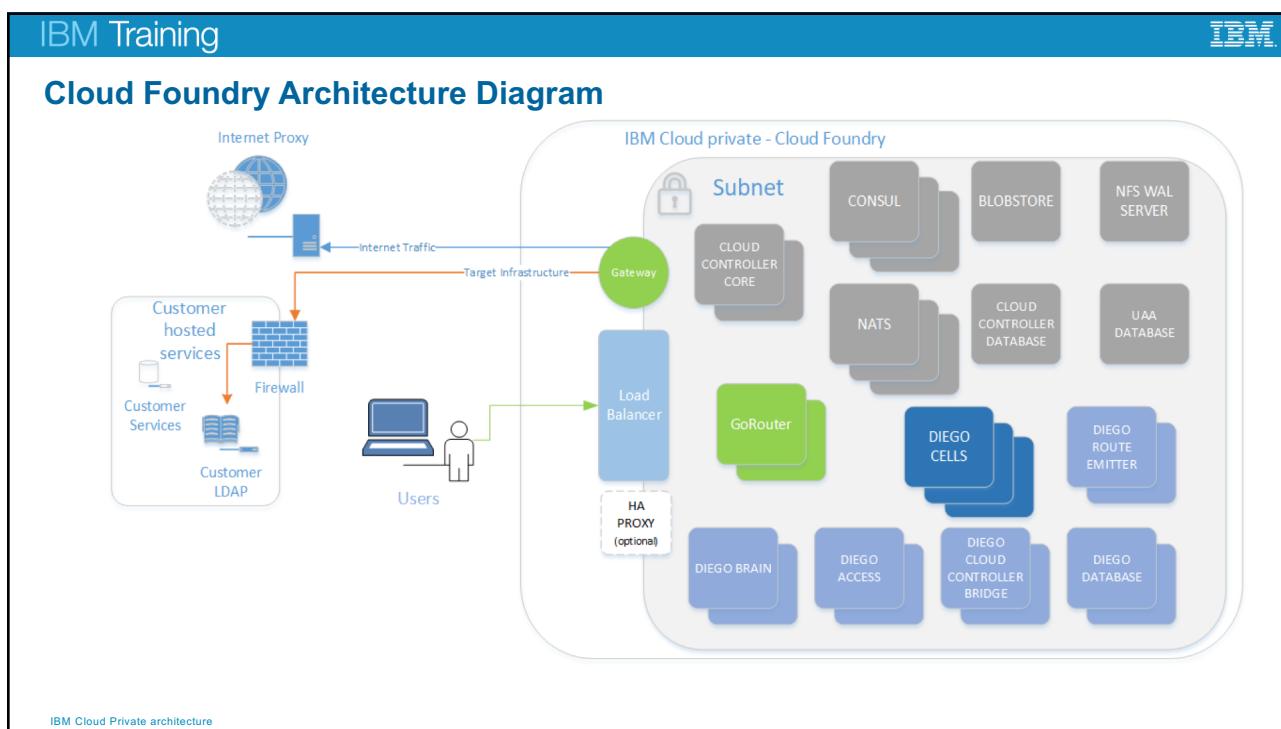
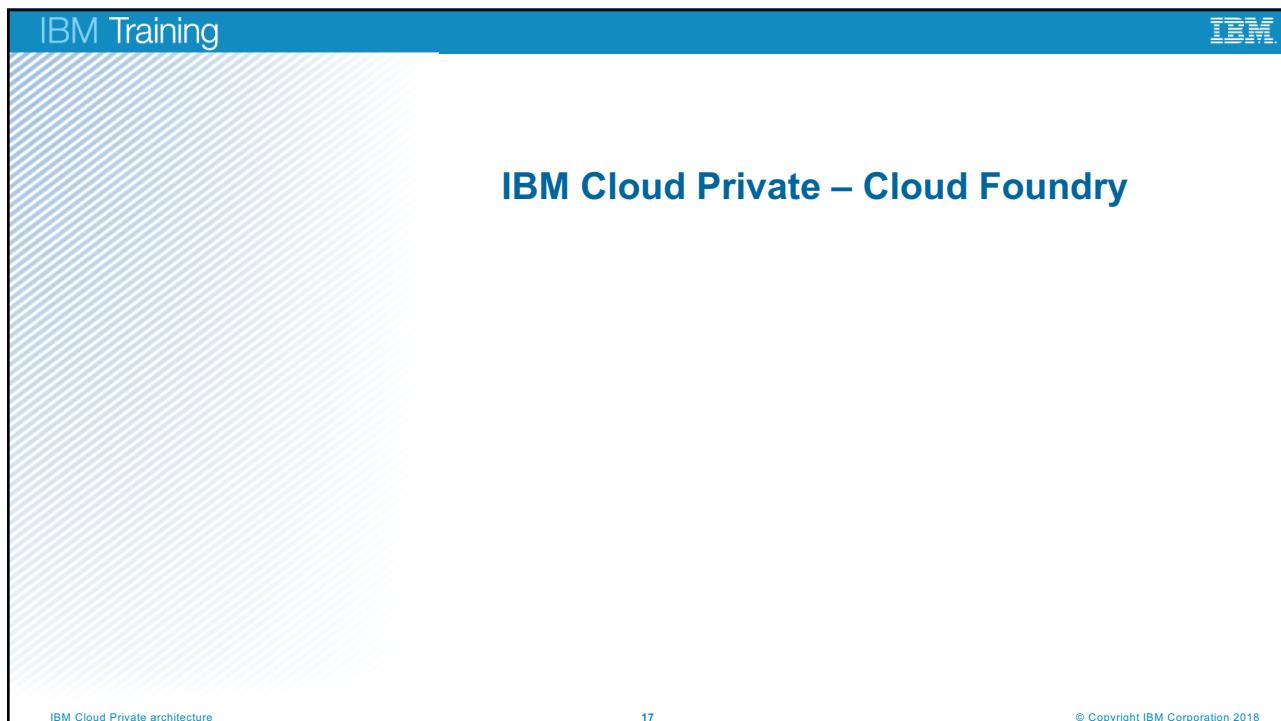
## Kubernetes Configuration topologies

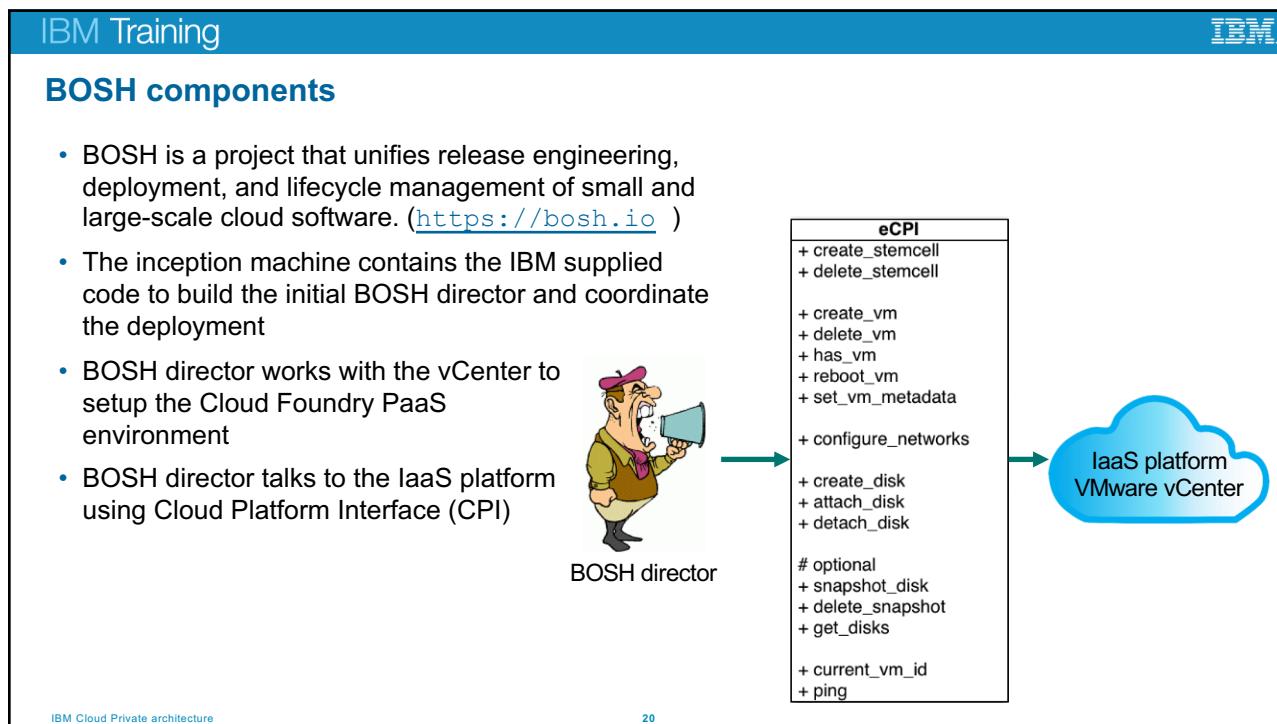
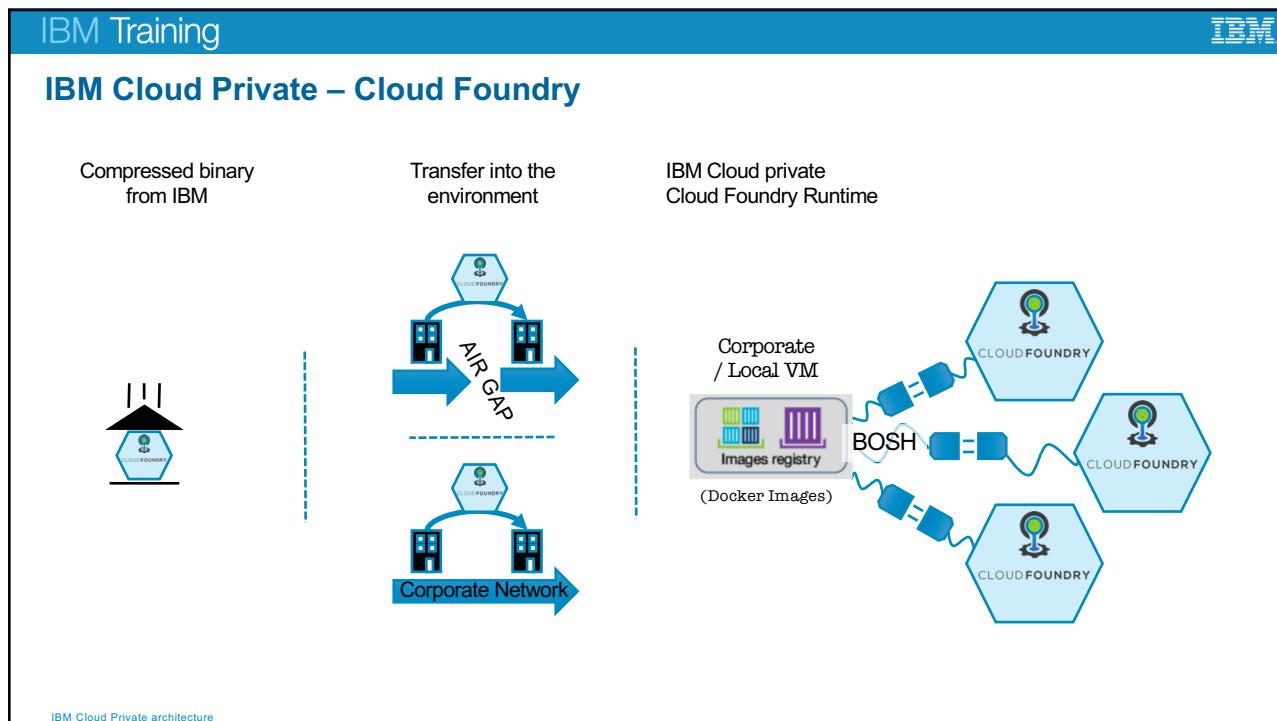
- Simple
  - Single machine install (master is a worker)
  - Great for testing and learning about the platform
- Standard
  - Single master (single master, 3 workers, 1 proxy)
  - Great for non-production testing environment
- High Availability
  - Multiple masters (3 masters, 3+ workers, 3 proxy)
  - Production installation

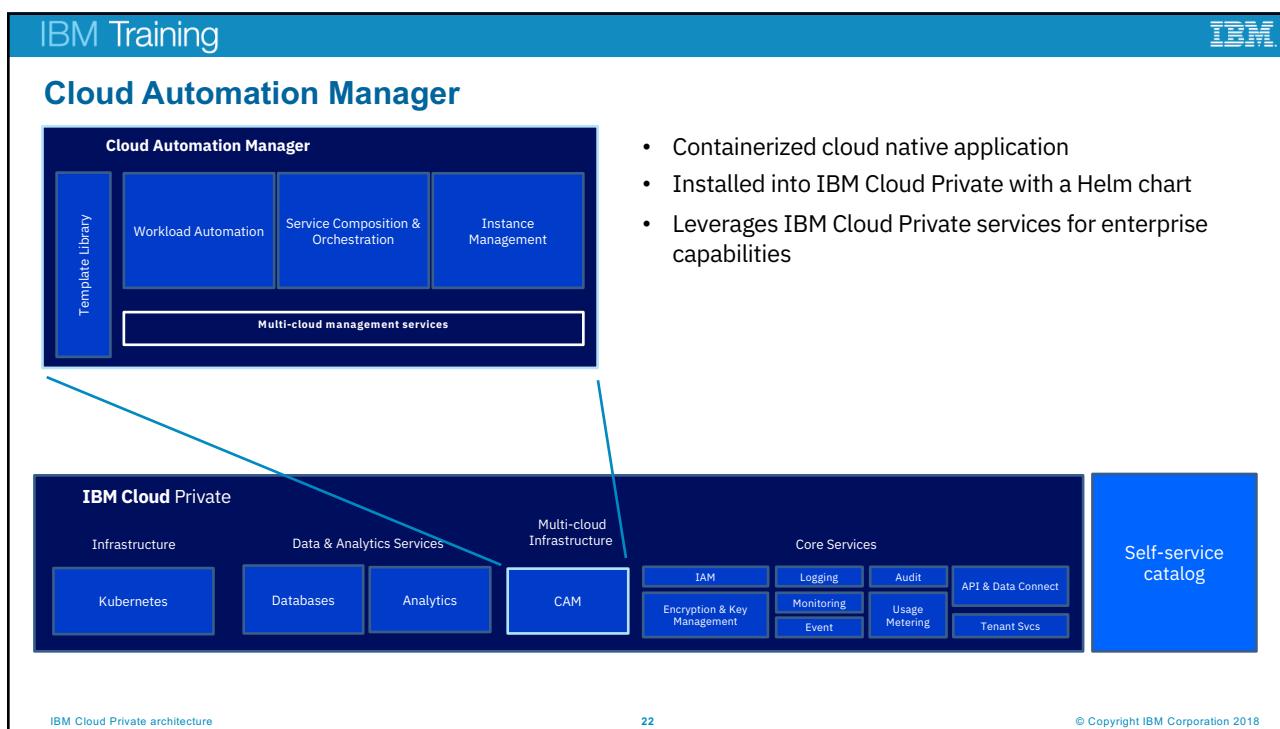
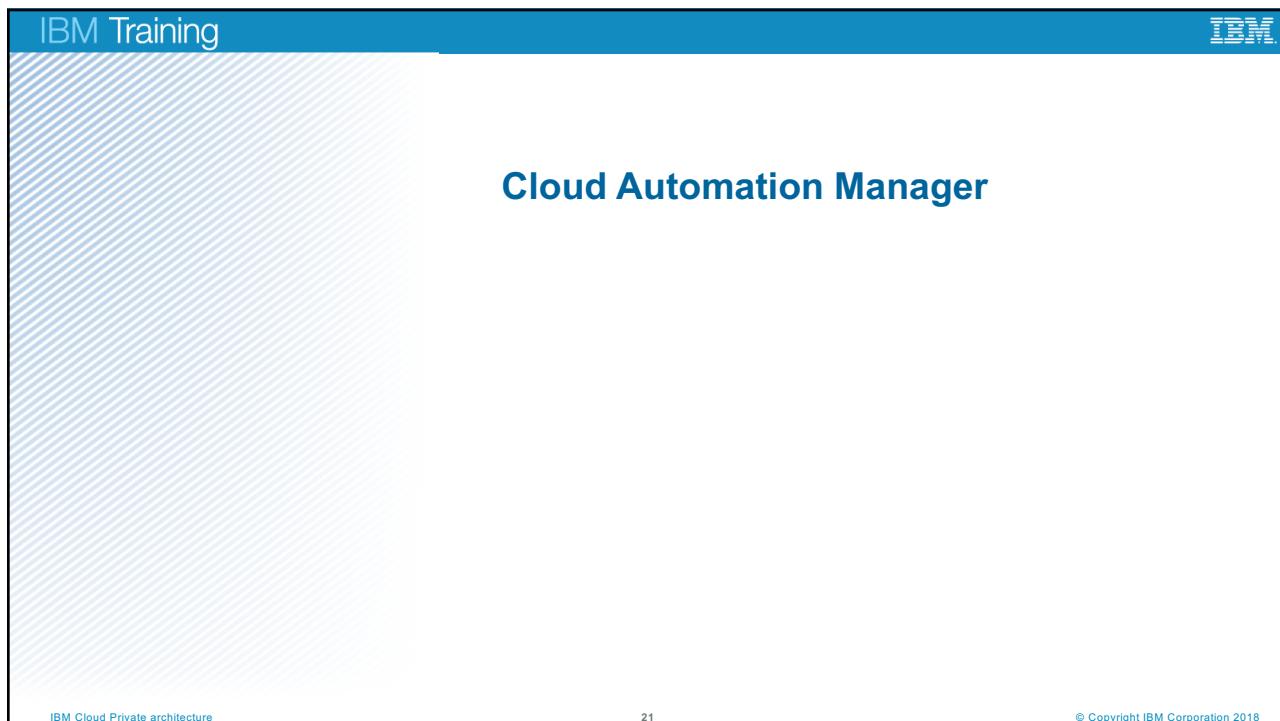


IBM Cloud Private architecture

© Copyright IBM Corporation 2018

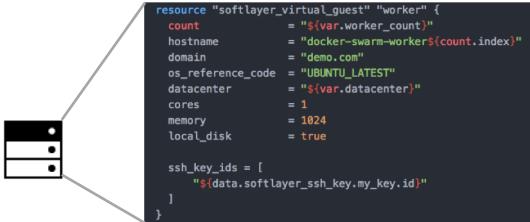






IBM Training IBM

## Terraform



```
resource "softlayer_virtual_guest" "worker" {  
    count      = "${var.worker_count}"  
    hostname   = "docker-swarm-worker${count.index}"  
    domain     = "demo.com"  
    os_reference_code = "UBUNTU_LATEST"  
    datacenter  = "${var.datacenter}"  
    cores       = 1  
    memory     = 1024  
    local_disk  = true  
  
    ssh_key_ids = [  
        "${data.softlayer_ssh_key.my_key.id}"  
    ]  
}
```

**Terraform** is an infrastructure as code software. It allows users to define a datacenter infrastructure in a high-level configuration language, from which it can create an execution plan to build the infrastructure in a service provider such as IBM SL , AWS, Google, MSFT as well as Open Stack & VMWare



IBMCLOUD Microsoft AWS openstack vmware



IBM Cloud Private architecture

IBM Training

IBM

## Introduction to Cloud Full Stack Development

Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

IBM Training

IBM

### Agenda

- Introduction
- Reference architecture
- Architectural components
- Development process

Introduction to IBM Cloud Full Stack Development

2

© Copyright IBM Corporation 2017

IBM Training

IBM

## Introduction

3

© Copyright IBM Corporation 2017

IBM Training

IBM

## Challenge

- How should I develop an application in IBM Cloud?
- This question can be challenging to answer
  - Every solution is custom and unique
  - IBM Cloud has many available services, but not all services are used by all applications
    - 80 / 20 rule
- There is an easier way
  - The basics of most applications running in IBM Cloud are the same
  - You don't need to learn all of IBM Cloud to implement applications
  - Learn the general architectural view and add other components/services as needed

Introduction to IBM Cloud Full Stack Development

4

© Copyright IBM Corporation 2017

IBM Training 

## Solution

- **Cloud Full Stack Development**
  - How to develop a typical application in IBM Cloud
    - End-to-end (a.k.a. full stack): From client GUIs through server app to backends
  - Reference architecture shows a typical, complete but basic application in IBM Cloud
  - Captures the sweet spot: 20% of IBM Cloud Public product and skills that are used 80% of the time
- Forms the common foundation for more advanced applications in IBM Cloud

Introduction to IBM Cloud Full Stack Development 5 © Copyright IBM Corporation 2017

IBM Training 



## Reference Architecture

6 © Copyright IBM Corporation 2017

**IBM Training**

**Cloud Full Stack Reference Architecture**

- Core architecture
  - For typical, basic applications in IBM Cloud
- Five layers
  - Clients
    - Web, mobile, CLI, etc.
  - API gateway (API Connect)
  - Business logic
  - Backends
    - Cloudant, on-prem resources, etc.
  - DevOps
- Three business logic architectures
  - Monolith
  - Dispatchers + Monolith
  - Microservices
- Two business logic technologies (in IBM Cloud Public)
  - Cloud Foundry
  - Docker containers in Container Service

The diagram illustrates the Cloud Full Stack Reference Architecture. At the top, three client boxes (Web App, Mobile App, CLI App) connect via 'Public REST Endpoints' to an 'API Connect' layer. This layer then connects via 'Private REST Endpoints' to a central 'Business Logic' box. Below 'Business Logic' are two 'Cloudant' databases. From each Cloudant database, a line goes to a 'Secure Gateway' (left) or 'IBM Cloud' (right). The 'Secure Gateway' leads to a 'Database of Record' (left) or an 'Adapter' (right), which then connects to a 'System of Record' (left) or a 'Data Center' (right). On the far left, a 'DevOps' box contains 'Continuous Delivery' and 'Monitoring' circles. A legend at the bottom identifies symbols: a circle for 'Client', a rectangle for 'Service', a cylinder for 'Data Store', and an oval for 'Adapter'.

Introduction to IBM Cloud Full Stack Development

7

© Copyright IBM Corporation 2017

**IBM Training**

**Business Logic Architecture: Monolith**

- Business logic architecture #1
  - Monolith
- REST API endpoints are all implemented by a single set of code running in a single process

The diagram shows a monolithic business logic architecture. It features a large, rounded rectangular box labeled 'Monolith' in the center. Inside this box, the words 'Business Logic' are written. Above the 'Monolith' box, there is a horizontal line with three circular endpoints, representing 'Private REST Endpoints'.

Introduction to IBM Cloud Full Stack Development

8

© Copyright IBM Corporation 2017

**IBM Training**

**Business Logic Architecture: Dispatchers + Monolith**

- Business logic architecture #2
  - Dispatchers + Monolith
- Each REST API endpoint is implemented by a separate dispatcher
  - I.E. Backend-for-frontend (BFF)
- All dispatchers route to a single set of code running in a single process

Introduction to IBM Cloud Full Stack Development

9

© Copyright IBM Corporation 2017

**IBM Training**

**Business Logic Architecture: Microservices**

- Business logic architecture #3
  - Microservices
- Each REST API endpoint is implemented by a separate dispatcher (BFF)
- Dispatchers route to a combination of business services
  - I.E. Single-task processes
- Also teaches microservices discovery
  - Netflix OSS, Amalgam8, ISTIO

Introduction to IBM Cloud Full Stack Development

10

© Copyright IBM Corporation 2017

IBM Training

IBM

The diagram illustrates a layered architectural model. At the top, three rectangular boxes represent external interfaces or gateways. Below them is a horizontal layer of four rectangular boxes, likely representing a service mesh or a set of microservices. This is followed by a layer of two circular nodes, possibly databases or message brokers. The bottom layer consists of two rectangular boxes, which could be application components or persistence layers.

**Architectural Components**

11 © Copyright IBM Corporation 2017

IBM Training

IBM

## Business Logic

- This is the custom code (Java, PHP, etc.) that makes your app unique
  - In IBM Cloud Public (or any PaaS), custom code is actually a much smaller part of the overall app
  - A lot of the behavior is implemented in (backing) services
    - What used to be the responsibility of middleware, such as the WebSphere stack products
- The app still starts with business logic
  - What do you want to do that makes this app unique?
- But let's first discuss the other parts of the architecture
  - How business logic fits into the rest of the architecture

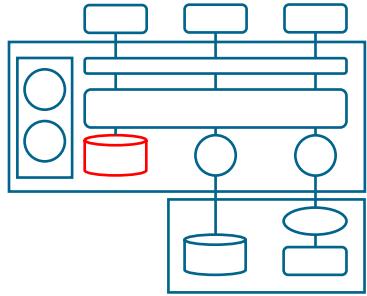
The same architectural diagram as above, but with a red rectangular box highlighting the middle layer of four rectangular boxes. This visual emphasizes that business logic is integrated within the core service mesh or microservices layer.

Introduction to IBM Cloud Full Stack Development 12 © Copyright IBM Corporation 2017

IBM Training 

## Backend: Cloudant NoSQL Database

- Pretty much every application needs to persist data
- NoSQL databases are more cloud-friendly than SQL databases
  - Horizontal scalability: better reliability and throughput
  - Documents: more natural than records, tables, and O/R mapping
- IBM Cloud catalog contains about 15 databases
  - An IBM Cloud developer doesn't need to learn them all
- Cloudant is a great general-purpose OLTP NoSQL database
  - Based on CouchDB, with lots of additional features
  - Available in IBM Cloud Public, Dedicated, and Private
- The first database cloud developers should learn is Cloudant

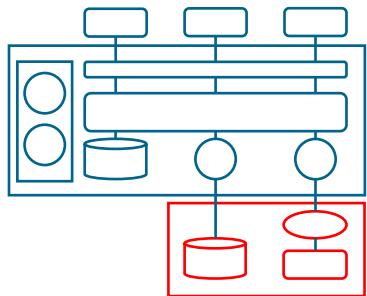


Introduction to IBM Cloud Full Stack Development 13 © Copyright IBM Corporation 2017

IBM Training 

## Data Center: Existing on-prem resources

- Many IBM and Enterprise customers have existing data centers
  - This on-premises technology is Traditional IT, not Cloud
- Existing systems that cannot easily be moved
  - Enterprise database of record
  - System of record
- Protected behind firewalls
  - Cannot be accessed publicly
  - Customer is unwilling to expose publicly
- IBM Cloud applications need to be able to access these existing systems
  - IBM Cloud Public and Dedicated are public, yet need private access



Introduction to IBM Cloud Full Stack Development 14 © Copyright IBM Corporation 2017

**IBM Training**

**Backend: Secure Gateway**

- IBM Cloud Public applications need to be able to access on-prem existing systems
  - IBM Cloud Public and Dedicated are public, yet need private access
- Secure Gateway service (IBM Cloud Public and Dedicated)
  - A quick, easy, and secure solution for connecting anything to anything
  - Gateway – Runs in IBM Cloud
  - Gateway client – Runs in the customer's data center (or other target)
  - Gateway establish a secure, persistent connection between client and IBM Cloud
  - Granular system-level access, not coarse network-level access like a VPN
- Use Secure Gateway to access on-prem systems and databases of record

Introduction to IBM Cloud Full Stack Development

15

© Copyright IBM Corporation 2017

**IBM Training**

**Clients: Web, Mobile, CLI, etc.**

- Clients enable users to access the server application
  - Clients provided by the server team or business partners
- Clients are multimodal – Different ways to access the same experience
  - Web – Clients that run in Web browsers – HTML 5, CSS 3, JavaScript
  - Mobile – Clients that run on phones and tablets – iOS, Android, etc.
  - CLI – Clients written by business partners, often complete applications
- Use clients to provide GUIs for server applications

Introduction to IBM Cloud Full Stack Development

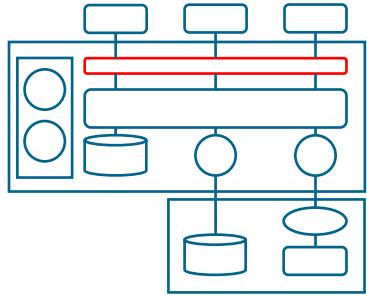
16

© Copyright IBM Corporation 2017

IBM Training 

## API Gateway: API Connect

- A gateway enables external clients to connect to internal server resources
  - Such connections need to be secured to block client attacks
- API Connect service
  - Gateway – Enforces policies that control access to services' APIs
  - Catalog – Lists available services, manages their lifecycle
  - Adapter – Connects gateway to services with inadequate APIs
- Use API Connect as an API gateway to control clients' access to services

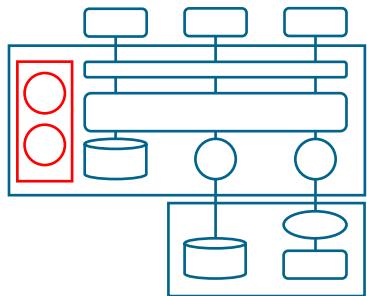


Introduction to IBM Cloud Full Stack Development 17 © Copyright IBM Corporation 2017

IBM Training 

## DevOps: Continuous Delivery and Monitoring

- DevOps ensures the application's quality
  - Delivered efficiently
  - Running properly
- Continuous Delivery service
  - A continuous integration/continuous delivery pipeline (CI/CD)
  - Each code change rapidly packaged, tested, and deployed into production
- Monitoring and Analytics service
  - Logs, runtimes, and services monitored for error conditions
  - Ensures an application successfully deployed into production is still running properly



Introduction to IBM Cloud Full Stack Development 18 © Copyright IBM Corporation 2017

IBM Training 

## Development Process

19 © Copyright IBM Corporation 2017

IBM Training 

### Overview and considerations

- Basic process
  - Phase 1: Monolith in Cloud Foundry
  - Phase 2: Dispatchers + Monolith in Cloud Foundry
  - Phase 3: Microservices in Cloud Foundry
  - Phase 4: Microservices in Docker containers
  - Phase 5: Microservices fabric and autoscaling
  - All phases: Add the rest of the architecture as needed
- Considerations
  - Monolith in Cloud Foundry: Simplest way to get started
    - Show how easy IBM Cloud is, show value quickly
  - Docker preferred?
    - If no opinion, Cloud Foundry is simpler
    - PaaS is simpler and provides greater value than CaaS
    - Postpones issues like microservices fabric
  - Microservices preferred?
    - If not, monolith is easier until the team gets too big
    - MVPs start as monoliths
    - Microservice architectures start as monoliths (or duoliths)

Introduction to IBM Cloud Full Stack Development 20 © Copyright IBM Corporation 2017

IBM Training 

## All phases: Add the rest of the architecture as needed

- As you're developing business logic
  - Regardless of architecture or technology
  - Add other architectural components as needed
  - Agile: Work on the other components in their own iterations
- When you need a GUI for testing and demos
  - Create a simple Web client
- When the business logic needs persistence
  - Add Cloudant
  - Make code modular, separate database/service instance per module
- When the business logic needs to connect to an on-prem resource
  - Create a gateway in Secure Gateway
  - Configure the gateway and client to connect to the resource(s)
- When connecting clients from an untrusted network
  - Expose the business logic's APIs via API Connect

Introduction to IBM Cloud Full Stack Development 21 © Copyright IBM Corporation 2017

IBM Training 

## Phase 1: Monolith in Cloud Foundry

- The simplest way to get started
- Monolith
  - What developers are already used to implementing
  - Least disruptive to their development teams and processes
- Cloud Foundry
  - Full Platform-as-a-Service (PaaS), not just Container-as-a-Service (CaaS)
  - Builds the runtime for you
  - Built-in service discovery
- Migrating to IBM Cloud
  - Migrating an existing application to IBM Cloud usually starts at this phase
  - The monolithic application needs to be cloud-ready (stateless, non-transactional, etc.)

Introduction to IBM Cloud Full Stack Development 22 © Copyright IBM Corporation 2017

IBM Training IBM

### Expose business logic as APIs

- Expose the business logic's functionality as APIs
  - RESTful endpoints (JSON/HTTP)
- Make the code modular
  - Separate modules for separate functionality
  - Each module exposes a single REST endpoint
  - For independent usage scenarios, create separate REST endpoints

The diagram illustrates a monolithic application structure. A large blue-outlined rectangle is labeled "Monolith" in its center. Inside the monolith, the text "Business Logic" is visible. Three red arrows originate from the word "APIs" at the top of the slide and point downwards to three small blue dots on the left edge, middle edge, and right edge of the monolith, which are labeled "Private REST Endpoints".

Introduction to IBM Cloud Full Stack Development      23      © Copyright IBM Corporation 2017

IBM Training IBM

### Phase 2: Dispatchers + Monolith in Cloud Foundry

- Refactor the dispatchers out of the monolith
  - Enables the dispatchers to be developed by the same teams developing the clients
  - Dispatchers can be developed and deployed independently of the monolith and each other
- Business functionality can continue to be developed as a monolith
- Cloud Foundry
  - Continues to be the simplest way to deploy runtimes
  - Simplest for both the monolith and the dispatchers

Introduction to IBM Cloud Full Stack Development      24      © Copyright IBM Corporation 2017

IBM Training 

## Phase 3: Microservices in Cloud Foundry

- Refactor the monolith into microservices
  - If the monolith code is modular, it will be easier to refactor into microservices
- Focus on developing good microservices
  - One business task per service
  - Business services implement requirements accurately
  - Advanced microservices fabric (autoscaling, circuit breaker, etc.) can come later
- Cloud Foundry
  - Simplest way to deploy each microservice
  - Built-in service discovery

Introduction to IBM Cloud Full Stack Development      25      © Copyright IBM Corporation 2017

IBM Training 

## Phase 4: Microservices in Docker containers

- Package the microservices as Docker containers
  - Each dispatcher and each business service in its own container
- Deploy the containers in IBM Cloud Container Service
  - May need service discovery beyond what Kubernetes provides
- Incremental step to container complexity
  - Microservices are already designed
  - Focus on building containers, deploying them, integrating them

Introduction to IBM Cloud Full Stack Development      26      © Copyright IBM Corporation 2017

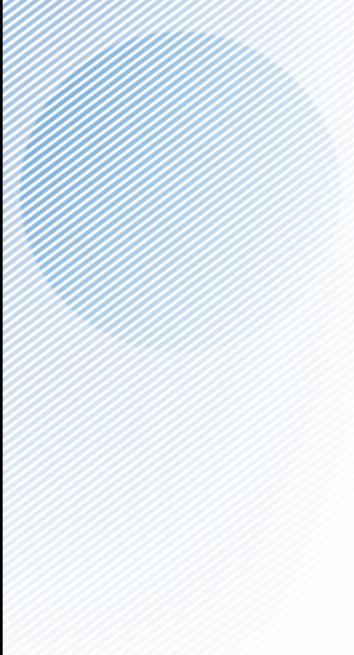
IBM Training 

## Phase 5: Microservices fabric and autoscaling

- Deploy microservices configured within a fabric for better QoS
  - Service registry and discovery
  - Circuit breaker, commands, and retry/alternate path
  - Monitoring, dashboards, and business transaction tracing
- Deploy microservices configured to run elastically
  - Horizontally scale based on client load
  - Crash recovery and health monitoring
- Focus on microservices packaged as Docker containers deployed in IBM Cloud Container Service
  - Can also apply to microservices deployed in Cloud Foundry
  - Different techniques needed for Cloud Foundry

Introduction to IBM Cloud Full Stack Development 27 © Copyright IBM Corporation 2017

IBM Training 



## Conclusion

28 © Copyright IBM Corporation 2017

IBM Training 

## Summary

We've now discussed

- Introduction
- Reference architecture
- Architectural components
- Development process

Introduction to IBM Cloud Full Stack Development      29      © Copyright IBM Corporation 2017

IBM Training IBM

## Introduction to Containers and Docker

 **LXC**     **docker**

© Copyright IBM Corporation 2016-18  
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

IBM Training IBM

### **Agenda**

- Why containers?
- Docker containers
- Conclusion

Introduction to Containers      2      © Copyright IBM Corporation 2016-18

IBM Training

IBM

## Why Containers?



3

© Copyright IBM Corporation 2016-18

IBM Training

IBM

## Containers

A standard way to **package** an application and all its **dependencies** so that it can be **moved** between environments and run **without change**

Containers work by **hiding the differences** between applications **inside** the container so that everything **outside** the container can be **standardized**

**Docker:** Created standard way to create images for Linux Containers

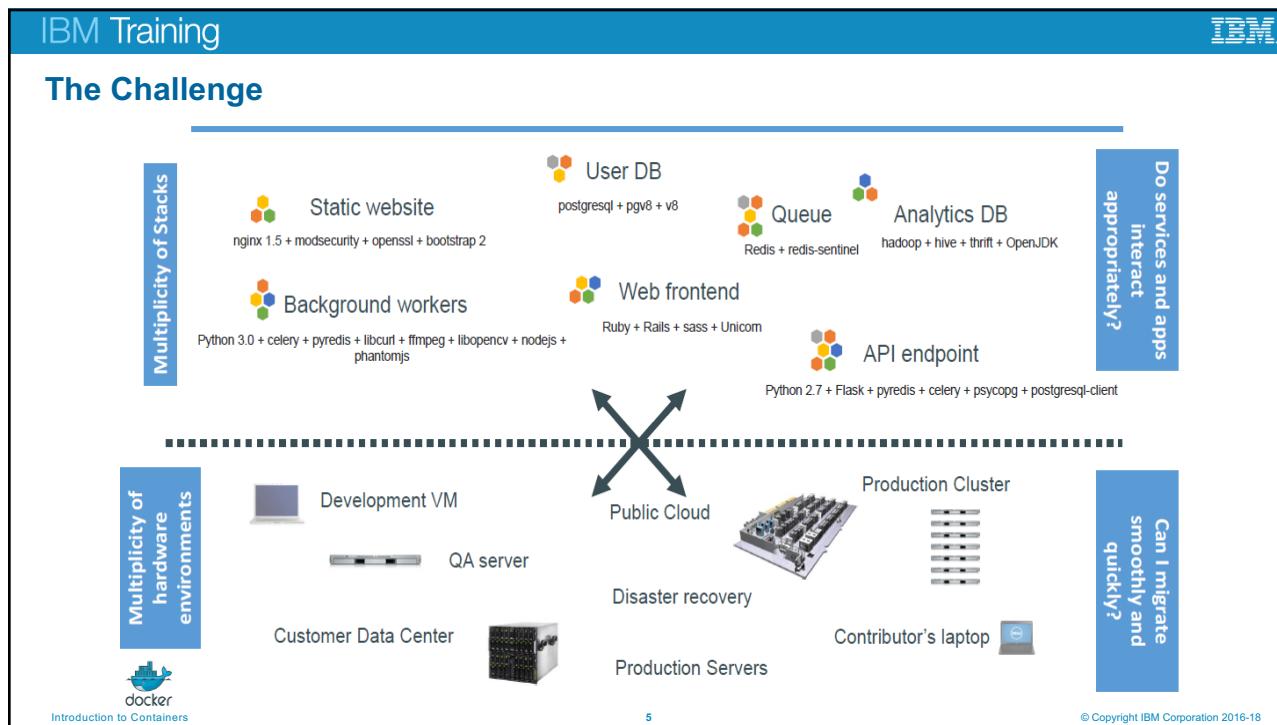


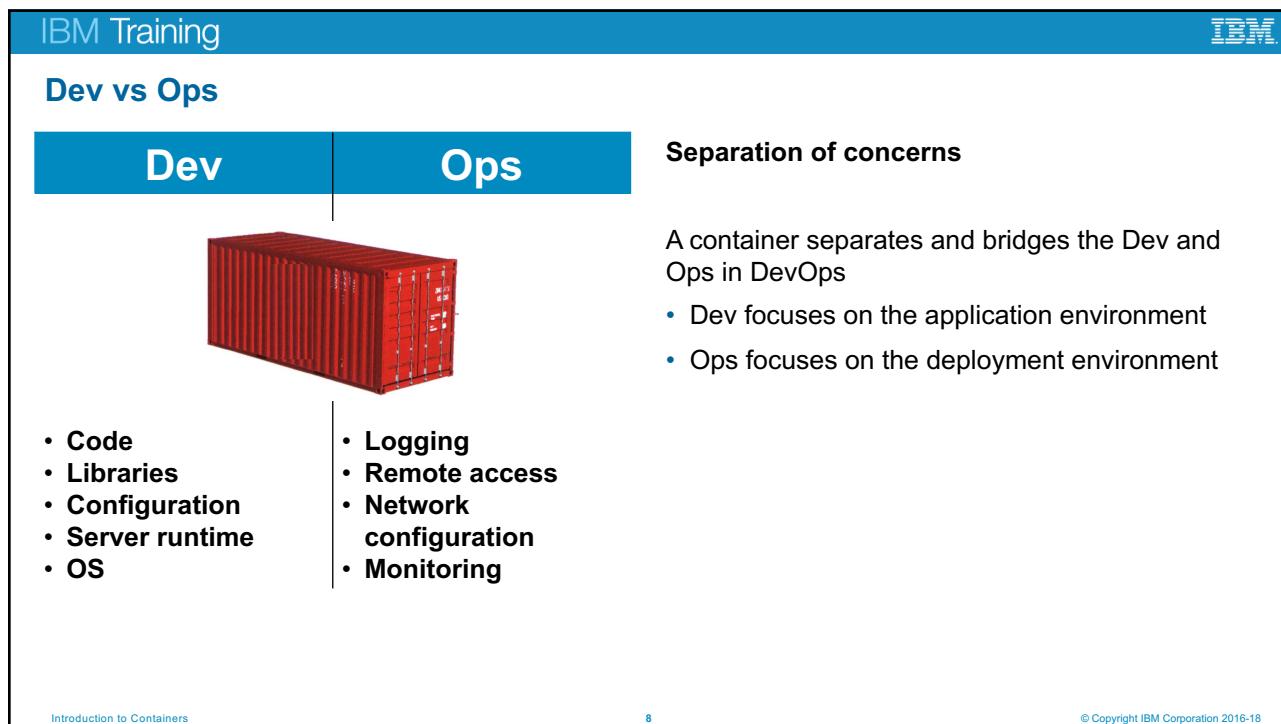
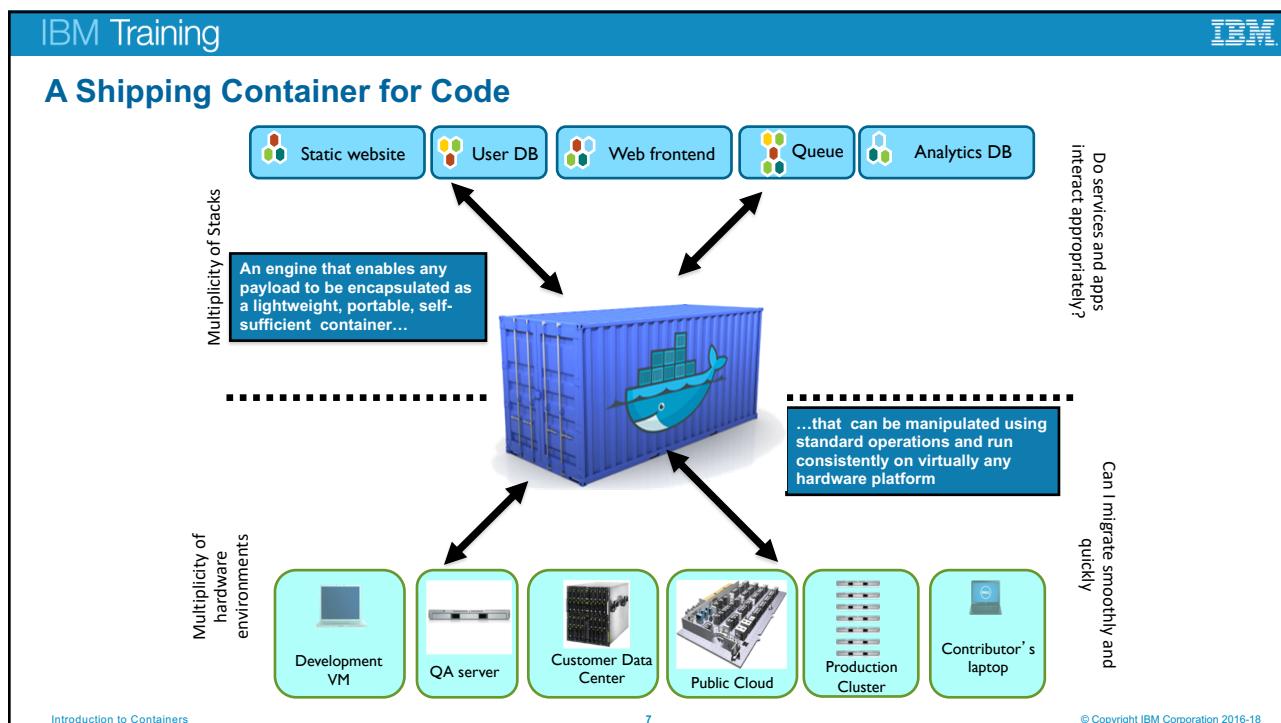
Linux Containers (LXC) details

- An isolated user space within a running Linux OS
- Shared kernel across containers
- Direct device access
- All packages and data in an isolated runtime, saved as a filesystem
- Resource management implemented with control groups (cgroups)
- Resource isolation through namespaces

4

© Copyright IBM Corporation 2016-18





**IBM Training**

## Everybody Loves Containers

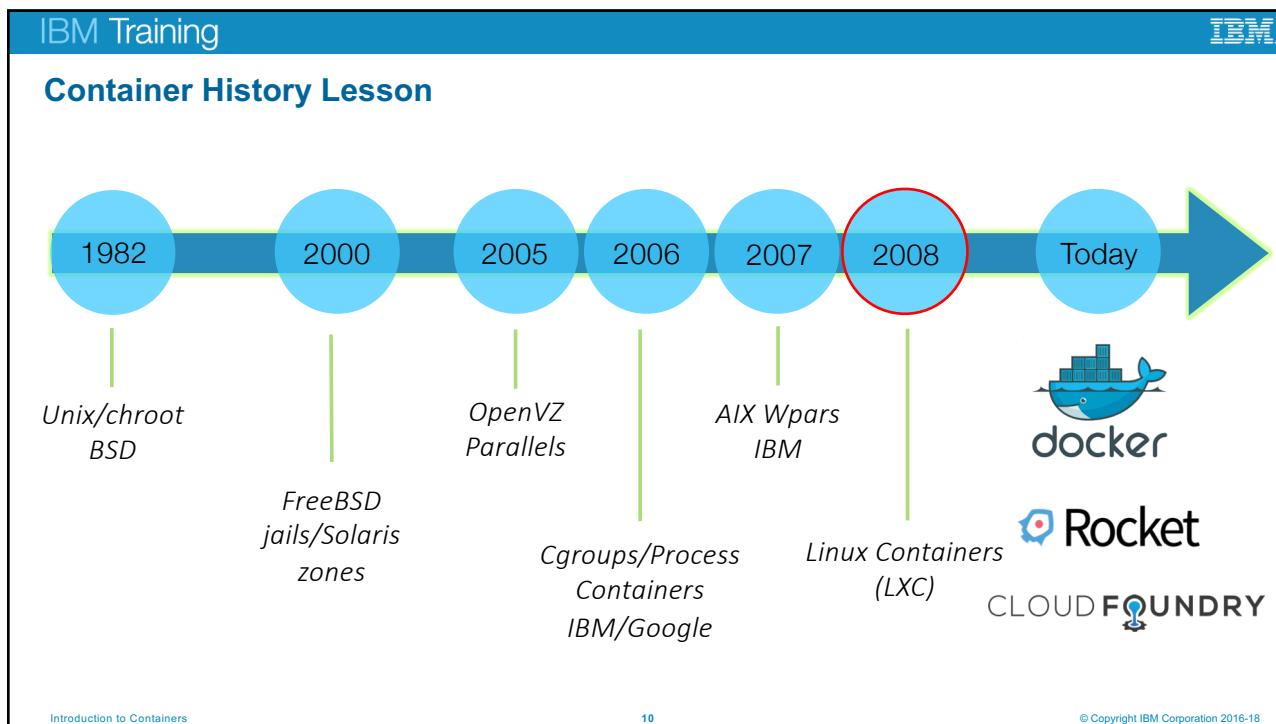



- A standard way to package an application and all its dependencies so that it can be moved between environments and run without changes
- Containers work by isolating the differences between applications inside the container so that everything outside the container can be standardized

Introduction to Containers

9

© Copyright IBM Corporation 2016-18



**IBM Training**

**Multiple (De facto) Container Standards**

- Docker
  - The most common standard, made Linux containers usable by the masses
- Rocket (rkt)
  - An emerging container standard from CoreOS, the company that developed etcd
- Garden
  - The format Cloud Foundry builds using buildpacks
- Open Container Initiative (OCI)
  - A Linux Foundation project developing a governed container standard

Docker logo: A blue whale carrying a stack of shipping containers.

Rocket logo: A white rocket ship with a red dot on its side.

Cloud Foundry logo: The words "CLOUD" and "FOUNDRY" in grey, with a small blue icon between them.

Open Container Initiative logo: A blue square with a white 'C' and 'I' inside, followed by the text "OPEN CONTAINER INITIATIVE".

Introduction to Containers 11 © Copyright IBM Corporation 2016-18

**IBM Training**

**Advantages of Containers**

- Containers are portable
  - Any platform with a container engine can run containers
- Containers are easy to manage
  - Container images are easy to share, download, and delete
    - Especially with Docker registries
  - Container instances are easy to create and delete
  - Each container instance is easy and fast to start and stop
- Containers provide “just enough” isolation
  - More lightweight than virtual machines
  - Processes share the operating system kernel but are segregated
- Containers use hardware more efficiently
  - Greater density than virtual machines
    - Especially Docker containers, which can share layers
- Containers are immutable
  - Container images are versions
  - Containers cannot (should not) be patched

The first container is red and empty. The second container is blue and features a white penguin icon. The third container is blue and features a white whale icon carrying a stack of shipping containers.

Introduction to Containers 12 © Copyright IBM Corporation 2016-18

IBM Training

## Docker Containers



13 © Copyright IBM Corporation 2016-18

IBM Training

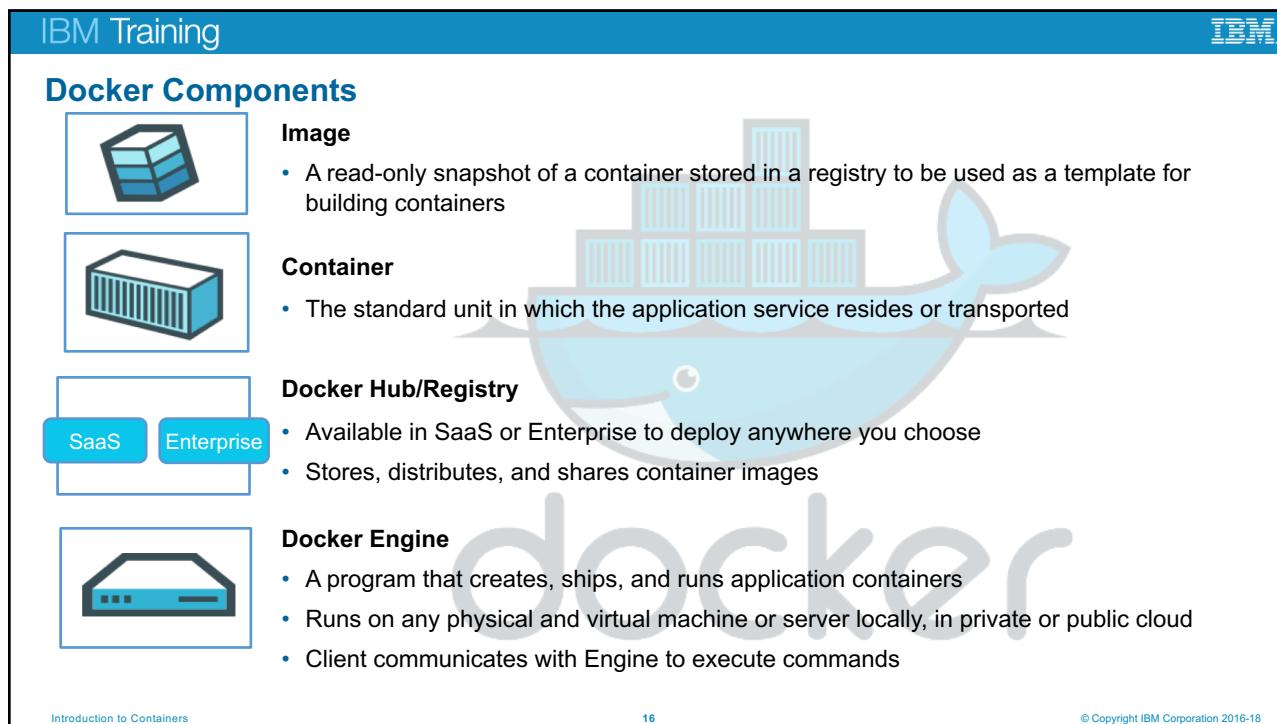
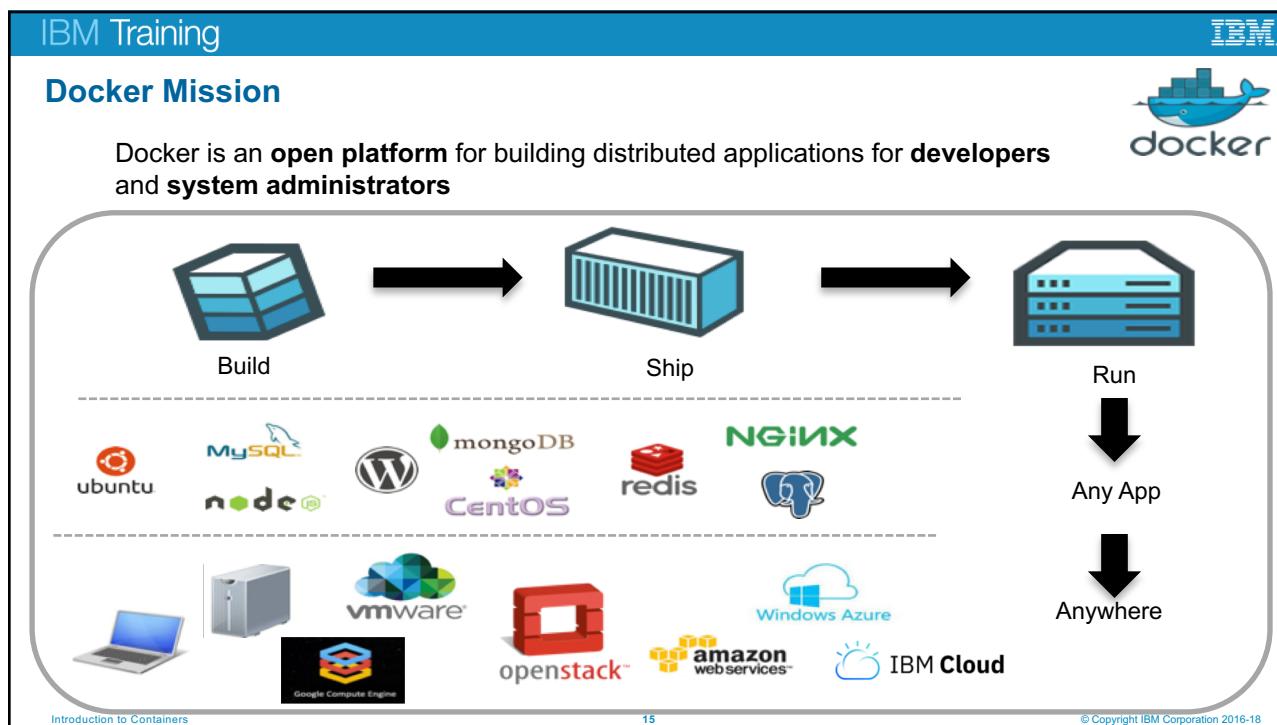
## Docker Adoption

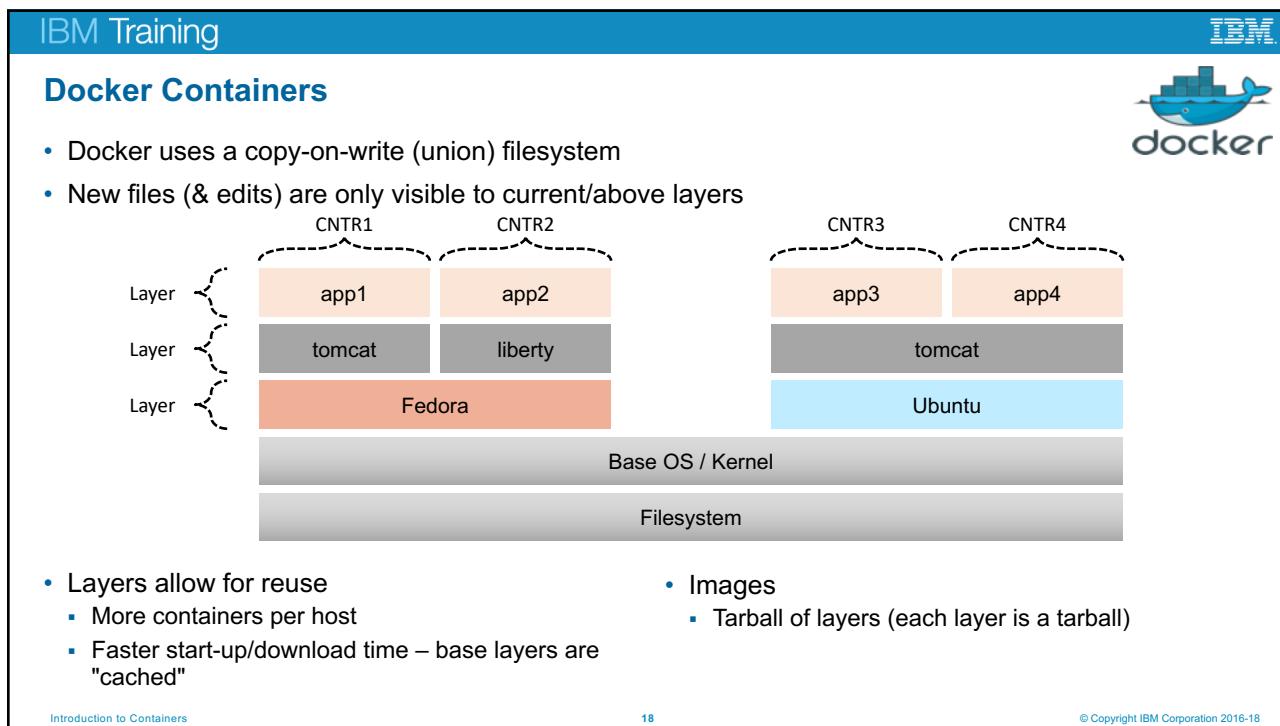
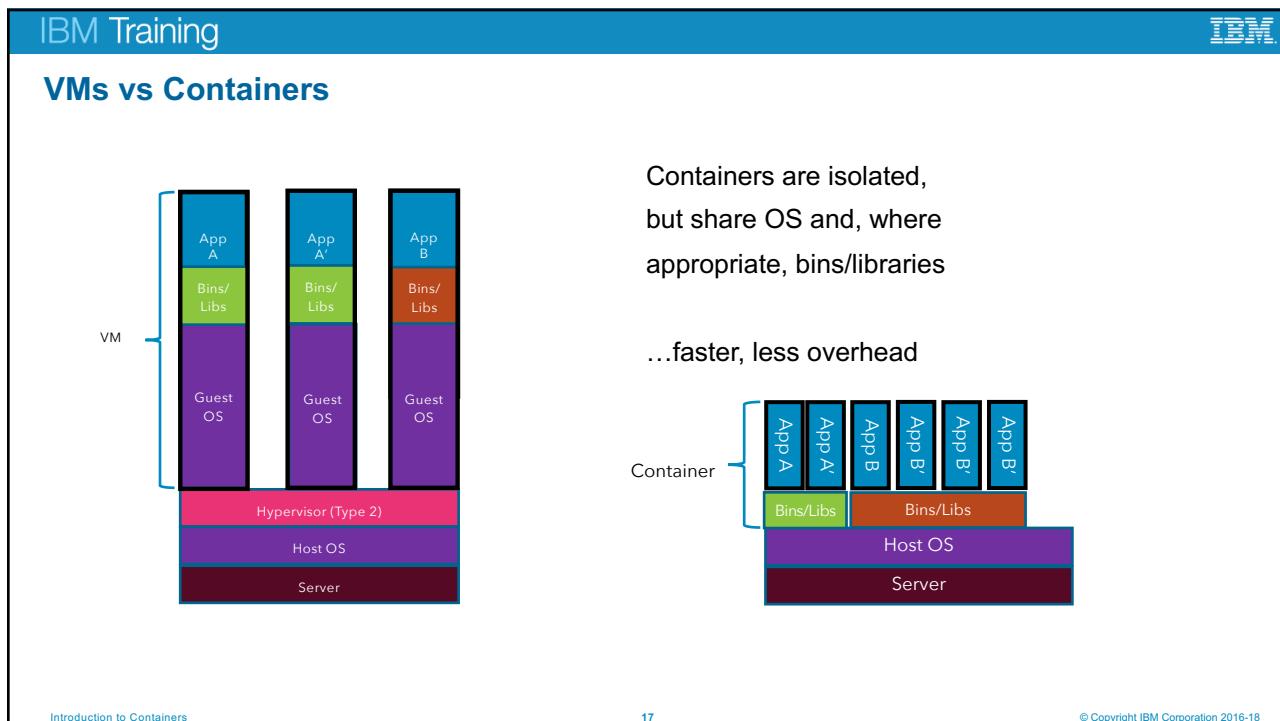
Docker enables application development **efficiency**, making deployment more **efficient**, eliminating vendor 'lock-in' with true **portability**



- Open software
  - Launched **March 2013**
  - 2.0+ billion downloads of Docker images
- Open contribution
  - 2000+ contributors
  - #2 most popular project
  - 185 community meet-up groups in 58 countries
- Open design
  - Contributors include IBM, Red Hat, Google, Microsoft, VMware, AWS, Rackspace, and others
- Open governance
  - Docker, the Open Container Initiative (OCI), and the Cloud Native Computing Foundation (CNCF) are jointly developing container standards

Introduction to Containers 14 © Copyright IBM Corporation 2016-18





IBM Training

## Why Customers are Interested in Docker Containers



- Containers are a critical foundation for hybrid apps
- Ship more software
  - Accelerate development and CI/CD pipelines by eliminating headaches of setting up environments and dealing with differences between environments
  - On average, Docker users ship software 7X more frequently
- Resource efficiency
  - Lightweight containers run on a single machine and share the same OS kernel while images are layered file systems sharing common files to make efficient use of RAM and disk and start instantly
- App portability
  - Isolated containers package the application, dependencies and configurations together
  - These containers can then seamlessly move across environments and infrastructures

Introduction to Containers 19 © Copyright IBM Corporation 2016-18

IBM Training

## Conclusion



20 © Copyright IBM Corporation 2016-18

IBM Training 

## Conclusion

- Why containers?
- Docker containers

Introduction to Containers 21 © Copyright IBM Corporation 2016-18

IBM Training 

## Resources

- Docker tutorial
  - <https://docs.docker.com/get-started/>
- The Evolution of Linux Containers and Their Future
  - <https://dzone.com/articles/evolution-of-linux-containers-future>
- rkt
  - <https://coreos.com/rkt>
- Cloud Foundry Garden container
  - <https://docs.cloudfoundry.org/concepts/architecture/garden.html>
- Open Container Initiative (OCI)
  - <https://www.opencontainers.org>
- Cloud Native Computing Foundation (CNCF)
  - <https://www.cncf.io>
- Demystifying the Open Container Initiative (OCI) Specifications
  - <https://blog.docker.com/2017/07/demystifying-open-container-initiative-oci-specifications>

Introduction to Containers 22 © Copyright IBM Corporation 2016-18

IBM Training

IBM

## Introduction to Kubernetes



**kubernetes**

© Copyright IBM Corporation 2017-18  
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

IBM Training

IBM

## Agenda

- Introduction
- Container orchestration
- Kubernetes
- Conclusion

Introduction to Kubernetes

2

© Copyright IBM Corporation 2017-18

IBM Training

IBM

## Introduction

3

© Copyright IBM Corporation 2017-18

IBM Training

IBM



Everyone's container journey starts with one container....

Introduction to Kubernetes

4

© Copyright IBM Corporation 2017-18

IBM Training IBM

At first the growth is easy to handle....

Introduction to Kubernetes 5 © Copyright IBM Corporation 2017-18

IBM Training IBM

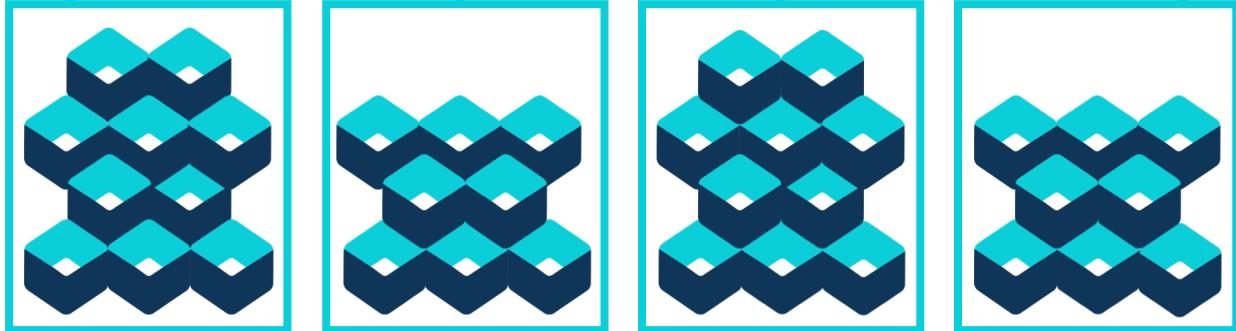
But soon it is overwhelming... chaos reigns

Introduction to Kubernetes 6 © Copyright IBM Corporation 2017-18

IBM Training

IBM

Regain control with Kubernetes



Introduction to Kubernetes

7

© Copyright IBM Corporation 2017-18

IBM Training

IBM

Container Orchestration



8

© Copyright IBM Corporation 2017-18

**IBM Training**

## More to Containers than just Docker

Serverless      APACHE OpenWhisk™

PaaS      CLOUD FOUNDRY

Container Orchestration      Apache MESOS      kubernetes

Container Engine      OPEN CONTAINER INITIATIVE      rkt      Swarm/ Swarm Mode

docker

Introduction to Kubernetes      9      © Copyright IBM Corporation 2017-18

**IBM Training**

## What is container orchestration?

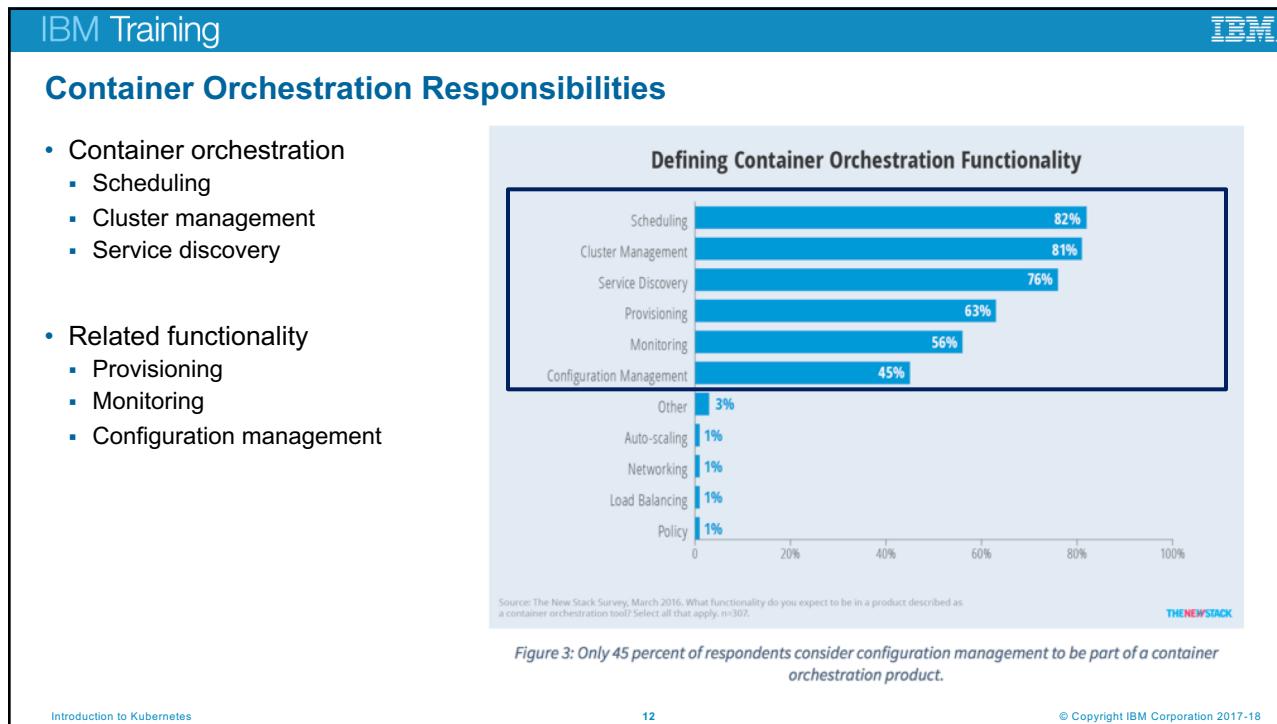
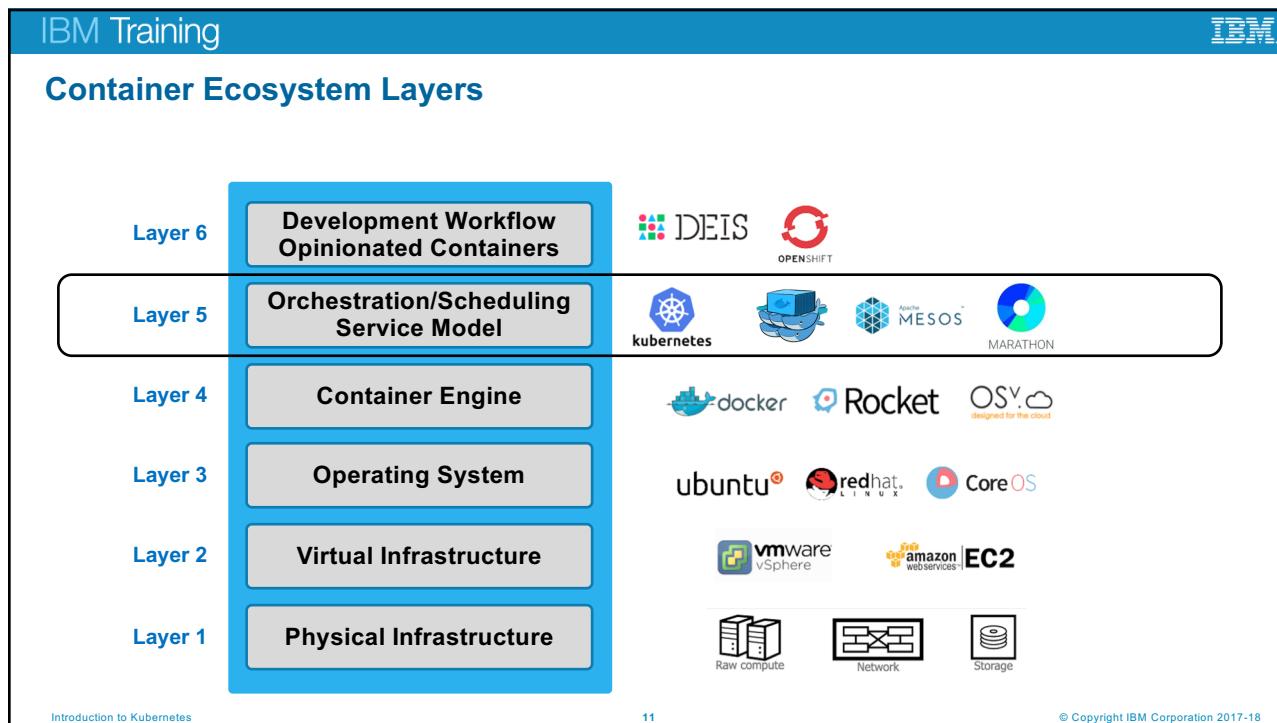
- Container orchestration
  - Manages the deployment, placement, and lifecycle of workload containers
- Cluster management
  - Federates multiple hosts into one target
- Scheduling
  - Distributes containers across nodes
- Service discovery
  - Knows where the containers are located
  - Distributes client requests across the containers
- Replication
  - Ensures the right number of nodes and containers
- Health management
  - Replaces unhealthy containers and nodes

**Container Orchestrator**

```

graph TD
    Manager[Manager<br/>Scheduler<br/>Replicator] --> ImageRepo[Image Repository]
    Manager --> DiscoveryDB[Discovery DB]
    Manager --- Node1[Node<br/>Daemon<br/>Containers]
    Manager --- Node2[Node<br/>Daemon<br/>Containers]
    Manager --- Node3[Node<br/>Daemon<br/>Containers]
    DiscoveryDB --> Node1
    DiscoveryDB --> Node2
    DiscoveryDB --> Node3
  
```

Introduction to Kubernetes      10      © Copyright IBM Corporation 2017-18



IBM Training 

## Container Orchestration in IBM Cloud

- In June 2015 when IBM Cloud Container Service launched, open source orchestration projects were not available for production
  - Kubernetes – July 21, 2015
  - Docker Swarm – November 3, 2015
  - Apache Mesos – July 27, 2016
- IBM developed custom orchestration to ensure anti co-location of containers within a group and container placement to the least utilized hosts
- IBM Research runs rigorous performance and scalability testing (<https://github.com/Open-IBeam/containers-os>) on various open source projects; then work with those communities to make improvements

Introduction to Kubernetes 13 © Copyright IBM Corporation 2017-18



The slide contains three logos side-by-side. From left to right: the Kubernetes logo (a blue hexagon with a white steering wheel icon), the Docker Swarm logo (a blue whale carrying a white cube), and the Apache Mesos logo (a blue hexagonal grid with the word "MESOS" next to it).

IBM Training 

## Kubernetes

14 © Copyright IBM Corporation 2017-18

**IBM Training**

## What is Kubernetes?

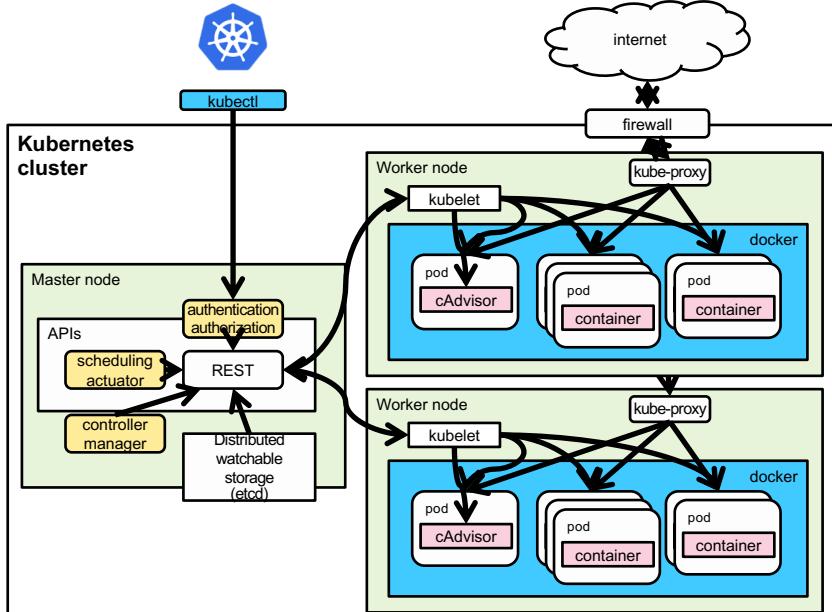


- Container orchestrator
  - Runs and manages containers
  - Unified API for deploying web applications, batch jobs, and databases
  - Maintains and tracks the global view of the cluster
  - Supports multiple cloud and bare-metal environments
- Manage applications, not machines
  - Rolling updates, canary deploys, and blue-green deployments
- Designed for extensibility
  - Rich ecosystem of plug-ins for scheduling, storage, networking
- Open source project managed by the Linux Foundation
  - Inspired and informed by Google's experiences and internal systems
  - 100% open source, written in Go

Introduction to Kubernetes      15      © Copyright IBM Corporation 2017-18

**IBM Training**

## Kubernetes Architecture



The diagram illustrates the Kubernetes architecture. It shows a central **Kubernetes cluster** containing a **Master node** and multiple **Worker nodes**. The Master node contains the **REST API**, **scheduling actuator**, **controller manager**, and **Distributed watchable storage (etcd)**. The REST API interacts with the scheduling actuator and controller manager. The scheduling actuator interacts with the REST API. The controller manager interacts with the REST API. The REST API interacts with the authentication/authorization module. The authentication/authorization module interacts with the REST API. The REST API interacts with the kubelet on the Worker nodes. The kubelet on the Worker nodes interacts with the pod, cAdvisor, and container. The pod, cAdvisor, and container interact with the docker. The docker interacts with the kube-proxy. The kube-proxy interacts with the Worker nodes. The Worker nodes interact with the internet through a firewall.

**Master nodes**

- Computer with processes that manage the cluster
- Multiple nodes for HA

**Worker nodes**

- Computers that host pods
- Pods host application containers

Introduction to Kubernetes      16      © Copyright IBM Corporation 2017-18

**IBM Training**

## Kubernetes Terminology: Topology

- Cluster
  - Collection of worker nodes managed by the same master node
  - Makes the worker nodes behave like one big computer
- Master node
  - Controls and manages the cluster
  - Scheduling and replication logic
  - Generally two or more master nodes for resiliency, but are not used for scaling out the cluster
- Worker node
  - Node where pods are run
  - Docker – Docker engine for running containers (including the kubelet)
  - kubelet
    - Kubernetes agent that accepts commands from the master
    - Manages pods in the node
  - cAdvisor – Container Advisor provides resources usage and performance statistics
  - kube-proxy – network proxy service responsible for routing activities for inbound or ingress traffic

Introduction to Kubernetes

17

© Copyright IBM Corporation 2017-18

**IBM Training**

## Kubernetes Terminology: Master Node Components

- Etcd
  - A highly-available key value store
  - All cluster data is stored here
- API Server
  - Exposes API for managing Kubernetes
  - Used by kubectl CLI
- Controller manager
  - Daemon that runs controllers, which are the background threads that handle routine tasks in the cluster
  - Node Controller – Responsible for noticing and responding when nodes go down
  - Replication Controller – Replaced by ReplicaSet
  - Endpoints Controller – Populates the Endpoints object (that is, joins services and pods)
  - Service Account & Token Controllers – Create default accounts and API access tokens for new namespaces
- Scheduler
  - Selects the worker node each pods runs in

Introduction to Kubernetes

18

© Copyright IBM Corporation 2017-18

**IBM Training**

**Kubernetes Architecture: Workloads**

- Container
  - Packaging of an app
- Pod
  - Unit of deployment
- Service
  - Fixed endpoint for 1+ pods

Master Node

App A Client

App A Service

App B Client

App B Service

Kubernetes components

Worker Node 1

Worker Node 2

Worker Node 3

App A

App B

Pod

© Copyright IBM Corporation 2017-18

**IBM Training**

**Kubernetes Terminology: Workloads**

- Container
  - Unit of packaging
- Pod
  - Smallest deployment unit in Kubernetes
  - A collection of containers that run on a worker node
  - Each pod has its own IP
  - A pod shares a PID namespace, network, and hostname
- Service
  - Collection of pods exposed as an endpoint
  - Information stored in the Kubernetes cluster state and networking info propagated to all worker nodes
  - Types
    - ClusterIP – Exposes the service on a cluster-internal IP
    - NodePort – Exposes the service on each Node's IP at a static port
    - LoadBalancer – Exposes the service externally using a cloud provider's load balancer
    - ExternalName – Maps the service to an external name (such as foo.bar.example.com)

Introduction to Kubernetes

20

© Copyright IBM Corporation 2017-18

**IBM Training**

## Kubernetes Terminology: Deployment

- Deployment
  - A set of pods to be deployed together, such as an application
  - Declarative: Revising a Deployment creates a ReplicaSet describing the desired state
  - Rollout: Deployment controller changes the actual state to the desired state at a controlled rate
  - Rollback: Each Deployment revision can be rolled back
  - Scale and autoscale: A Deployment can be scaled
- ReplicaSet
  - The next-generation ReplicationController
  - A set of pod templates that describe a set of pod replicas
  - Uses a template that describes specifically what each pod should contain
  - Ensures that a specified number of pod replicas are running at any given time

Introduction to Kubernetes

21

© Copyright IBM Corporation 2017-18

**IBM Training**

## Kubernetes Autoscaling

- Horizontal Pod Autoscaler (HPA)
  - Automatically scales the number of pods in a replication controller, deployment, or replica set
  - Matches the observed average CPU utilization to the specified target
  - Fetches metrics in two different ways: direct Heapster access and REST client access
  - Kubernetes Heapster enables container cluster monitoring and performance analysis
  - Default config: query every 30 sec, maintain 10% tolerance, wait 3 min after scale-up, wait 5 min after scale-down

```
$ kubectl autoscale deployment <deployment-name>
--min=3 --max=10 --cpu-percent=50
```

- Creates a horizontal pod autoscaler
  - An HPA instance that autoscales the specified deployment
  - Maintains between 3 and 10 replicas of the pods controlled by the deployment
  - Maintains an average CPU utilization across all pods of 50%

Introduction to Kubernetes

22

© Copyright IBM Corporation 2017-18

IBM Training 

## Kubernetes Terminology: Naming



- Name
  - Each resource object by type has a unique name
- Namespace
  - Resource isolation: Each namespace is a virtual cluster within the physical cluster
    - Resource objects are scoped within namespaces
    - Low-level resources are not in namespaces: nodes, persistent volumes, and namespaces themselves
    - Names of resources need to be unique within a namespace, but not across namespaces
  - Resource quotas: Namespaces can divide cluster resources
  - Initial namespaces
    - default – The default namespace for objects with no other namespace
    - kube-system – The namespace for objects created by the Kubernetes system
- Resource Quota
  - Limits resource consumption per namespace
  - Limit can be number of resource objects by type (pods, services, etc.)
  - Limit can be total amount of compute resources (CPU, memory, etc.)
  - Overcommit is allowed; contention is handled on a first-come, first-served basis

Introduction to Kubernetes 23 © Copyright IBM Corporation 2017-18

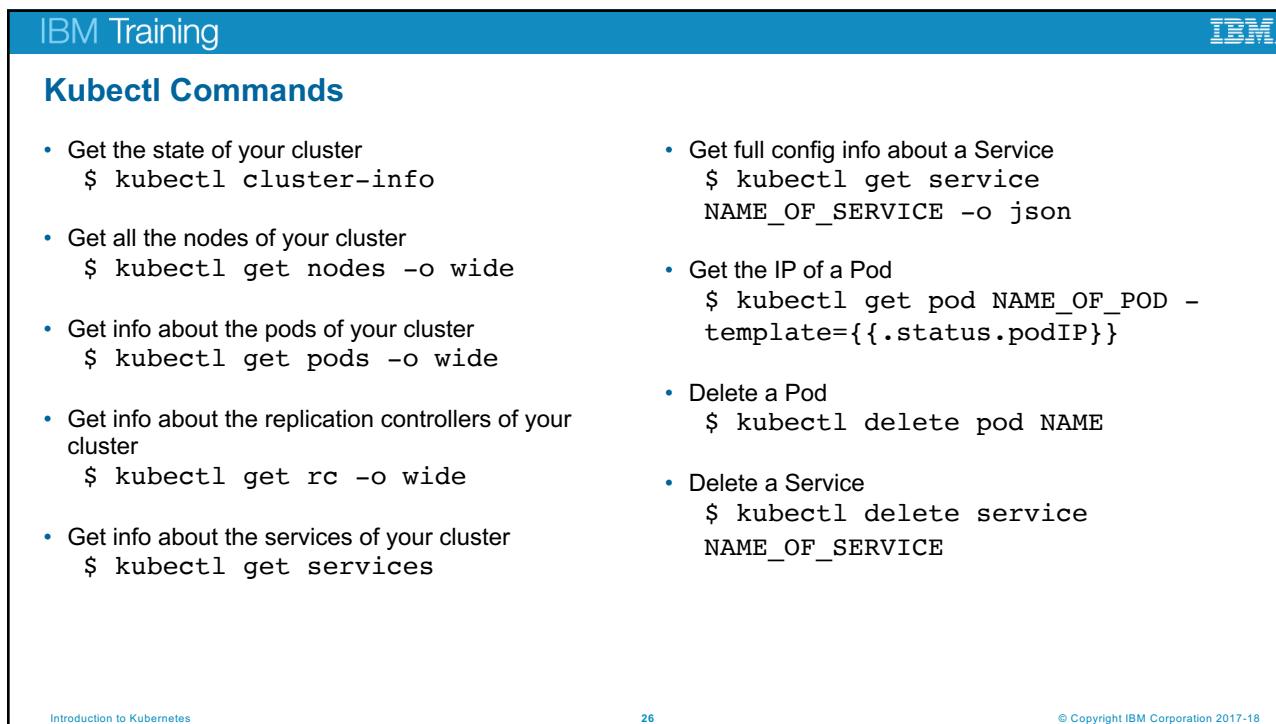
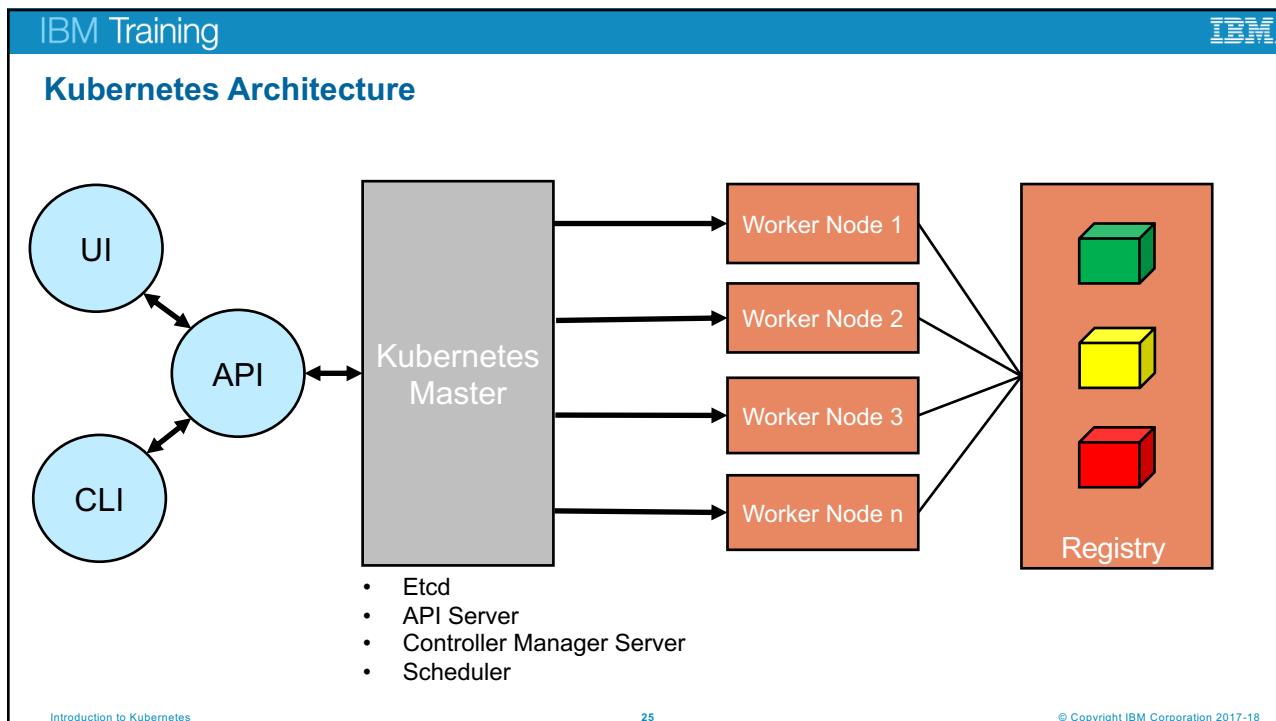
IBM Training 

## Kubernetes Configuring Resources and Containers



- Label
  - Metadata assigned to Kubernetes resources (pods, services, etc.)
  - Key-value pairs for identification
  - Critical to Kubernetes as it relies on querying the cluster for resources that have certain labels
- Selector
  - An expression that matches labels to identify related resources
- ConfigMap
  - Configuration values to be used by containers in a pod
  - Stores configuration outside of the container image, making containers more reusable
- Secrets
  - Sensitive info that containers need to read or consume
  - Encrypted in special volumes mounted automatically

Introduction to Kubernetes 24 © Copyright IBM Corporation 2017-18



IBM Training

IBM

## Conclusion

27

© Copyright IBM Corporation 2017-18

IBM Training

IBM

## Conclusion

- Container orchestration
- Kubernetes

Introduction to Kubernetes

28

© Copyright IBM Corporation 2017-18

IBM Training 

## Resources

- Kubernetes tutorial
  - <https://kubernetes.io/docs/tutorials/kubernetes-basics/>
- Introduction to container orchestration
  - <https://www.exoscale.ch/syslog/2016/07/26/container-orch/>
- TNS Research: The Present State of Container Orchestration
  - <https://thenewstack.io/tns-research-present-state-container-orchestration/>
- Large-scale cluster management at Google with Borg
  - <https://research.google.com/pubs/pub43438.html>

Introduction to Kubernetes      29      © Copyright IBM Corporation 2017-18

IBM Training

IBM

## Kubernetes in IBM Cloud Private

© Copyright IBM Corporation 2018  
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

1

IBM Training

IBM

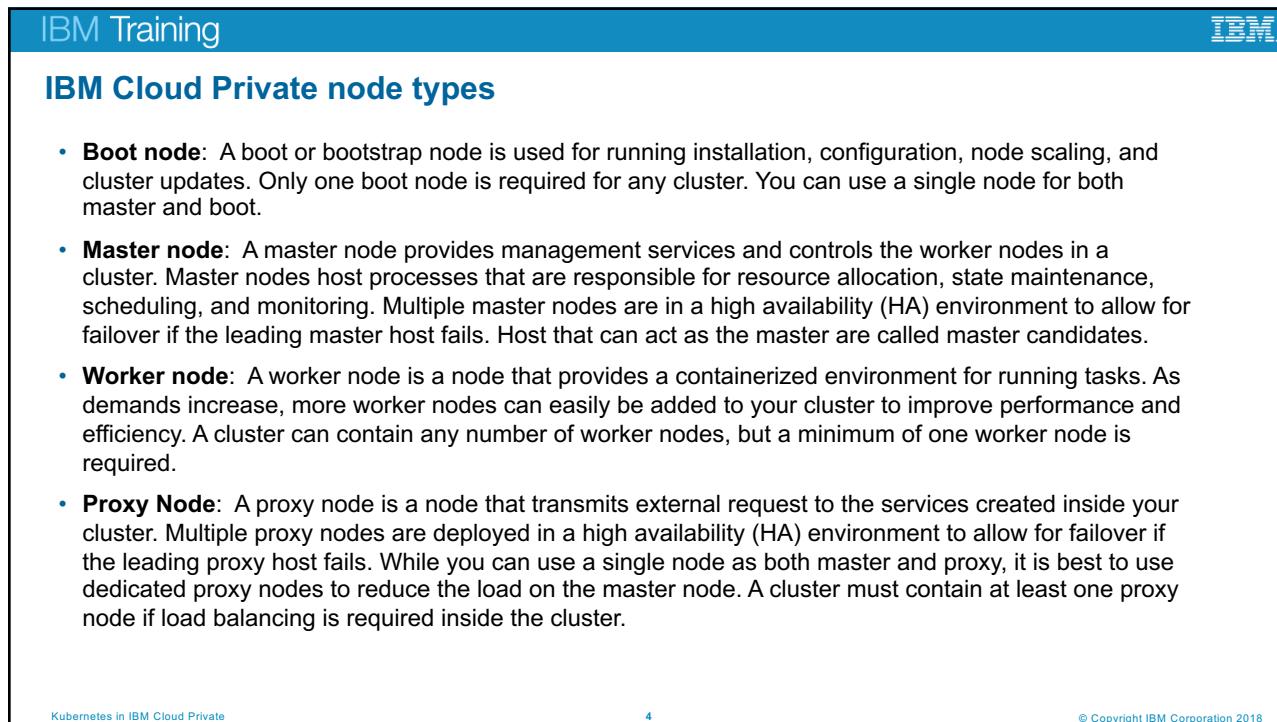
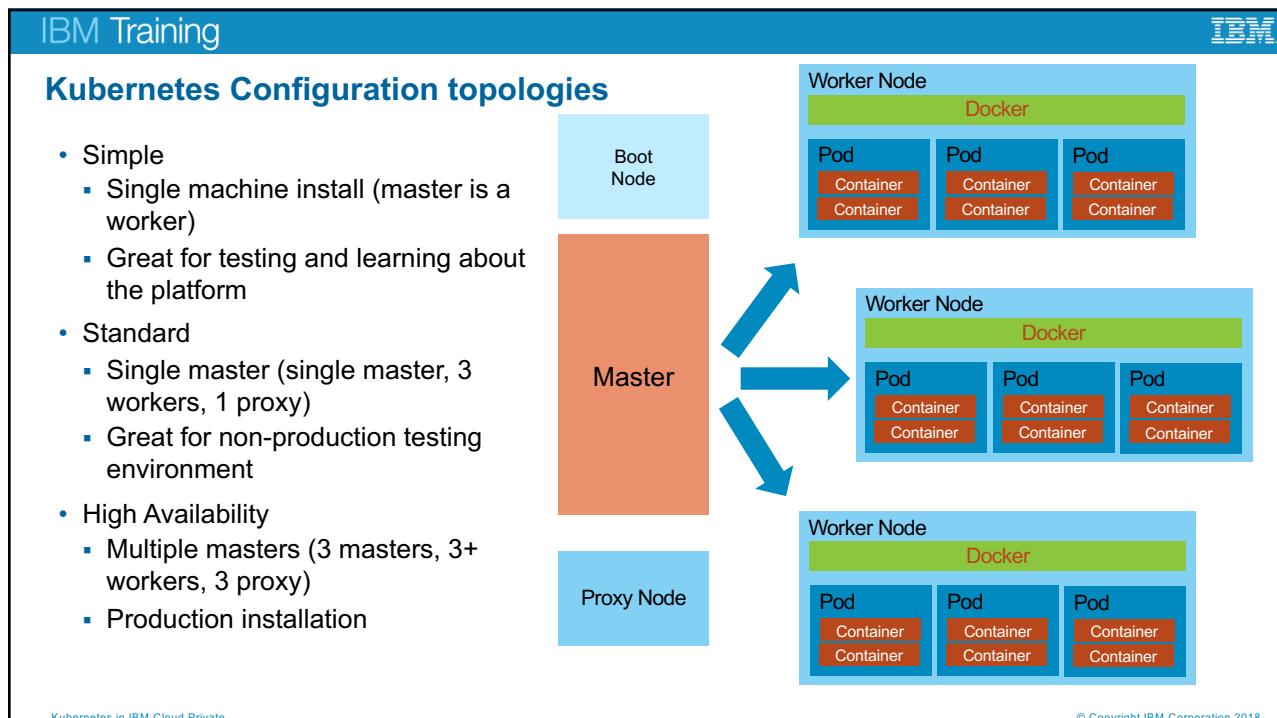
### Objectives

- Explain IBM Cloud Private implementation of Kubernetes
- List major components of IBM Cloud Private implementation of Kubernetes (calico, authentication, user interface, etc)

Kubernetes in IBM Cloud Private

2

© Copyright IBM Corporation 2018

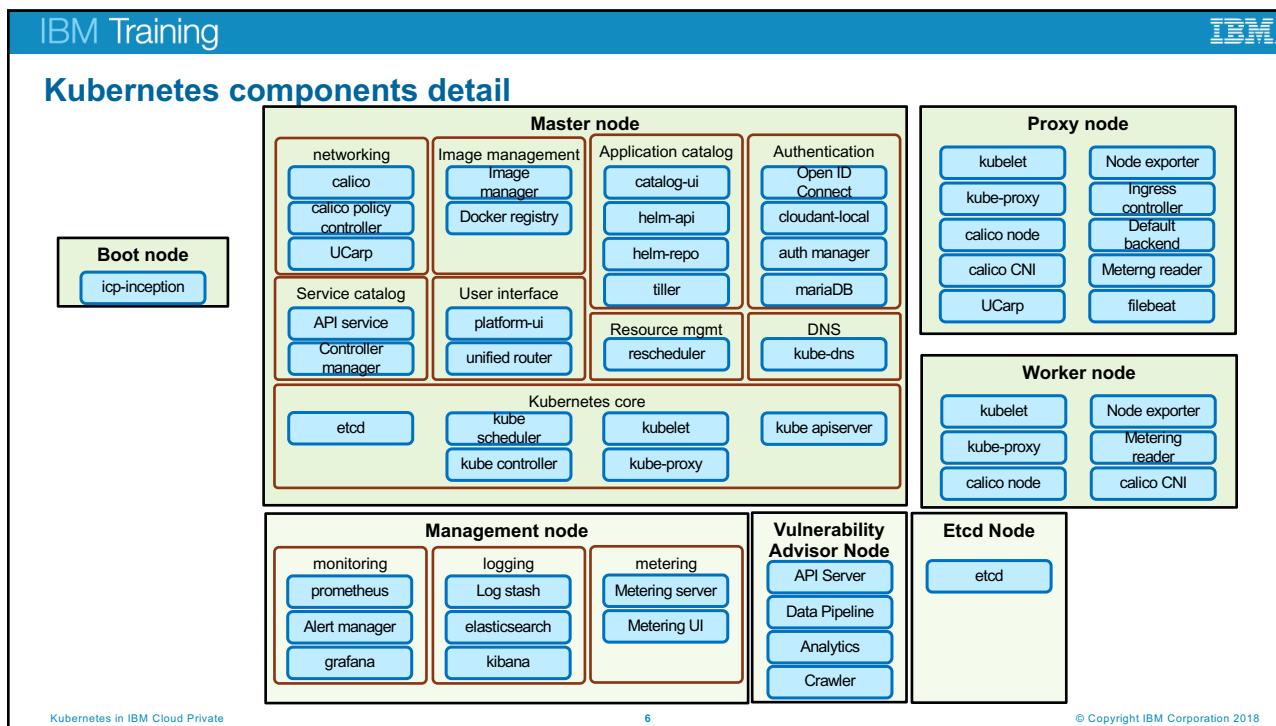


**IBM Training** **IBM**

## IBM Cloud Private node types – optional nodes

- Management node:** A management node is an optional node that only hosts management services such as monitoring, metering, and logging. By configuring dedicated management nodes, you can prevent the master node from becoming overloaded. You can enable the management node only during IBM Cloud Private installation.
- VA node:** A VA (Vulnerability Advisor) node is an optional node that is used for running the Vulnerability Advisor services. Vulnerability Advisor services are resource intensive. If you use the Vulnerability Advisor service, specify a dedicated VA node.
- Etcd node:** An etcd node is an optional node that is used for running the etcd distributed key value store. Configuring an etcd node in an IBM Cloud Private cluster that has many nodes, such as 100 or more, helps to improve the etcd performance.

Kubernetes in IBM Cloud Private 5 © Copyright IBM Corporation 2018



## IBM Training



### Kubernetes master components

Master components provide the cluster's control plane and global decisions about the cluster and detecting and responding to cluster events

- **Etcd:** A strong, consistent, and highly-available key value store which Kubernetes uses for persistent storage of all of its API objects.
- **Kubelet:** The primary “node agent” that runs on each node.
- **Kube Proxy:** The Kubernetes network proxy runs on each node.
- **Kube Scheduler:** A policy-rich, topology-aware, workload-specific function that significantly impacts availability, performance, and capacity.
- **Kube Control Manager:** A daemon that embeds the core control loops shipped with K8s. In K8s, a controller is a control loop that watches the shared state of the cluster through the apiserver and makes changes attempting to move the current state towards the desired state.
- **Kube apiserver:** Validates and configures data for the api objects which include pods, services, replication controllers, and others. The API Server services REST operations and provides the frontend to the cluster's shared state through which all other components interact.

## IBM Training



### Network services components

- **DNS:** (kube-dns, Cluster DNS) K8s DNS schedules a DNS pod and service on the cluster, and configures the kubelets to tell individual containers to use the DNS service's IP to resolve DNS names. Every service defined in the cluster (including the DNS server itself) is assigned a DNS name. By default, a client pod's DNS search list will include the pod's own namespace and the cluster's default domain.
- **UCarp:** UCarp allows a couple of hosts to share common virtual IP (or floating IP) addresses in order to provide automatic failover.
- **Calico:** A new approach to virtual networking and network security for containers, VMs, and bare metal services, that provides a rich set of security enforcement capabilities running on top of a highly scalable and efficient virtual network
  - The calico/node Docker container runs on the Kubernetes master and each Kubernetes node in the cluster
  - The calico-cni plug-in integrates directly with the Kubernetes kubelet process on each node to discover which pods have been created, and adds them to Calico networking
  - The calico/kube-policy-controller container runs as a pod on top of Kubernetes and implements the NetworkPolicy API

IBM Training IBM

## Authentication

- **Authentication Manager:** Provides an HTTP API for managing users. Protocols are implemented in a RESTful manner. Open ID Connect is used for authentication.
- **Open ID Connect:** WebSphere Liberty based authentication server that implements Open ID Connect protocol.
- **MariaDB:** An open source relational database made by the original developers of MySQL.

Kubernetes in IBM Cloud Private 9 © Copyright IBM Corporation 2018

IBM Training IBM

## Images and Registries

- You create a Docker image and push it to a registry before referring to it in a Kubernetes pod
- There will likely be many registries used in your deployment

The diagram illustrates a workflow for managing Docker images across different environments. It starts with a 'Docker public registry' containing a Docker image icon. This image is pushed to a 'Docker private registry' located at 'acme.com'. From the private registry, the image is pulled into a 'Developer' environment, which contains a developer icon. The developer environment then pushes the image to a 'UAT' (User Acceptance Testing) environment, which contains a UAT icon. The UAT environment then promotes the image to a 'Production' environment, which contains a production icon. The Production environment also contains a 'Promote' icon. Arrows indicate the direction of the 'Push/Pull' and 'Pull to deploy' operations between the stages.

Kubernetes in IBM Cloud Private 10 © Copyright IBM Corporation 2018

**IBM Training**

## User Interfaces

- Cluster Management Console: (ICP component) Use to manage, monitor, and troubleshoot your applications and cluster from a single, centralized, and secure management console.
- Platform-UI:** Can use to deploy containerized applications to a Kubernetes cluster, troubleshoot your containerized application, and manage the cluster itself along with its attendant resources.
- kubectl:** A command-line interface for running commands against Kubernetes clusters.

Kubernetes in IBM Cloud Private      11      © Copyright IBM Corporation 2018

**IBM Training**

## Application Center components

- Catalog-UI:** provides a centralized location from which you can browse, and install packages in your cluster.
- Helm:** A tool for managing Kubernetes charts. Charts are packages of pre-configured Kubernetes resources.
- Helm Repository:** A Helm chart repository is a location where packaged charts can be stored and shared.
- Tiller:** Runs inside of the cluster, and manages releases (installations) of your charts.

Kubernetes in IBM Cloud Private      12      © Copyright IBM Corporation 2018

**IBM Training** 

## Ingress resources

- **Ingress:** Typically, services and pods have IPs only routable by the cluster network. All traffic that ends up at an edge router is either dropped or forwarded elsewhere.
  - An Ingress is a collection of rules that allow inbound connections to reach the cluster services.
  - It can be configured to give services externally-reachable URLs, load balance traffic, terminate SSL, offer name based virtual hosting etc.
  - Users request ingress by POSTing the Ingress resource to the API server
- **Ingress Controller:** Responsible for fulfilling the Ingress, usually with a load balancer, though it may also configure your edge router or additional frontends to help handle the traffic in an HA manner.

Kubernetes in IBM Cloud Private 13 © Copyright IBM Corporation 2018

**IBM Training** 

## Logging components

The easiest and most embraced logging method for containerized applications is to write to standard out and standard error

- **Filebeat:** A log data shipper for local files. Filebeat monitors the log directories or specific log files, tails the files, and forwards them either to [Elasticsearch](#) and/or [Logstash](#) for indexing.
- **Elasticsearch:** An open source full-text search engine based on Lucene. It provides HTTP web interface and schema-free JSON documents.
- **Log stash:** A open source tool for collecting, parsing, and storing logs for future use.
- **Heapster:** The Kubernetes network proxy runs on each node.
- **Kibana:** An open source data visualization plugin for Elasticsearch. Users can create bar, line and scatter plots, or pie charts and maps on top of large volumes of data.

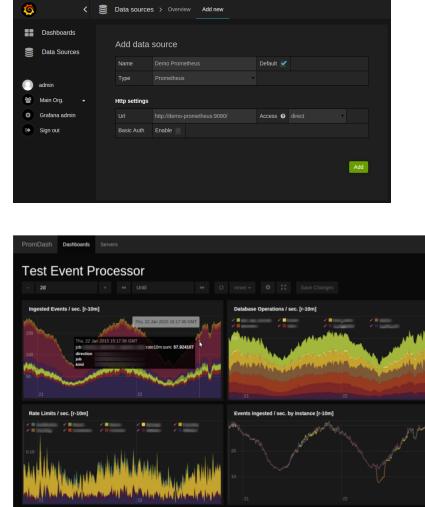


Kubernetes in IBM Cloud Private 14 © Copyright IBM Corporation 2018

**IBM Training**

## Monitoring: Prometheus and Grafana

- **Prometheus:** An open-source systems monitoring and alerting toolkit originally built at [SoundCloud](#). Since its inception in 2012, many companies and organizations have adopted Prometheus, and the project has a very active developer and user [community](#). It is now a standalone open source project and maintained independently of any company.
- **Grafana:** An open-source, general purpose dashboard and graph composer, which runs as a web application.



Kubernetes in IBM Cloud Private      15      © Copyright IBM Corporation 2018

**IBM Training**

## Persistent storage components

- Traditionally Containers: stateless, ephemeral in nature
  - Storage exists within the container
  - The container goes away and so goes the storage
- Some applications desire state and thus persistent storage:
  - Specific aspects of configuration
  - Database (structured and unstructured)
  - Application data (website definitions, etc.)
- Storage must be universally accessible across the K8s environment
- ICP Persistent Storage Support: HostPath, NFS, GlusterFS, vSphereVolume
  - hNote: Reclaim policy and access modes and behaviors can vary
- Access Modes:
  - ReadWriteOnce – the volume can be mounted as read-write by a single node
  - ReadOnlyMany – the volume can be mounted read-only by many nodes
  - ReadWriteMany – the volume can be mounted as read-write by many nodes

Kubernetes in IBM Cloud Private      16      © Copyright IBM Corporation 2018

IBM Training IBM

## Introduction to Helm



**kubernetes**

Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

IBM Training IBM

### Agenda

- What is Helm?
- Using Helm
- Developing charts
- Developing templates
- Conclusion
- Resources

Introduction to Helm 2 © Copyright IBM Corporation 2017

IBM Training 

## What is Helm?

3 © Copyright IBM Corporation 2017

IBM Training 

### Helm – A package manager for Kubernetes

- What is a package manager?
  - Automates the process of installing, configuring, upgrading, and removing computer programs
  - Examples: Red Hat Package Manager (RPM), Homebrew (macOS), Chocolatey (Windows)
- Helm enables multiple Kubernetes resources to be created with a single command
  - Deploying an application often involves creating and configuring multiple resources
  - A Helm chart defines multiple resources as a set
- An application in Kubernetes typically consists of (at least) two resource types
  - Deployment – Describes a set of pods to be deployed together
  - Services – Endpoints for accessing the APIs in those pods
  - Could also include ConfigMaps, Secrets, Ingress, etc.
- A default chart for an application consists of a deployment template and a service template
  - The chart creates all of these resources in a Kubernetes cluster as a set
  - Rather than manually having to create each one separately via kubectl

Introduction to Helm 4 © Copyright IBM Corporation 2017

**IBM Training**

## Helm Terminology

- Helm – The CLI
  - Helm installs charts into Kubernetes, creating a new release for each installation
  - To find new charts, search Helm chart repositories
- Chart – The application package
  - Templates for a set of resources necessary to run an application
  - The chart includes a values file that configures the resources
- Repository – The library
  - Storage for Helm charts
  - stable – The namespace of the hub for official charts
- Release – The application runtime
  - An instance of a chart running in a Kubernetes cluster
  - The same chart installed multiple times creates many releases
- Tiller – The server-side engine
  - Helm templating engine, runs in a pod in a Kubernetes cluster
  - Tiller processes the chart to generate the resource manifests, then installs the release into the cluster
  - Tiller stores each release as a Kubernetes config map

Introduction to Helm

5

© Copyright IBM Corporation 2017

**IBM Training**

## Advantages of Using Helm

- Deploy all of the resources for an application with a single command
  - Makes deployment easy and repeatable
  - `$ helm install <chart>`
- Separates configuration settings from manifest formats
  - Edit the values without changing the rest of the manifest
  - `values.yaml` – Update to deploy the application differently
- Upgrade a running release to a new chart version
  - `$ helm upgrade <release> <chart>`
- Rollback a running release to a previous revision
  - `$ helm rollback <release> <revision>`
- Delete a running release
  - `$ helm delete <release>`

Introduction to Helm

6

© Copyright IBM Corporation 2017

IBM Training 

## Installing Helm

- Helm runs as a CLI client
  - Typically installed on your laptop
- See Installing Helm
  - [https://docs.helm.sh/using\\_helm/#installing-helm](https://docs.helm.sh/using_helm/#installing-helm)
- Options for installing Helm
  1. Download the release, including the binary
    - <https://github.com/kubernetes/helm/releases>
  2. Homebrew on MacOS
    - brew install kubernetes-helm
  3. Installer script
    - curl https://raw.githubusercontent.com/kubernetes/helm/master/scripts/get > get\_helm.sh

Introduction to Helm 7 © Copyright IBM Corporation 2017

IBM Training 

## Using Helm

8 © Copyright IBM Corporation 2017

**IBM Training** 

## Helm Commands

- Install Tiller  
    \$ helm init
- Create a chart  
    \$ helm create <chart>
- List the repositories  
    \$ helm repo list
- Search for a chart  
    \$ helm search <keyword>
- Info about a chart  
    \$ helm inspect <chart>
- Deploy a chart (creates a release)  
    \$ helm install <chart>
- List all releases  
    \$ helm list --all
- Get the status of a release  
    \$ helm status <release>
- Get the details about a release  
    \$ helm get <release>
- Upgrade a release  
    \$ helm upgrade <release> <chart>
- Rollback a release  
    \$ helm rollback <release> <revision>
- Delete a release  
    \$ helm delete <release>

Introduction to Helm 9 © Copyright IBM Corporation 2017

**IBM Training** 

## Working with Repositories

```
$ helm repo list
NAME          URL
stable        https://kubernetes-charts.storage.googleapis.com/
```

```
$ helm search jenkins
NAME          VERSION      DESCRIPTION
stable/jenkins  0.1.14     A Jenkins Helm chart for Kubernetes.
```

```
$ helm repo add my-charts https://my-charts.storage.googleapis.com
$ helm repo list
NAME          URL
stable        https://kubernetes-charts.storage.googleapis.com/
my-charts    https://my-charts.storage.googleapis.com
```

Introduction to Helm 10 © Copyright IBM Corporation 2017

**IBM Training**

**Installing an Application**

- To deploy an application into Kubernetes, install that application's Helm chart

```
$ helm search mysql
NAME      VERSION  DESCRIPTION
stable/mysql  0.1.1   Chart for MySQL

$ helm install stable/mysql
Fetched stable/mysql to mysql-0.1.1.tgz
NAME: loping-toad
LAST DEPLOYED: Thu Oct 20 14:54:24 2016
NAMESPACE: default
STATUS: DEPLOYED

RESOURCES:
=> v1/Secret
NAME          TYPE    DATA  AGE
loping-toad-mysql  Opaque  2    3s

=> v1/Service
NAME          CLUSTER-IP      EXTERNAL-IP  PORT(S)  AGE
loping-toad-mysql  192.168.1.5  <none>     3306/TCP  3s

=> extensions/Deployment
NAME          DESIRED  CURRENT  UP-TO-DATE  AVAILABLE  AGE
loping-toad-mysql  1        0        0           0          3s

=> v1/PersistentVolumeClaim
NAME          STATUS  VOLUME  CAPACITY  ACCESSMODES  AGE
loping-toad-mysql  Pending
```

- Install output
  - Details about the release
  - Details about its resources
- Chart
  - stable/mysql
- Release name
  - loping-toad (auto generated)
- Resources
  - Four total, one of each type
  - All named loping-toad-mysql
  - Secret
  - Service
  - Deployment
  - PersistentVolumeClaim

Introduction to Helm      11      © Copyright IBM Corporation 2017

**IBM Training**

**Overriding Values**

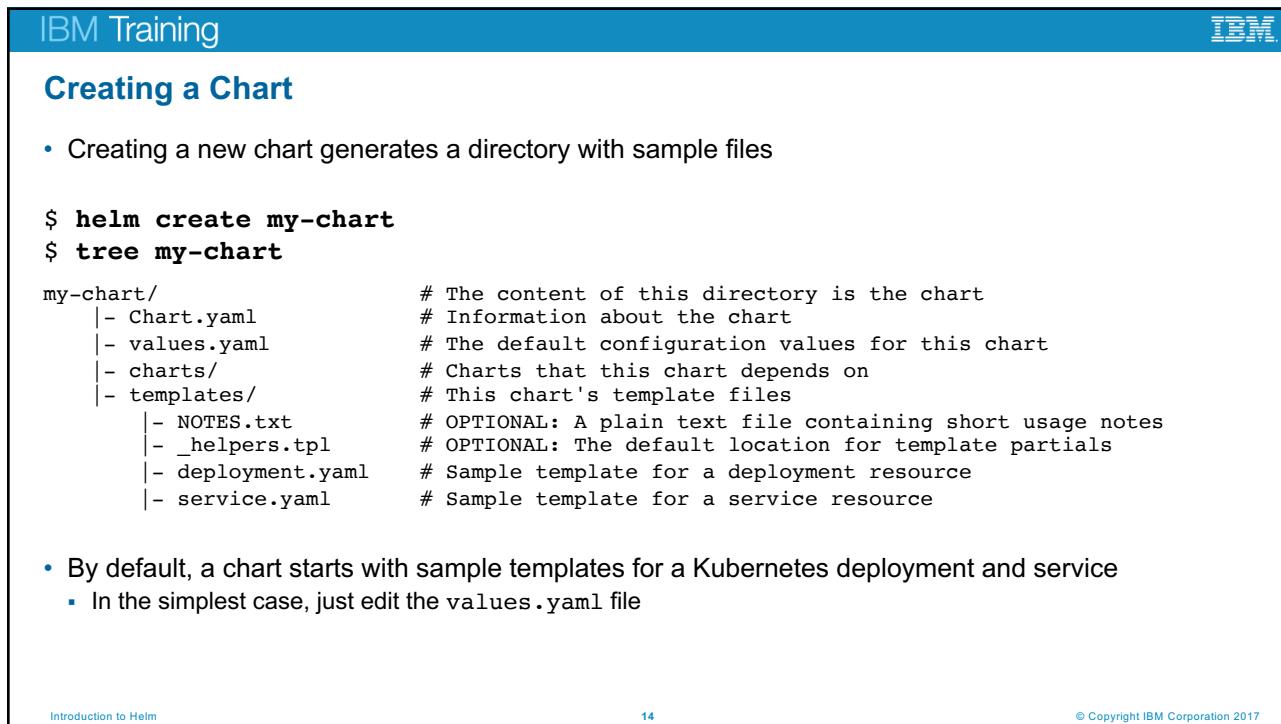
- Default values are stored in the chart  
`<chart-path>/values.yaml`
- Helm CLI uses Kubernetes CLI's configuration to connect to your current cluster  
`~/.kube/config`  
`$ kubectl config view`
- To specify a release's name, use the `name` flag  
`$ helm install --name CustomerDB stable/mysql`
- To deploy the release into a particular Kubernetes namespace, use the `namespace` flag  
`$ helm install --namespace ordering-system stable/mysql`
- To override an individual value, use the `set` flag  
`$ helm install --set user.name='student',user.password='passw0rd' stable/mysql`
- To override values with a values file, use the `values` or `f` flag  
`$ helm install --values myvalues.yaml stable/mysql`

Introduction to Helm      12      © Copyright IBM Corporation 2017



This slide is titled "Developing Charts". It features a large blue header bar with the text "IBM Training" on the left and the IBM logo on the right. The main content area has a light gray diagonal striped background. The title "Developing Charts" is centered in a large, bold, dark blue font.

13 © Copyright IBM Corporation 2017



This slide is titled "Creating a Chart". It features a large blue header bar with the text "IBM Training" on the left and the IBM logo on the right. The main content area contains a bulleted list and a code snippet.

- Creating a new chart generates a directory with sample files

```
$ helm create my-chart
$ tree my-chart
my-chart/
  |- Chart.yaml          # The content of this directory is the chart
  |- values.yaml         # Information about the chart
  |- charts/              # The default configuration values for this chart
  |- templates/           # Charts that this chart depends on
  |  |- NOTES.txt          # This chart's template files
  |  |- _helpers.tpl        # OPTIONAL: A plain text file containing short usage notes
  |  |- deployment.yaml     # OPTIONAL: The default location for template partials
  |  |- service.yaml        # Sample template for a deployment resource
  |  |- deployment.yaml     # Sample template for a service resource
```

- By default, a chart starts with sample templates for a Kubernetes deployment and service
  - In the simplest case, just edit the `values.yaml` file

Introduction to Helm 14 © Copyright IBM Corporation 2017

**IBM Training**

**How Install Uses Charts**

- The main step of installing a chart is rendering its templates
- How Helm installs a chart
  1. User runs an install in the Helm CLI  
\$ helm install myapp
  2. Helm CLI loads the chart into Tiller
  3. **Tiller renders the chart templates**
  4. Tiller loads the resulting resources into Kubernetes
  5. Tiller returns the release data to the client
  6. The client exits
- Rendering the templates
  - Each template generates a Kubernetes resource manifest file (yaml)
  - Tiller runs each of the template files, generating the resource files
- Tiller then loads the resources—as described by the manifests—into the Kubernetes cluster

Introduction to Helm      15      © Copyright IBM Corporation 2017

**IBM Training**

**Chart Lifecycle Hooks**

<p><b>Hooks</b></p> <ul style="list-style-type: none"> <li>• <b>pre-install</b> <ul style="list-style-type: none"> <li>▪ Executes after templates are rendered</li> <li>▪ Before any resources are created in Kubernetes</li> </ul> </li> <li>• <b>post-install</b> <ul style="list-style-type: none"> <li>▪ Executes after all resources are loaded into Kubernetes</li> </ul> </li> <li>• <b>pre-delete</b> <ul style="list-style-type: none"> <li>▪ Executes before any resources are deleted from Kubernetes</li> </ul> </li> <li>• <b>post-delete</b> <ul style="list-style-type: none"> <li>▪ Executes after all of the release's resources have been deleted</li> </ul> </li> <li>• <b>pre-upgrade</b> <ul style="list-style-type: none"> <li>▪ Executes after templates are rendered</li> <li>▪ Before any resources are loaded into Kubernetes</li> </ul> </li> <li>• <b>post-upgrade</b> <ul style="list-style-type: none"> <li>▪ Executes after all resources have been upgraded</li> </ul> </li> <li>• <b>pre-rollback</b> <ul style="list-style-type: none"> <li>▪ Executes after templates are rendered</li> <li>▪ Before any resources have been rolled back</li> </ul> </li> <li>• <b>post-rollback</b> <ul style="list-style-type: none"> <li>▪ Executes after all resources have been modified</li> </ul> </li> </ul>	<p><b>Hooks in the Helm Install Lifecycle</b></p> <ol style="list-style-type: none"> <li>1. User runs an install in the Helm CLI</li> <li>2. Helm CLI loads the chart into Tiller</li> <li>3. Tiller renders the chart templates</li> <li>4. <b>Tiller executes the pre-install hooks</b></li> <li>5. Tiller loads the resulting resources into Kubernetes</li> <li>6. <b>Tiller executes the post-install hook</b></li> <li>7. Tiller returns the release data to the client</li> <li>8. The client exits</li> </ol> <ul style="list-style-type: none"> <li>• A hook can be any Kubernetes resource           <ul style="list-style-type: none"> <li>▪ A hook is often a Kubernetes job</li> <li>▪ Goes in the templates directory</li> </ul> </li> </ul>
---	--

Introduction to Helm      16      © Copyright IBM Corporation 2017

IBM Training 

## Sharing Charts

- A chart is a directory
  - Easy for a Helm client to use the chart directories on the same computer
  - Difficult to share with other users on other computers
- Packaging a chart
  - Bundle `Chart.yaml` and related files into a tar file

```
$ helm package <chart-path>          # Bundles chart directory into a tar file  
$ helm install <chart-name>.tgz      # Installs the chart in the chart file
```
- Chart repository
  - HTTP server that houses an `index.yaml` file and optionally some packaged charts
  - Server can be any HTTP server that can serve YAML and tar files and can answer GET requests
    - Ex: Google Cloud Storage (GCS) bucket, Amazon S3 bucket, Github Pages, or even create your own web server
  - To add a chart to the repository, copy it to the directory and regenerate the index

```
$ helm repo index <charts-path>      # Generates an index of the charts in the repo
```

Introduction to Helm

17

© Copyright IBM Corporation 2017

IBM Training 

## Developing Templates

18

© Copyright IBM Corporation 2017

**IBM Training**

**Creating Templates**

- The main aspect of implementing a chart is implementing its templates
- A related task: Create and populate the files that contain the settings used by the templates
  - These settings files, particularly `values.yaml`, define the chart's API
  - The settings files list the variables the templates can use, therefore the only values worth changing
- Examples of chart templates can be found in <https://github.com/kubernetes/charts/>
  - Each file is a Golang template
  - Includes functions from the Sprig template library
  - A template can create the manifest for any type of Kubernetes resource
- Each file in a chart's `templates` directory is expected to be a template
  - Expected to generate a Kubernetes resource manifest
  - Filename can be anything, should describe the resource it defines
  - Exception: The notes file (i.e. `NOTES.txt`) provides instructions to the chart's users
  - Exception: Files whose names begin with an underscore (e.g. `_helpers.tpl`) are expected to contain partials

Introduction to Helm      19      © Copyright IBM Corporation 2017

**IBM Training**

**Chart Template for Deployment Manifest**

**Kubernetes Deployment Manifest**

```
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.7.9
          ports:
            - containerPort: 80
```

**Helm Deployment Template**

```
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: {{ template "fullname" . }}
  labels:
    app: {{ template "name" . }}
    chart: {{ .Chart.Name }}-{{ .Chart.Version }}
    heritage: {{ .Release.Service }}
    release: {{ .Release.Name }}
spec:
  replicas: {{ .Values.replicaCount }}
  template:
    metadata:
      {{- if .Values.podAnnotations }}
        annotations:
          {{ toYaml $.Values.podAnnotations | indent 8 }}
      {{- end }}
      labels:
        app: {{ template "name" . }}
        release: {{ .Release.Name }}
    spec:
      containers:
        - name: {{ template "name" . }}
          image: "{{ .Values.image.repository }}:{{ .Values.image.tag }}"
          imagePullPolicy: {{ .Values.image.pullPolicy }}
          ports:
            - name: http
              containerPort: 80
              protocol: TCP
      ...
```

Introduction to Helm      20      © Copyright IBM Corporation 2017

**IBM Training** **IBM**

## Chart Template for Service Manifest

**Kubernetes Service Manifest**

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

**Helm Service Template**

```
apiVersion: v1
kind: Service
metadata:
{{- if .Values.service.annotations }}
  annotations:
{{ toYaml .Values.service.annotations | indent 4 }}
{{- end }}
  name: {{ template "fullname" . }}
  labels:
    app: {{ template "name" . }}
    chart: {{ .Chart.Name }}-{{ .Chart.Version }}
    heritage: {{ .Release.Service }}
    release: {{ .Release.Name }}
spec:
  selector:
    app: {{ template "name" . }}
    release: {{ .Release.Name }}
  ports:
    - name: http
      protocol: TCP
      port: {{ .Values.service.port }}
      targetPort: http
{{- if (and (eq .Values.service.type "NodePort") ...) }}
      nodePort: {{ .Values.service.nodePort }}
{{- end }} ...
...
```

Introduction to Helm      21      © Copyright IBM Corporation 2017

**IBM Training** **IBM**

## Values YAML – A Chart's API

**Values (values.yaml)**

```
replicaCount: 1
restartPolicy: Never
# Evaluated by the post-install hook
sleepyTime: "10"
index: >-
  <h1>Hello</h1>
  <p>This is a test</p>
image:
  repository: nginx
  tag: 1.11.0
  pullPolicy: IfNotPresent
service:
  annotations: {}
  clusterIP: ""
  externalIPs: []
  loadBalancerIP: ""
  loadBalancerSourceRanges: []
  type: ClusterIP
  port: 8888
  nodePort: ""
podAnnotations: {}
resources: {}
nodeSelector: {}
```

**Helm Deployment Template**

```
...
spec:
  replicas: {{ .Values.replicaCount }}
  template:
    metadata:
{{- if .Values.podAnnotations }}
      annotations:
{{ toYaml .Values.podAnnotations | indent 8 }}
{{- end }} ...
...
```

**Helm Service Template**

```
...
spec:
  ports:
    - name: http
      protocol: TCP
      port: {{ .Values.service.port }}
      targetPort: http
{{- if (and (eq .Values.service.type "NodePort") ...) }}
      nodePort: {{ .Values.service.nodePort }}
{{- end }} ...
...
```

Introduction to Helm      22      © Copyright IBM Corporation 2017

**IBM Training** 

## Chart YAML – A Chart's Meta Information

<b>Chart (Chart.yaml)</b> <pre> name: nginx description: A basic NGINX HTTP server version: 0.1.0 keywords:   - http   - nginx   - www   - web home: https://github.com/kubernetes/helm sources:   - https://hub.docker.com/_/nginx/ maintainers:   - name: technosophos     email: mbutcher@deis.com </pre>	<b>Helm Template</b> <pre> ... metadata: {{- if .Values.service.annotations --}}   annotations: {{ toYaml .Values.service.annotations   indent 4 }} {{- end --}}   name: {{ template "fullname" . }}   labels:     app: {{ template "name" . }}     chart: {{ .Chart.Name }}-{{ .Chart.Version }}     heritage: {{ .Release.Service }}     release: {{ .Release.Name }} ... </pre>
--	--

Introduction to Helm 23 © Copyright IBM Corporation 2017

**IBM Training** 

## Chart Template Helpers – More Default Settings

<b>Helpers (templates/_helpers.tpl)</b> <pre> {{/* vim: set filetype=mustache: */}} {{/* Expand the name of the chart. */}} {{- define "name" -}} {{- default .Chart.Name .Values.nameOverride   trunc 63   trimSuffix "-" -}} {{- end -}} {{/* Create a default fully qualified app name. We truncate at 63 chars because . . . */}} {{- define "fullname" -}} {{- \$name := default .Chart.Name .Values.nameOverride -}} {{- printf "%s-%s" .Release.Name \$name   trunc 63   trimSuffix "-" -}} {{- end -}} </pre>	<b>Helm Template</b> <pre> ... metadata:   name: {{ template "fullname" . }}   labels:     app: {{ template "name" . }}     chart: {{ .Chart.Name }}-{{ .Chart.Version }}     heritage: {{ .Release.Service }}     release: {{ .Release.Name }} ... </pre>
---	--

Introduction to Helm 24 © Copyright IBM Corporation 2017

IBM Training 

## Chart Predefined Values – More Default Settings

<h3>Predefined Values</h3> <ul style="list-style-type: none"> <li>• <b>Release</b> – Information about the release being created           <ul style="list-style-type: none"> <li>▪ <code>Release.Name</code> – The name of the release (not the chart)</li> <li>▪ <code>Release.Service</code> – The service that conducted the release               <ul style="list-style-type: none"> <li>- Usually this is Tiller</li> </ul> </li> <li>▪ <code>Release.Revision</code> – The revision number               <ul style="list-style-type: none"> <li>- It begins at 1, and increments with each helm upgrade</li> </ul> </li> <li>▪ Lots of other Release values</li> </ul> </li> <li>• <b>Chart</b> – The contents of the <code>Chart.yaml</code> <ul style="list-style-type: none"> <li>▪ <code>Chart.Name</code> – The chart name</li> <li>▪ <code>Chart.Version</code> – The chart version</li> <li>▪ <code>Chart.Maintainers</code> – The maintainers</li> <li>▪ Etc.</li> </ul> </li> <li>• <b>Files</b> – Map of all non-special files in the chart</li> <li>• <b>Capabilities</b> – Map of info about Kubernetes and Helm           <ul style="list-style-type: none"> <li>▪ <code>Capabilities.KubeVersion</code> – Version of Kubernetes</li> <li>▪ <code>Capabilities.TillerVersion</code> – Version of Tiller</li> <li>▪ <code>Capabilities.APIVersions</code> – Kubernetes API versions</li> </ul> </li> <li>• <b>Template</b> – Information about the current template</li> </ul>	<h3>Helm Template</h3> <pre> . . . metadata: {{- if .Values.service.annotations --}}   annotations: {{ toYaml .Values.service.annotations   indent 4 }} {{- end --}}   name: {{ template "fullname" . }}   labels:     app: {{ template "name" . }}     chart: {{ .Chart.Name }}-{{ .Chart.Version }}     heritage: {{ .Release.Service }}     release: {{ .Release.Name }} . . . </pre>
---	--

Introduction to Helm 25 © Copyright IBM Corporation 2017

IBM Training 

## Conclusion

26 © Copyright IBM Corporation 2017

IBM Training 

## Conclusion

- What is Helm?
- Using Helm
- Developing charts
- Developing templates

Introduction to Helm 27 © Copyright IBM Corporation 2017

IBM Training 

## Resources – Introduction

- Helm - The Kubernetes Package Manager
  - <https://helm.sh>
  - <https://docs.helm.sh>
  - <https://github.com/kubernetes/helm>
  - <https://github.com/kubernetes/helm/blob/master/docs/index.md>
- Taking the Helm: Delivering Kubernetes-Native Applications by Michelle Noorali (KubeCon 2016)
  - <https://www.youtube.com/watch?v=zBc1goRfk3k>
- Installing Helm
  - [https://docs.helm.sh/using\\_helm/#installing-helm](https://docs.helm.sh/using_helm/#installing-helm)

Introduction to Helm 28 © Copyright IBM Corporation 2017

IBM Training 

## Resources – Developing Charts

- Helm examples
  - <https://github.com/kubernetes/helm/tree/master/docs/examples>
- Stable Helm charts
  - <https://github.com/kubernetes/charts/tree/master/stable>
- Golang templates
  - <https://golang.org/pkg/text/template>
- Sprig template library
  - <https://godoc.org/github.com/Masterminds/sprig>
- Getting Started Authoring Helm Charts
  - <https://deis.com/blog/2016/getting-started-authoring-helm-charts>
- How to Create Your First Helm Chart
  - <https://docs.bitnami.com/kubernetes/how-to/create-your-first-helm-chart>
- Packaged Kubernetes Deployments – Writing a Helm Chart
  - <https://www.influxdata.com/packaged-kubernetes-deployments-writing-helm-chart>

Introduction to Helm      29      © Copyright IBM Corporation 2017

IBM Training

IBM

## Helm in IBM Cloud Private

Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

IBM Training

IBM

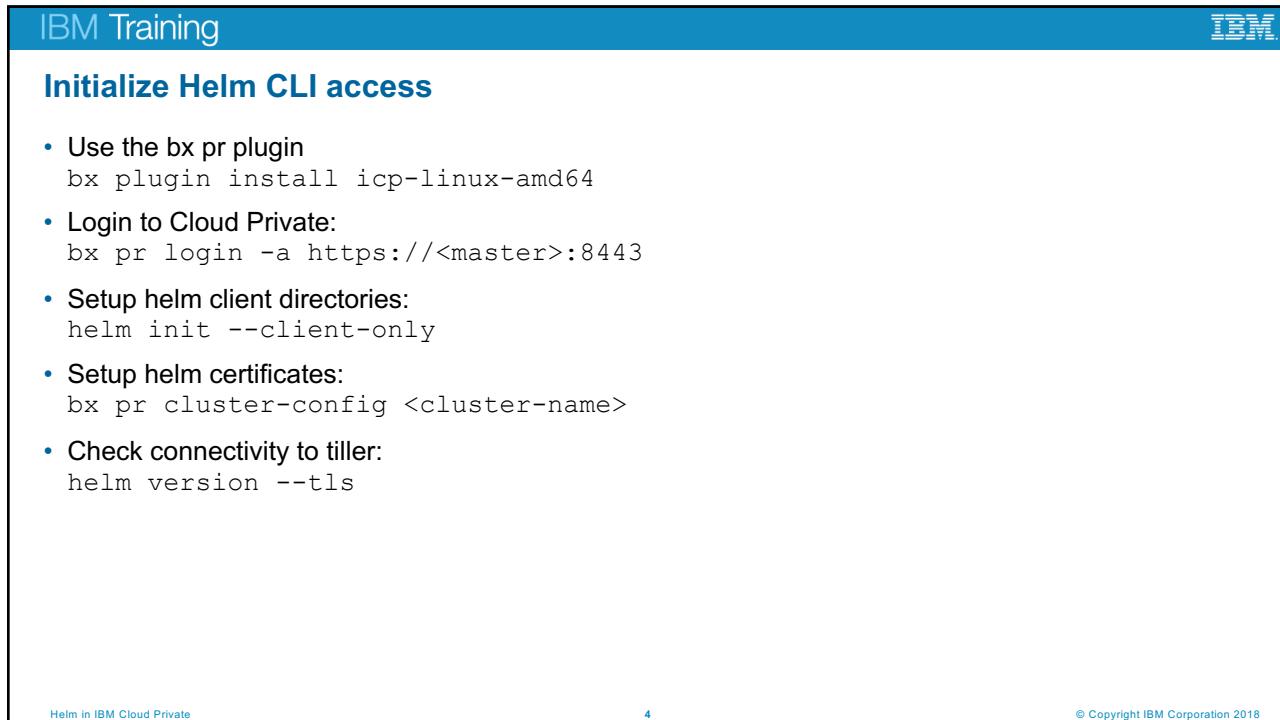
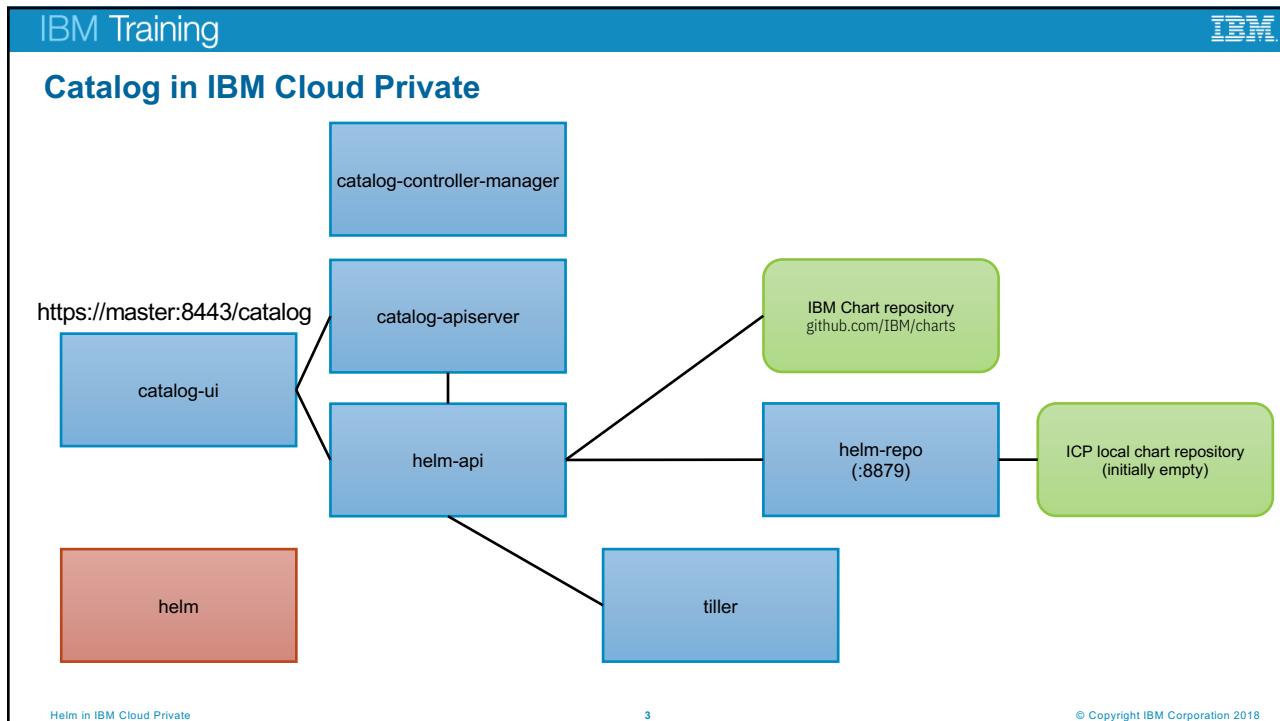
### Objectives

- List IBM Cloud Private catalog feature and usage

Helm in IBM Cloud Private

2

© Copyright IBM Corporation 2018



**IBM Training**

## Helm repositories

- Manage > Helm Repositories

The screenshot shows the 'Helm repositories' page in the IBM Cloud Private interface. At the top, there's a navigation bar with 'IBM Cloud Private', 'Create resource', 'Docs', 'Support', and a user icon. Below the navigation is a search bar and a table titled 'Repositories'. The table has columns for 'NAME' and 'URL'. It contains two entries:

NAME	URL	ACTION
ibm-charts	<a href="https://raw.githubusercontent.com/IBM/charts/master/repo/stable/">https://raw.githubusercontent.com/IBM/charts/master/repo/stable/</a>	⋮
local-charts	<a href="https://icp-management-ingress:8443/helm-repo/charts">https://icp-management-ingress:8443/helm-repo/charts</a>	⋮

At the bottom right of the page is a large blue circular button with a white arrow pointing right.

Helm in IBM Cloud Private 5 © Copyright IBM Corporation 2018

**IBM Training**

## Catalog entries

**Toolchain & Runtimes**

- UrbanCode Deploy
- Microclimate
- Microservice Builder
- Jenkins (open source)
- IBM WebSphere Liberty (MicroProfile, Web Profile, JEE Profile)
- Open Liberty (open source)
- IBM SDK for Node.js
- Swift runtime (open source)
- Nginx (open source)

**Logging & Monitoring Services**

- ELK (open source)
- Prometheus & Grafana (open source)

**App Modernization Tooling**

- IBM Transformation Advisor

**Multi-cloud Management**

- IBM Cloud Automation Manager

**Digital Business Automation**

- IBM Operational Decision Manager
- IBM Operational Decision Manager for Developers

**Mobile**

- IBM Mobile Foundation

**Data Services**

- IBM Db2 Direct Advanced Edition / AESE with Data Server Manager
- IBM Db2 Dev-C
- IBM Data Server Manager (for Db2 Dev-C)
- IBM Db2 Warehouse Enterprise
- IBM Db2 Warehouse Dev-C
- IBM Cloudant Developer Edition
- MongoDB (open source)
- PostgreSQL (open source)
- MariaDB (open source)
- Galera clustering with MariaDB (open source)
- Redis HA Topology (open source)

**Messaging**

- IBM MQ Advanced
- IBM MQ Advanced for Developers
- Rabbit MQ (open source)

**Integration**

- IBM Integration Bus
- IBM Integration Bus for Developers
- IBM DataPower Gateway Virtual Edition
- IBM DataPower Gateway for Developers
- IBM API Connect Professional
- IBM API Connect Enterprise

**Watson**

- IBM Watson Compare & Comply: Element Classification

**Data Science and Business Analytics**

- IBM Data Science Experience Local
- IBM Data Science Experience Developer Edition
- IBM Watson Explorer Deep Analytics Edition

**Data Governance and Integration**

- IBM InfoSphere Information Server for Evaluation

**Management**

- IBM Netcool - integration (Probe for ICP Services – Logging events & Monitoring alerts)

**Connectivity**

- IBM Voice Gateway Developer Trial

**Tooling**

- Web Terminal (open source)
- Skydive – network analyzer (open source)

**HPC / HPDA**

- IBM Spectrum LSF Community Edition
- IBM Spectrum Symphony Community Edition
- IBM Spectrum Conductor Tech Preview

Helm in IBM Cloud Private 6 © Copyright IBM Corporation 2018

IBM Training

IBM

## Helm catalog

- Catalog > Helm Charts

IBM Cloud Private

Create resource Docs Support

### Catalog

Search items Filter

Deploy your applications and install software packages

Chart Name	Description	Version	Chart Type
ibm-cam-prod	IBM Cloud Automation Manager	1.0.0	ibm-charts
ibm-cloudant-dev	Cloudant for Linux.	1.0.0	ibm-charts
ibm-datapower-dev	IBM DataPower Gateway	1.0.0	ibm-charts
ibm-db2oltp-dev	IBM Db2 Developer-C Edition 11.1.3.3	1.0.0	ibm-charts
ibm-db2warehouse-dev	Db2 Warehouse Developer-C for Non-Production v2.5.0	1.0.0	ibm-charts
ibm-dsm-dev	IBM Data Server Manager Developer-C Edition. Note that there can only be one	1.0.0	ibm-charts
ibm-dsx-dev	IBM Data Science Experience (DSX) Developer Edition brings together best of	1.0.0	ibm-charts
ibm-galera-mariadb-dev	Galera Cluster is a multi-master solution for MariaDB which provides an easy-to-	1.0.0	ibm-charts
ibm-icplogging	Log storage and search management solution	1.0.0	ibm-charts

Helm in IBM Cloud Private 7 © Copyright IBM Corporation 2018

IBM Training

IBM

## Building and deploying applications with Kubernetes

© Copyright IBM Corporation 2018  
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

1

IBM Training

IBM

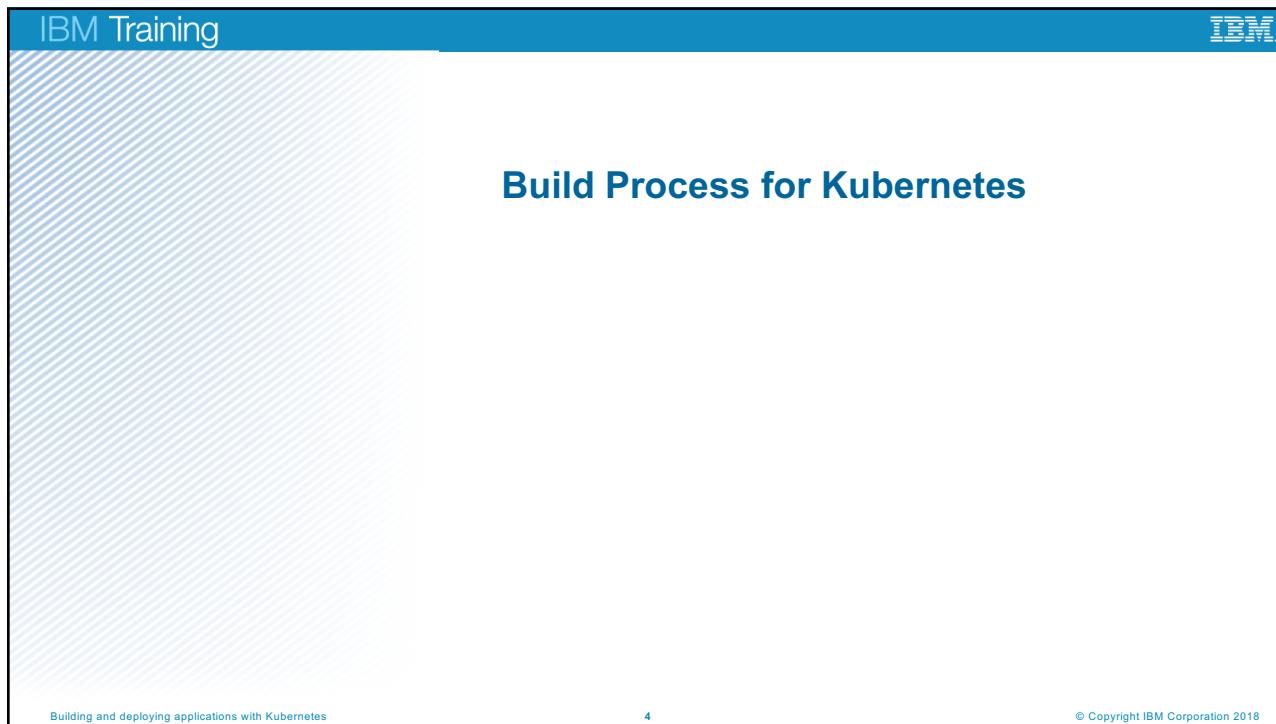
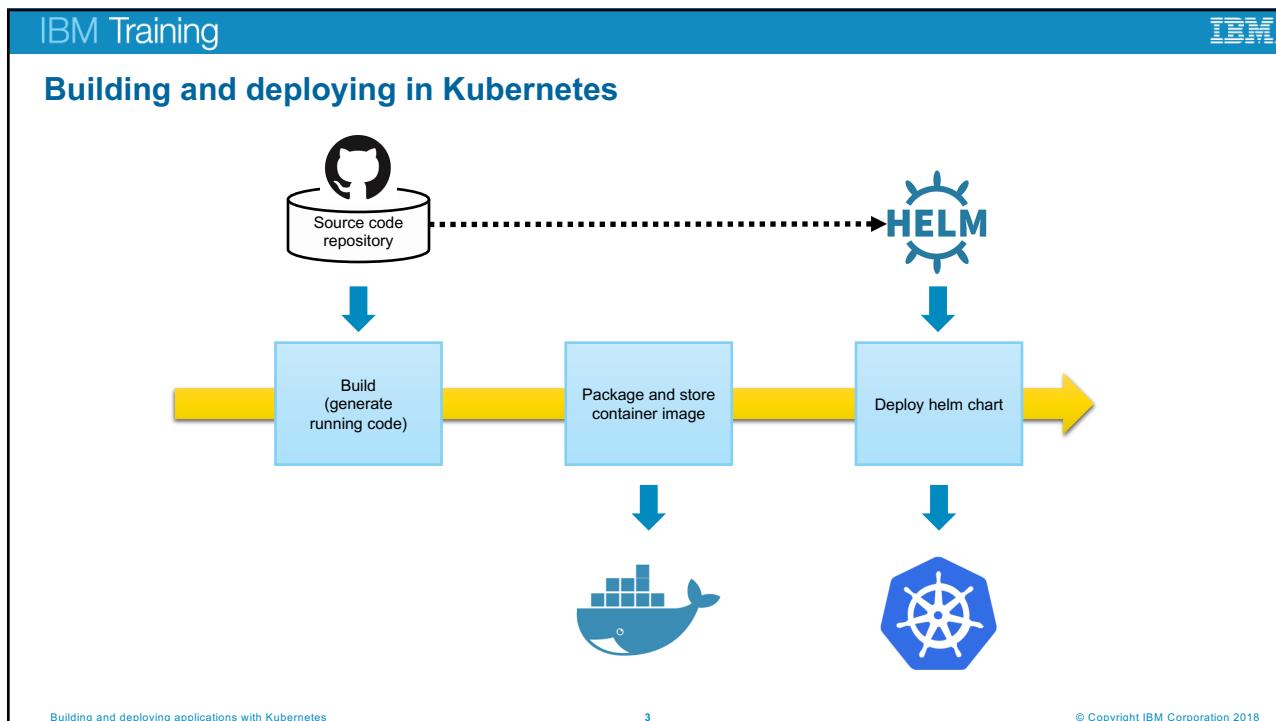
### Objectives

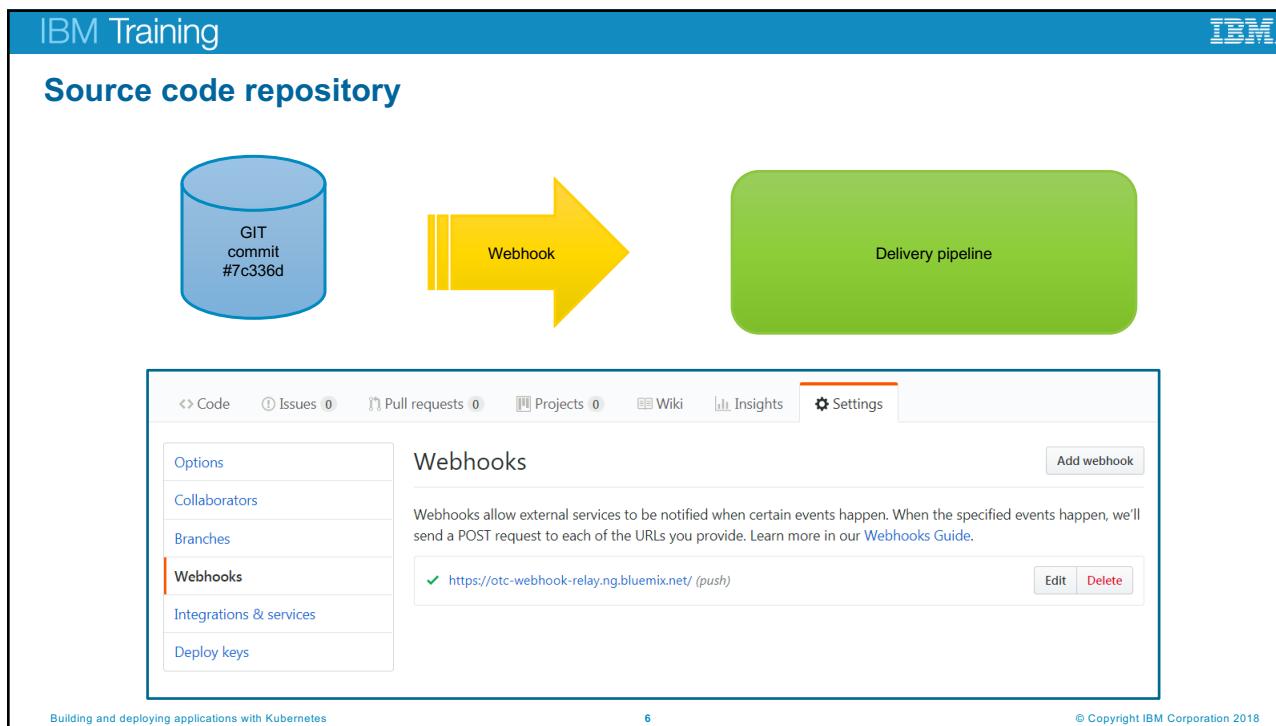
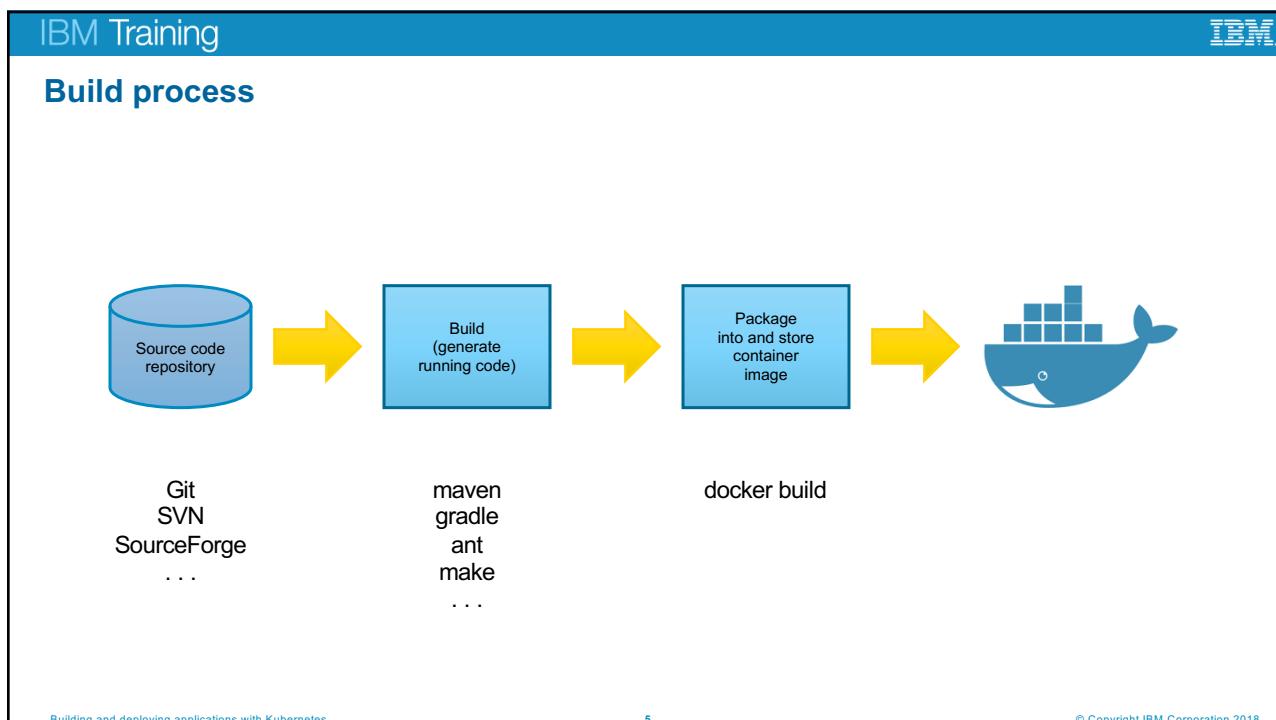
- List the steps to build an application
- Identify steps to re-create and store the container image for an application
- Build a helm chart for an application
- Describe the process to redeploy a helm chart

Building and deploying applications with Kubernetes

2

© Copyright IBM Corporation 2018





## IBM Training

### Build running code

The diagram illustrates three popular build tools:

- Apache Ant:** Represented by a purple mountain icon. It includes a yellow box labeled "build.xml" and a blue box containing the commands "ant compile" and "ant jar".
- Gradle:** Represented by a black elephant icon. It includes a yellow box labeled "gradle.build" and a blue box containing the command "gradlew build".
- Maven:** Represented by a colorful "Maven" logo. It includes a yellow box labeled "pom.xml" and a blue box containing the commands "mvn compile" and "mvn package".

Building and deploying applications with Kubernetes

7

© Copyright IBM Corporation 2018

## IBM Training

### Package container image

The diagram shows the components of a Dockerfile:

- A blue box labeled "Dockerfile".
- A blue box containing the command "docker build -t <image> <path>".
- A small blue icon of a ship with a stack of shipping containers.

Building and deploying applications with Kubernetes

8

© Copyright IBM Corporation 2018

IBM Training 

## Typical Dockerfile flow

- Load an existing image  
FROM
- Copy or add files into the image  
ADD/COPY
- Set an environment variables  
ENV
- Run commands to configure the image  
RUN
- Specify startup command and arguments  
ENTRYPOINT/CMD
- Set port to be available (mapped)  
EXPOSE
- Set initial working directory  
WORKDIR

```
FROM mysql

ADD scripts/load-data.sh load-data.sh
ADD scripts/load-data.sql load-data.sql
RUN chmod u+x load-data.sh

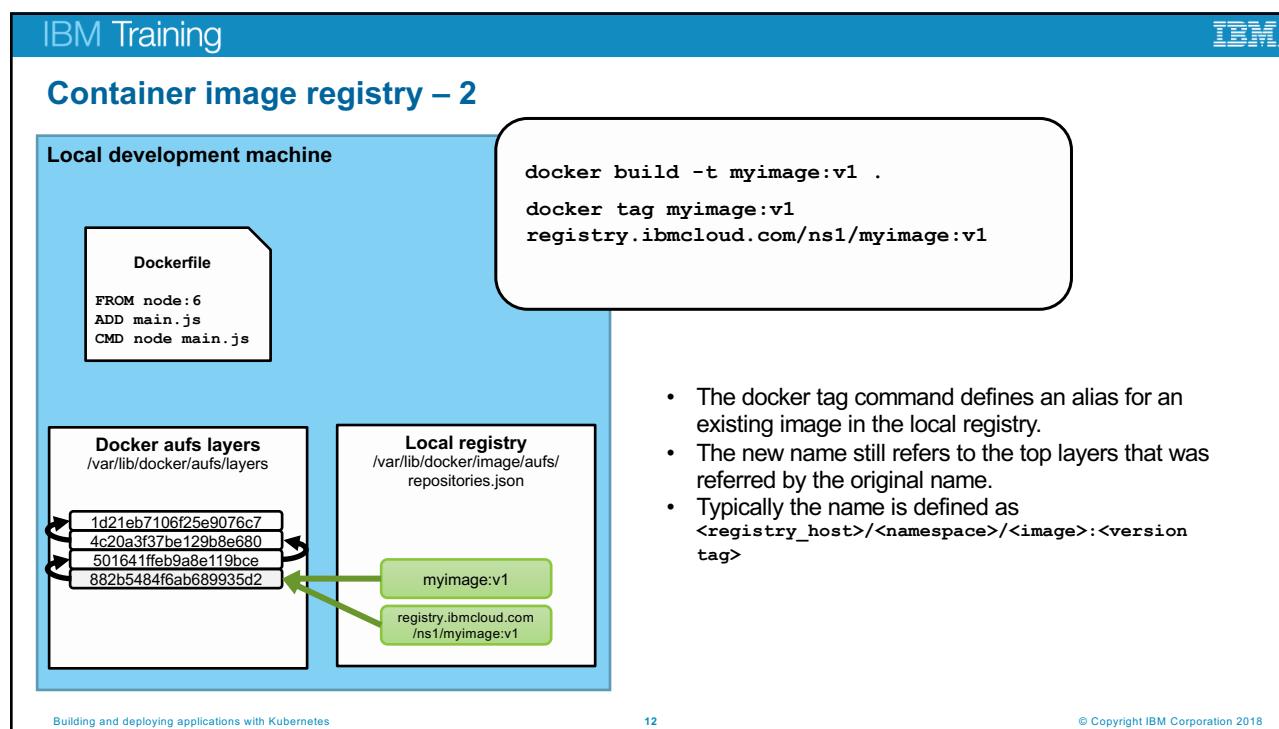
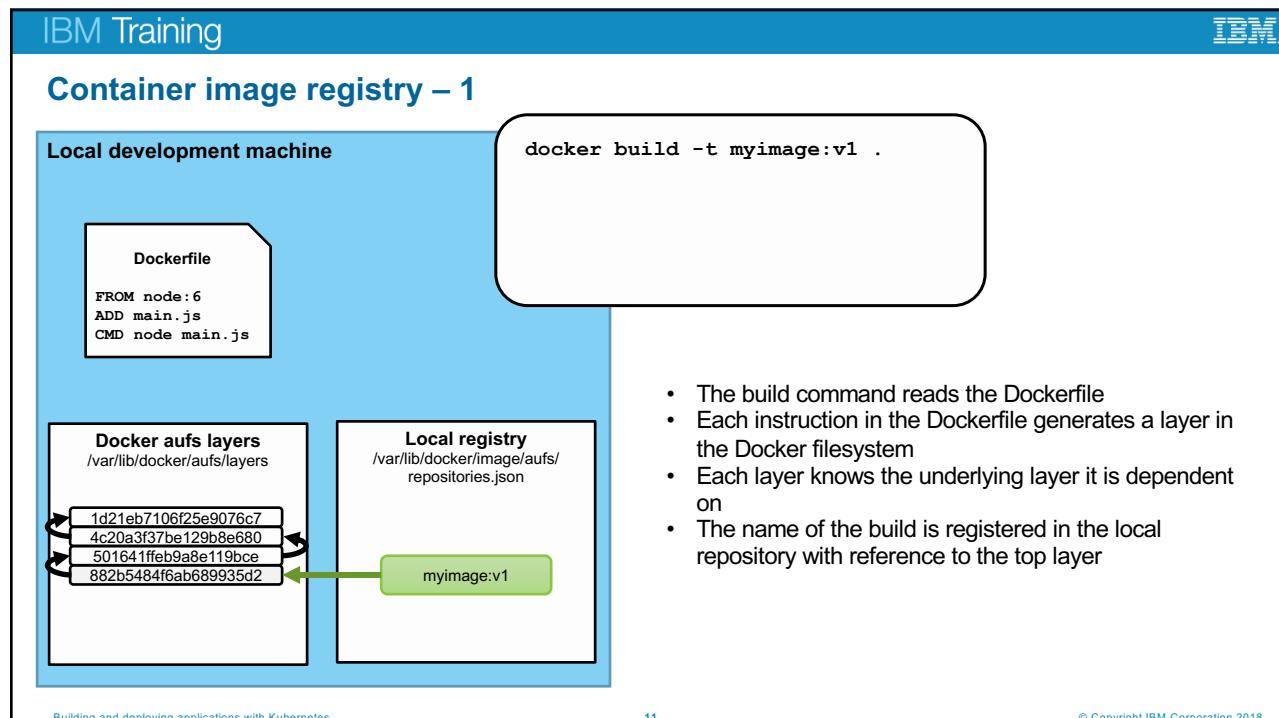
CMD ["mysqld"]
EXPOSE 3306
```

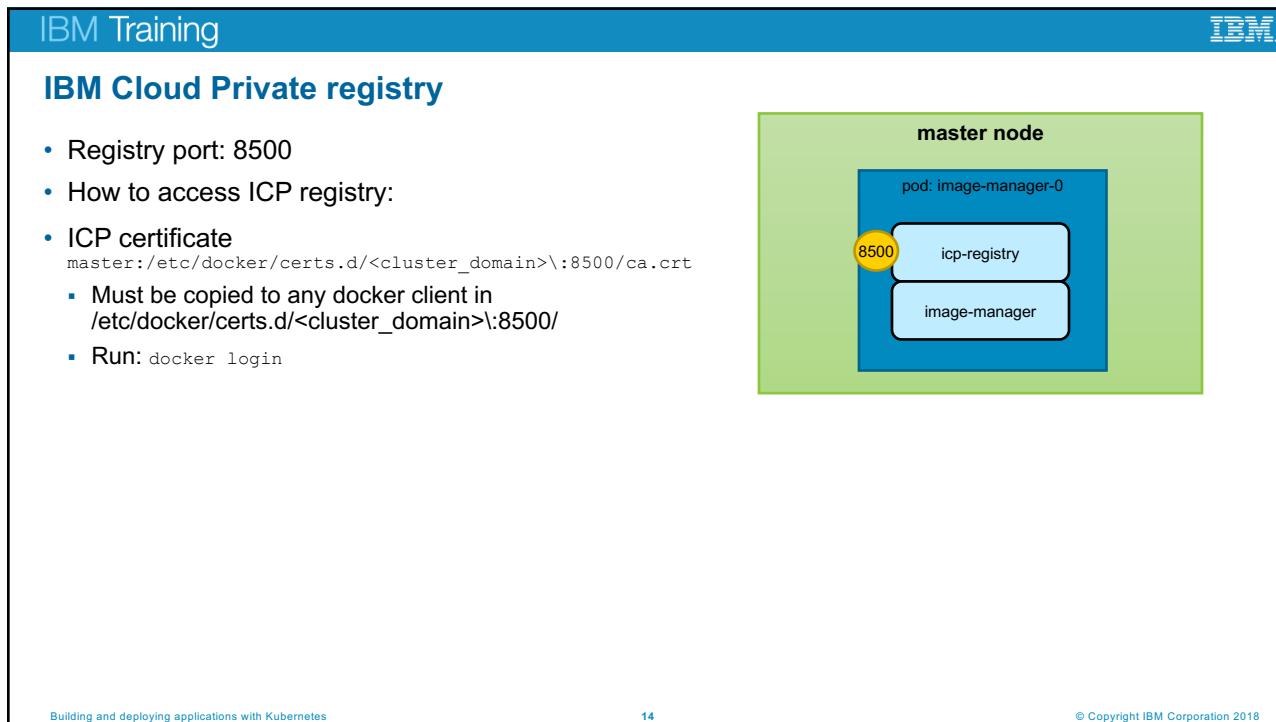
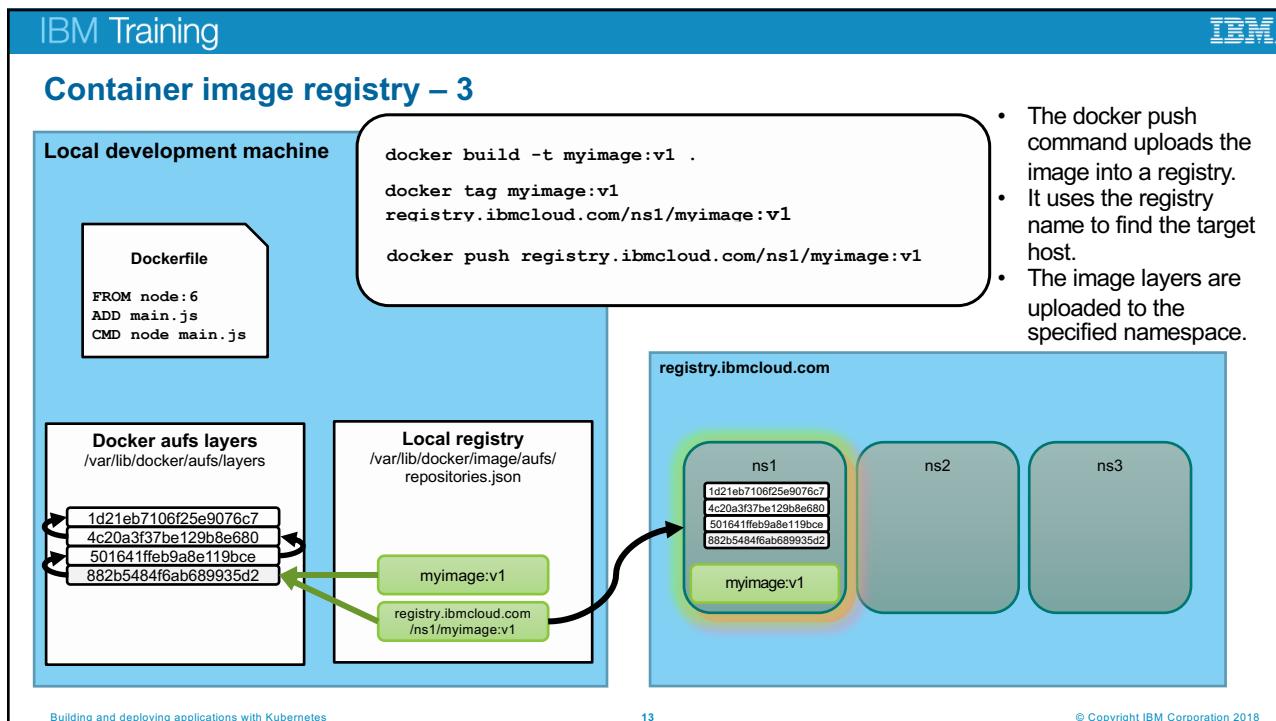
Building and deploying applications with Kubernetes 9 © Copyright IBM Corporation 2018

IBM Training 

## Working with a container image

Building and deploying applications with Kubernetes 10 © Copyright IBM Corporation 2018





## Using IBM Cloud Private registry

- Pulling image out from a private (not global registry entry)
- ImagePullSecrets
  - Create secret:  
kubectl create secret docker-registry myregistrykey  
--docker-server=<cluster\_CA\_domain>:8500  
--docker-username=<user\_name>  
--docker-password=<user\_password>  
--docker-email=<user\_email>
  - Use in template:  
template:  
spec:  
imagePullSecrets:  
- name: myregistrykey

master node

pod: image-manager-0

8500 icp-registry

image-manager

Building and deploying applications with Kubernetes

15

© Copyright IBM Corporation 2018

## Using Helm charts for deployment

Building and deploying applications with Kubernetes

16

© Copyright IBM Corporation 2018

## IBM Training

### Helm chart components

```
helm create <chart_name>
```

```
localuser@ubuntu-base:~$ tree catalog
catalog
├── charts
│   └── Chart.yaml
└── templates
    ├── deployment.yaml
    │   └── helpers.tpl
    ├── ingress.yaml
    └── NOTES.txt
    └── service.yaml
    └── values.yaml
2 directories, 7 files
```

- The `create` sub-command build the directory structure
- `Chart.yaml` and `values.yaml` are used for variable substitution
- `templates` directory contains resources that this Helm chart built
- `templates` also contains lifecycle scripts (will be discussed later)

Building and deploying applications with Kubernetes 17 © Copyright IBM Corporation 2018

## IBM Training

### Helm chart life cycle

```
graph LR; chart_tgz[chart.tgz] -- install --> release_v1[release-v1]; release_v1 -- upgrade --> release_v2[release-v2]; release_v2 -- rollback --> release_v1; release_v1 -- delete --> chart_tgz
```

The diagram illustrates the Helm chart life cycle. It shows three main components: `chart.tgz` (blue box), `release-v1` (orange box), and `release-v2` (orange box). The flow of the life cycle is as follows:

- `install`: `chart.tgz` leads to `release-v1`.
- `upgrade`: `release-v1` leads to `release-v2`.
- `rollback`: `release-v2` leads back to `release-v1`.
- `delete`: `release-v1` leads back to `chart.tgz`.

Building and deploying applications with Kubernetes 18 © Copyright IBM Corporation 2018

**IBM Training** **IBM**

## Chart Lifecycle Hooks

- pre-install
  - Executes after templates are rendered
  - Before any resources are created
- post-install
  - Executes after all resources are loaded
- pre-delete
  - Executes before any resources are deleted
- post-delete
  - Executes after all of the release's resources have been deleted
- pre-upgrade
  - Executes after templates are rendered
  - Before any resources are loaded
- post-upgrade
  - Executes after all resources have been upgraded
- pre-rollback
  - Executes after templates are rendered
  - Before any resources have been rolled back
- post-rollback
  - Executes after all resources have been modified

```
apiVersion: batch/v1
kind: Job
metadata:
  name: "{{ .Release.Name }}-initialize-tables"
  annotations:
    # This is what defines this resource as a hook. Without this line, the
    # job is considered part of the release.
    "helm.sh/hook": post-install
    "helm.sh/hook-weight": "30"
spec:
  template:
    metadata:
      name: {{ .Release.Name }}-initialize-tables
    spec:
      restartPolicy: Never
      containers:
        - name: initialize-tables
          image: "ibm-cloud-academy/ibmcloudcli:v01"
          command: [ "/bin/sh", "-c" ]
          args: ["git clone https://github.com/ibm-cloud-academy/tools;cd tools;./initdb.sh"]
          imagePullPolicy: IfNotPresent
```

Building and deploying applications with Kubernetes 19 © Copyright IBM Corporation 2018

**IBM Training** **IBM**

## Microservice in Kubernetes – resources

The Helm chart must contain the templates for each of these Kubernetes resources

Building and deploying applications with Kubernetes 20 © Copyright IBM Corporation 2018

IBM Training IBM

## Deploying a helm chart

Building and deploying applications with Kubernetes 21 © Copyright IBM Corporation 2018

IBM Training IBM

## Checking chart

The **lint** sub-command checks and verifies the chart definitions.

```
helm lint <chart_name>
```

Options

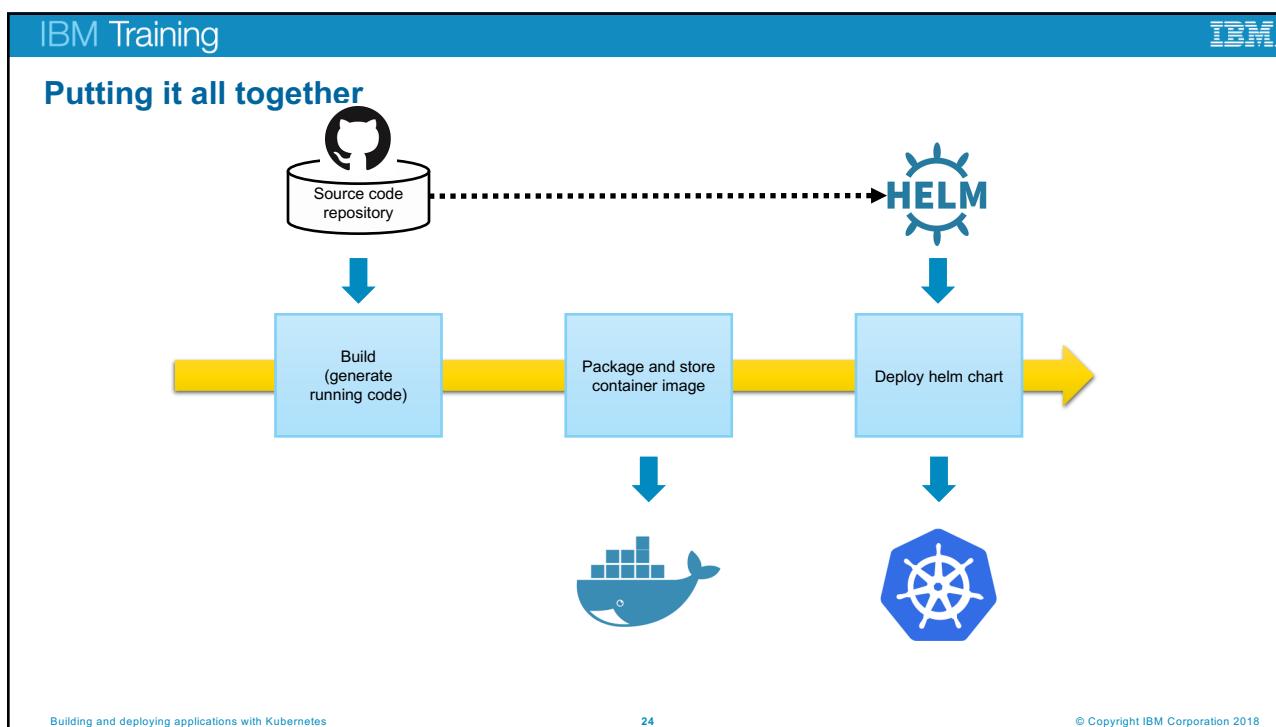
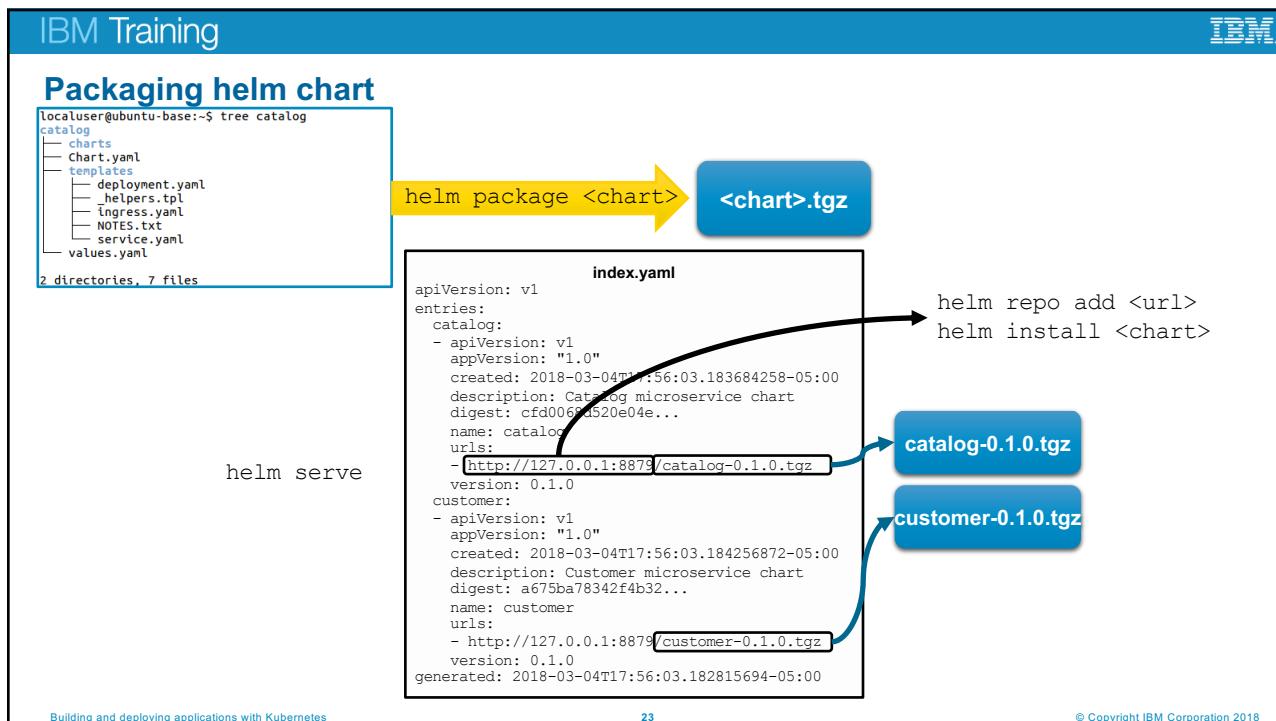
```
--namespace string      namespace to install the release into (only used if -  
--set stringArray      set values on the command line (can specify multiple  
--strict               fail on lint warnings  
-f, --values valueFiles  specify values in a YAML file (can specify multiple)
```

Options inherited from parent commands

```
--debug                enable verbose output  
--home string           location of your Helm config. Overrides $HELM_HOM  
--host string            address of Tiller. Overrides $HELM_HOST  
--kube-context string   name of the kubeconfig context to use  
--tiller-namespace string  namespace of Tiller (default "kube-system")
```

localuser@ubuntu-base:~\$ helm lint catalog  
==> Linting catalog  
[INFO] Chart.yaml: icon is recommended  
1 chart(s) linted, no failures

Building and deploying applications with Kubernetes 22 © Copyright IBM Corporation 2018



IBM Training

IBM

## DevOps Concepts

© Copyright IBM Corporation 2017  
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

IBM Training

IBM

### Agenda

- What is DevOps
- Definitions - CI / CD
- Zero Downtime Deployment
- DevOps in Twelve-Factor Apps
- Implementing DevOps

DevOps Concepts

2

© Copyright IBM Corporation 2017

IBM Training

IBM

## What is DevOps?

3

© Copyright IBM Corporation 2017

IBM Training

IBM

### Definition

**DevOps** (*n.*) –

“DevOps is a philosophy, a **cultural shift** that **merges operations with development** and demands a **linked toolchain** of technologies to facilitate collaborative change. DevOps toolchains ... can include dozens of non-collaborative tools, making the task of automation a technically complex and arduous one.”

Gartner

The diagram illustrates the Gartner view of DevOps as a linked toolchain. It features three green rectangular boxes labeled "Customer", "Development", and "Operations". Between the "Customer" and "Development" boxes is a light blue circle labeled "Agile". Between the "Development" and "Operations" boxes is another light blue circle labeled "DevOps". The word "Gartner" is written in the top right corner of the slide area.

Customer      Agile      Development      DevOps      Operations

DevOps Concepts

4

© Copyright IBM Corporation 2017

**IBM Training**

**Every organization has critical business apps**

Trading halted for half a day on the biggest US exchange for financial options following an outage caused by software problems.

Airline canceled more than 700 flights and another 765 flights are delayed due to a software outage – Blamed ticketing partner while the real problem was on their end

Software problem led to two days of downtime at the largest bank in Europe has tarnished their image as the most reliable banking website.

Not surprisingly, many angry customers poured out their wrath via social networking after the largest video streaming company had a software outage for more than 20 hours

and in today's competitive world, those apps can never go DOWN!!!

DevOps Concepts

5

© Copyright IBM Corporation 2017

**IBM Training**

**Evolving perspective on DevOps**

Large Enterprises Traditional Workloads (historical view)

Large Enterprises with Traditional Workloads Moving to Cloud

Newer companies with Cloud-Native Workloads

DevOps Integrated Team. Full Lifecycle

DevOps Concepts

6

© Copyright IBM Corporation 2017

**IBM Training**

## Newer approaches becoming mainstream

**2015**

Large Enterprises experimenting with full integrated DevOps lifecycle for cloud-native workloads with small innovation teams.

How to scale to the enterprise?

**2016 / 2017**

Several testimonials of large enterprises applying integrated DevOps to real projects.

Message expanding beyond just Continuous Integration / Continuous Delivery to also include reliability, security and insights.

Containers and microservice architectures more visible

DevOps toolchains enter mainstream

ElectricCloud What We Offer Success Stories Plugins Resources

Plug in the tools you already use.

- Across LOBs, Shared Services and 3rd Parties
- Tested end-to-end
- All dependencies are satisfied

Deliver High Quality Working Software Faster

- No security flaws
- No legal flaws
- Minimum defects
- All levels of testing done
- Code reviewed and source controlled

- How fast! ASAP!

Pipeline must have 16 gates

- Source code version control
- Optimum branching strategy
- Static analysis
- > 100% coverage
- Vulnerability scan
- Open source scan
- Artifact version control
- Auto promotion
- Intelligent triggers
- Integration testing
- Performance testing
- Build, Deploy, Testing automated for every commit
- Automated deployment
- Zero downtime release
- Feature Toggle

© Copyright IBM Corporation 2017

DevOps Concepts

**IBM Training**

## Definitions - CI / CD

8

© Copyright IBM Corporation 2017

**IBM Training** **IBM**

## Definitions

- **CI / CD**
  - Acronym for **Continuous Integration** plus **Continuous Deployment**
- **Continuous integration**
  - Creating packaged builds of the latest code changes as soon as they're available
- **Continuous deployment**
  - Progressing each new packaged build through the deployment lifecycle stages as rapidly as possible, resulting in that package getting into production
- **Continuous delivery**
  - Continuous integration + continuous deployment
- **Delivery pipeline**
  - An automated sequence of steps to perform CI / CD

DevOps Concepts 9 © Copyright IBM Corporation 2017

**IBM Training** **IBM**

## Continuous Integration

Frequently performing all of these steps in sequence:

- **Development**
  - Rapidly implementing changes in small, tested batches
- **SCM (Source Code Management)**
  - Merging changes from multiple developers
  - Tools like GitHub, SVN, etc
- **Build**
  - Creating new deployment artifacts
  - Tools like Jenkins, Gradle, Maven, etc
- **Package**
  - Installing builds into runtimes
  - Releasing runtimes as immutable images
  - Cloud Foundry push, Building container images

DevOps Concepts 10 © Copyright IBM Corporation 2017

**IBM Training**

## Continuous integration – package step

In the package step, an immutable image is created

- An immutable image does not get changed – only used for deploying instances
- If it needs to change, you delete it and create a new one

**Examples**

- Docker containers
  - A container image is created from a Dockerfile
  - After that, it is only deployed as a container, not changed
  - If you need to make changes, make a new container image by creating a new build and running the Dockerfile again
- Cloud Foundry
  - When an app is pushed to Cloud Foundry, the buildpack creates a droplet – from that CF creates Garden container instances that run the app
  - If you need to make changes, create a new build and push it

The diagram illustrates the Docker build process and the Cloud Foundry deployment architecture. On the left, a 'Dockerfile' icon leads to a 'Docker Image' icon via a 'docker build' arrow. On the right, a 'Cloud Foundry' box contains a 'Cloud Foundry Cloud Controller' box, a 'Buildpack' box, and a 'Diego' box. Arrows show the flow from 'cf push' to the controller, which then interacts with the buildpack and diego components. Below the controller are four 'Cell / Container' boxes.

DevOps Concepts      11      © Copyright IBM Corporation 2017

**IBM Training**

## Continuous Deployment

Rapidly progressing the latest packaged build through the test lifecycle stages and into production

- Deploy to Test
  - Perform functional testing
  - Automated use of test tools
- Deploy to Stage
  - Rehearse production deployment
  - Perform integration testing
- Deploy to Prod
  - Make the build available to users

The diagram shows a 'Release Train' moving through 'Test', 'Stage', and 'Prod' environments. It starts with a 'Package Repo' leading to a 'Deploy' step, then moving through 'Test', 'Stage', and finally 'Prod'. Above the train, a 'Configure' section lists: Applications, Middleware, Databases, APIs. To the right, a 'Rules for promotion' section shows three circular arrows labeled 'Automated Test Deployments' pointing from Test to Stage, Stage to Prod, and Prod back to Test. A final 'Release' step is at the end.

DevOps Concepts      12      © Copyright IBM Corporation 2017

