

IBM Training 

Zero Downtime Deployment

Deploying a new version of an application to replace an old version without incurring a service outage

- Characteristics of zero downtime deployment
 - Application is always available, first old version, then new version
 - User does not experience an outage (a.k.a. downtime)
 - Old version and new version are both deployed at the same time – traffic is directed to both
 - User's only indication of new deployment is improved functionality
- Necessitated by DevOps' practice of frequent redeployment
 - Unacceptable for each new deployment to cause an outage

DevOps Concepts 17 © Copyright IBM Corporation 2017

IBM Training 

Implementing Zero Downtime Deployment – Blue Green Deployment

- Some challenges of continuous deployment
 - Application availability during cut-over
 - Rapidly reverting to prior version if necessary
- A zero-downtime deployment technique
 - Also known as Red Black Deployment (Netflix term)
- Two nearly identical production environments, called Blue and Green
 1. Deploy v1 (blue environment), which users use
 2. Deploy v2 (green environment)
 3. Run smoke tests and automated tests to verify the new environment before taking it live
 4. Shift client load from Blue to Green
 5. Delete Blue (or keep as a hot backup)
- Kubernetes incorporates “rolling updates” as its default method of deployment
 - Updates are incrementally rolled out pod by pod, so some pods have the “old” version and some have the “new” version – traffic can be managed between these pods
 - Updates can be rolled back if there is a problem

“Blue green deployments”
<https://www.ng.bluemix.net/docs/manageapps/updapps.html>

DevOps Concepts 18 © Copyright IBM Corporation 2017

IBM Training

IBM

DevOps in Twelve-Factor Apps

19 © Copyright IBM Corporation 2017

IBM Training

IBM

The Twelve-Factor App

- A set of best practices for creating applications
 - Implementing, deploying, monitoring, and managing
- Typical modern applications
 - Deployed in the cloud
 - Accessible as web applications that deliver software-as-a-service
- Can be applied to any application
 - Implemented in any programming language
 - Using any backing services (database, messaging, caching, and so on)
- Addresses common problems
 - The dynamics of the organic growth of an app over time
 - The dynamics of collaboration between developers
 - Avoiding the cost of software erosion
 - Systemic problems in modern application development
- Provides a shared vocabulary for addressing these problems

Sources: <http://www.12factor.net>
<http://www.clearlytech.com/2014/01/04/12-factor-apps-plain-english/>

20 © Copyright IBM Corporation 2017

IBM Training **IBM**

Twelve Factors related to DevOps

- I. **Codebase:** One codebase that is tracked in revision control, with many deployments
- II. **Dependencies:** Explicitly declare and isolate dependencies
- III. **Configuration:** Store Configuration in the environment
- IV. **Backing services:** Treat backing services as attached resources
- V. **Build, release, run:** Strictly separate build and run stages
- VI. **Processes:** Execute the app as one or more stateless processes
- VII. **Port binding:** Export services with port binding
- VIII. **Concurrency:** Scale out using the process model
- IX. **Disposability:** Maximize robustness with fast startup and efficient shutdown
- X. **Development and production parity:** Keep development, staging, and production as similar as possible
- XI. **Logs:** Treat logs as event streams
- XII. **Admin processes:** Run administrative and management tasks as one-off processes

DevOps Concepts 21 © Copyright IBM Corporation 2017

IBM Training **IBM**

Twelve Factors related to DevOps

- I. **Codebase:** One codebase that is tracked in revision control, with many deployments
- II. **Dependencies:** Explicitly declare and isolate dependencies
- III. **Configuration:** Store Configuration in the environment
- IV. **Backing services:** Treat backing services as attached resources
- V. **Build, release, run:** Strictly separate build and run stages
- VI. **Processes:** Execute the app as one or more stateless processes
- VII. **Port binding:** Export services with port binding
- VIII. **Concurrency:** Scale out using the process model
- IX. **Disposability:** Maximize robustness with fast startup and efficient shutdown
- X. **Development and production parity:** Keep development, staging, and production as similar as possible
- XI. **Logs:** Treat logs as event streams
- XII. **Admin processes:** Run administrative and management tasks as one-off processes

The diagram illustrates how various DevOps practices relate to a delivery pipeline. Arrows point from specific factors to callout boxes detailing the pipeline's behavior:

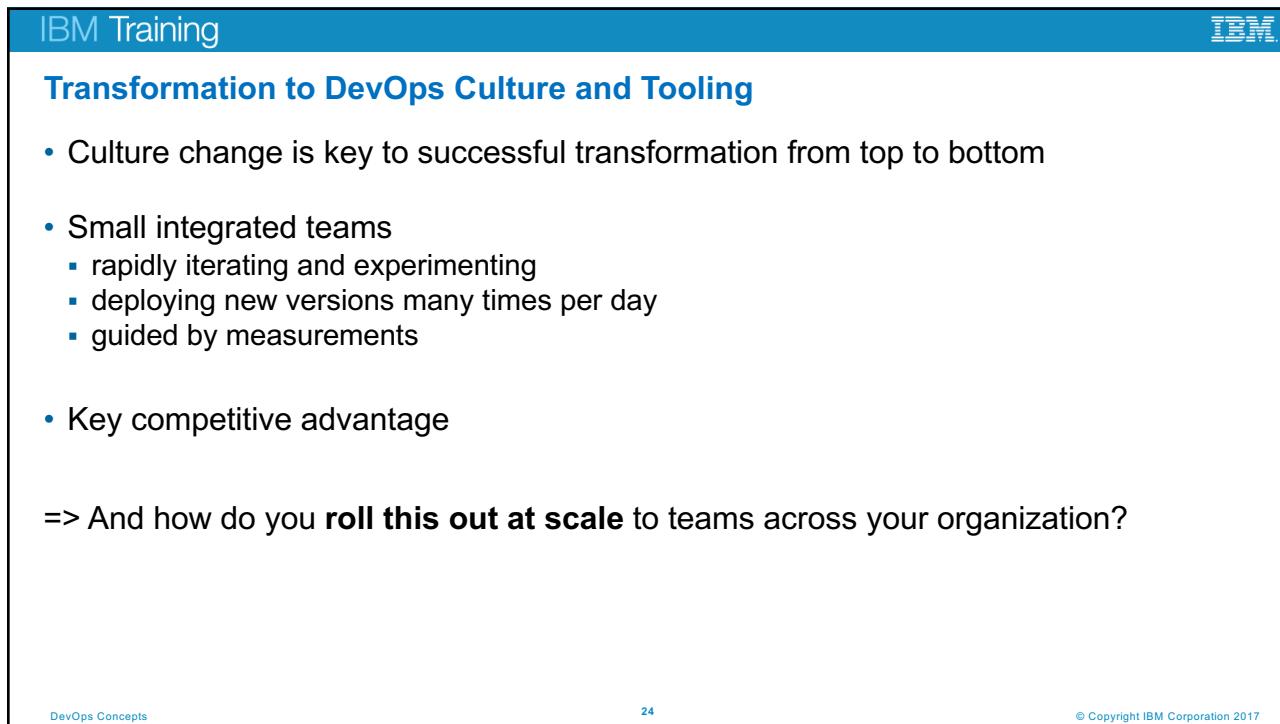
- Codebase:** Points to a box stating "Delivery pipeline is triggered by SCM such as Git repos".
- Build, release, run:** Points to a box stating "Delivery pipeline uses environment variables, both system and user".
- Processes:** Points to a box stating "Delivery pipeline builds deployment artifacts, packages immutable images, and deploys to lifecycle stages (e.g. test, stage, prod)".
- Development and production parity:** Points to a box stating "Delivery pipeline deploys same immutable image to each lifecycle stage".

DevOps Concepts 22 © Copyright IBM Corporation 2017



The slide features a blue header bar with "IBM Training" on the left and the IBM logo on the right. The main content area has a light gray diagonal striped background. The title "Implementing DevOps" is centered in a large, bold, dark blue font. At the bottom center, there is small text: "23" on the left, "© Copyright IBM Corporation 2017" on the right, and a small blue link "DevOps Concepts" at the very bottom.

Implementing DevOps



The slide features a blue header bar with "IBM Training" on the left and the IBM logo on the right. The main content area contains the title "Transformation to DevOps Culture and Tooling" in a bold, dark blue font. Below it is a bulleted list of transformation steps:

- Culture change is key to successful transformation from top to bottom
- Small integrated teams
 - rapidly iterating and experimenting
 - deploying new versions many times per day
 - guided by measurements
- Key competitive advantage

A question follows: "**=> And how do you roll this out at scale to teams across your organization?**" At the bottom, there is small text: "DevOps Concepts" on the left, "24" in the center, and "© Copyright IBM Corporation 2017" on the right.

DevOps Concepts
24
© Copyright IBM Corporation 2017

IBM Training

Implementing DevOps is not all about technology

You need new practices and you need tools

Method
IBM Cloud Garage Method

Integrated Toolchain
IBM Cloud Continuous Delivery

Innovation Platform
IBM Cloud

DevOps Concepts 25 © Copyright IBM Corporation 2017

IBM Training

Method - IBM Cloud Garage Method

ibm.com/cloud/garage

Combining Industry Best Practices for Design Thinking, Lean Startup, Agile Development, DevOps and Cloud to build and deliver innovative solutions

Reference architectures show how services work together, with sample implementations

Toolchains support the practices and architectures

Accelerate delivery of innovation to the market

Practices
IBM's Garage Method combines the best practices from Design Thinking, Agile development, Lean Startup, and DevOps to build innovative solutions.

Culture Think Code Deliver Run Manage Learn

Putting it all together

Architectures & Technologies
Get started with your favorite languages and runtimes, and deploy to the cloud within minutes. Incrementally add capabilities to build enterprise scale applications using our reference architectures.

Practices & Toolchains
Transform your teams and make them more productive with best practices and open toolchains that span the entire DevOps lifecycle.

The Garage
Get hands-on experience at one of many Garage locations.

Cloud Adoption & Transformation Framework
Coming soon

Build and scale with examples and proven architectures

Architecture Center Proven architectures at enterprise scale	Cognitive	Microservices	Data and analytics	Mobile
IoT	Virtualization	Hybrid Cloud	Security	Service management
DevOps				

DevOps Concepts 26 © Copyright IBM Corporation 2017

IBM Training

Integrated Toolchain - Built on the Open Toolchain foundation

A sample open toolchain for building, and deploying and managing three microservices

Toolchains provide an integrated set of tools that support the best practices to build, deploy and manage your apps

You can create toolchains that include IBM Cloud services, open source tools, and third-party tools that make development and operations repeatable and easier to manage

Rapidly instantiate new toolchains from templates to on-board new teams quickly

DevOps Concepts 27 © Copyright IBM Corporation 2017

IBM Training

Innovation Platform - IBM Cloud

DevOps Tooling

Your Own Hosted Apps / Services

Catalog of Services that Extend Apps' Functionality

- Web
- Data
- Mobile
- Cognitive
- Analytics
- IoT
- Security
- Yours

Integration & API Mgmt

Flexible Compute Options to Run Apps

- Virtual Machines (openstack)
- Containers (docker)
- Instant Runtimes (Cloud Foundry)
- Serverless

Platform Deployment Options to meet Workload Requirements

- Public
- Dedicated
- IBM Cloud Private

IBM Cloud Infrastructure Your Data Center

DevOps Concepts 28 © Copyright IBM Corporation 2017

IBM Training

IBM

Conclusion

29

© Copyright IBM Corporation 2017

IBM Training

IBM

Summary

- What is DevOps
- Definitions - CI / CD
- Zero Downtime Deployment
- DevOps in Twelve-Factor Apps
- Implementing DevOps

DevOps Concepts

30

© Copyright IBM Corporation 2017

IBM Training

IBM

Jenkins in IBM Cloud Private

© Copyright IBM Corporation 2018
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

1

IBM Training

IBM

Objectives

- Describe Jenkins components and architecture
- Describe how the Jenkins interact with Git
- Explain what a Jenkins pipeline is
- Build pipeline for build and deployment in IBM Cloud Private

Jenkins in ICP

2

© Copyright IBM Corporation 2018

IBM Training IBM

Jenkins components

Jenkins in ICP 3 © Copyright IBM Corporation 2018

IBM Training IBM

Jenkins components

The diagram illustrates the Jenkins master-slave architecture. It shows a single master node on the left and three slave nodes on the right. The master node contains a cartoon character of a man in a tuxedo holding a white mug, with the word "master" written above it. To its right is a blue Jenkins logo. Three lines connect the master node to the first slave node. The slave nodes are represented by three stacked rectangles. Each slave node contains a smaller master icon, a Jenkins logo, and three yellow rectangular boxes labeled "executor".

Jenkins in ICP 4 © Copyright IBM Corporation 2018

IBM Training **IBM**

Jenkins plugins

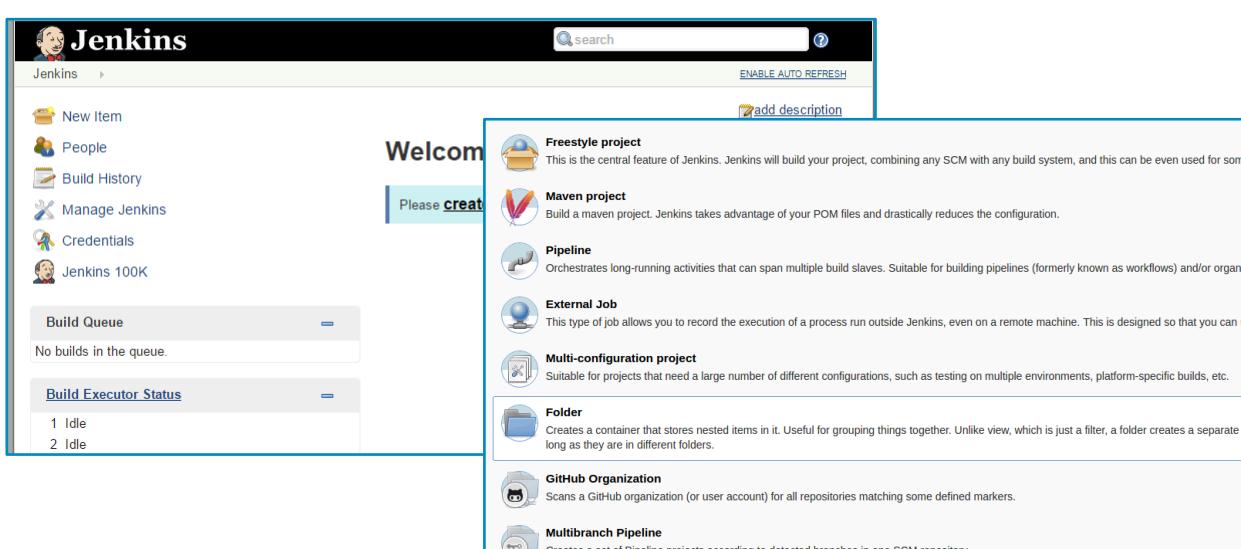


- SCM plugin**
Interact with SCM repository to check for updates and invokes pipelines
- Build action plugin**
Deployment tools
Build actions and notification
- Administration plugin**
Controlling agents, user security
CLI extensions, cluster management
- User interface plugin**
Modify the Jenkins Web UI widgets
Custom column displays
- Platform plugin**
Development target platform that provide custom build processing: iOS, Android, Ruby, .NET, etc

Jenkins in ICP 5 © Copyright IBM Corporation 2018

IBM Training **IBM**

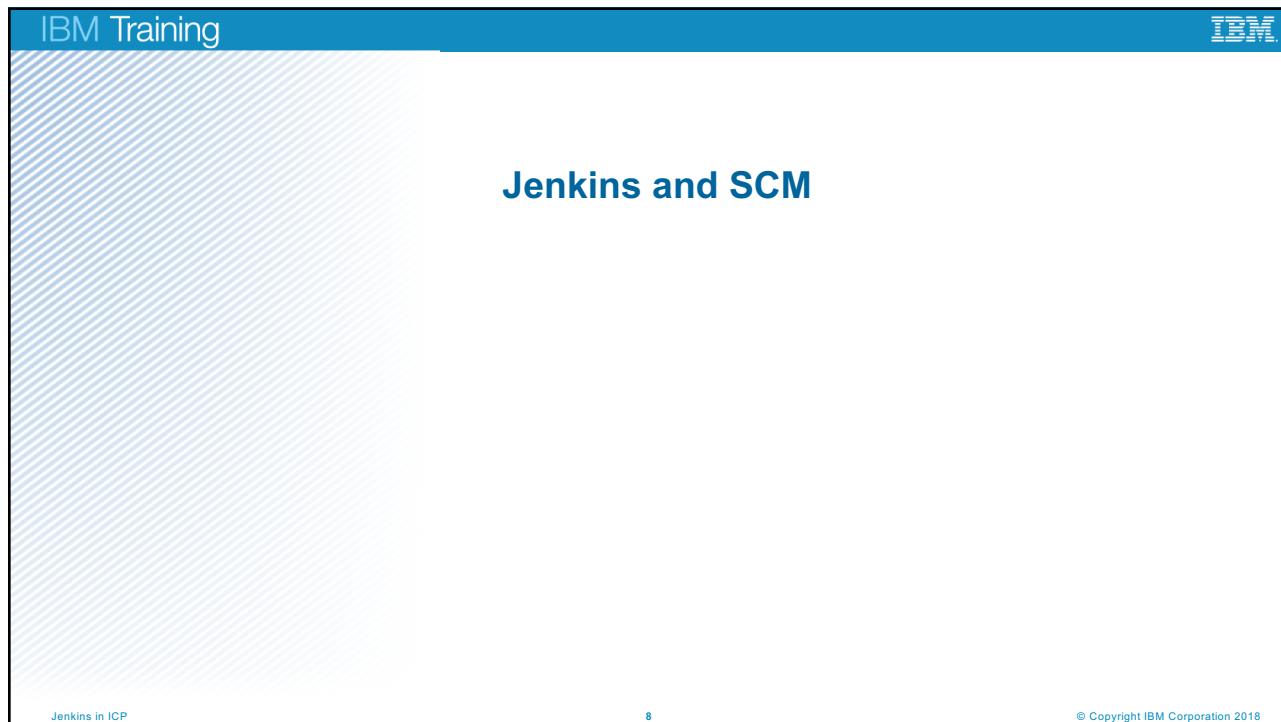
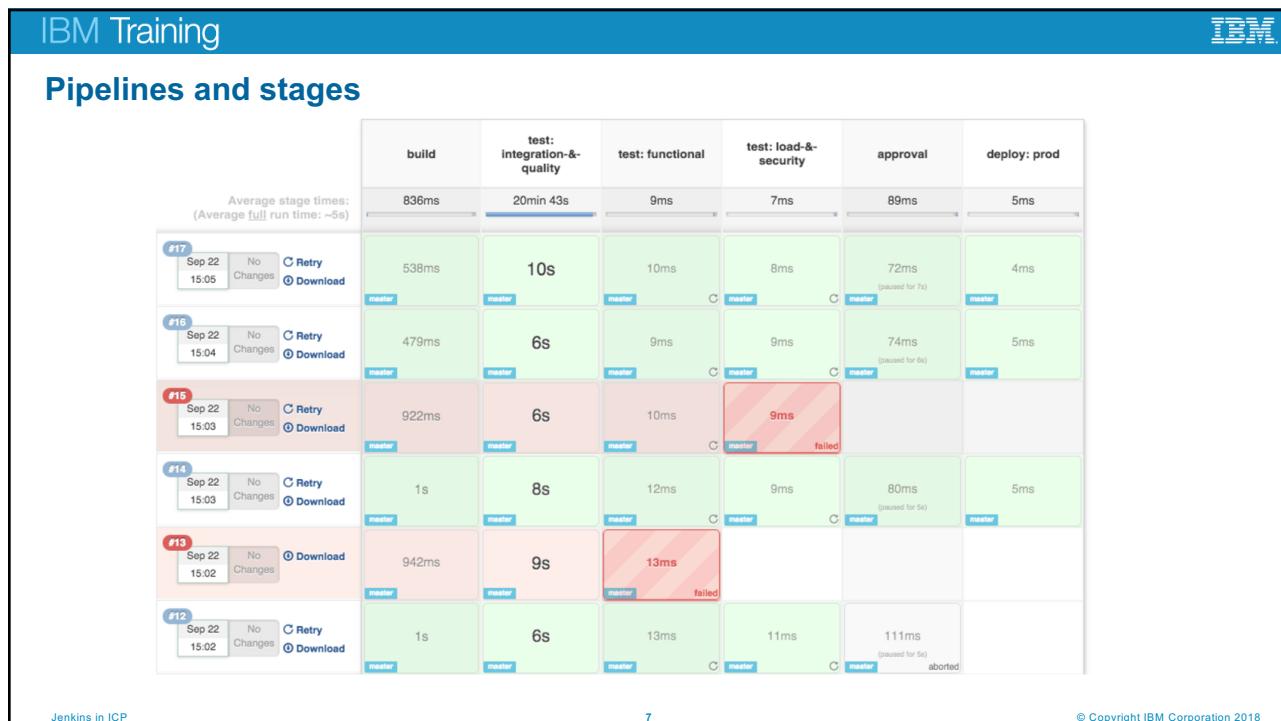
Jenkins Web UI



The screenshot shows the Jenkins Web UI homepage. The left sidebar includes links for New Item, People, Build History, Manage Jenkins, Credentials, and Jenkins 100K. Below these are sections for Build Queue (No builds in the queue) and Build Executor Status (1 Idle, 2 Idle). The main content area features a "Welcome" message with a "Please create" button. To the right, there is a list of project types with icons and descriptions:

- Freestyle project**
This is the central feature of Jenkins. Jenkins will build your project, combining any SCM with any build system, and this can be even used for som
- Maven project**
Build a maven project. Jenkins takes advantage of your POM files and drastically reduces the configuration.
- Pipeline**
Orchestrates long-running activities that can span multiple build slaves. Suitable for building pipelines (formerly known as workflows) and/or organi
- External Job**
This type of job allows you to record the execution of a process run outside Jenkins, even on a remote machine. This is designed so that you can u
- Multi-configuration project**
Suitable for projects that need a large number of different configurations, such as testing on multiple environments, platform-specific builds, etc.
- Folder**
Creates a container that stores nested items in it. Useful for grouping things together. Unlike view, which is just a filter, a folder creates a separate
- GitHub Organization**
Scans a GitHub organization (or user account) for all repositories matching some defined markers.
- Multibranch Pipeline**
Creates a set of Pipeline projects according to detected branches in one SCM repository.

Jenkins in ICP 6 © Copyright IBM Corporation 2018



IBM Training IBM

Jenkins SCM plugin

- Allows Jenkins to interact with SCM
 - Check for any changes in SCM
 - Pull SCM content
 - Get change log from SCM
- Available for GIT, github, BitBucket, SVN and others
- Provides a trigger Webhook for SVN to trigger checking

Jenkins in ICP 9 © Copyright IBM Corporation 2018

IBM Training IBM

Jenkins SCM interaction modes

Poll
Did changes occurs?
Invoke pipeline on changes

Post-commit hook
invoke Webhook

Jenkins check for changes
and invoke pipeline

Jenkins in ICP 10 © Copyright IBM Corporation 2018

IBM Training

Jenkins poll to SCM

Build Triggers

- Build after other projects are built
- Build periodically
- Poll SCM

Schedule

Cron-like time specification

No schedules so will only run due to SCM changes

Ignore post-commit hooks

Disable this project

Save Apply (e.g., from scripts)

Git Polling Log

```

Started on Apr 25, 2018 5:46:42 PM
Using strategy: Default
[poll] Last Built Revision: Revision 3fff56db222819402837ee30b8902a77cbfd08bf
(refs/remotes/origin/master)
> git --version # timeout=10
> git ls-remote -h https://github.com/ibm-cloud-academy/icp-jenkins-helm-bluecompute #
timeout=10
Found 1 remote heads on https://github.com/ibm-cloud-academy/icp-jenkins-helm-
bluecompute
[poll] Latest remote head revision on refs/heads/master is:
86ea5ea5bf21e2742284ea785dd833a999eaff91
Done. Took 1.4 sec
Changes found

```

Build History

#	Date
4	Apr 25, 2018 5:46 PM
3	Apr 25, 2018 4:03 PM

11 © Copyright IBM Corporation 2018

IBM Training

Jenkins and GIT

Build Triggers

- Build after other projects are built
- Build periodically
- Poll SCM

Schedule

No schedules

Ignore post-commit hooks

Disable this project

Save Apply (e.g., from scripts)

Git Polling Log

```

Started on Apr 25, 2018 5:46:42 PM
Using strategy: Default
[poll] Last Built Revision: Revision 3fff56db222819402837ee30b8902a77cbfd08bf
(refs/remotes/origin/master)
> git --version # timeout=10
> git ls-remote -h https://github.com/ibm-cloud-academy/icp-jenkins-helm-bluecompute #
timeout=10
Found 1 remote heads on https://github.com/ibm-cloud-academy/icp-jenkins-helm-
bluecompute
[poll] Latest remote head revision on refs/heads/master is:
86ea5ea5bf21e2742284ea785dd833a999eaff91
Done. Took 1.4 sec
Changes found

```

File Edit View Search Terminal Help

```

localuser@bmcloudacademy:~/icp-jenkins-helm-bluecompute/.git/hooks$ vi post-commit
localuser@bmcloudacademy:~/icp-jenkins-helm-bluecompute/.git/hooks$ cat post-commit
#!/bin/sh
curl http://proxy:30472/git/notifyCommit?url=https://github.com/lbm-cloud-academy/icp-jenkins-helm-bluecompute
localuser@bmcloudacademy:~/icp-jenkins-helm-bluecompute/.git/hooks$ 

```

12 © Copyright IBM Corporation 2018

1. The post-commit hook in .git/hooks is invoked when git push command is issued.
2. The post-commit triggers Jenkins webhook to check the SCM (web hook is created when Poll SCM box is checked)
3. As this is triggered by post-commit hook, changes that is found invokes the corresponding build pipeline

IBM Training IBM

Pipeline definition

```
{  
    node ('mypod') {  
        checkout scm  
        container('cli') {  
            stage('Build Gradle Project')  
            sh """  
                ....  
            stage('Build Docker Image')  
            sh """  
                ....  
        }  
    }  
}
```

Jenkins in ICP 13 © Copyright IBM Corporation 2018

IBM Training IBM

Kubernetes plugin and pipelines

IBM Training **IBM**

Jenkins Kubernetes plugin

Kubernetes plugin allows running of Jenkins pipeline in a Kubernetes cluster
The pipeline would run on a slave pod which is built based on a podTemplate
The slave pod is controlled by a Jenkins executor (which is a JNLP process)

Jenkins in ICP 15 © Copyright IBM Corporation 2018

IBM Training **IBM**

Pipeline podTemplate

```
podTemplate(
    label: 'mypod',
    volumes: [
        hostPathVolume(hostPath: '/var/run/docker.sock', mountPath: '/var/run/docker.sock'),
        secretVolume(secretName: 'icpadmin', mountPath: '/var/run/secrets/registry-account'),
        configMapVolume(configMapName: 'icpconfig', mountPath: '/var/run/configs/registry-config')
    ],
    containers: [
        containerTemplate(name: 'gradle', image: 'ibmcase/gradle:jdk8-alpine', ttyEnabled: true, command: 'cat'),
        containerTemplate(name: 'kubectl', image: 'ibmcloudacademy/k8s-icp:v1.0', ttyEnabled: true, command: 'cat'),
        containerTemplate(name: 'docker', image: 'docker:17.06.1-ce', ttyEnabled: true, command: 'cat')
    ]
)
```

Two snippets of YAML code are shown on the right, corresponding to the volumes defined in the podTemplate:

```
apiVersion: v1
kind: Secret
metadata:
  name: icpadmin
type: Opaque
data:
  username: <user-encoded>
  password: <password-encoded>
```

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: icpconfig
data:
  namespace: default
  registry: cloudcluster.icp:8500
```

Jenkins in ICP 16 © Copyright IBM Corporation 2018

IBM Training IBM

Jenkins pipeline in Kubernetes

```

{
  node ('mypod') {
    checkout scm
    container('cli') {
      stage('Build Gradle Project') {
        sh """
          #!/bin/sh
          gradle build -x test
        """
      }
    }
    container('docker') {
      stage ('Build Docker Image') {
        sh """
          #!/bin/bash
          NAMESPACE=`cat /var/run/configs/registry-account/namespace` 
          REGISTRY='cat /var/run/configs/registry-account/registry'
          cd docker
          docker build -t \${REGISTRY}/\${NAMESPACE}/bluecompute-customer:\${env.BUILD_NUMBER} .
          docker login -u=\${DOCKER_USER} -p=\${DOCKER_PASSWORD} \${REGISTRY}
          docker push \${REGISTRY}/\${NAMESPACE}/bluecompute-customer:\${env.BUILD_NUMBER}
        """
      }
    }
  }
}

```

Jenkins in ICP 17 © Copyright IBM Corporation 2018

IBM Training IBM

Jenkins in IBM Cloud Private

Jenkins in ICP 18 © Copyright IBM Corporation 2018

IBM Training

Jenkins in IBM Cloud Private

The screenshot shows the IBM Cloud Catalog interface. On the left, there's a sidebar titled "Repositories" with a search bar and dropdowns for "NAME" and "URL". Below it is a table with three rows: "ibm-charts", "local-charts", and "stable". The "stable" row has a URL link: "https://github.com/kubernetes/charts/tree/master/stable/jenkins". On the right, there's a large orange box labeled "Jenkins master" containing a green cylinder icon labeled "Persistence Disk". A red box highlights the "jenkins" entry in the catalog search results.

<https://github.com/kubernetes/charts/tree/master/stable/jenkins>

Jenkins in ICP 19 © Copyright IBM Corporation 2018

IBM Training

Jenkins pipeline access to IBM Cloud Private repository

A Jenkins pipeline script stage is shown:

```
stage ('Push Docker Image to Registry') {
    sh """
    #!/bin/bash
    NAMESPACE=`cat /var/run/configs/registry-config/namespace`
    REGISTRY=`cat /var/run/configs/registry-config/registry`
    DOCKER_USER=`cat /var/run/secrets/registry-account/username`
    DOCKER_PASSWORD=`cat /var/run/secrets/registry-account/password`
    docker login -u=\${DOCKER_USER} -p=\${DOCKER_PASSWORD} \${REGISTRY}
    docker push \${REGISTRY}/\${NAMESPACE}/bluecompute-customer:\${env.BUILD_NUMBER}
    """
}
```

A blue box labeled "Push image to ICP" is placed over the "push" command in the script.

A terminal window shows the creation of a secret named "tcpregistry" in a namespace "tcp" using the command:

```
kubectl create secret docker-registry tcpregistry --docker-server=tcp://cloudcluster.tcp:8500 --docker-username=admin --docker-password=password --namespace=tcp --from=service=tcpadmin@cloudcluster.tcp
```

A blue box labeled "Pull image from ICP" is placed over the "pullSecret" field in the "values.yaml" file.

A snippet of a "deployment.yaml" file is shown:

```
58 {{ if .Values.image.pullSecret }}
59   imagePullSecrets:
60     - name: {{ .Values.image.pullSecret }}
61 {{ end }}
```

A snippet of a "values.yaml" file is shown:

```
10   image:
11     repository: ibmcse/bluecompute-customer
12     tag: latest
13     pullPolicy: Always
14     pullSecret: tcpregistry
```

A red box highlights the "pullSecret" field in the "values.yaml" file.

Jenkins in ICP 20 © Copyright IBM Corporation 2018

Jenkins deploying to IBM Cloud Private

```
container('kubectl') {
    stage('Deploy new Docker Image') {
        sh """
        #!/bin/bash
        set +e
        NAMESPACE=`cat /var/run/configs/registry-config/namespace`
        REGISTRY=`cat /var/run/configs/registry-config/registry`
        DOCKER_USER=`cat /var/run/secrets/registry-account/username`
        DOCKER_PASSWORD=`cat /var/run/secrets/registry-account/password`
        wget --no-check-certificate https://master:8443/api/cli/icp-linux-amd64
        bx plugin install icp-linux-amd64
        bx pr login -a https://master:8443 --skip-ssl-validation \
            -u ${DOCKER_USER} -p ${DOCKER_PASSWORD} -c id-cloudcluster-account
        helm init --client-only
        bx pr cluster-config cloudcluster
        helm repo add bluecompute \
            https://raw.githubusercontent.com/ibm-cloud-academy/icp-jenkins-helm-bluecompute/master/charts
        helm install --tls -n bluecompute-customer \
            --set image.repository=${REGISTRY}/${NAMESPACE}/bluecompute-customer \
            --set image.tag=${env.BUILD_NUMBER} bluecompute/customer
        """
    }
}
```

IBM Training

IBM

Monitoring IBM Cloud Private

Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

IBM Training

IBM

Agenda

- Monitoring approaches in IBM Cloud Private
- What is available out of the box
- Cloud Service Management and Operations

Monitoring IBM Cloud Private

2

© Copyright IBM Corporation 2018

IBM Training **IBM**

Who monitors?

Primary Role



Todd
Operations / Admin
Responsible for infrastructure, security, and management of the environment

Todd has confidence in the private cloud platform. He knows corporate data is secure and in compliance with government regulations and/or industry standards. He has control and visibility of the running applications and can quickly update workloads and the underlying platform using continuous delivery techniques, without requiring downtime. Application data is automatically managed for backup, recovery, and disaster failover into alternate locations.

Todd receives incident alerts
He ensures monitoring is in place for the application components

Secondary Role



Jane
Enterprise Developer
Responsible for modernizing existing applications and creating new Cloud Native Workloads

Jane can immediately leverage CI/CD to rapidly develop, test and deliver her applications on the IBM cloud platform. She has a core set of services she needs, including data and application runtime services, with prescriptive guidance for moving heritage workloads to the cloud and/or creating new cloud native applications. She can also safely integrate off-premises cloud services into her application.

Jane leverages build to manage so that Todd can deploy "health" tests against the APIs. She is concerned with her applications performance and trends

Monitoring IBM Cloud Private

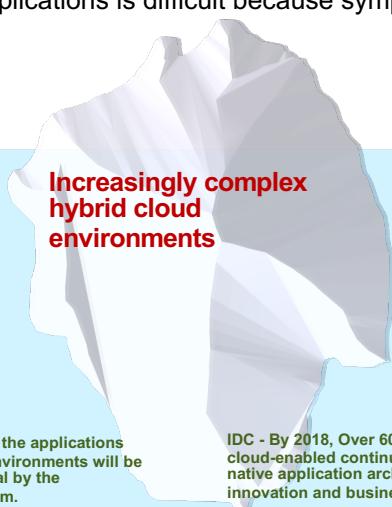
3

© Copyright IBM Corporation 2018

IBM Training **IBM**

Challenges with today's IT landscapes – yet another reality!

...Managing today's applications is difficult because symptoms may not be visible



Mean Time to Repair for application-related problems is often **3-6 hours**

User calls to applications support are the enterprise's "heads up" for an application problem **35% of the time**

Gartner – By 2018, 50% of the applications running in public cloud environments will be considered mission-critical by the organizations that use them.

Increasingly complex hybrid cloud environments

Issues that go beyond Level 1 support often involve at least **3-4 people** taking up an average of **5-7 total man hours**

58% of survey participants said their enterprise had between **6 and 40 different monitoring tools**

IDC - By 2018, Over 60% of new apps will use cloud-enabled continuous delivery and cloud-native application architectures to enable faster innovation and business agility.

IDC FutureScape: Worldwide Cloud 2016 Predictions -Mastering the Raw Material of Digital Transformation, IDC, Nov. 3, 2015
Monitoring IBM Cloud Private

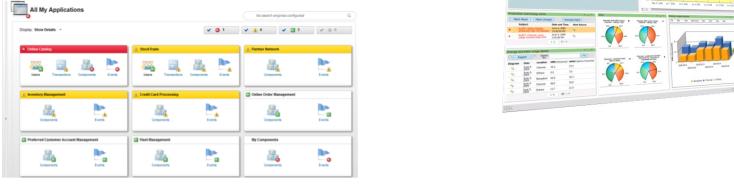
4

Gartner Predicts 2016: Cloud Computing to Drive Digital Business 12/8/15 G00292047
© Copyright IBM Corporation 2018

IBM Training 

Primary purpose for monitoring and logging

- Overall system health & availability
 - Understand, maintain, prevent, diagnose issues related to availability. What's broken and why?
- Business relevant data
 - Understand application & infrastructure usage patterns
 - Improve customer experience, test and measure hypothesis (A/B testing for example)
 - Optimize resources for greater economic efficiency



Monitoring IBM Cloud Private 5 © Copyright IBM Corporation 2018

IBM Training 

Traditional and Cloud-native

- Cloud-native has the same monitoring characteristics as traditional (non-Cloud-native)
- Infrastructure monitoring covering hosts and the network
- Application monitoring covering service level data like errors or response times
- Real-user monitoring covering end-user response times and errors as well behavioral data
- Log monitoring

Monitoring IBM Cloud Private 6 © Copyright IBM Corporation 2018

IBM Training



Monitoring and logging

- Monitoring helps identify and solve immediate problems such as:
 - Code problem
 - Deployment problem
 - Response time issue
 - Application availability
 - Platform issue
 - Typically numeric time series data used to record state that changes frequently
- Logging is used
 - Log files
 - Used to record important events, such as server start/stop, authentication results, a new version of the service is deployed, unexpected errors, application start/stop, configuration changes, etc.
 - Critical for effective debugging
 - To identify longer term trends and problems
 - Resource issues over time
 - Does an application consume more memory at certain times of day or season?
 - Is more/less hardware needed based on resource consumption usage?
 - Are workloads meeting Service Level Agreements?

IBM Training



Operational Visibility

- Alerts
 - Used to automatically detect changes in monitored data relative to a threshold or other condition
 - CPU >90% for 3 minutes, disk volume > 70% full, response latency > 100ms for 50% of requests, etc.
 - Not all alerts may trigger a notification
- Notifications
 - Communicate an important alert has been triggered and needs attention
 - Could be text, Slack post, PagerDuty, sirens, flashing-lights, etc.
- Dashboards
 - Typically a web app used to visualize a summary of the most important data relative to its users
 - Usually have ability to filter and/or ability to do ad-hoc queries
 - Very common to have multiple dashboards for different audiences (executives, offering managers, SRE/ops, developers, etc.)
 - Very unlikely to have a once-size-fits all dashboard

IBM Training

Differences between traditional and Cloud-Native

Traditional	Cloud-Native
Centralized operations staff with monitoring and system health responsibility	Development teams have end-to-end responsibility (or share with operations)
Fewer and more controlled releases	Many releases, sometimes daily/hourly
Few language / technology choices	Polyglot language / technology choices (for example, databases)
Anomalies detected from small number of monitored systems	High number of inter-related components making anomaly detection difficult
Monitor and manage at infrastructure level	Monitor and manage at application (workload) level
Pro-active long term capacity planning	Real-time capacity planning with auto-scaling
Fix failing components	Replace failing components

Monitoring IBM Cloud Private

9

© Copyright IBM Corporation 2018

IBM Training

IBM Cloud Private Monitoring and Logging

1. The ICP dashboard shows the status, metrics and log entries for all the workloads & Infrastructure.

```

graph TD
    Host[Hosting Infrastructure (Virtual servers)] --> CF[Cloud Foundry]
    Host --> KH[Kubernetes Heapster]
    CF --> DD[Default Dashboards]
    KH --> DD
    subgraph ICP [IBM Cloud Private]
        DD
    end
  
```

Monitoring IBM Cloud Private

10

© Copyright IBM Corporation 2018

IBM Training

IBM Cloud Private Monitoring and Logging

1. The ICP dashboard shows the status, metrics and log entries for all the workloads & Infrastructure.
2. An Elasticsearch stack collects logs from all the workloads and Infrastructure

Monitoring IBM Cloud Private

11

© Copyright IBM Corporation 2018

IBM Training

IBM Cloud Private Monitoring and Logging

1. The ICP dashboard shows the status, metrics and log entries for all the workloads & Infrastructure.
2. An Elasticsearch stack collects logs from all the workloads and Infrastructure
3. A Prometheus stack collects logs from the Kubernetes workload

Monitoring IBM Cloud Private

12

© Copyright IBM Corporation 2018

IBM Training

IBM Cloud Private Monitoring and Logging

1. The ICP dashboard shows the status, metrics and log entries for all the workloads & Infrastructure.
2. An Elasticsearch stack collects logs from all the workloads and Infrastructure
3. A Prometheus stack collects logs from the Kubernetes workload
4. Customizable Grafana and Kibana dashboards are available for use

```

graph TD
    HI[Hosting Infrastructure (Virtual servers)] --> CF[Cloud Foundry]
    CF --> KH[Kubernetes Heapster]
    KH --> ELK[ELK Stack]
    ELK --> P[Prometheus]
    P --> DD[Default Dashboards]
    P --> ED[Extendable Dashboards & Alerts]
    
```

ICP

Monitoring IBM Cloud Private

© Copyright IBM Corporation 2018

IBM Training

IBM Cloud Private Monitoring and Logging

1. The ICP dashboard shows the status, metrics and log entries for all the workloads & Infrastructure
2. An Elasticsearch stack collects logs from all the workloads and Infrastructure
3. A Prometheus stack collects logs from the Kubernetes workload
4. Customizable Grafana and Kibana dashboards are available for use
5. Prometheus can send alerts to external system (email, webhooks)
 - Dashboards are view only
 - APIs available to manage

```

graph TD
    HI[Hosting Infrastructure (Virtual servers)] --> CF[Cloud Foundry]
    CF --> KH[Kubernetes Heapster]
    KH --> ELK[ELK Stack]
    ELK --> P[Prometheus]
    P --> DD[Default Dashboards]
    P --> ED[Extendable Dashboards & Alerts]
    P --> A[Alerts]
    
```

ICP

Monitoring IBM Cloud Private

© Copyright IBM Corporation 2018

IBM Training

Monitoring Components

- Prometheus - <https://prometheus.io/>
 - Systems and service monitoring system
 - Cloud Native Computing Foundation (CNCF) project
 - Collects metrics from configured targets at given intervals, evaluates rule expressions, displays the results, and can trigger alerts if some condition is observed to be true
- Grafana - <https://grafana.com/grafana>
 - Open source metric analytics and visualization suite
 - Rich display features (compared to Kibana)
 - Commonly used for visualizing time series data for infrastructure and application analytics
 - No data searching and exploring (unlike Kibana)
- Grafana dashboards are used
 - Prometheus has limited dashboarding capabilities of its own
- Prometheus and the Grafana Dashboards are installed ICP using Helm charts

IBM



Prometheus



Grafana

Monitoring IBM Cloud Private

15

© Copyright IBM Corporation 2018

IBM Training

Logging Components

- Elasticsearch stack * - <https://www.elastic.co/elk-stack>
- Log Management system which is composed of:
 - Elasticsearch
 - Distributed, JSON-based search and analytics engine designed for horizontal scalability, maximum reliability, and easy management
 - Logstash
 - Dynamic data collection pipeline with an extensible plugin ecosystem and strong Elasticsearch synergy
 - Kibana
 - Extensible user interface for configuring and managing all aspects of the Elastic Stack
 - Comprehensive log analytics dashboard to visualize and build dashboards for log data
 - Beats
 - Platform for lightweight shippers that send data from edge machines to Logstash and Elasticsearch
- Packaged as a series of Helm charts
 - Deployed automatically when ICP is installed
 - Can be customized and/or opted-out of

* Previously known as ELK Stack

Monitoring IBM Cloud Private

16

© Copyright IBM Corporation 2018

IBM Training

Monitoring an IBM Cloud Private Deployment

- ICP Prometheus display
 - Data is there but hard to absorb

Monitoring IBM Cloud Private

17

© Copyright IBM Corporation 2018

IBM Training

Monitoring an IBM Cloud Private Deployment

- ICP dashboard shows a high level overview of resource usage and status

Monitoring IBM Cloud Private

18

© Copyright IBM Corporation 2018

IBM Training

Monitoring an IBM Cloud Private Deployment

- ICP-provided Grafana dashboard shows more detail of resource usage and status

The dashboard displays various metrics for the IBM Cloud Private cluster. Key sections include:

- Cluster Summary:** Shows Youngest Node Update (6.32 hour), Total memory (15.51 GiB), Available Memory (5.16 GiB), Memory Free (33%), ICP Total CPU 15 Minutes Average (16%), Avg. Machine Count (2), and Active Pods (70).
- Memory by node:** A heatmap showing memory usage across nodes.
- Top 5 Containers by CPU:** A list of containers with their CPU usage.
- Container CPU Utilization:** A line chart showing CPU usage over time for various containers.
- Container Memory Usage:** A line chart showing memory usage over time for various containers.
- Pods CPU Utilization:** A line chart showing CPU usage for individual pods.
- Pods Memory Usage:** A line chart showing memory usage for individual pods.

Monitoring IBM Cloud Private 19 © Copyright IBM Corporation 2018

IBM Training

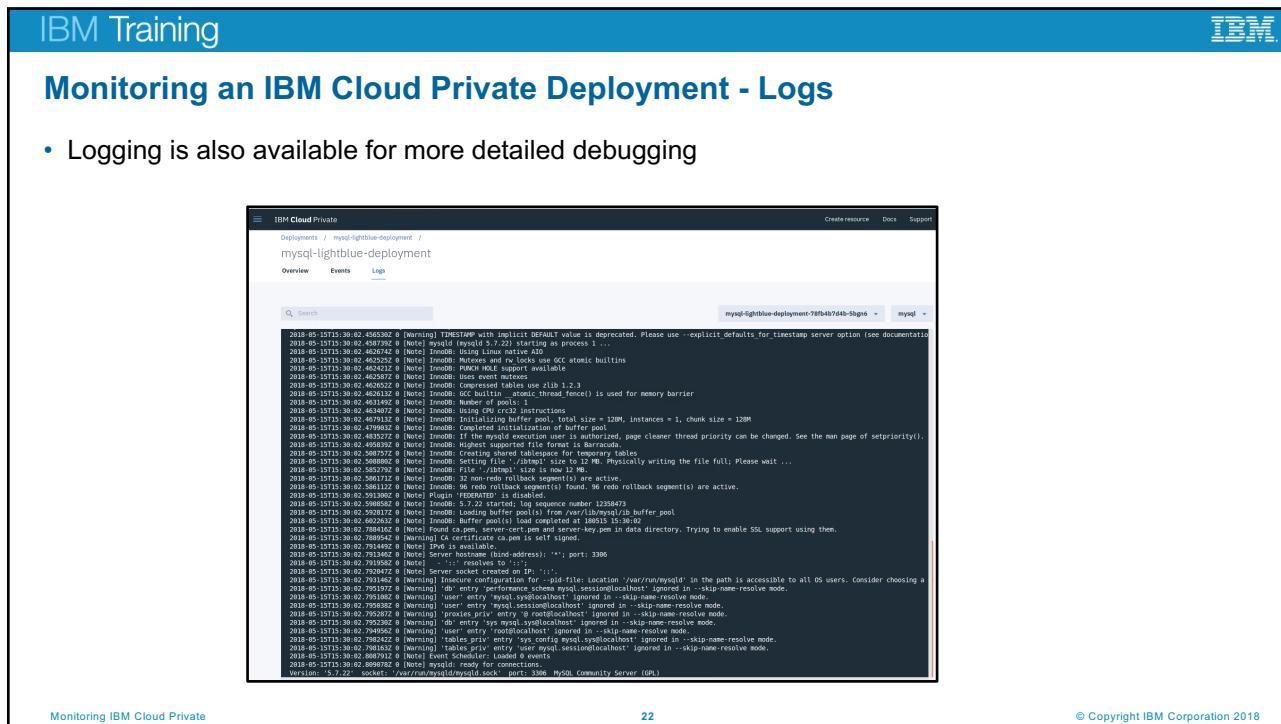
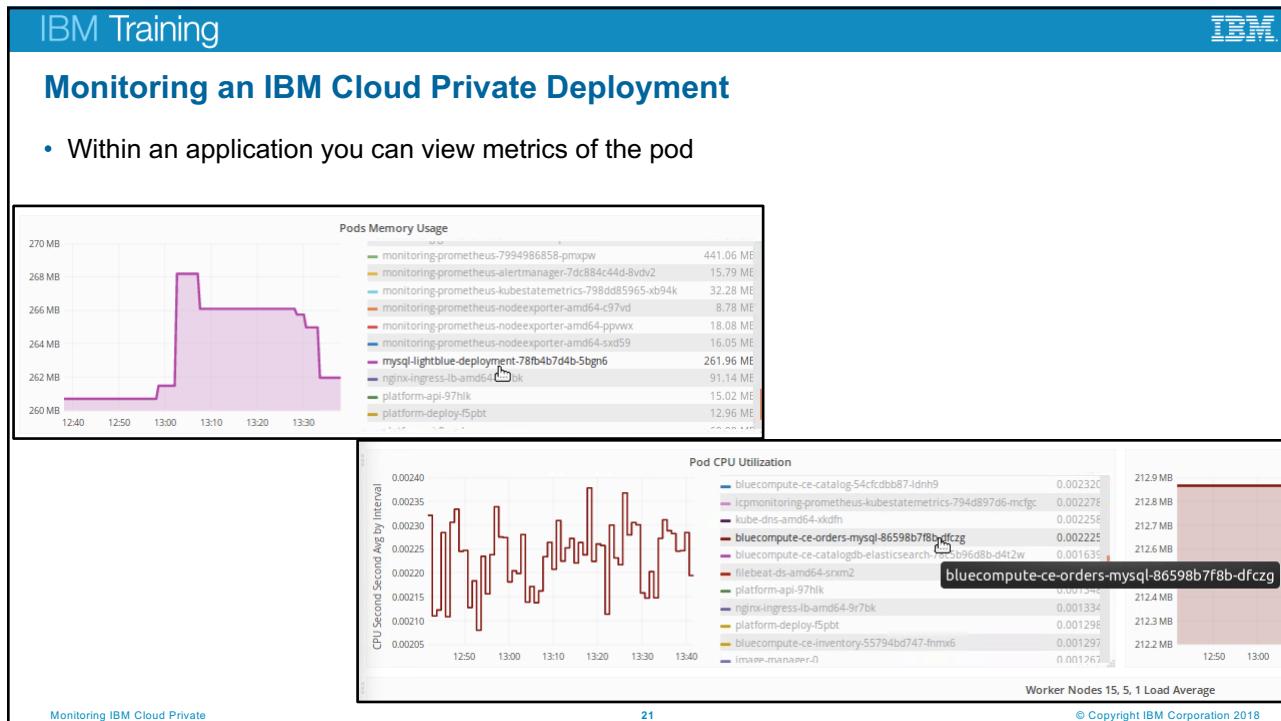
Monitoring an IBM Cloud Private Deployment

- Kubernetes Dashboard shows a variety of cluster metrics

The Kubernetes Dashboard provides a comprehensive view of the cluster's health and resource consumption. Key features include:

- Network I/O pressure:** Three donut charts showing Cluster memory usage (70%, 10.84 GiB / 15.51 GiB), Cluster CPU usage (15.36%, 0.92 cores / 6.00 cores), and Cluster filesystem usage (12.79%, 43.76 GiB / 342.16 GiB).
- Pods CPU usage (5m avg):** A line chart showing average CPU usage for individual pods over a 5-minute interval.
- Containers CPU usage (5m avg):** A line chart showing average CPU usage for individual containers over a 5-minute interval.
- System services CPU usage:** A line chart showing average CPU usage for system services over a 5-minute interval.

Monitoring IBM Cloud Private 20 © Copyright IBM Corporation 2018



IBM Training

Monitoring an IBM Cloud Private Deployment - Alerts

- Organizations want to monitor and alert on Out of Bounds conditions for the applications and infrastructure

The screenshot shows the Alertmanager interface with two main sections of alerts:

- HighCPUUsage:** Three alerts from 201143, 2018-05-15, triggered by 'service='backend'' and 'job='kubeentes-service-endpoints'', with instances '10.10.1.30:9100', '10.10.1.10:9100', and '10.10.1.20:9100'.
- NodeMemoryUsage:** Two alerts from 201143, 2018-05-15, triggered by 'severity='page'' and 'job='kubeentes-service-endpoints'', with instances '10.10.1.10:9100' and '10.10.1.30:9100'. These alerts involve 'kubernetes_name="monitoring-prometheus-nodeexporter"' and 'app="monitoring-prometheus"'.

At the bottom left: Monitoring IBM Cloud Private

At the bottom right: © Copyright IBM Corporation 2018

IBM Training

Managing an IBM Cloud Private Deployment – Enterprise View

- Enterprises want to consolidate operational and monitoring resources
 - Prometheus can send alerts to external enterprise management system such as IBM Application Performance Management (APM)
- Monitoring and logging part of some enterprises larger operational objective

...

The diagram illustrates the monitoring architecture within the ICP (IBM Cloud Private) environment:

- Default Dashboards** (blue box) are connected to **Kubernetes Heapster** and **Cloud Foundry**.
- ELK Stack** (red box) provides data to both **Kubernetes Heapster** and **Cloud Foundry**.
- Prometheus** (red box) also provides data to both **Kubernetes Heapster** and **Cloud Foundry**.
- Extendable Dashboards & Alerts (Grafana/Kibana)** (orange box) receives data from **Kubernetes Heapster**, **Cloud Foundry**, and **Alerts** (red box).
- Alerts** (red box) receive input from both **Prometheus** and the **ELK Stack**.

At the bottom: Hosting Infrastructure (Virtual servers)

At the bottom left: Monitoring IBM Cloud Private

At the bottom right: © Copyright IBM Corporation 2018

IBM Training **IBM**

Cloud Service Management and Operations

- What is it?
 - All the activities an organization does to plan, design, deliver, operate, & control the IT & cloud services it offers
 - Includes the operational aspects of applications and services
 - Once pushed to production, an application must be managed and monitored to ensure availability and performance according to service level agreements (SLAs) or service level objectives (SLOs)
- Why is it important?
 - Poor application performance will drive customers away
 - As organizations move workloads to cloud faster and faster it's important to know how the workload is performing against customer demand
 - The ability to rapidly respond and correct anomalies is impossible without:
 - Proper monitoring of workload and associated infrastructure
 - Incident notification and management
 - Performance Management data provides the ability to trend the application performance and understand its growth and seasonal patterns

Monitoring IBM Cloud Private 25 © Copyright IBM Corporation 2018

IBM Training **IBM**

CSMO Incident Management Tool Chain

- Incident management
 - The practice of restoring a damaged service to health as quickly as possible
- Incident management tool chain
 - Set of tools used to aid incident management
- Different levels based on enterprise need
- First level is monitoring and logging

The diagram illustrates the CSMO Incident Management Tool Chain. It starts with a central box labeled "CSMO Incident Management Tool Chain" containing several components: Monitoring, Event Management, Analyze, Plan, and Execute. To the left, there is a box labeled "Our focus" with an orange arrow pointing to the Monitoring and Logging components. Monitoring and Logging feed into Event Management. Event Management feeds into Analyze, which then feeds into Plan and Execute. Plan and Execute both feed into Ticketing & Trending. Ticketing & Trending feeds into Collaboration. Collaboration feeds into Runbooks. Runbooks feeds into Dashboards & Reporting. There are also direct connections from Monitoring to Dashboards & Reporting, and from Event Management to Dashboards & Reporting.

CSMO Incident Management Tool Chain

Monitoring

Event Management

Analyze

Plan

Execute

Ticketing & Trending

Collaboration

Runbooks

Dashboards & Reporting

Monitoring

Event Management

Analyze

Plan

Execute

Ticketing & Trending

Collaboration

Runbooks

Dashboards & Reporting

Our focus

Monitoring IBM Cloud Private 26 © Copyright IBM Corporation 2018

IBM Training

IBM

Unit Summary

- Monitoring approaches in IBM Cloud Private
- What is available out of the box
- Cloud Service Management and Operations

Monitoring IBM Cloud Private

27

© Copyright IBM Corporation 2018

IBM Training

IBM

Microservices Application Architecture

© Copyright IBM Corporation 2016-18
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

IBM Training

IBM

Prerequisites

- Developing cloud applications
 - Cloud-native application development
 - Twelve-Factor Apps
- Deploying cloud applications
 - Docker and Kubernetes in IBM Cloud Container Service
 - Cloud Foundry in IBM Cloud
 - Swarm in Docker Enterprise Edition for IBM Cloud
 - OpenWhisk in IBM Cloud Functions
- IBM Cloud capabilities
 - IBM Cloud architecture
 - WebSphere Liberty for microservices using Java, Java EE, and MicroProfile

Microservices Application Architecture

2

© Copyright IBM Corporation 2016-18

Unit objectives

- List the benefits of cloud native applications and the microservices application architecture
- Describe how the microservices application architecture compares with existing application architectures
- Show how an application can be designed as a set of microservices

Lesson 1 Introduction to microservices

The curse of the unmaintainable application

Back in 2005, the applications for each of Expedia's travel sites were too difficult to modify

"Each site was an interconnected tangle of code. If programmers wanted to add the ability to handle credit card billing in Germany, for example, that might end up causing errors in the flight-search feature. As a result, changes were rare: A site would be updated only a couple times a year."

Don't let this happen to you (any more)!

"Getting to Know You: Expedia has bet everything on understanding the psyche of the modern traveler" in *Bloomberg Business Week*, February 25, 2016
<http://www.bloomberg.com/news/features/2016-02-25/expedia-thinks-it-can-help-you-find-the-dream-vacation-you-didn-t-know-you-wanted>

Cloud native

An **application architecture** designed to use the *strengths* and accommodate the *challenges* of a standardized cloud environment, including:

- Elastic scaling
- Immutable deployment
- Disposable instances
- Less predictable infrastructure



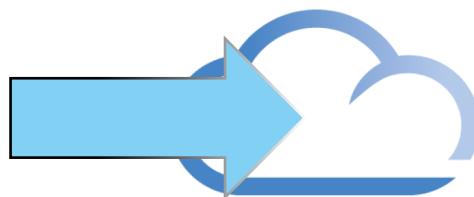
What it means to be cloud native

- Clean contract with underlying OS to ensure maximum portability
- Scale elastically without significant changes to tooling, architecture, or development practices
- Resilient to inevitable failures in the infrastructure and application
- Instrumented to provide both technical and business insight
- Use cloud services such as storage, queuing, and caching
- Rapid and repeatable deployments to maximize agility
- Automated setup to minimize time and cost for new developers

Microservices

An **engineering approach** focused on **decomposing** an application into **single-function** modules with **well defined interfaces** which are **independently deployed** and operated by **small teams** who own the **entire lifecycle** of the service

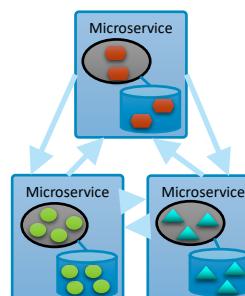
Microservices accelerate delivery by **minimizing communication** and coordination between people while **reducing the scope** and risk of change



IBM Training **IBM**

Microservices: Using technology more efficiently

- Microservices is an application architectural style
 - The application is composed of microservice components
- Microservice, a component in this architecture
 - Each is a miniature application
 - Each is focused on one task, a business capability
 - The Single Responsibility principle
 - Each can be deployed and updated independently
 - They are loosely coupled
 - Each has a well-defined interface
 - REST APIs

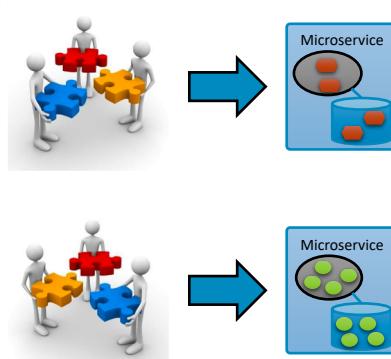


Microservices Application Architecture 9 © Copyright IBM Corporation 2016-18

IBM Training **IBM**

Microservices: Making developers more efficient

- Each microservice is developed and deployed by a small team
 - The team owns the entire lifecycle of the service
- Microservices accelerate delivery
 - Minimize communication and coordination between people
 - “Less time in meetings, more time coding”
 - Reducing the scope and risk of change
- Microservices facilitate agile development
 - They make their development teams loosely coupled



Microservices Application Architecture 10 © Copyright IBM Corporation 2016-18

IBM Training IBM

Microservices are an evolution

Evolution of architectural styles

- **Monolithic**
One large application that does everything
- **Microservices**
Several smaller applications that each does part of the whole

Evolution of service orientation

- **SOA**
Focused on reuse, technical integration issues, technical APIs
- **Microservices**
Focused on functional decomposition, business capabilities, business APIs

Microservices Application Architecture 11 © Copyright IBM Corporation 2016-18

IBM Training IBM

Microservices architecture

- Microservices architecture is about breaking down large silo applications into more manageable, fully decoupled pieces
- A *microservice* is a granular, decoupled component within a broader application

The diagram illustrates the shift from a monolithic application to a microservices architecture. On the left, a dashed box labeled 'Monolithic application' contains a single large gray rectangle labeled 'Silo'. On the right, a larger dashed box labeled 'Microservices application' contains three components: a top section with bulleted text ('Agility', 'Scalability', 'Resilience') and two bottom sections, each labeled 'Microservice (component)'.

Microservices Application Architecture 12 © Copyright IBM Corporation 2016-18

IBM Training **IBM**

Sample application that uses microservices

- Airline reservation application
 - Seven services (in this example)
- Each service includes
 - Logging
 - Metrics
 - Health check
 - Service endpoint
 - Service registry
 - Service management
- What do the different colors of the tiles mean (red, blue, and green)?
- The tiles are in stacks, some higher than others. What does the height mean?

Microservices Application Architecture 13 © Copyright IBM Corporation 2016-18

IBM Training **IBM**

Key tenets of a microservices architecture

1. Large monoliths are broken down into many small services
 - Each service runs in its own process
 - One service per container
2. Services are optimized for a single function
 - There is only one business function per service
 - The Single Responsibility Principle
 - A microservice should have one, and only one, reason to change
3. Communication via REST API and message brokers
 - Avoid tight coupling introduced by communication through a database
4. Per-service continuous integration and continuous deployment (CI/CD)
 - Services evolve at different rates
 - You let the system evolve but set architectural principles to guide that evolution
5. Per-service high availability (HA) and clustering decisions
 - One size or scaling policy is not appropriate for all
 - Not all services need to scale; others require autoscaling up to large numbers

Microservices Application Architecture 14 © Copyright IBM Corporation 2016-18

IBM Training 

Advantages of microservices

- Developed independently
 - Limited, explicit dependencies on other services
- Developed by a single team
 - The team is small
 - All team members can understand the entire code base
- Developed on its own timetable
 - New versions delivered independently of other services
- Polyglot: Each can be developed in a different language
 - Select the best language
- Manages its own data
 - Select the best technology and schema
- Scales and fails independently
 - Isolates problems

Microservices Application Architecture

15

© Copyright IBM Corporation 2016-18

IBM Training 

Comparing monolithic and microservices architectures

Category	Monolithic architecture	Microservices architecture
Architecture	Built as a single logical executable	Built as a suite of small services
Modularity	Based on language features	Based on business capabilities
Agility	Changes require building and deploying a new version of the entire application	Changes can be applied to each service independently
Scaling	Entire application scaled when only one part is the bottleneck	Each service scaled independently when needed
Implementation	Typically entirely developed in one programming language	Each service can be developed in a different programming language
Maintainability	Large code base is intimidating to new developers	Smaller code bases easier to manage
Deployment	Complex deployments with maintenance windows and scheduled downtimes	Simple deployment as each service can be deployed individually, with minimal downtime

Microservices Application Architecture

16

© Copyright IBM Corporation 2016-18

Technology advances have made microservices possible

- Ease and feasibility of distributing components
 - Internet, intranet, or network maturity
 - RESTful API conventions or *perceived* simplicity, and lightweight messaging
- Ease and simplicity of hosting
 - Lightweight runtimes
 - Examples: Node.js and WebSphere Liberty
 - Simplified infrastructure
 - OS virtualization (hypervisors), containerization (Docker), infrastructure as a service (cloud infrastructure)
 - Workload virtualization (examples: Cloud Foundry, Kubernetes, OpenWhisk, Swarm)
 - Platform as a service
 - Autoscaling, SLA management, messaging, caching, build management
- Agile development methods
 - Examples: IBM Cloud Garage Method, XP, TDD, Scrum, Continuous Delivery
 - Standardized code management, such as GitHub

Microservice challenges

- Greater operational complexity because there are more moving parts to monitor and manage
- Developers must have significant operational skills (DevOps)
- Service interfaces and versions
- Duplication of effort across service implementations
- Extra complexity of creating a distributed system
 - Network latency
 - Fault tolerance
 - Serialization
- Designing decoupled, non-transactional systems is difficult
- Locating service instances
- Maintaining availability and consistency with partitioned data
- End-to-end testing

IBM Training IBM

Lesson 2

Application architecture evolution

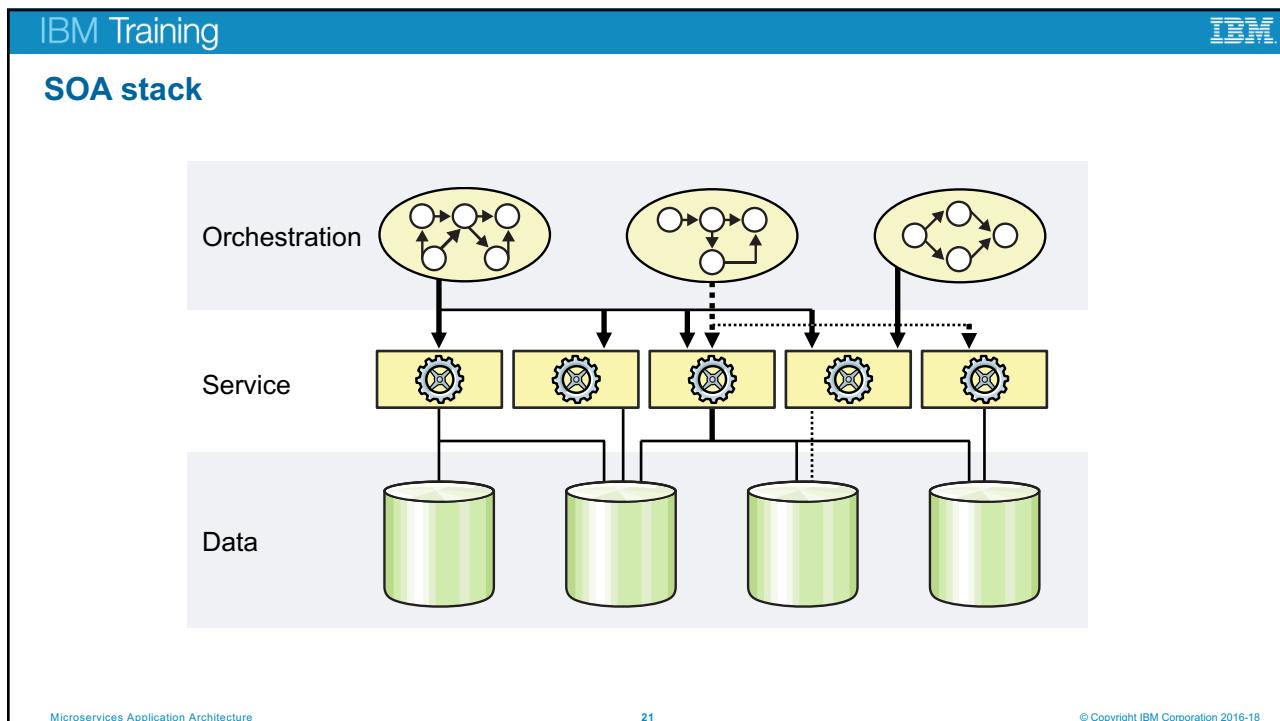
Microservices Application Architecture 19 © Copyright IBM Corporation 2016-18

IBM Training IBM

Layered application architecture

The diagram illustrates a layered application architecture. At the top, a vertical stack of four layers is labeled "Monolithic application": View (white), Application model (white), Domain model (yellow), and Integration (light green). A horizontal line connects the bottom of this stack to a horizontal cylinder below. The cylinder represents the "Database" layer, which is divided into several segments of varying widths, all colored light green.

Microservices Application Architecture 20 © Copyright IBM Corporation 2016-18



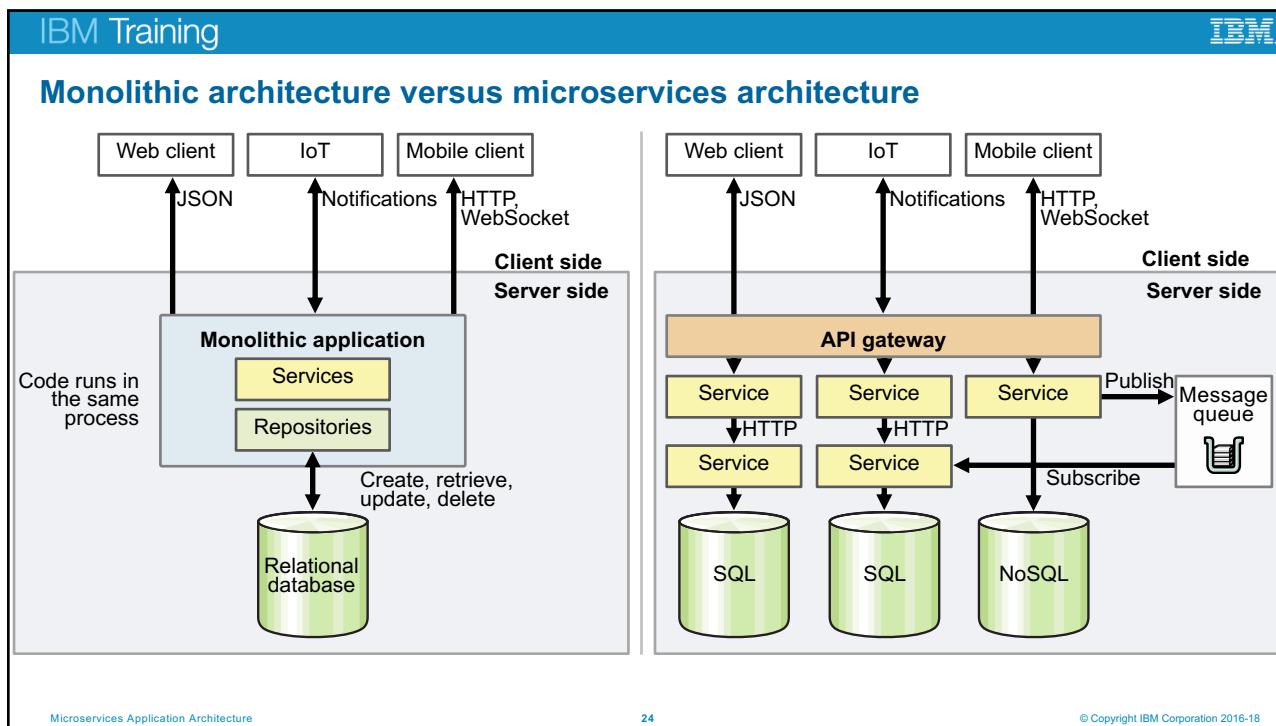
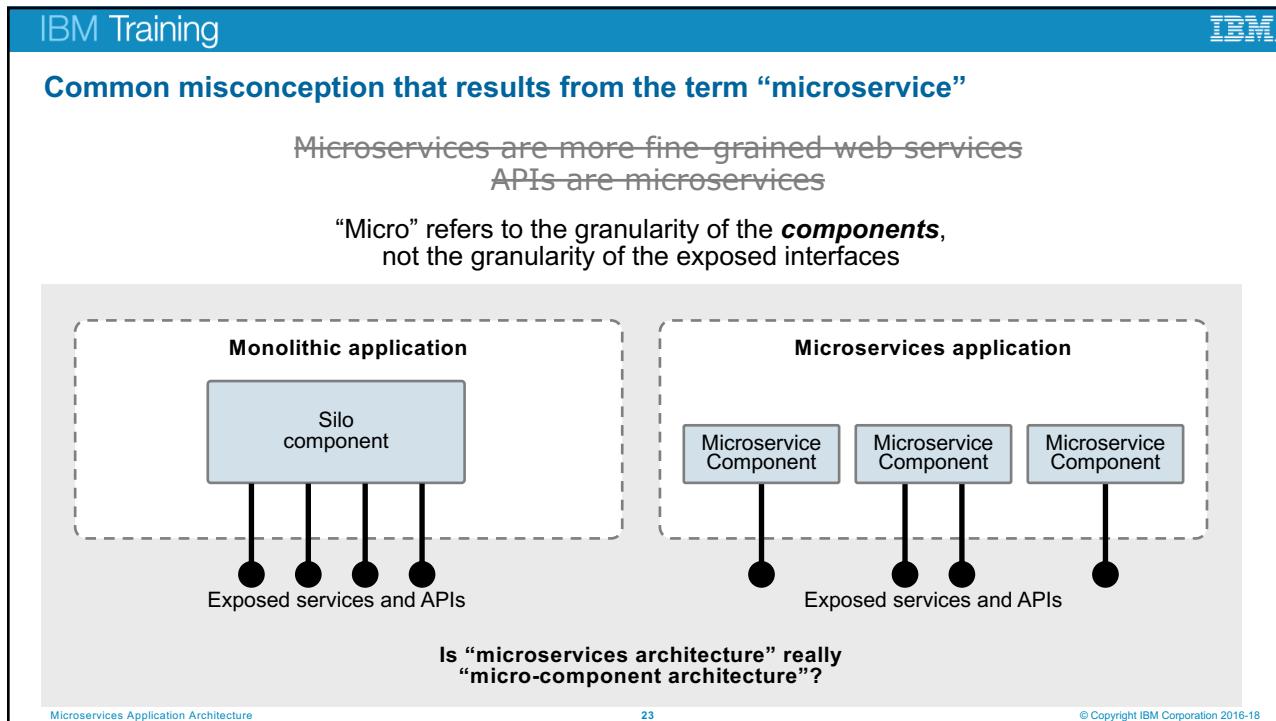
IBM Training IBM

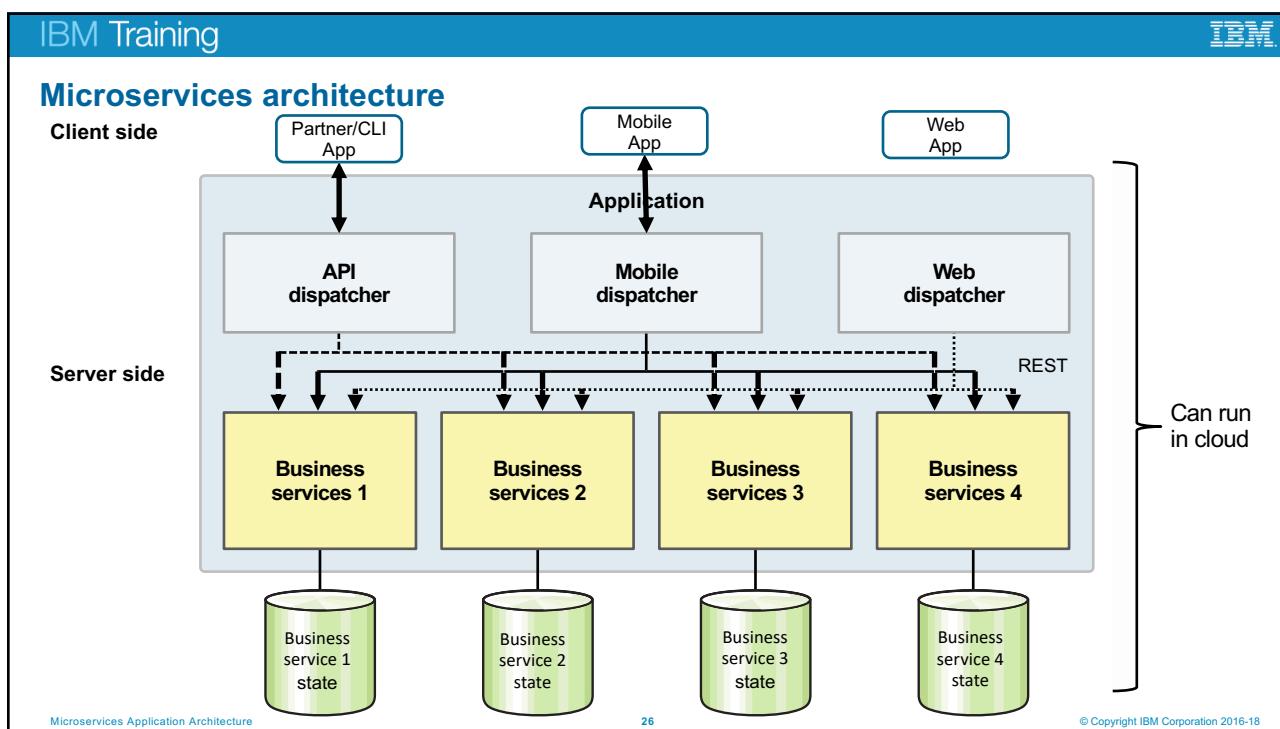
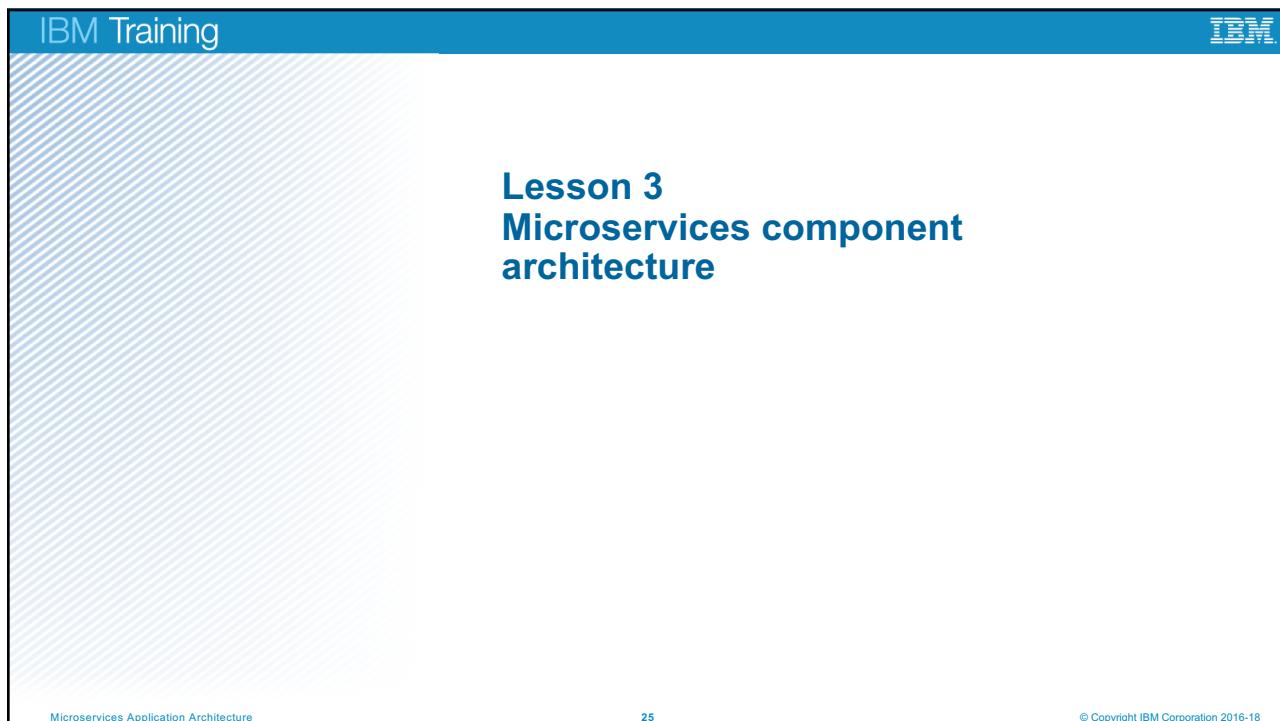
Microservices and SOA

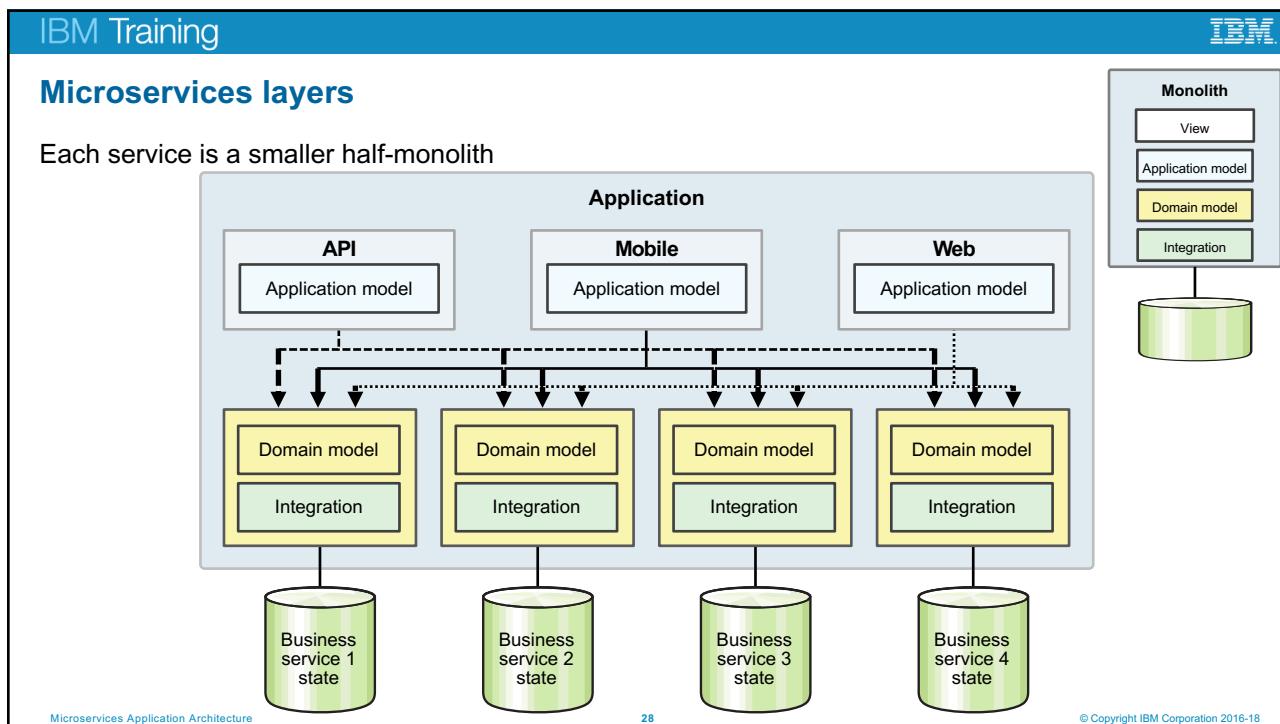
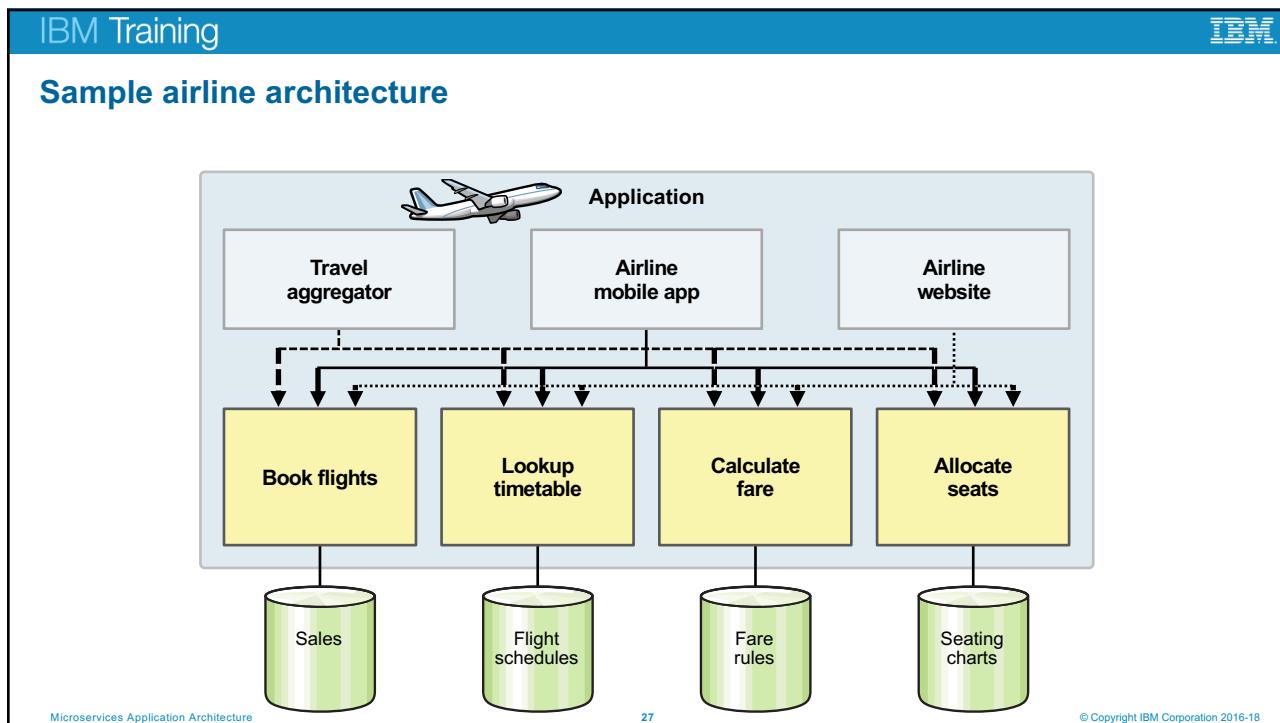
Both SOA and microservices deal with a system of services that communicate over a network
But there are differences

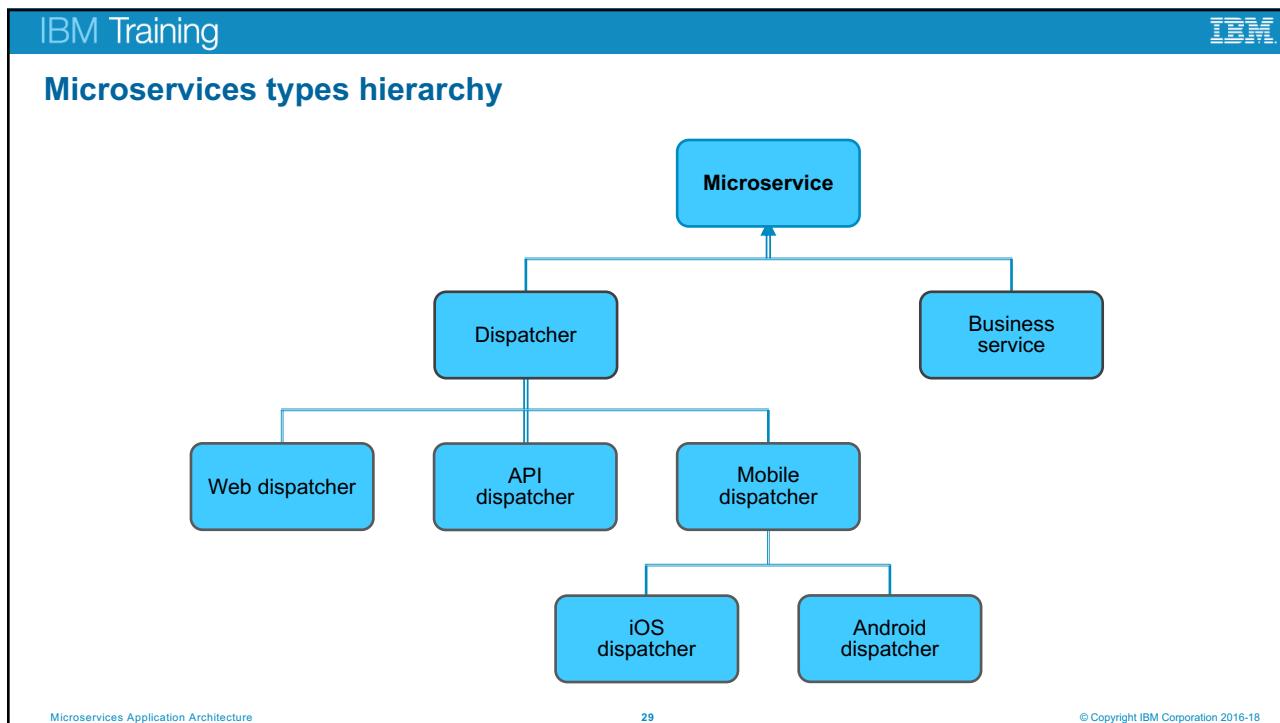
- The focus of SOA is on reuse
 - This tends to align with a centrally funded model
 - SOA services tend to be “servants of many masters”
 - This means that a change to a SOA service might impact multiple consumers
- The focus of microservices is on breaking down a potentially monolithic application into smaller, more manageable components
 - With the objective of more flexible, decoupled, faster development
 - Challenges here relate to needing good DevOps, management views, and controls

Microservices Application Architecture 22 © Copyright IBM Corporation 2016-18









IBM Training IBM

Language decisions

- Dispatchers are most often written in Node.js
 - Better fidelity with the clients
 - Especially clients that run JavaScript
 - iOS teams might want to use the Swift server-side runtime
 - I/O intensive, supports numerous concurrent clients
- Business services are most often written in Java
 - Better for CPU-intensive tasks
 - Better connectivity to external systems
- The team selects the language
 - Choose the language that best fits the job

Microservices Application Architecture 30 © Copyright IBM Corporation 2016-18

IBM Training **IBM**

Back ends for front ends

- Each dispatcher is a back end for an external front end, typically a UI
- The same team develops each back end and front end pair
 - Use the same or compatible languages in the pair

```

graph LR
    subgraph API_team [API team]
        direction TB
        API_dispatcher[API dispatcher] --- REST --- CLI_app[CLI app]
        CLI_app --- REST --- API_dispatcher
    end
    subgraph Mobile_team [Mobile team]
        direction TB
        Mobile_dispatcher[Mobile dispatcher] --- REST --- Mobile_app[Mobile app]
        Mobile_app --- REST --- Mobile_dispatcher
    end
    subgraph Web_team [Web team]
        direction TB
        Web_dispatcher[Web dispatcher] --- REST --- Web_JavaScript[Web JavaScript]
        Web_JavaScript --- REST --- Web_dispatcher
    end

```

Microservices Application Architecture 31 © Copyright IBM Corporation 2016-18

IBM Training **IBM**

Business service microservices dependencies: Typical

- Business services can delegate to other business services
 - Try to avoid circular dependencies
- Be careful that each service still implements a complete task
 - They are not separate layers

```

graph TD
    BS1[Business service 1] --> BS3[Business service 3]
    BS2[Business service 2] --> BS4[Business service 4]
    BS2 --> BS5[Business service 5]
    BS3 --> BS6[Business service 6]
    BS3 --> BS7[Business service 7]

```

Microservices Application Architecture 32 © Copyright IBM Corporation 2016-18

IBM Training IBM

Business service microservices dependencies: Death Star

- Practically every service delegates to or coordinates with every other service
- Difficult to deploy and maintain

The diagram illustrates a fully connected network of five business services, labeled Business service 1 through Business service 4 and Business service n. Every service is interconnected with every other service, creating a complex web of dependencies. The connections are represented by double-headed arrows between all pairs of services.

Microservices Application Architecture © Copyright IBM Corporation 2016-18

IBM Training IBM

Conclusion

Microservices Application Architecture © Copyright IBM Corporation 2016-18

IBM Training 

Unit summary

- List the benefits cloud native applications and the microservices application architecture
- Describe how the microservices application architecture compares with existing application architectures
- Show how an application can be designed as a set of microservices

Microservices Application Architecture 35 © Copyright IBM Corporation 2016-18

IBM Training 

Next

- More on microservices
 - Microservice development
 - Service mesh
- Deploying cloud applications
 - Docker and Kubernetes in IBM Cloud Container Service
 - Cloud Foundry in IBM Cloud
 - Swarm in Docker Enterprise Edition for IBM Cloud
 - OpenWhisk in IBM Cloud Functions
- IBM Cloud capabilities
 - IBM Cloud architecture
 - WebSphere Liberty for microservices using Java, Java EE, and MicroProfile
 - Cloud data services
 - Integration with systems of record
- DevOps
 - Continuous delivery of microservices

Microservices Application Architecture 36 © Copyright IBM Corporation 2016-18

Resources

- *Microservices from Theory to Practice: Creating Applications in IBM Bluemix Using the Microservices Approach*, IBM Redbook, 26 August 2015, <http://bit.ly/ibmredmicro>
- "Microservices, SOA, and APIs: Friends or enemies?" by Kim J. Clark, *IBM developerWorks*, 21 January 2016
- "Microservices.TV" <https://developer.ibm.com/tv/microservices/>
- "Microservices in action, Part 1: Introduction to microservices" by Rick Osowski, *IBM developerWorks*, 26 June 2015
- "Introduction to Microservices and Cloud Native Application Architecture" by David Currie
- "Microservices" by James Lewis and Martin Fowler, <http://bit.ly/microsvc>
- *Building Microservices: Designing Fine-Grained Systems* by Sam Newman, O'Reilly Media, February 2015, <http://bit.ly/oreillymicro>
- "Netflix Hystrix: How It Works" <https://github.com/Netflix/Hystrix/>
- Spring Cloud <https://cloud.spring.io>
- *Release It!: Design and Deploy Production-Ready Software* by Michael T. Nygard, Pragmatic Bookshelf, April 2007
- *Domain-Driven Design: Tackling Complexity in the Heart of Software* by Eric Evans, Addison-Wesley Professional, August 2003
- *Enterprise Integration Patterns* by Gregor Hohpe and Bobby Woolf, Addison-Wesley Professional, October 2003

IBM Training

IBM

IBM Cloud Data Services Overview

© Copyright IBM Corporation 2016 - 2018
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

IBM Training

IBM

Topics

- ▶ Where and how will data be used?
 - Why use NoSQL?
 - Attributes of NoSQL databases
 - Implementation: Cloudant

IBM Cloud Data Services Overview

2

© Copyright IBM Corporation 2016 - 2018

IBM Training

Polyglot Persistence

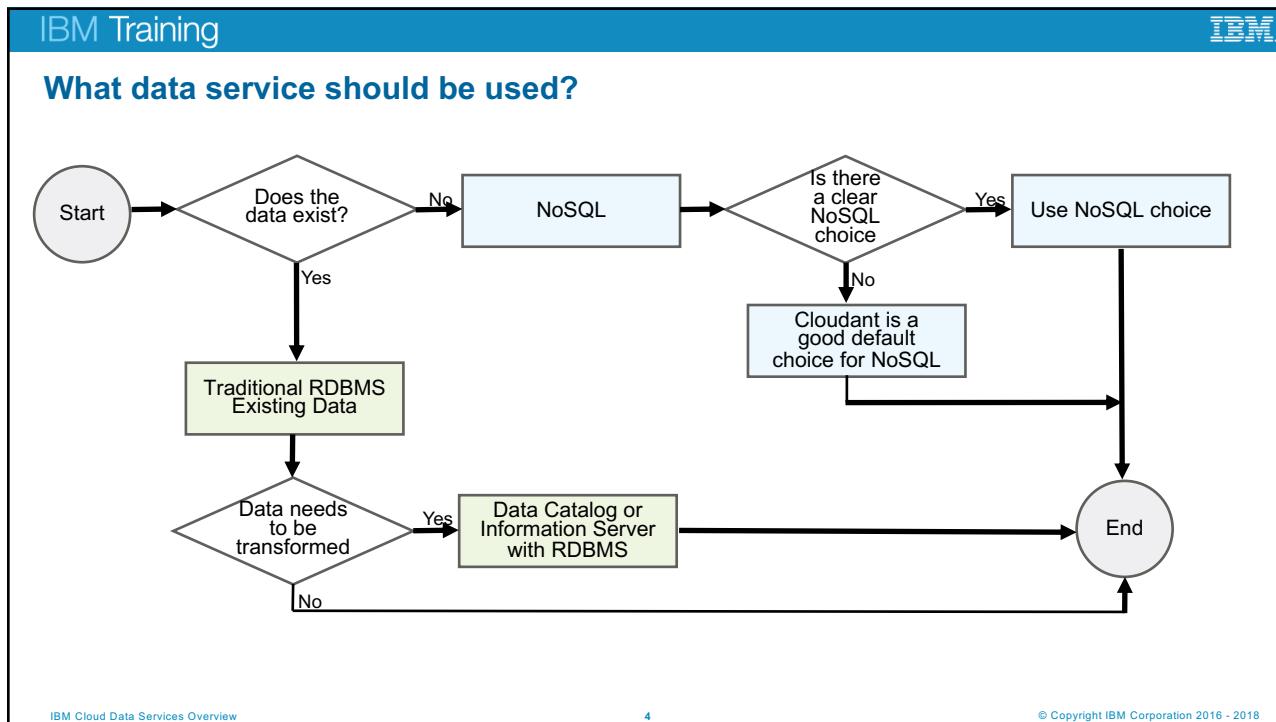
- A basic assumption is that your application will probably have to store some data persistently
 - For longer than the length of a single user session
- At one time, it was a simple assumption that Web programs would run on a relational database
 - Multiple relational database options exist for IBM Cloud Public
- The assumption of using a SQL database is challenged by the data management options collectively called “NoSQL”
- These data stores are grouped into four basic types
 - Key-value
 - Document
 - Column-family
 - Graph



IBM Cloud Data Services Overview

3

© Copyright IBM Corporation 2016 - 2018

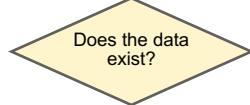


IBM Training 

Ask the right questions about the data

- Is there existing data in a relational database management system (RDBMS)?
- For cloud native development with new data don't create a new RDBMS
 - Initially it might seem to be easier and quicker
 - Can hinder rapid development
 - Traditional RDBMS requires structure, a schema
- What does the data "look like"
- Existing data probably already lives in an RDBMS
 - That doesn't mean you need to use it "as is"
 - There are services for data transformation and caching
- NoSQL databases provide:
 - Faster prototyping
 - No need for strict data typing, schema-less
 - No need to throw away non-conforming data

IBM Cloud Data Services Overview 5 © Copyright IBM Corporation 2016 - 2018



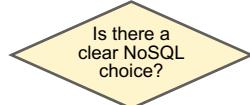
Does the data exist?

IBM Training 

Choosing NoSQL over RDBMS is not the final answer

- Look to see what type of problem you're trying to solve. Let the structure of your data and your queries define your technology choice.
- Simple key-value cache?
 - Use Redis
- Storing documents?
 - Use Cloudant
- Representing graph data?
 - Use Compose for JanusGraph
- Cloudant is a good default NoSQL database

IBM Cloud Data Services Overview 6 © Copyright IBM Corporation 2016 - 2018



Is there a clear NoSQL choice?

IBM Training

Options to work with SQL data in Cloud

- Use a provider with SQL data in cloud:
 - Compose databases
 - Db2, mysql, postgres, etc
 - Run database container with persistent volume
- Connect to data source outside the cloud
 - MQ
 - Message Hub
 - Kafka
- Transform and cleanse data to cloud
 - InfoSphere Information Server

The screenshot shows a grid of service cards:

- ibm-db2oltp-dev**: IBM Db2 Developer-C Edition 11.1.3.3. (Icon: Database)
- ibm-mariadb-dev**: MariaDB is developed as open source software and as... (Icon: Database)
- ibm-mqadvanced-server-dev**: IBM MQ queue manager. (Icon: Queue)
- ibm-rabbitmq-dev**: Open source message broker software that implements the Advanced... (Icon: RabbitMQ)
- ibm-eventstreams-dev**: Kafka is an open source stream processing platform used... (Icon: Kafka)
- ibm-ilsee-eval**: IBM InfoSphere Information Server for Evaluation v11.7 (Evaluation) (Icon: Server)

At the bottom left: IBM Cloud Data Services Overview. At the bottom right: © Copyright IBM Corporation 2018.

IBM Training

Topics

- Where and how will data be used?
- Why use NoSQL?
- Attributes of NoSQL databases
- Implementation: Cloudant

At the bottom left: IBM Cloud Data Services Overview. At the bottom right: © Copyright IBM Corporation 2016 - 2018.

IBM Training

Why use NoSQL databases

- Some data does not need to be normalized or categorized
 - CLOB, BLOB, reports, customer reviews, and so on
- If the data would be stored as a single unit, it is a waste of effort to write it to columns and tables of an RDBMS
 - A NoSQL database can handle this type of data much more rapidly
- Horizontal scaling
 - Query massive amount of data
 - Parallel processing is a common capability
- Designed to be reliable on unreliable hardware
 - NoSQL databases are distributed

IBM Cloud Data Services Overview 9 © Copyright IBM Corporation 2016 - 2018

IBM Training

NoSQL Landscape

Database Type	Examples	When best used
Key-Value	Redis, Memcached, DynamoDB	Storing interaction data, user preferences, simple shopping carts. Not great at set operations.
Document	Cloudant, MongoDB	Event logging, content management, analytics. Not great at complex queries.
Column-Family	Cassandra, Db2 Warehouse	Event logging, counters, blogs. Not compatible with ACID transactions.
Graph	GraphDB, JanusGraph, Neo4J, OrientDB	Social Networks, Location-based Services. Not optimal for certain types of bulk updates.



ibm-redis-ha-dev
Highly available Redis cluster with multiple sentinels and standbys.
[ibm-charts](#)



ibm-mongodb-dev
NoSQL document oriented database that stores JSON like documents...
[ibm-charts](#)



ibm-db2warehouse-dev
Db2 Warehouse Developer-C for Non-Production v2.5.0
[ibm-charts](#)



cassandra
Apache Cassandra is a free and open source distributed...
[incubator](#)



neo4j
Neo4j is the world's leading graph database.
[Incubator](#)



couchdb
A database featuring seamless multi master sync, that scales...
[incubator](#)

IBM Cloud Data Services Overview 10 © Copyright IBM Corporation 2016 - 2018

IBM Training

Topics

- Where and how will data be used?
- Why use NoSQL?

▶ Attributes of NoSQL databases

- Implementation: Cloudant

IBM Cloud Data Services Overview 11 © Copyright IBM Corporation 2016 - 2018

IBM Training

CAP Theorem

RDBMS

Consistency

Availability

Partition Tolerance

Redis
MongoDB
HBase

Cloudant
CouchDB
Cassandra

CA

CP

AP

CAP Theorem says you can only two:

- Consistency
- Availability
- Partition Tolerance

Horizontal replicas can become partitioned

- Consequence: Choose always consistent or always available (can't have both!)
- Configurable in many NoSQL databases

Eventually consistent

- To maintain availability, partitioned replicas lose synchronization

IBM Cloud Data Services Overview 12 © Copyright IBM Corporation 2016 - 2018

IBM Training **IBM**

NoSQL and eventual consistency

- Eventual consistency
 - Many NoSQL databases use replicas (horizontal scaling) to maintain high availability
 - The network between database replicas can become partitioned
 - CAP Theorem: When partitioned, either consistency or availability must be sacrificed
 - To maintain availability, most databases are configured to sacrifice consistency
 - When the partitioning is resolved, synchronization reestablishes consistency
- Consequences
 - Transactions are not ACID across replicas
 - Updates are inconsistent until they replicate (usually milliseconds)
 - Applications must tolerate the possibility that two reads may produce inconsistent results
- Be aware of the limitations of your particular database choice
 - Don't assume that your data is always consistent!

IBM Cloud Data Services Overview 13 © Copyright IBM Corporation 2016 - 2018

IBM Training **IBM**

Eventual consistency works well

- Eventual consistency sounds problematic
 - Many business domains are eventually consistent (hotels, airlines, and so on)
- How is conflict resolution handled?
 - "Last writer wins" can overwrite data, resolution is non trivial
 - Reconciliation can be achieved by time stamps (epochs, vector clocks)
- CouchDB (Cloudant) uses Multi-Version Concurrency Control (MVCC)
 - No locking required
 - Each read request sees the most recent snapshot of the database
- Database detects update conflicts
 - Application or user can reconcile the conflicts

IBM Cloud Data Services Overview 14 © Copyright IBM Corporation 2016 - 2018

IBM Training 

Understanding NoSQL data retrieval

- There are several approaches for accessing data
 - These are not necessarily specific to NoSQL
 - Understanding these options can help in deciding which NoSQL to use
- Four common approaches
 1. SQL-style query
 2. Look-up by ID
 3. Object navigation
 4. Map-reduce

IBM Cloud Data Services Overview  15 © Copyright IBM Corporation 2016 - 2018

IBM Training 

SQL-style queries

- If you have always used SQL queries, your natural tendency is to use them for your new cloud app

```
SELECT * FROM customers WHERE postal_code = '51922'
```

- Database considers all elements, filters out those that don't match

• Don't Do It 

- Sequential queries perform poorly with most NoSQL databases
- NoSQL databases provide APIs for other methods to retrieve your data much more efficiently

IBM Cloud Data Services Overview  16 © Copyright IBM Corporation 2016 - 2018

IBM Training

Lookup by ID

- Do you already know the key or ID of the data you need?
 - Primary key access
 - Typical for relational databases
 - Available for most NoSQL databases

Yes  <ul style="list-style-type: none"> • Key Value Pair Retrieval <ul style="list-style-type: none"> ▪ Provide the key and retrieve the values ▪ A key-value database fits this scenario well 	No  <ul style="list-style-type: none"> • This requires another approach for accessing the data <ul style="list-style-type: none"> ▪ Object navigation ▪ Map-reduce • Many NoSQL databases provide built-in map-reduce functionality
------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

IBM Cloud Data Services Overview

17

© Copyright IBM Corporation 2016 - 2018

IBM Training

Key-Value Store Databases

- Values are stored in relation to unique keys
- Values are retrieved by supplying their key


```
String key = "location";
Object value = find(key);
```
- Value data usually has no known structure
 - Any kind of data can be stored as a value
 - Very flexible
- Examples
 - Compose for Redis
 - Memcached Cloud
 - Amazon DynamoDB

Key	Value
firstName	Bugs
lastName	Bunny
location	Earth

IBM Cloud Data Services Overview

18

© Copyright IBM Corporation 2016 - 2018

IBM Training

Object Navigation

- Entity-Relationship-Attribute (ERA)
 - Each entity has attributes
 - Entities are connected by relationships
 - a.k.a. Nodes-Edges-Properties
- Implemented in graph databases
 - Examples: Compose for JanusGraph, Neo4J, OrientDB
 - Most NoSQL databases have this capability
- Object-oriented
 - Entities are objects
 - Relationships are links
 - Attributes are instance variables
- Retrieval done by following relationships
 - “Give me the members of the Chess Club”
 - Follows links rather than searching multiple tables – more efficient

Entity
Relationship
Attributes

Entity: Id: 2, Name: Bob, Age: 22
Relationship: Id: 3, Type: Group, Name: Chess
Attributes: Id: 1, Name: Alice, Age: 18

Labels: knows, Members, is_member
Since: 2001/10/03, 2001/10/04, 2005/7/01
Ids: 100, 101, 102, 103, 104, 105

IBM Training

Why use a graph database?

- Relationships are important in graph databases
 - Allow more complex relationships than relational databases
 - Allow for dynamic relationships
 - Can describe relationships in greater detail
- Graph databases can be schema-less
 - Allows flexibility in data typing
- Can apply Graph Theory to databases
 - Discover disjoint sets within the data
 - Find minimal routes between nodes

IBM Training **IBM**

Property Graphs

- Contains vertices and edges with properties
- Vertex represents an entity
 - A vertex has a label
 - Vertices have unique IDs
 - Vertices have properties
- Edge represents a directional relationship
 - Edges have labels
 - Each edge has a unique ID
 - Edges have properties
- Graph computing makes a distinction between the structure of the graph and the process of navigating the graph
 - JanusGraph based on TinkerPop3 and the traversal language is Gremlin

Graph Computing

Structure + Process

Graph

Traversal

Traversing

IBM Cloud Data Services Overview 21 © Copyright IBM Corporation 2016 - 2018

IBM Training **IBM**

The JanusGraph data browser

```
def g=ConfiguredGraphFactory.open("example").traversal();def saturn=g.V().has("name", "saturn");saturn.v
```

Filter:

Vertices: 12

IBM Cloud Data Services Overview 22 © Copyright IBM Corporation 2018

IBM Training

Map-Reduce

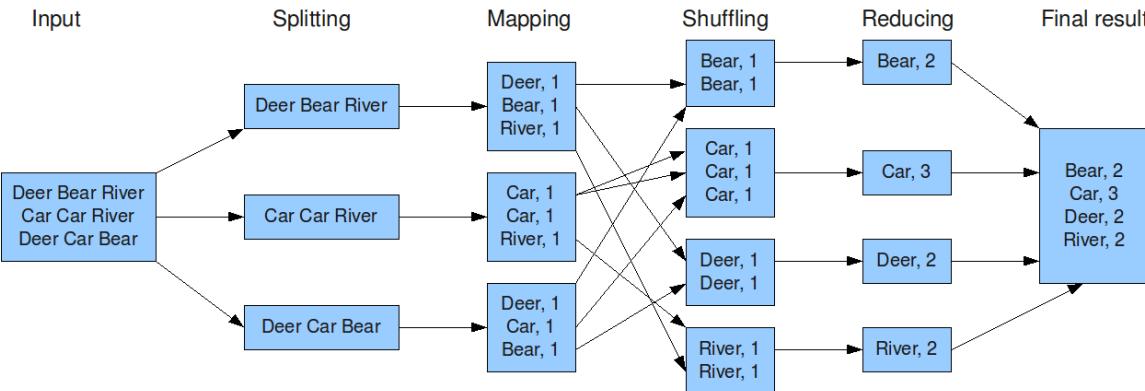
- Programming model designed to exploit parallel processing opportunities of clusters
 - Invented at Google Research
 - Horizontal scalability – job can be split across multiple servers
 - Hadoop is most well known implementation
 - Many NoSQL databases provide this capability
- Comprised of 2 major steps:
 1. Map
 - Converts a set of data into intermediate key / value pairs
 2. Reduce
 - Takes the output of Map step as input
 - Combines output of many map steps into a smaller set of key / value pairs

 hadoop

IBM Cloud Data Services Overview 23 © Copyright IBM Corporation 2016 - 2018

IBM Training

Consider this Example – How many times is each word used?



```

graph LR
    Input[Deer Bear River  
Car Car River  
Deer Car Bear] --> Splitting[Deer Bear River  
Car Car River  
Deer Car Bear]
    Splitting --> Mapping[Deer, 1  
Bear, 1  
River, 1  
Car, 1  
Car, 1  
River, 1  
Deer, 1  
Car, 1  
Bear, 1]
    Mapping --> Shuffling[Bear, 1  
Bear, 1  
Car, 1  
Car, 1  
Car, 1  
Deer, 1  
Deer, 1  
River, 1  
River, 1]
    Shuffling --> Reducing[Bear, 2  
Car, 3  
Deer, 2  
River, 2]
    Reducing --> FinalResult[Bear, 2  
Car, 3  
Deer, 2  
River, 2]
  
```

Note that this method could be used to answer other questions:
 What word is used most, least?

IBM Cloud Data Services Overview 24 © Copyright IBM Corporation 2016 - 2018

IBM Training IBM

Topics

- Where and how will data be used?
- Why use NoSQL?
- Attributes of NoSQL databases

▶ Implementation: Cloudant

IBM Cloud Data Services Overview 25 © Copyright IBM Corporation 2016 - 2018

IBM Training IBM

Introduction to Cloudant

- A fully managed NoSQL Database as a Service
- Transactional JSON *document* database with RESTful API
- Can spread data across data centers and devices for scale plus high availability (HA)
- Ideal for apps that require these features:
 - Massive, elastic scalability
 - High availability
 - Geolocation services
 - Full-text search
 - Occasionally connected users

 **Cloudant**



IBM Cloud Data Services Overview 26 © Copyright IBM Corporation 2016 - 2018

IBM Training

Cloudant database replication

- Bi-directional (push/pull == sync)
- Continuous or point-in-time
- Filterable (replicate subset of docs in a database)
- Durable (replication will pick up where it left off)
- Allows for eventual consistency
- Separate clusters
 - In different regions, active-active
 - Can be used to cache data on a phone

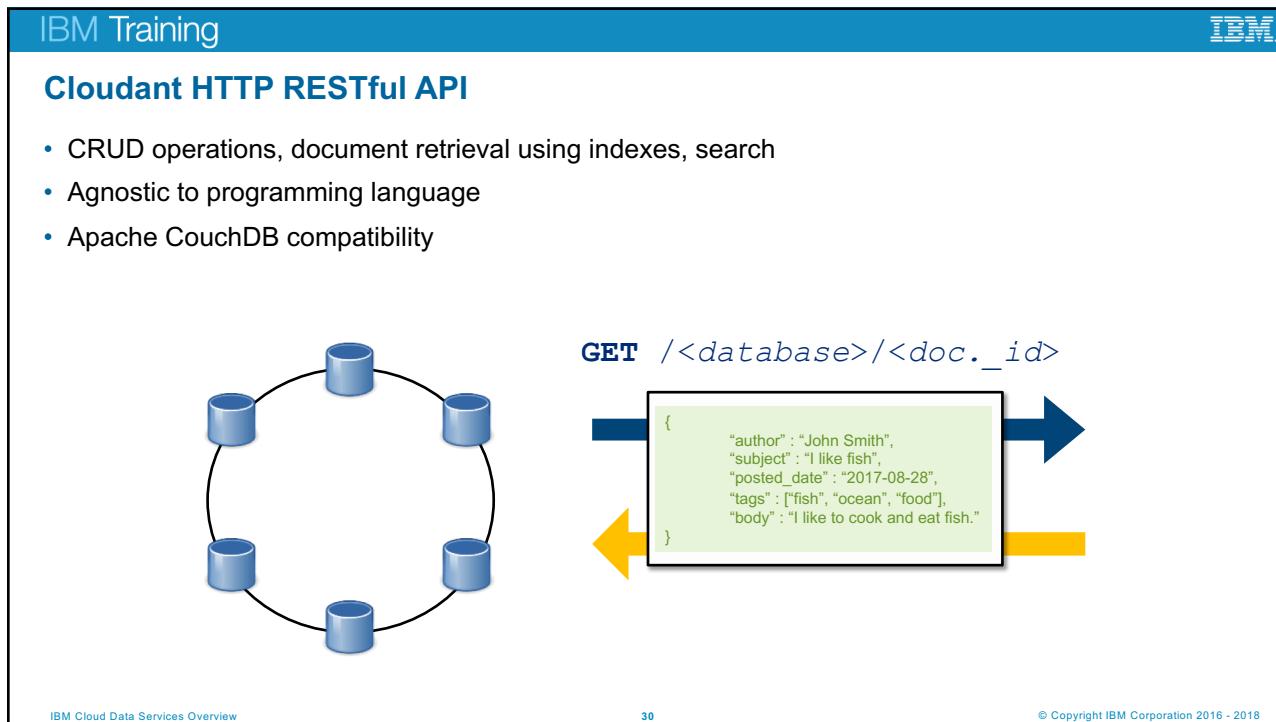
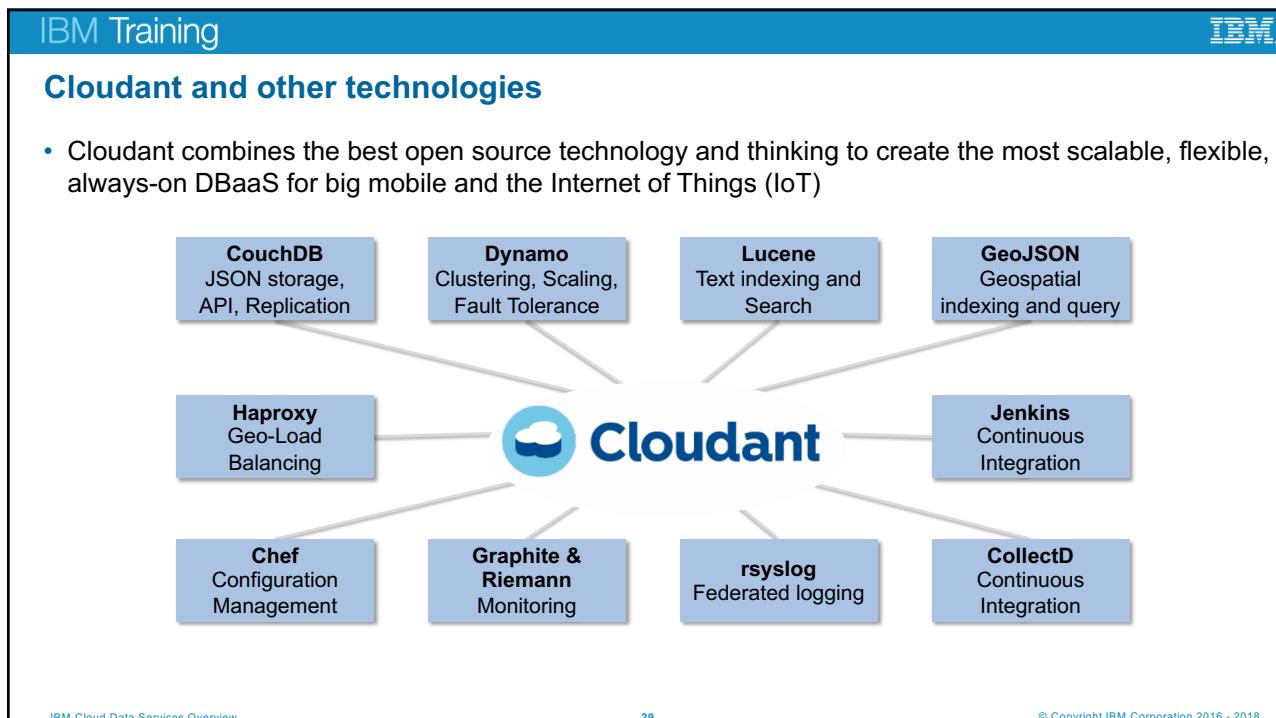
IBM Cloud Data Services Overview 27 © Copyright IBM Corporation 2016 - 2018

IBM Training

Cloudant API and Query Options

CRUD	Primary Index	Secondary Indices	Search	Cloudant Geospatial	Cloudant Query
JSON Documents	<ul style="list-style-type: none"> • Direct doc lookup by _id • Use when you want a single document and can find by the doc _id 	<ul style="list-style-type: none"> • Exists OOTB • Stored in a b-tree • Primary key → doc._id • Use when you can find documents based on their _id • Pull back a range of keys 	<ul style="list-style-type: none"> • Built using MapReduce • Stored in a B-tree • Key → user-defined field(s) • Use when you need to analyze data or get a range of keys • Ex: count data fields, sum/average numeric results, advanced stats, group by date, etc. 	<ul style="list-style-type: none"> • Built using Lucene • FTI: Any or all fields can be indexed • Ad-hoc queries • Find docs based on their contents • Can do groups, facets, and basic geo queries (bbox & sort by distance) 	<ul style="list-style-type: none"> • Stored in R* tree • Lat/Long coordinates in GeoJSON • Complex geometries (polygon, circularstring, etc.) • Advanced relations (intersect, overlaps, etc.) • Mongo-style querying • Not a 1:1 mapping • Built natively in Erlang • Ad-hoc queries • Lots of operators (>, <, IN, OR, AND, etc.) • Intuitive for people coming from Mongo or SQL backgrounds

IBM Cloud Data Services Overview 28 © Copyright IBM Corporation 2016 - 2018



IBM Training **IBM**

Cloudant API: Inserting a document - HTTP PUT or POST

- using **POST** with `_id` in the document body:

```
POST https://<username>.cloudant.com/authors
```

```
{
  "_id": "101",
  "name": "Mary Smith",
  "agent": "John Reid",
  "telephone": "512-555-1212"
}
```

```
{
  "ok": "true",
  "id": "101",
  "rev": "1-0af5e..."
}
```

- using **PUT** with `_id` in the URL:

```
PUT https://<username>.cloudant.com/authors/101
```

```
{
  "name": "Mary Smith",
  "agent": "John Reid",
  "telephone": "512-555-1212"
}
```

```
{
  "ok": "true",
  "id": "101",
  "rev": "1-0af5e..."
}
```

IBM Cloud Data Services Overview 31 © Copyright IBM Corporation 2016 - 2018

IBM Training **IBM**

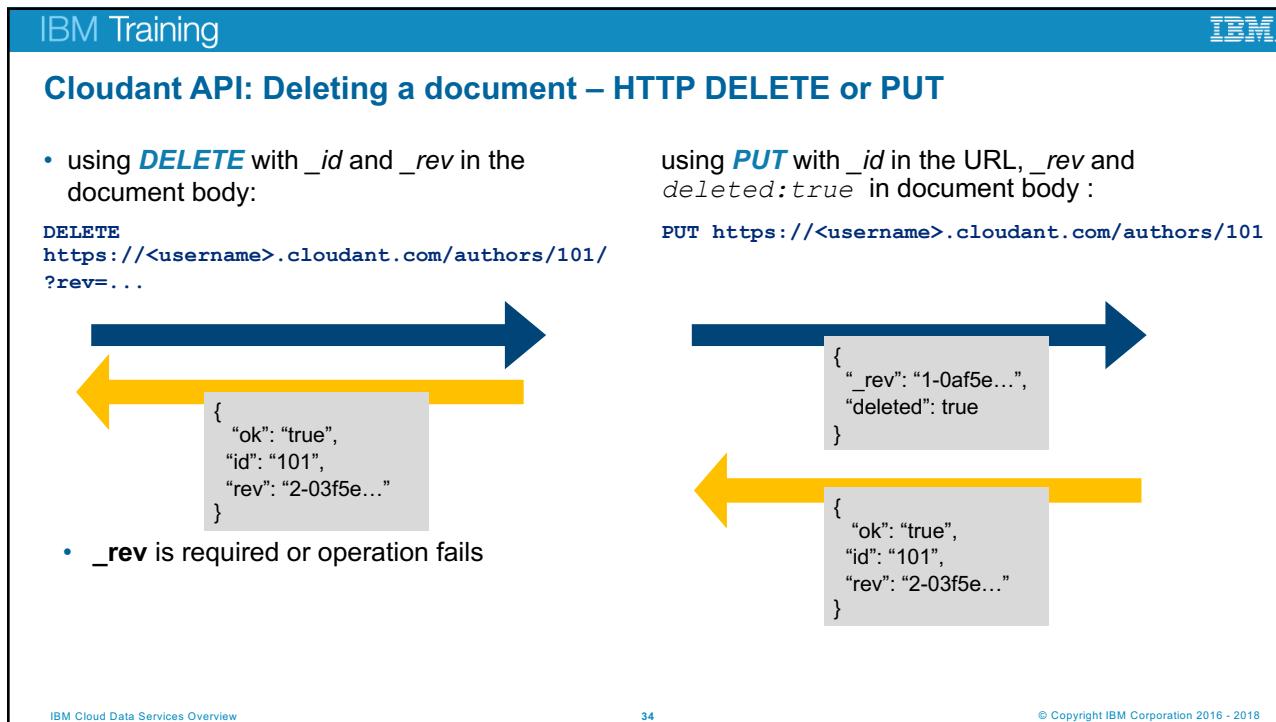
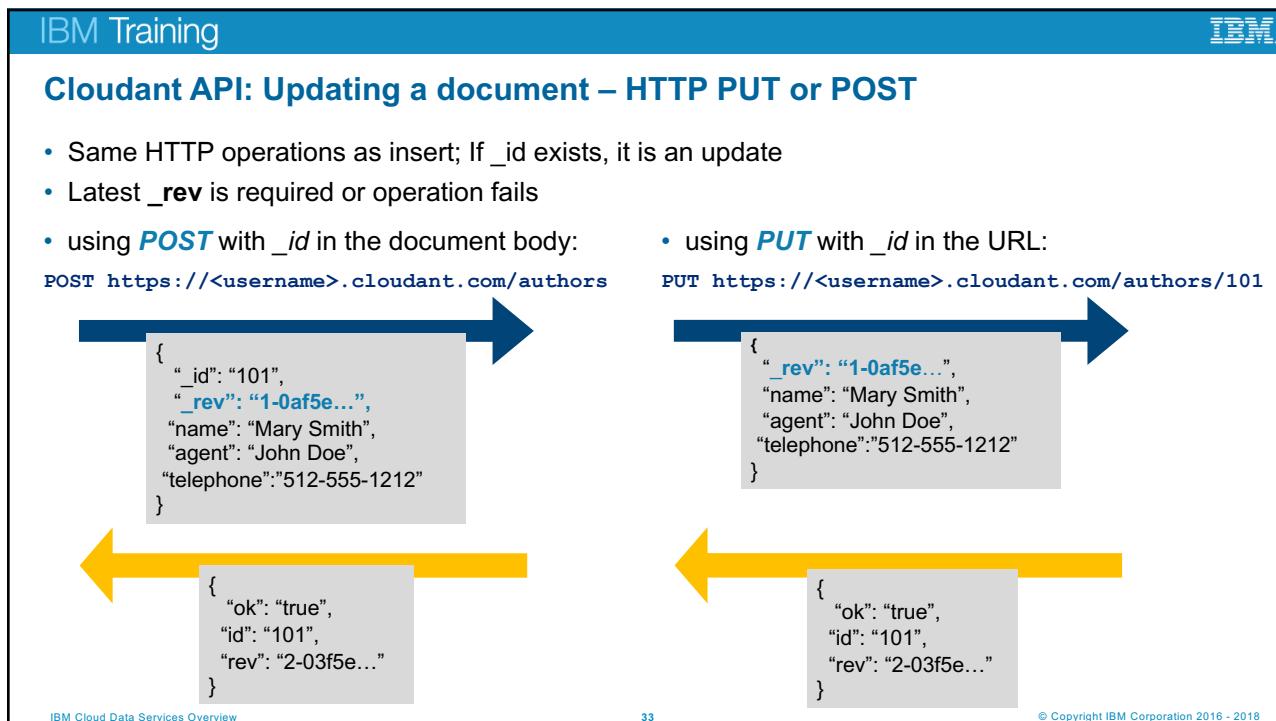
Cloudant API: Returning a single document

- HTTP GET with database name and `_id` of document
- For example, to get document with `_id` **100** from the **authors** database:

```
GET https://<username>.cloudant.com/authors/100
```

```
{
  "_id": "100",
  "rev": "1-044d6f...",
  "name": "John Smith",
  "agent": "Mary Reid",
  "telephone": "512-555-1212"
}
```

IBM Cloud Data Services Overview 32 © Copyright IBM Corporation 2016 - 2018



IBM Training 

Cloudant documents: `_id` and `_rev`

- Each document has an `_id` (ID) field that is unique per database
 - Any string can be supplied as an `_id`, but it is recommended that you allow Cloudant to generate a UUID (universally unique identifier)
- There is also a unique `_rev` (revision number) field per document
 - Generated by an md5 hash of the transport representation of the document
 - N -prefix reflects the number of times this document has been updated
 - Updates to existing documents must provide the latest `_rev` value
 - Otherwise the update request is rejected

```
{  
  "_id": "7f123e23a328bd50ee123cd35452ae47",  
  "_rev": "2-3123414209",  
  "title": "IBM Cloudant Redbook",  
  "author": "Christopher Bienko"  
}
```

IBM Cloud Data Services Overview  35 © Copyright IBM Corporation 2016 - 2018

IBM Training 

Summary

- Think about how and where data is used
- Select the best cloud database service
- NoSQL databases do not all have the same capabilities
- Cloudant is a good choice for NoSQL

IBM Cloud Data Services Overview  36 © Copyright IBM Corporation 2016 - 2018

IBM Training 

References (1 of 2)

Cloudant

- <https://cloudant.com/learning-center/>

“RDBMS to NoSQL: Reviewing Some Next-Generation Non-Relational Database's”

- <http://liacs.leidenuniv.nl/~stefanovtp/courses/StudentenSeminarium/Papers/DB/3.IJAEST-Vol-No-11-Issue-No-1-RDBMS-to-NoSQL-Reviewing-Some-Next-Generation-Non-Relational-Database's-015-030.pdf>

Eventual consistency

- <http://guide.couchdb.org/draft/consistency.html>

“NoSQL Distilled”

- <http://www.amazon.com/dp/0321826620>

Extract, Transform, and Load (ETL)

- https://en.wikipedia.org/wiki/Extract,_transform,_load

Graph database introduction

- <https://www.compose.com/articles/graph-101-getting-started-with-graphs/>

IBM Cloud Data Services Overview 37 © Copyright IBM Corporation 2016 - 2018

IBM Training 

References (2 of 2)

JanusGraph concepts

- <https://help.compose.com/docs/janusgraph-concepts>

Introduction to graph concepts

- <https://tinkerpop.apache.org/docs/3.2.3/reference/#intro>

Compose for JanusGraph on IBM Cloud

- <https://console.bluemix.net/docs/services/ComposeForJanusGraph/index.html#getting-started-with-compose-for-janusgraph>

IBM Cloud Data Services Overview 38 © Copyright IBM Corporation 2016 - 2018

IBM Training

IBM

Developing Microservices

© Copyright IBM Corporation 2016-18
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

IBM Training

IBM

Prerequisites

- Microservices
 - Microservice architecture

Developing Microservices

2

© Copyright IBM Corporation 2016-18

IBM Training 

Unit objectives

- Explain how to use microservices to design for failure
- Describe how microservices connect with one another
- List techniques for refactoring existing applications into microservices

Developing Microservices 3 © Copyright IBM Corporation 2016-18

IBM Training 

Developing microservices

Developing Microservices 4 © Copyright IBM Corporation 2016-18

IBM Training 

Developing a microservices application

1. Identify a set of independent business tasks
 - Build initial microservices around those tasks
 - Teams that can work independently
 - Components that can be deployed, scale, and fail independently
2. Design for failure
 - Use microservices to make the application more robust
3. Design for scale
 - Service discovery
 - For example, Eureka
 - Configuration repositories
 - For example Zookeeper, Archaius
 - Common logging

- A service mesh is very helpful for steps 2 and 3

Developing Microservices 5 © Copyright IBM Corporation 2016-18

IBM Training 

Design for failure

- Any service can fail
 - In a monolith, when one part stops working, it all stops working
 - Application must keep working and stay responsive
- Employ patterns for resiliency
 - **Service Registry:** Dynamic listing of service instances
 - **Circuit Breaker:** Block calls to a service that's not working
 - **Bulkhead:** Separate connection pools for separate resources
 - **Command:** Make requests easy to retry, throttle, and monitor
- Test for failure
 - Purposely, randomly cause problems
 - Example: Istio fault injection, Netflix OSS Simian Army framework



Developing Microservices 6 © Copyright IBM Corporation 2016-18

IBM Training 

Service registry

- How service clients find service providers
 - Each service provider is a microservice instance
 - Client can use any instance of a microservice class
 - The pool of instances and their locations can change:
 - Elasticity
 - Health management
 - Load distribution
- Service registry
 - As service instances are started, they are listed in the registry
 - The registry knows the location of each service, even after it moves
 - The registry can distribute load across the instances
- Examples
 - Eureka, Consul
 - KubeDNS
 - CloudFoundry GO Router

Developing Microservices 7 © Copyright IBM Corporation 2016-18

IBM Training 

Service discovery in IBM Cloud compute models

- The Cloud Foundry router is the service registry
 - Each microservice class is a runtime with a route
 - Cloud Controller manages microservice instances
 - Router distributes load across the microservice instances
 - Blue-green deployment supports multiple versions of a microservice
- Kubernetes clients find the microservices in pods using Kube DNS and services
 - DNS enables finding services within namespaces
 - Each pod replica registers its IP address with its service
 - The service distributes load across the pod replicas
 - Istio Pilot provides platform agnostic service discovery interface
- For microservices deployed across regions, use an external registry
 - Examples:
 - Eureka and Zookeeper

Developing Microservices 8 © Copyright IBM Corporation 2016-18

IBM Training

Circuit breaker

- A service consumer is only as reliable as its service provider
- Avoid invoking an unreliable service instance
 - Service might throw an error
 - Service might time out
 - Network might fail
- Circuit breaker
 - Consumer can invoke reliably
 - Invokes unreliable service
 - Fails fast when the service isn't working

Developing Microservices

9

© Copyright IBM Corporation 2016-18

IBM Training

Command and circuit breaker

- Circuit breakers are typically used with command objects
- What happens if a circuit breaker fails?
 - How do you retry using a different service instance?
 - Every consumer shouldn't have to implement retry logic
- Command object
 - Wrap a service request as a command
 - A failed command can be retried
 - A command can implement a Plan B
 - Commands can be queued for throttling
 - Enable metrics to be measured for monitoring
- Example: Netflix Hystrix
 - Framework for making microservice invocations fault and latency tolerant
 - Classes and code for commands, circuit breakers, and bulkheads
 - Monitoring dashboard to view services' availability and performance
- Example: Istio Envoy failure handler

Developing Microservices

10

© Copyright IBM Corporation 2016-18

IBM Training IBM

Bulkhead

No bulkhead

- Single connection pool shared to connect to multiple external systems
- If one system blocks connections
 - All connections block on the one system
 - No way to connect to the working systems

Bulkhead

- Separate connection pool for each external system
- If one system blocks connections
 - All its connections block
 - Meanwhile, connections to the working systems still work
- Example: Hystrix or Istio

No bulkhead

```
graph TD; UserAction[User action] --> SharedPool[Shared connection pool]; SharedPool --> Backend1[Backend system 1]; SharedPool --> Backend2[Backend system 2];
```

Bulkhead

```
graph TD; UserAction1[User action] --> ConnPool1[Connection pool 1]; ConnPool1 --> Backend1[Backend system 1]; UserAction2[User action] --> ConnPool2[Connection pool 2]; ConnPool2 --> Backend2[Backend system 2];
```

Developing Microservices 11 © Copyright IBM Corporation 2016-18

IBM Training IBM

Microservices integration

Developing Microservices 12 © Copyright IBM Corporation 2016-18

IBM Training **IBM**

Communication between services

- Communication between services should be language-neutral
 - Services can be written in different languages
 - Cloud Foundry currently supports only inbound HTTP
- Most often REST synchronous protocol
 - JSON, HTTP
- For asynchronous protocol
 - Messaging system, typically AMQP or Kafka
 - JSON message payloads
 - Service activators to perform command messages synchronously
- Avoid alternative invocation styles
 - Leads to an explosion of functions and invocation styles

Developing Microservices 13 © Copyright IBM Corporation 2016-18

IBM Training **IBM**

Asynchronous communication

Asynchronous communication can make microservices more robust

- Requester doesn't have to block while provider runs
- Different requester instance can handle response
- Messaging system holds action and result

The diagram illustrates two communication models:

- Synchronous:** Shows a Business service requester and a Business service provider. A horizontal arrow labeled "Request" points from the requester to the provider, and another labeled "Response" points back from the provider to the requester, both using the "HTTP" protocol.
- Asynchronous:** Shows a Business service requester 1, a Business service requester 2, and a Business service provider. Requests from both requesters go to a central Messaging system. The Messaging system contains a "Request queue" and a "Response queue". Responses from the provider are stored in the Response queue. Arrows indicate the flow from requesters to the queue, from the queue to the provider, and from the provider back to the queue, labeled "Request" and "Response" respectively.

Developing Microservices 14 © Copyright IBM Corporation 2016-18

IBM Training **IBM**

Microservices intercommunication

- Lightweight protocols
 - REST such as JSON and HTTP
 - Messaging such as Kafka
- The aim is complete decoupling
 - Messaging wherever possible
 - Service registry or discovery
 - Load balancing
 - Circuit breaker patterns

Microservices application

Microservice component

Microservice component

Microservice component

REST / HTTP

Messaging

REST / HTTP

Developing Microservices 15 © Copyright IBM Corporation 2016-18

IBM Training **IBM**

Message Hub service in IBM Cloud

Hub for asynchronously connecting services inside IBM Cloud or beyond

Applications that are connected to events that happen in other IBM Cloud services, or from beyond the cloud

MQ, and other on-premises data sources

Spark

Microservices enable applications to evolve

Open protocols support polyglot runtimes, application-controlled behavior, and reactive scale

- Python
- Java
- Sinatra
- Ruby
- node.js

Insights from the data you already have

Data needs to be streamed from anywhere to one or many analytics engines

Developing Microservices 16 © Copyright IBM Corporation 2016-18

IBM Training 

Refactoring to microservices

Developing Microservices 17 © Copyright IBM Corporation 2016-18

IBM Training 

Evolving to microservices

- Agile development: What's the simplest thing that could possibly work?
 - A monolith is simpler
- Start with a monolith
 - Make it modular, and plan that modules can become microservices
 - Each module should be a vertical slice, a mini-app with its own data
- Start with a minimally viable product (MVP)
 - MVP is a module – improvements to existing features go in the same module
 - Implement new functional tasks in new modules
 - When a module implements more than one task, refactor into separate modules
- Separate modules into microservices as needed
 - Multiple teams want to work independently
 - Monolithic code base is becoming unwieldy to maintain
 - Modules need to deploy, scale, or fail independently

Developing Microservices 18 © Copyright IBM Corporation 2016-18

IBM Training **IBM**

Refactoring an existing application

- Places to refactor
 - Each REST service is potentially a microservice
 - Each SOAP web service or EJB is potentially a microservice
 - Especially stateless session beans
 - Redesign function-oriented interfaces to asset-oriented interfaces
 - Make the interface RESTful, such as using command objects
 - Use domain-driven design to discover your assets, which might be microservices
- Refactor the data for the microservices code
 - Each microservice manages its own data
 - A set of tables that is only used by one module
 - A set of tables that is only used at a particular time, like a daily status summary
 - Refactor the tables so that each table is used by only one module
 - Optimize for query time, not storage efficiency

Developing Microservices 19 © Copyright IBM Corporation 2016-18

IBM Training **IBM**

Repackaging an existing application

- Split up EAR files
 - Divide along functional lines
 - Package independent WAR files
 - Might require minor code changes
- Deploy WAR files separately
 - Its own Liberty for Java instant runtime
 - Its own Liberty Docker container
- Build, deploy, and manage separately
 - Use independent DevOps pipelines for each WAR file
 - Scale each separately and manage independently

Before

```

graph TD
    EAR[EAR] --- WARA[WAR A]
    EAR --- WARB[WAR B]
    EAR --- WARC[WAR C]
    EAR --- AppServer[App server]
  
```

After

```

graph LR
    WARA --- AppServerA[App server]
    WARB --- AppServerB[App server]
    WARC --- AppServerC[App server]
  
```

Developing Microservices 20 © Copyright IBM Corporation 2016-18

IBM Training

Refactoring data for microservices

- Master data management
 - Form a single, consistent view of widely used data entities
 - Develop microservices for working with that
- Code is storing blobs in your SQL database
 - Store those objects in a NoSQL database instead
 - Key-value store
 - Redis
- Active Record pattern: Flat objects unrelated to others
 - Islands of data unrelated to other data
 - Use a document model database
 - Cloudant or Mongo

Developing Microservices

21

© Copyright IBM Corporation 2016-18

IBM Training

Configuration for microservices

Make configuration part of your DevOps process

- Treat infrastructure as code, software-defined data center
- All application deployment must be automated
- Start small with simple scripts and build from there

Developing Microservices

22

© Copyright IBM Corporation 2016-18

IBM Training 

Conclusion

Developing Microservices 23 © Copyright IBM Corporation 2016-18

IBM Training 

Unit summary

- Explain how to use microservices to design for failure
- Describe how microservices connect with one another
- List techniques for refactoring existing applications into microservices

Developing Microservices 24 © Copyright IBM Corporation 2016-18

IBM Training 

Next

- More on microservices
 - Service mesh
- Deploying cloud applications
 - Docker and Kubernetes in IBM Cloud Container Service
 - Cloud Foundry in IBM Cloud
 - Swarm in Docker Enterprise Edition for IBM Cloud
 - OpenWhisk in IBM Cloud Functions
- IBM Cloud capabilities
 - IBM Cloud architecture
 - WebSphere Liberty for microservices using Java, Java EE, and MicroProfile
 - Cloud data services
 - Integration with systems of record
- DevOps
 - Continuous delivery of microservices

Developing Microservices 25 © Copyright IBM Corporation 2016-18

IBM Training 

Resources

- *Microservices from Theory to Practice: Creating Applications in IBM Bluemix Using the Microservices Approach*, IBM Redbook, 26 August 2015, <http://bit.ly/ibmredmicro>
- “Microservices, SOA, and APIs: Friends or enemies?” by Kim J. Clark, *IBM developerWorks*, 21 January 2016
- “Microservices.TV” <https://developer.ibm.com/tv/microservices/>
- “Microservices in action, Part 1: Introduction to microservices” by Rick Osowski, *IBM developerWorks*, 26 June 2015
- “Introduction to Microservices and Cloud Native Application Architecture” by David Currie
- “Microservices” by James Lewis and Martin Fowler, <http://bit.ly/microsvc>
- *Building Microservices: Designing Fine-Grained Systems* by Sam Newman, O'Reilly Media, February 2015, <http://bit.ly/oreillymicro>
- “Netflix Hystrix: How It Works” <https://github.com/Netflix/Hystrix/>
- Spring Cloud <https://cloud.spring.io>
- *Release It!: Design and Deploy Production-Ready Software* by Michael T. Nygard, Pragmatic Bookshelf, April 2007
- *Domain-Driven Design: Tackling Complexity in the Heart of Software* by Eric Evans, Addison-Wesley Professional, August 2003
- *Enterprise Integration Patterns* by Gregor Hohpe and Bobby Woolf, Addison-Wesley Professional, October 2003

Developing Microservices 26 © Copyright IBM Corporation 2016-18

IBM Training

IBM

Service Mesh

© Copyright IBM Corporation 2016-18
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

IBM Training

IBM

Prerequisites

- Microservices
 - Microservice architecture
 - Microservice development

Service Mesh 2 © Copyright IBM Corporation 2016-18

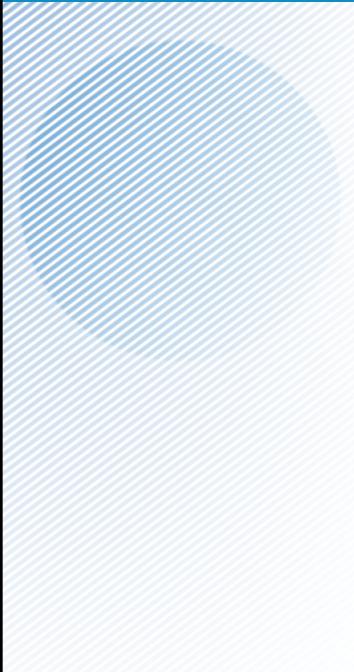
IBM Training 

Agenda

- Why service mesh?
- Common functions of a service mesh
- Services mesh implementations

Service Mesh 3 © Copyright IBM Corporation 2016-18

IBM Training 



Why Service Mesh?

4 © Copyright IBM Corporation 2016-18

IBM Training

Why Service Mesh?

- In monolithic applications, the location of resources used are:
 - Well-known
 - Relatively static
 - Found in configuration files or hard-coded
- This is not true for a microservices implementation
 - Required services could be anywhere
 - Cloud providers, public or private, datacenter, VMs, containers, servers
 - Service instances have dynamic locations
 - Come and go because of elasticity, failures, upgrades

ToonClips.com #6212 service@toonclips.com

© Copyright IBM Corporation 2016-18

IBM Training

The complexity tradeoff

Faster Independent Deployments

Hailo

Increased Operational Complexity

Service Mesh

6

© Copyright IBM Corporation 2016-18

IBM Training IBM

What are the challenges with microservices?

- How can I find the service I need?
 - What if it moves or fails?
- How can I scale my services?
 - Spread my load against multiple instances?
 - Can I control how I spread the load?
- How can I test new versions of services?
 - Without impacting existing users
- How can I test against failures?
 - How does the failure of one service affect the others?
 - Do my services have reasonable timeouts?
 - Can I avoid cascading failures?



Service Mesh 7 © Copyright IBM Corporation 2016-18

IBM Training IBM

A service mesh can help

- These are common challenges for everyone creating microservices
- Don't reinvent the wheel
- Variety of solutions using open source tools



Service Mesh 8 © Copyright IBM Corporation 2016-18

IBM Training IBM

A service mesh can help

- How can I find the service I need?
 - What if it moves or fails?
- How can I scale my services?
 - Spread my load against multiple instances?
 - Can I control how I spread the load?
- How can I test new versions of services?
 - Without impacting existing users
- How can I test against failures?
 - How does the failure of one service affect the others?
 - Do my services have reasonable timeouts?
 - Can I avoid cascading failures?

Service Mesh 9 © Copyright IBM Corporation 2016-18

IBM Training IBM

Common Functions of a Service Mesh

10 © Copyright IBM Corporation 2016-18

IBM Training

What is a service mesh?

- A network for services, not bytes
 - A domain specific router for microservice interactions
 - Makes service-to-service communication work better
- Distributed security
 - Service-to-service authentication and encryption – TLS
- Intelligent traffic management
 - Resiliency – Failure handling: circuit breaker, deadlines, retry
 - Routing – Routing rules, service discovery, load balancing, latency
 - Policy enforcement – Shared configuration, uniform communication
- Visibility
 - Distributed tracing – Track each request through the microservices
 - Instrumentation – Metrics and monitoring
 - Dashboards – See the microservices instances and the traffic through them

Service Mesh 11 © Copyright IBM Corporation 2016-18

IBM Training

Service Registry

- Database containing the network locations of service instances
- Should be highly available
- Instances self-register or are registered through a 3rd party registrar
- Heartbeat confirms instances are still alive

Examples:

- Netflix Eureka
- etcd (Used by Kubernetes & Cloud Foundry)
- Consul
- Apache Zookeeper

```

graph TD
    API[REST API] -- "Self-registration" --> ServiceInstanceA[SERVICE INSTANCE A]
    ServiceInstanceA -- "register(serviceName, \"10.4.3.1:8756\")\nheartbeat()\nunregister()" --> Registry[SERVICE REGISTRY]
  
```

The diagram illustrates the self-registration process. A green hexagonal box labeled 'SERVICE INSTANCE A' contains a 'REST API' icon and the IP address '10.4.3.1:8756'. An arrow labeled 'Self-registration' points from this box to a yellow rectangular box labeled 'SERVICE REGISTRY'. Below the arrow, the code 'register("serviceName, "10.4.3.1:8756") heartbeat() unregister()' is shown.

Source: dzone.com

Service Mesh 12 © Copyright IBM Corporation 2016-18

IBM Training

Service Discovery / Service Proxy

- How client services find & connect to provider services
- Each provider service provider could have multiple instances
- The pool of instances and their locations can change:
 - Elasticity
 - Health management
 - Load distribution
- 2 Models
 - Client-side
 - Server-side

The diagram shows a central **Service Registry** containing three entries: **SERVICE INSTANCE A**, **SERVICE INSTANCE B**, and **SERVICE INSTANCE C**. A **Service Proxy** is positioned between the client and the registry. The client has two options: it can either use a **Client library** to directly query the Service Registry, or it can use an **app** with a **sidecar** microservice. The sidecar interacts with the Service Registry to get the location of the provider services.

IBM Training

Client-Side Discovery

- Client code is responsible for:
 - Finding locations of available service instances
 - Load balancing logic for requests
- Client queries a service registry for location of service instances
- Client uses its own load balancing algorithm to select which instance to request
- Pros
 - Less components
 - Client can make application-specific load distribution decisions
- Cons
 - Client load distribution logic needs to be developed for every language used
 - Logic changes need to be in code

Examples:

- Netflix Eureka Service Registry
- Netflix Ribbon IPC Client

The diagram illustrates client-side discovery. A central **SERVICE REGISTRY** contains three entries: **SERVICE INSTANCE A**, **SERVICE INSTANCE B**, and **SERVICE INSTANCE C**. Three **Registry-aware HTTP Client**s query the registry. Each client then connects to its respective service instance via a **REST API**. The IP addresses of the instances are shown: **10.4.3.1:8756** for instance A, **10.4.3.89:4545** for instance B, and **10.4.3.20:333** for instance C.

IBM Training

Server-Side Discovery

- Client makes request to a load balancer
 - Load balancer queries service registry and routes requests to instances
 - Distribution logic in load balancer
- Pros
 - Client does not need discovery & load balancing logic
- Cons
 - Unless the load balancer is provided by the mesh, it is another component to manage

Examples:

- AWS Elastic Load Balancer (ELB)
- NGINX / NGINX Plus Http Server & Load Balancer
- Istio

Source: dzone.com

Service Mesh

15

© Copyright IBM Corporation 2016-18

IBM Training

Service discovery in IBM Cloud compute models

- The Cloud Controller and Router in Cloud Foundry provide service registry and service discovery
 - Each microservice class is a runtime with a route
 - Cloud Controller manages microservice instances
 - Router distributes load across the microservice instances
 - Blue-green deployment supports multiple versions of a microservice
- Kubernetes clients find the microservices in pods using Kube DNS and services
 - DNS enables finding services within namespaces
 - Each pod replica registers its IP address with its service
 - The service distributes load across the pod replicas

Service Mesh

16

© Copyright IBM Corporation 2016-18

IBM Training

Automated Testing

- Test for failure
 - Purposely, randomly cause problems
- Employ patterns for resiliency
 - Service Registry: Dynamic listing of service instances
 - Circuit Breaker: Block calls to a service that's not working
 - Bulkhead: Separate connection pools for separate resources
 - Command: Make requests easy to retry, throttle, and monitor

Examples:

- Netflix Chaos Monkey, Simian Army
- Istio



Service Mesh

17

© Copyright IBM Corporation 2016-18

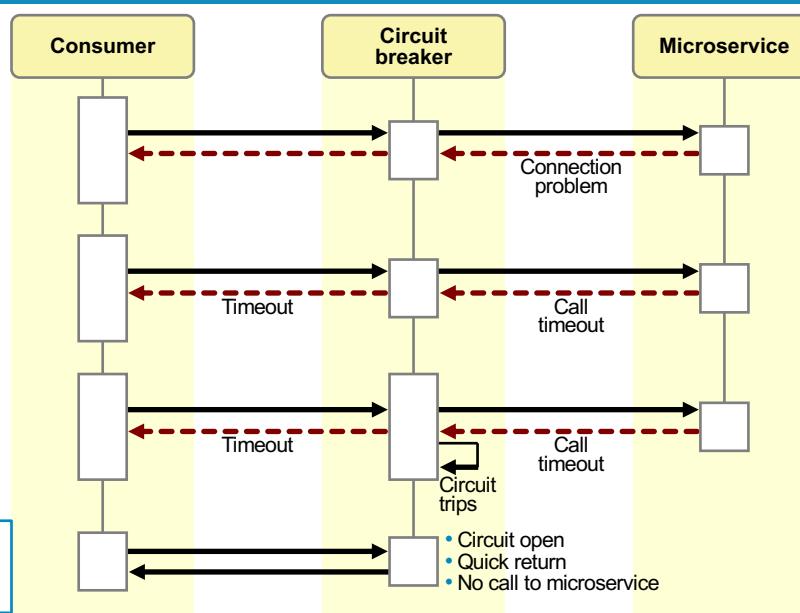
IBM Training

Circuit breaker

- A service consumer is only as reliable as its service provider
- Avoid invoking an unreliable service instance:
 - Service might throw an error
 - Service might time out
 - Network might fail
- Circuit breaker
 - Consumer can invoke reliably
 - Invokes unreliable service
 - Fails fast when the service isn't working

Example:

- Netflix Hystrix
- Istio



Service Mesh

18

© Copyright IBM Corporation 2016-18

IBM Training

IBM

Bulkhead

- No bulkhead
 - Single connection pool shared to connect to multiple external systems
 - If one system blocks connections
 - All connections block on the one system
 - No way to connect to the working systems
- Bulkhead
 - Separate connection pool for each external system
 - If one system blocks connections
 - All its connections block
 - Meanwhile, connections to the working systems still work

Example:

- Netflix Hystrix
- Istio

The diagram illustrates two connection architectures:

- No bulkhead:** A single "User action" box leads to a "Shared connection pool", which then branches to "Legacy system 1" and "Legacy system 2". If "Legacy system 1" blocks, both connections are lost.
- Bulkhead:** A "User action" box branches directly to two separate "Connection pool 1" and "Connection pool 2", each connected to a "Legacy system 1" or "Legacy system 2". Even if one connection pool is blocked, the other continues to work.

IBM Training

IBM

Services Mesh Implementations

Netflix OSS
Istio

20

© Copyright IBM Corporation 2016-18

IBM Training 



21 © Copyright IBM Corporation 2016-18

IBM Training 

What is Netflix OSS?

- **Netflix OSS** – A software-based control plane that is portable across clouds, can be run the same locally on a laptop or in any number of public Cloud providers.
- Netflix is the “poster child” of microservices
 - Netflix OSS is open source of many of the tools and methods they have used to build and improve their site
 - Over 33 root-level projects (and more actively maintained)
 - Most are Java libraries – Class structure could restrict you to only running your code in their framework
- Netflix OSS components available via <https://netflix.github.io>

Service Mesh 22 © Copyright IBM Corporation 2016-18

IBM Training

Major Netflix OSS Projects

- Eureka, Ribbon, Zuul – Service Registry / Discovery / Proxy
 - Client-side discovery – logic built into application libraries
 - Many applications automatically integrate with Eureka (or are enabled to through the use of [Spring Cloud](#) enhancements)
- Chaos Monkey, Simian Army – Failure Testing
 - Chaos Monkey was the first project
 - Simian Army includes Chaos Monkey and an “army” of other monkeys
 - Latency Monkey, Conformity Monkey, Doctor Monkey, Janitor Monkey, Security Monkey, and more
 - Programmatically control what parts of their infrastructure are unstable, unavailable, and unreliable
 - See how this impacts overall system health.
 - Original Chaos Monkey was built to the AWS APIs
 - Other parts of the Simian Army not written to specific provider APIs

Service Mesh 23 © Copyright IBM Corporation 2016-18

IBM Training

Netflix OSS Circuit Breaker – Hystrix, Turbine, Hystrix Dashboard

- Java-based implementation of the Circuit Breaker pattern
- Provides reliability beyond individual service calls & visibility into hot-spots / bottlenecks
- Isolates latency and increases fault tolerance at runtime
- Turbine aggregates clustered services into a single stream for visibility through Hystrix Dashboards

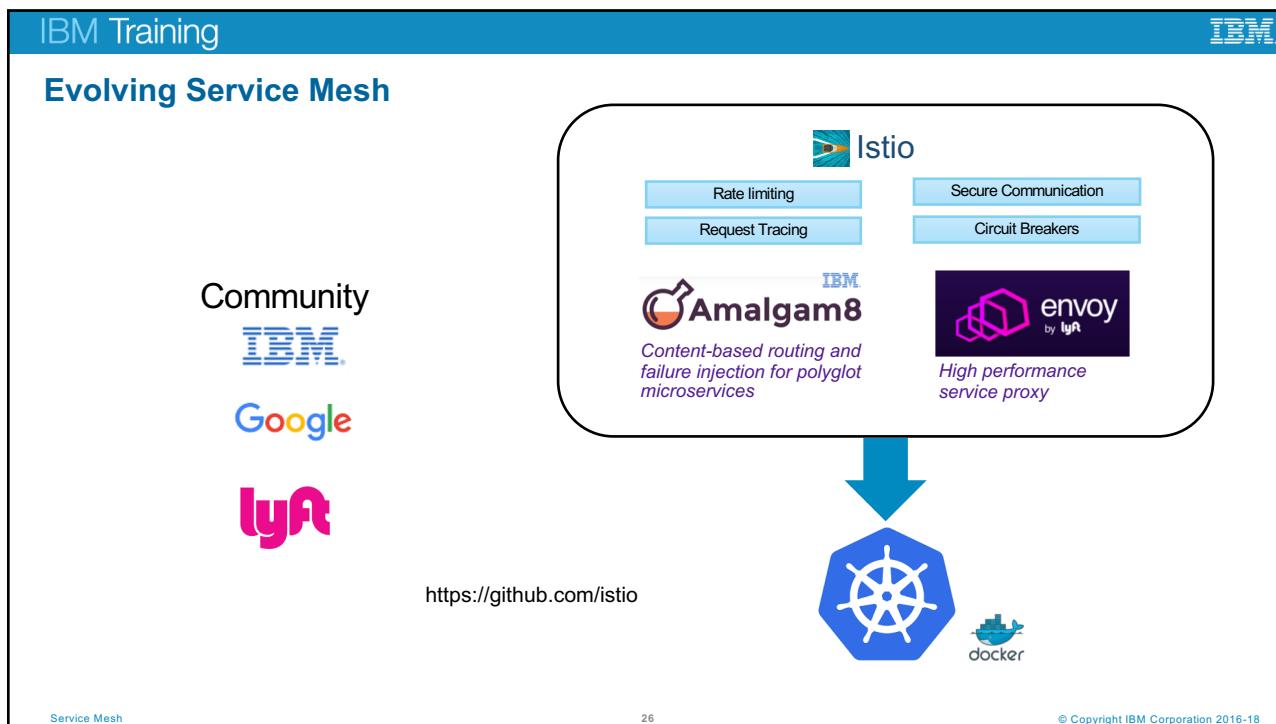
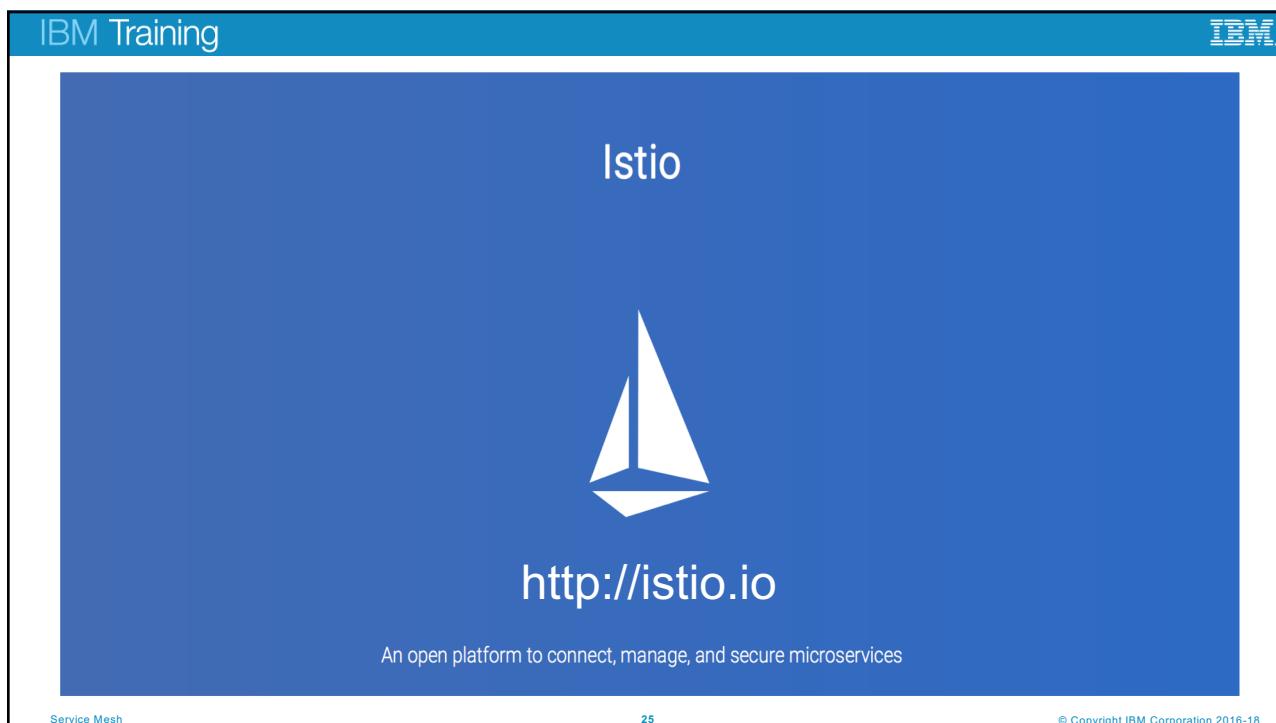
The diagram illustrates the data flow from multiple application instances to a central Turbine Stream Aggregator, which then feeds into Hystrix Dashboards for monitoring individual hosts or clusters. The dashboard provides real-time insights into service performance and reliability.

Hystrix Dashboard Cluster View Metrics:

- Subscribers: 200,545
- Hosts: 54.0/s
- Cluster: 20,056.0/s
- Circuit Closed
- Request rate: 370 1ms 4ms 90th 99th 99.9th
- Latency Percentiles: 10ms 44ms 61ms
- Error percentage of last 10 seconds: 0 %
- Hosts Median Mean: 370 1ms 4ms
- 2 minutes of request rate to show relative changes in traffic
- hosts reporting from cluster
- Rolling 10 second counters with 1 second granularity
- Successes: 200,545
- Short-circuited (rejected): 0
- Thread-pool Rejections: 94
- Failures/Exceptions: 0

Source: techblog.netflix.com

Service Mesh 24 © Copyright IBM Corporation 2016-18



IBM Training IBM

Intelligent Routing and Load Balancing

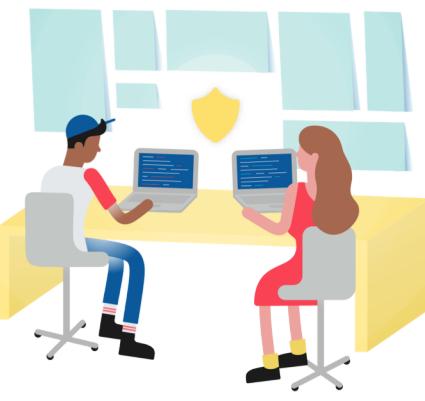
Control traffic between services with dynamic route configuration, conduct A/B tests, release canaries, and gradually upgrade versions using red/black deployments



Service Mesh 27 © Copyright IBM Corporation 2016-18

IBM Training IBM

Resilience Across Languages and Platforms

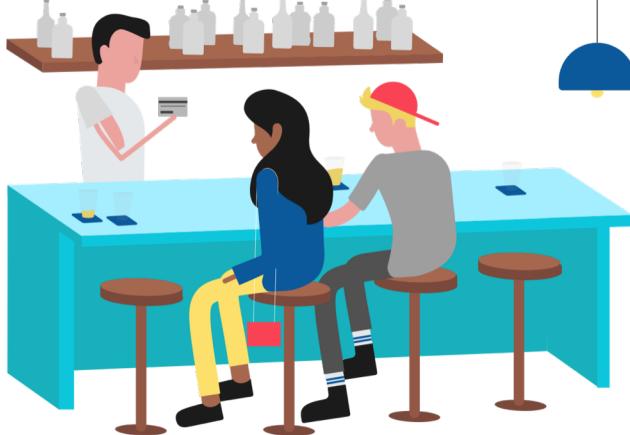


Increase reliability by shielding applications from flaky networks and cascading failures in adverse conditions

Service Mesh 28 © Copyright IBM Corporation 2016-18

IBM Training IBM

Fleet Wide Policy Enforcement

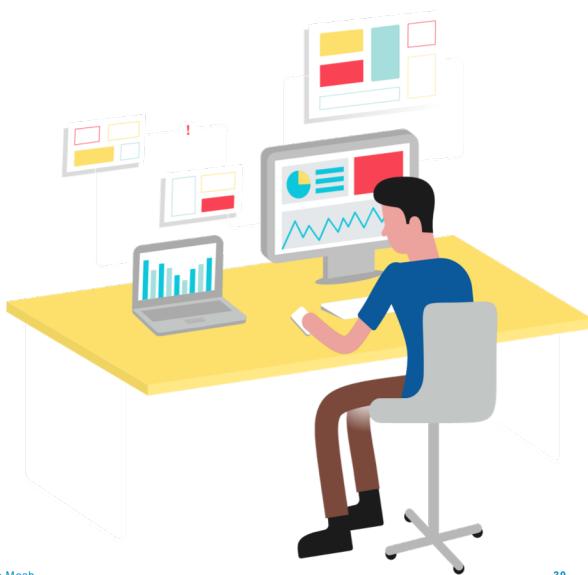


Apply organizational policy to the interaction between services, ensure access policies are enforced and resources are fairly distributed among consumers

Service Mesh 29 © Copyright IBM Corporation 2016-18

IBM Training IBM

In-Depth Telemetry and Reporting



Understand the dependencies between services, the nature and flow of traffic between them and quickly identify issues with distributed tracing

Service Mesh 30 © Copyright IBM Corporation 2016-18

IBM Training

Managing services: Proxy and sidecar

- A service mesh manages the interactions between services
 - It uses proxies, usually deployed as sidecars in pods
- Proxy – Controls access to another object
 - A service mesh uses proxies between services (and also clients)
 - Enables the mesh to manage interactions
- Pod – Kubernetes deploys a set of containers as a pod
 - All of the containers in a pod are hosted on the same node
 - A pod can manage the network access to the containers in the pod
 - A pod typically hosts a single container, but can host multiple as a unit
- Sidecar – Adds behavior to a container without changing it
 - Sidecar and service behave as a single enhanced unit
 - Pod hosts them as a single unit
- Mesh Member – Access to the service goes through the proxy
 - Deploys the proxy as a sidecar in the same pod
 - Configures the pod so that all network access goes through the proxy

The diagram illustrates three components:

- Proxy:** Shows a blue arrow pointing to a blue box labeled "proxy", which then points to a green box labeled "service".
- Sidecar:** Shows a "Pod" containing a "sidecar" (blue box) and a "service" (green box). They are connected by a double-headed arrow.
- Mesh Member:** Shows a "Pod" containing a "proxy" (blue box), which is connected to a "service" (green box) by a double-headed arrow. A blue arrow enters the pod from the left and exits to the right, indicating the proxy handles all network traffic.

Service Mesh 31 © Copyright IBM Corporation 2016-18

IBM Training

Mesh Request Flow

The diagram shows the flow of a request from the Internet through an Ingress Proxy (Kubernetes ingress controller) to a Pod containing Service A, then to another Pod containing Service B, and finally to External Services.

Inbound features (Service A):

- Service authentication
- Authorization
- Rate limits
- Load shedding
- Telemetry
- Request Tracing
- Fault Injection

Outbound features (Service B):

- Service authentication
- Load balancing
- Circuit breaker and retry
- Fine-grained routing
- Telemetry
- Request Tracing
- Fault Injection

Service Mesh 32 © Copyright IBM Corporation 2016-18

IBM Training

Envoy: Istio's proxy of choice



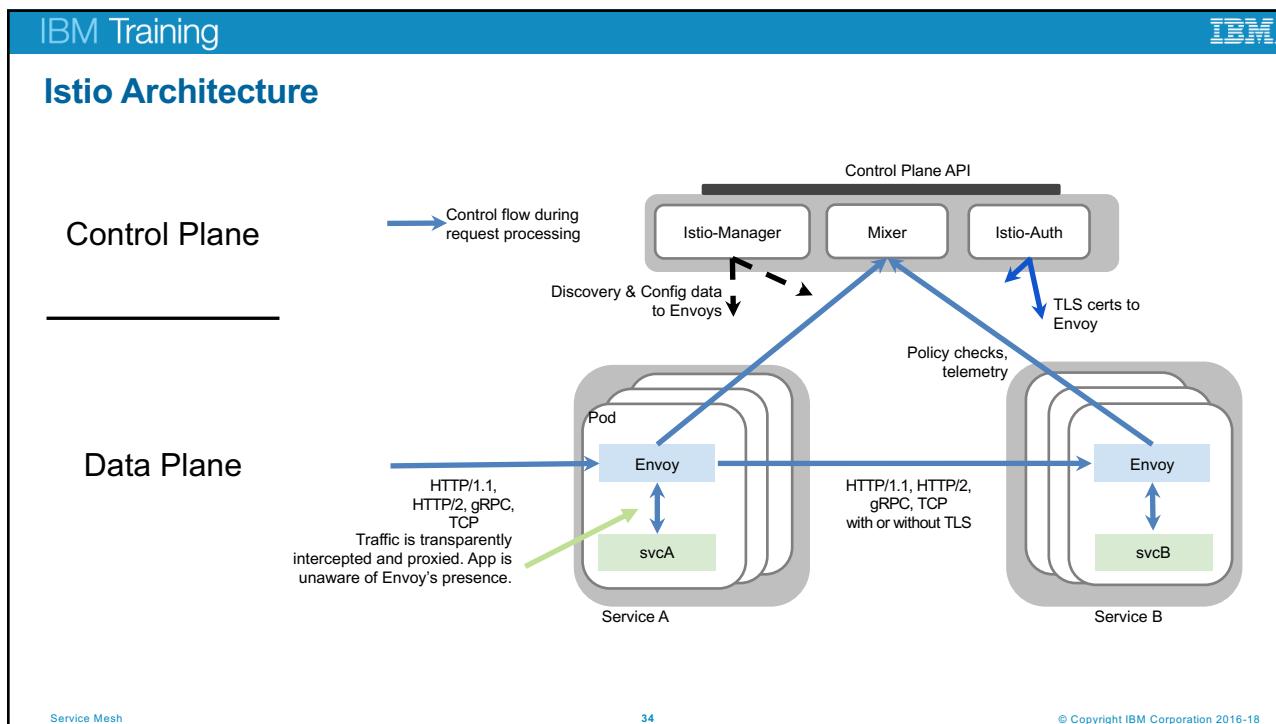
Goodies:

- A C++ based L4/L7 proxy
- Low memory footprint
- Battle-tested @ Lyft
 - 100+ services
 - 10,000+ VMs
 - 2M req/s
- Now a Cloud Native Computing Foundation (CNCF) open source project
 - Not unilaterally controlled by Lyft or Istio

Istio's contributions:

- Transparent proxying w/SO_ORIGINAL_DST
- Traffic routing and splitting
- Request tracing using Zipkin
- Fault injection

Service Mesh 33 © Copyright IBM Corporation 2016-18



IBM Training

Traffic Splitting

```
// A simple traffic splitting rule
destination: serviceB.example.cluster.local
match:
  source: serviceA.example.cluster.local
route:
- tags:
    version: v1.5
    env: us-prod
    weight: 99
- tags:
    version: v2.0-alpha
    env: us-staging
    weight: 1
```

Traffic control is decoupled from infrastructure scaling

Service Mesh

35

© Copyright IBM Corporation 2016-18

IBM Training

Traffic Steering

```
// Content-based traffic steering rule
destination: serviceB.example.cluster.local
match:
  httpHeaders:
    user-agent:
      regex: ^(.*)?(iPhone)(;.*)?$
precedence: 2
route:
- tags:
    version: canary
```

Content-based traffic steering

Service Mesh

36

© Copyright IBM Corporation 2016-18

IBM Training **IBM**

Resiliency

Istio adds fault tolerance to your application without any changes to code

```
// Circuit breakers
destination: serviceB.example.cluster.local
policy:
- tags:
  version: v1
circuitBreaker:
  simpleCb:
    maxConnections: 100
    httpMaxRequests: 1000
    httpMaxRequestsPerConnection: 10
    httpConsecutiveErrors: 7
    sleepWindow: 15m
    httpDetectionInterval: 5m
```

Resilience features

- ❖ Timeouts
- ❖ Retries with timeout budget
- ❖ Circuit breakers
- ❖ Health checks
- ❖ AZ-aware load balancing w/ automatic failover
- ❖ Control connection pool size and request load
- ❖ Systematic fault injection

Service Mesh 37 © Copyright IBM Corporation 2016-18

IBM Training **IBM**

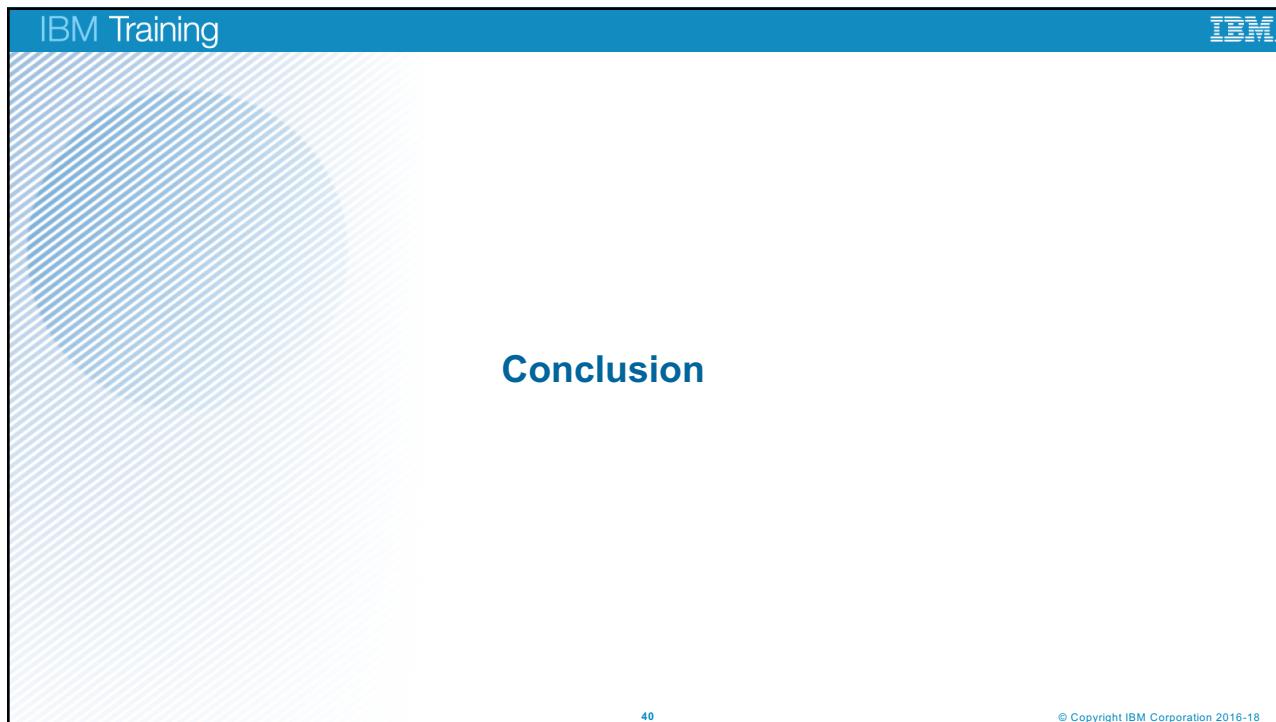
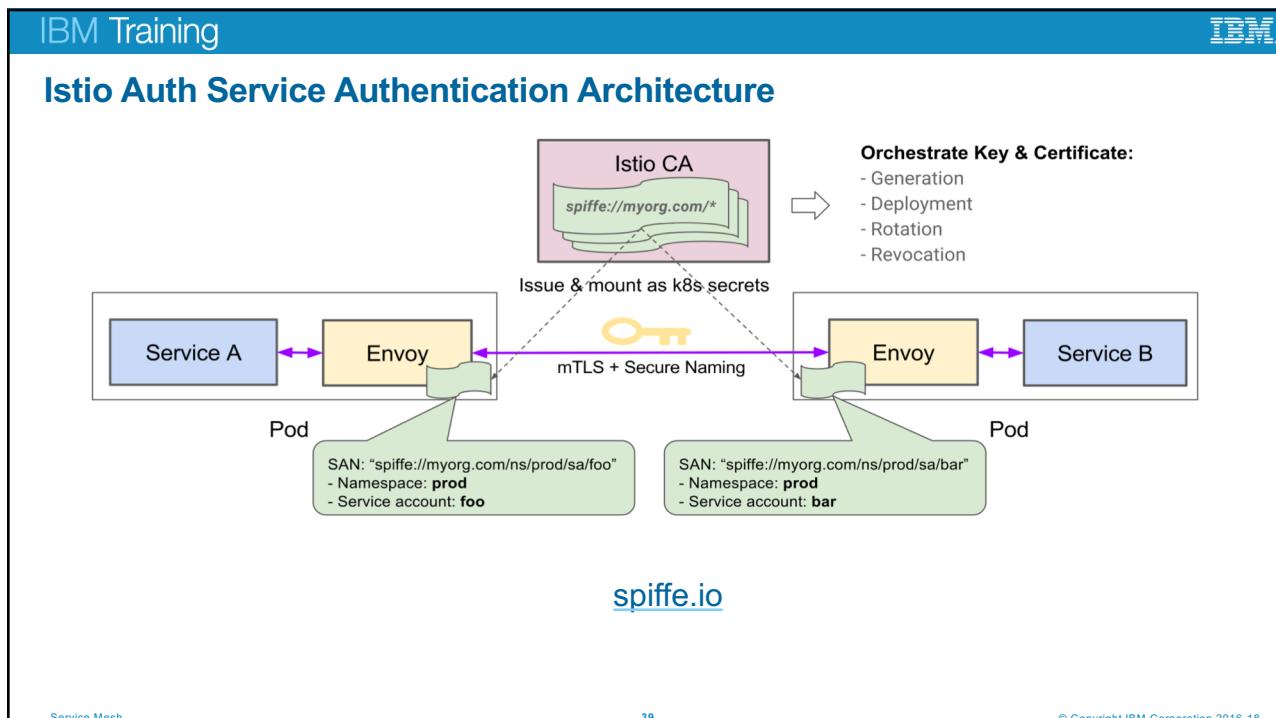
Resiliency Testing

- Systematic fault injection to identify weaknesses in failure recovery policies
 - HTTP/gRPC error codes
 - Delay injection

```

graph LR
    subgraph Service_A [Service A]
        direction TB
        EnvoyA[Envoy] --- svcA[svcA]
    end
    subgraph Service_B [Service B]
        direction TB
        EnvoyB[Envoy] --- svcB[svcB]
    end
    subgraph Service_C [Service C]
        direction TB
        EnvoyC[Envoy] --- svcC[svcC]
    end
    EnvoyA -- "Timeout: 100ms  
Retries: 3  
300ms" --> EnvoyB
    EnvoyB -- "Timeout: 200ms  
Retries: 2  
400ms" --> EnvoyC
  
```

Service Mesh 38 © Copyright IBM Corporation 2016-18



IBM Training 

Conclusion

- Why service mesh?
- Common functions of a service mesh
- Services mesh implementations

Service Mesh 41 © Copyright IBM Corporation 2016-18

IBM Training 

Next

- Deploying cloud applications
 - Docker and Kubernetes in IBM Cloud Container Service
 - Cloud Foundry in IBM Cloud
 - Swarm in Docker Enterprise Edition for IBM Cloud
 - OpenWhisk in IBM Cloud Functions
- IBM Cloud capabilities
 - IBM Cloud architecture
 - WebSphere Liberty for microservices using Java, Java EE, and MicroProfile
 - Cloud data services
 - Integration with systems of record
- DevOps
 - Continuous delivery of microservices

Service Mesh 42 © Copyright IBM Corporation 2016-18

IBM Training 

Resources

- IBM, Google and Lyft give microservices a ride on the Istio Service Mesh
 - <https://developer.ibm.com/dwblog/2017/istio/>
- Istio
 - <https://istio.io>
 - <https://github.com/istio>
- developerWorks Container Orchestration Journeys
 - <http://ibm.biz/kube-journeys>
 - <https://developer.ibm.com/code/journey/manage-microservices-traffic-using-istio>
- Microservices
 - <http://kasunpanorama.blogspot.com/2015/11/microservices-in-practice.html>
- Service discovery
 - <https://dzone.com/articles/service-discovery-in-a-microservices-architecture>
- Service mesh
 - <https://blog.buoyant.io/2017/04/25/whats-a-service-mesh-and-why-do-i-need-one/>
 - <https://content.pivotal.io/blog/pivotal-and-istio-advancing-the-ecosystem-for-microservices-in-the-enterprise>
- Sidecars
 - <https://blog.buoyant.io/2016/06/17/small-memory-jvm-techniques-for-microservice-sidecars/>

Service Mesh 43 © Copyright IBM Corporation 2016-18

IBM Training 

Microservices traffic management and telemetry using Istio

Optimizing and visualizing the network in your
microservice application

© Copyright IBM Corporation 2018
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

IBM Training 

Topics

- Why Istio?
- Architecture
- Traffic management
 - Rule configuration
 - Traffic routing
 - Bookinfo example
- Policies and telemetry
- Debugging Istio

Microservices traffic management and telemetry using Istio 2 © Copyright IBM Corporation 2018

IBM Training

Why Istio

3 © Copyright IBM Corporation 2018

IBM Training

What is a service mesh?

- Service mesh
 - The network of microservices that make up an application and the interactions between them
 - A network of services, not bytes
- Why do we need a service mesh?
 - As a mesh gets bigger, it's harder to understand and manage
 - Manage the mesh as a whole, not as its parts
 - Provide behavioral insights and operational control
- New requirements
 - Quality of service:* discovery, load balancing, failure recovery, rate limiting, metrics, and monitoring
 - Continuous deployment:* canary rollouts, A/B testing, staged rollouts with %-based traffic splits, etc.
 - Security:* end-to-end authentication and access control

© Copyright IBM Corporation 2018

IBM Training 

What is Istio?

- An open platform to connect, secure, control and observe microservices
- A language-neutral service mesh for modern cloud-native applications
- Istio provides an easy way to create a network of deployed services
 - Load balancing, service-to-service authentication, monitoring, and more
 - Does not require any changes to the service code
- Istio uses a special service proxy throughout your environment
 - Intercepts all network communication between microservices
 - Configured and managed using Istio's control plane functionality

Microservices traffic management and telemetry using Istio

5

© Copyright IBM Corporation 2018

IBM Training 

Why use Istio?

Intelligent Routing and Load Balancing
Control traffic between services with dynamic route configuration, conduct A/B tests, release canaries, and gradually upgrade versions using red/black deployments

Resilience Across Languages and Platforms
Increase reliability by shielding applications from flaky networks and cascading failures in adverse conditions

Fleet Wide Policy Enforcement
Apply organizational policy to the interaction between services, ensure access policies are enforced and resources are fairly distributed among consumers

In-Depth Telemetry and Reporting
Understand the dependencies between services, the nature and flow of traffic between them and quickly identify issues with distributed tracing

Microservices traffic management and telemetry using Istio

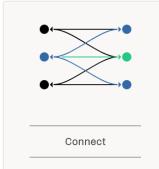
6

© Copyright IBM Corporation 2018

IBM Training IBM

Core capabilities

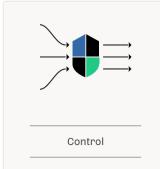
- Traffic Management**
 - Intelligently control the flow of traffic and API calls between services
 - Make calls more reliable and make the network more robust in the face of adverse conditions
- Service Identity and Security**
 - Provide services in the mesh with a verifiable identity
 - Provide the ability to protect service traffic as it flows over networks of varying degrees of trustability
- Policy Enforcement**
 - Apply organizational policy to the interaction between services
 - Ensure access policies are enforced and that resources are fairly distributed among consumers
 - Make policy changes by changing the configuration of the mesh, not the application code
- Telemetry**
 - Gain understanding of the dependencies between services and the nature and flow of traffic between them
 - Provide the ability to quickly identify issues



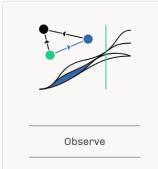
Connect



Secure



Control



Observe

Microservices traffic management and telemetry using Istio

7

© Copyright IBM Corporation 2018

IBM Training IBM

Istio Architecture

8

© Copyright IBM Corporation 2018

IBM Training

Istio advantages

- Language-neutral**
 - Istio implementation is performed outside of the application container
 - Does not matter what programming language the application is written in
- Transparent/Non-invasive**
 - Istio intercepts network traffic outside of the application
 - Masking the application network connection which is not detected by the application

Microservices traffic management and telemetry using Istio

9

© Copyright IBM Corporation 2018

IBM Training

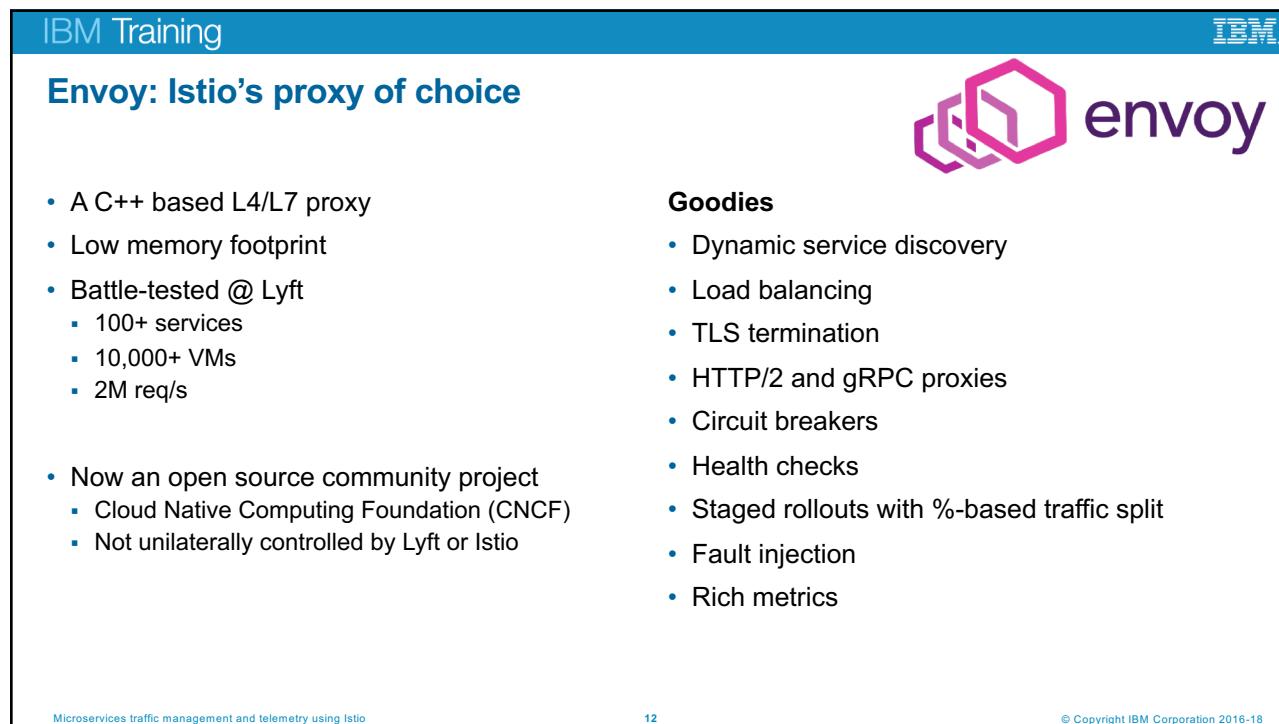
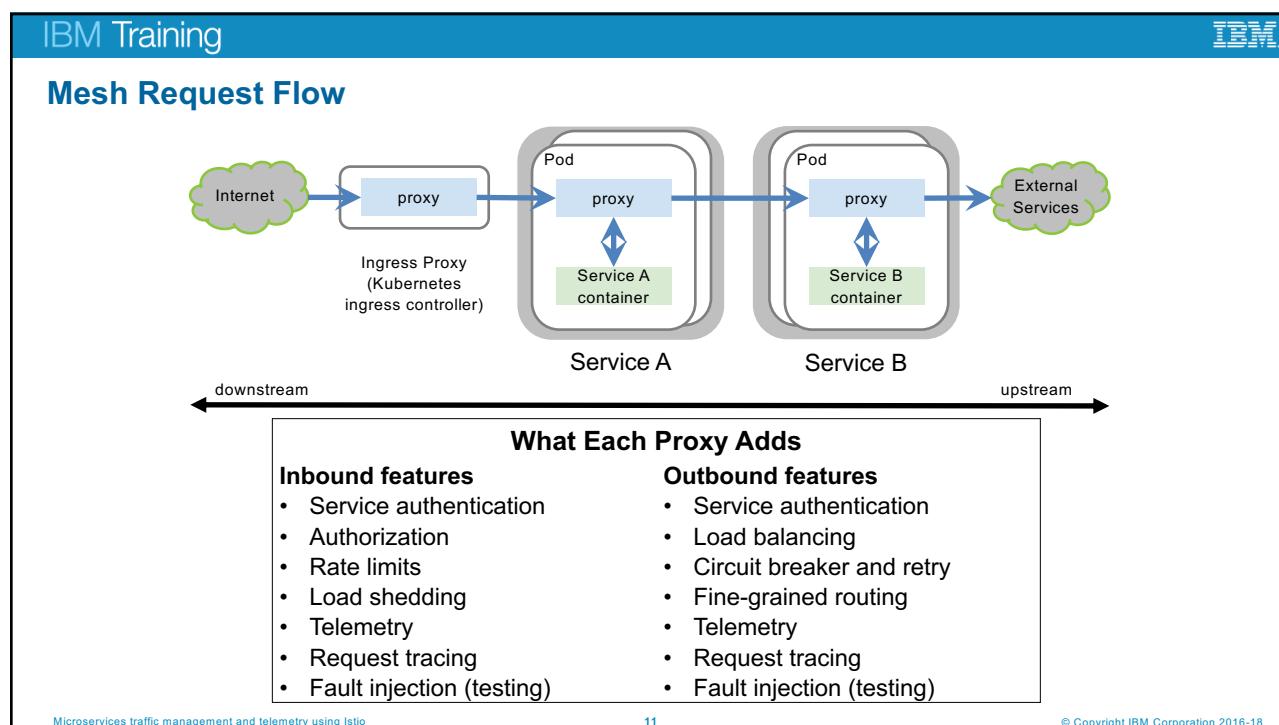
Managing services: Proxy and sidecar

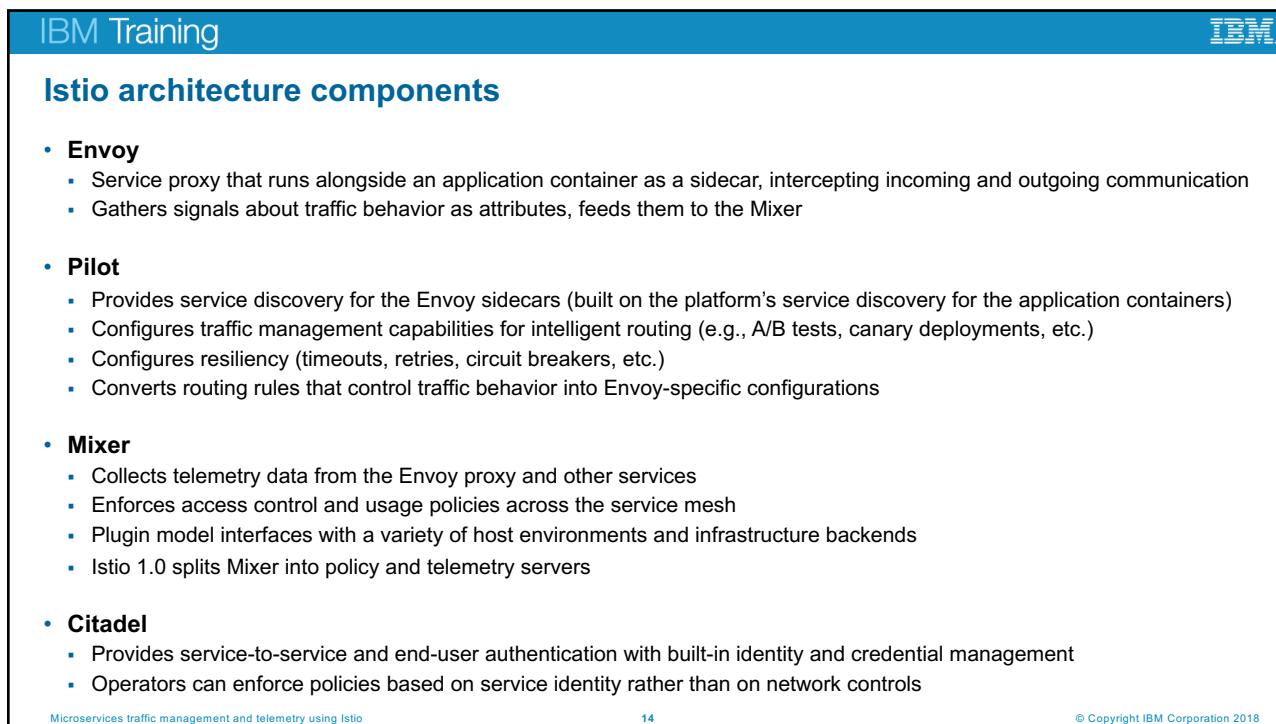
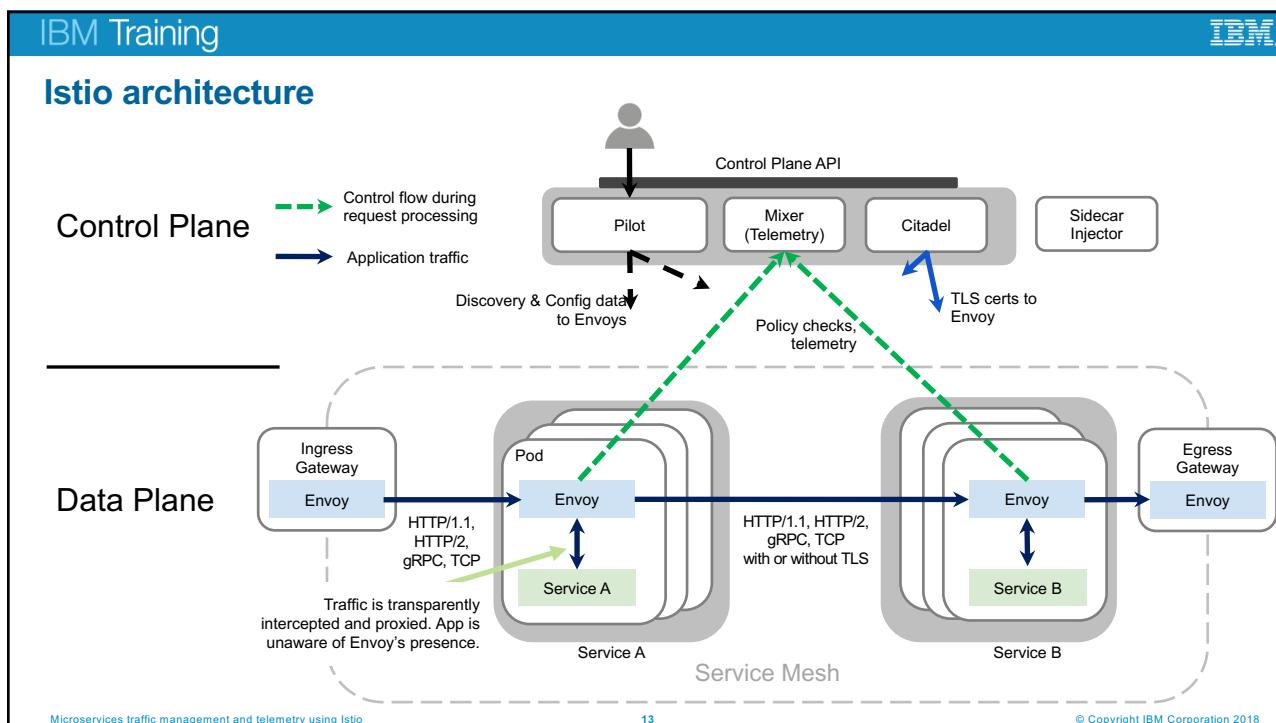
- A service mesh manages the interactions between services
 - It uses *proxies*, usually deployed as *sidecars* in *pods*
- Proxy** – Controls access to another object
 - A service mesh uses proxies between services (and also clients)
 - Enables the mesh to manage interactions
- Pod** – Kubernetes deploys a set of containers as a pod
 - All of the containers in a pod are hosted on the same node
 - A pod can manage the network access to the containers in the pod
 - A pod typically hosts a single container, but can host multiple as a unit
- Sidecar** – Adds behavior to a container without changing it
 - Sidecar and service behave as a single enhanced unit
 - Pod hosts them as a single unit
- Mesh Member** – Controlled access to the service
 - Deploys the proxy as a *sidecar* in the same *pod*
 - Configures the pod so that all network access goes through the proxy

Microservices traffic management and telemetry using Istio

10

© Copyright IBM Corporation 2016-18





IBM Training IBM

Additional Istio architecture components

- **Ingress gateway**
 - Handles incoming traffic to the mesh from external clients
- **Egress gateway**
 - Handles outgoing traffic from the mesh to external services
- **Sidecar injector**
 - Functionality to add Envoy proxies to deployed application containers

Microservices traffic management and telemetry using Istio 15 © Copyright IBM Corporation 2018

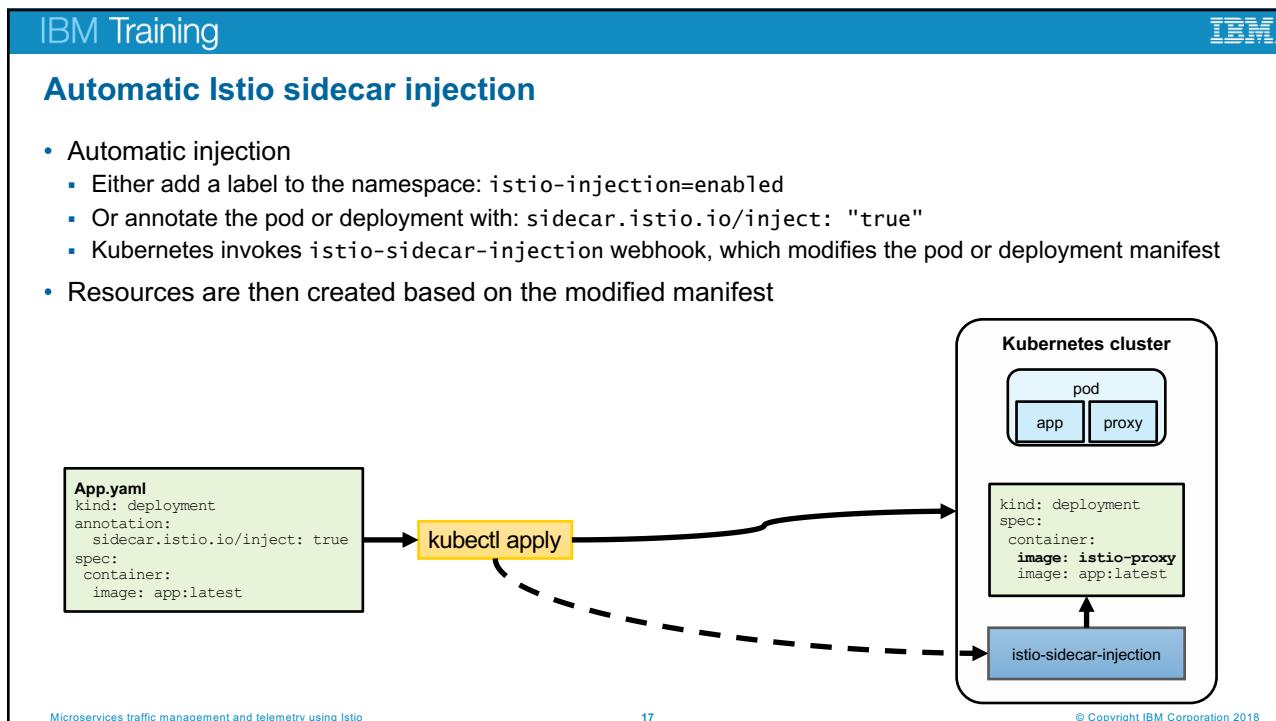
IBM Training IBM

Manual Istio sidecar injection

- To configure the service mesh, Istio must add its Envoy proxy to each of the services it will manage
 - This addition is performed by configuring Kubernetes using YAML file injection
- Manual injection
 - Run the command `istioctl kube-inject` on the YAML file
 - Resources are then created based on the modified manifest
 - `istioctl kube-inject -f app.yaml | kubectl apply -f -`

```
graph LR; A[App.yaml<br/>kind: deployment<br/>spec:<br/>container:<br/>image: app:latest] --> B[istioctl kube-inject]; B --> C["kind: deployment<br/>spec:<br/>container:<br/>image: istio-proxy<br/>image: app:latest"]; C --> D[kubectl apply]; D --> E["Kubernetes cluster<br/>pod<br/>app proxy"]; style A fill:#e0f2e0,stroke:#0070C0,stroke-width:1px; style B fill:#FFF,stroke:#0070C0,stroke-width:2px; style C fill:#FFF,stroke:#0070C0,stroke-width:2px; style D fill:#FFF,stroke:#0070C0,stroke-width:2px; style E fill:#FFF,stroke:#0070C0,stroke-width:1px;
```

Microservices traffic management and telemetry using Istio 16 © Copyright IBM Corporation 2018



Istio interfaces

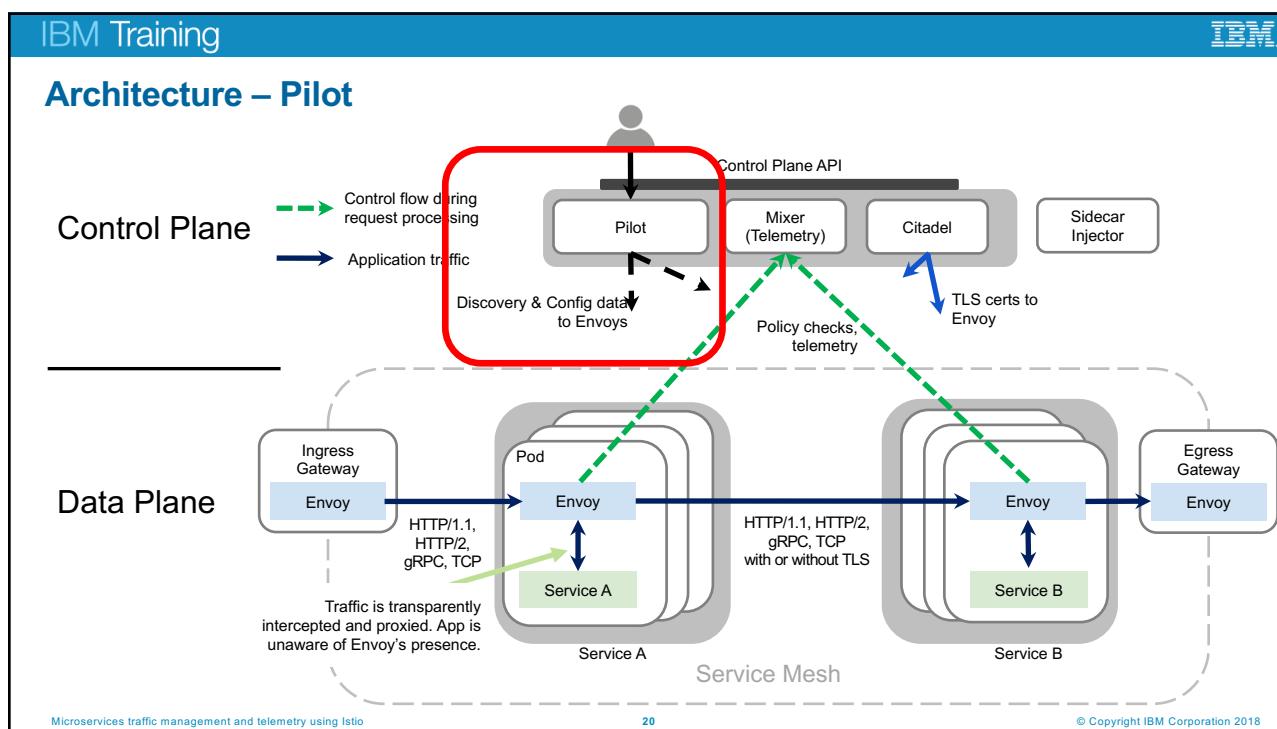
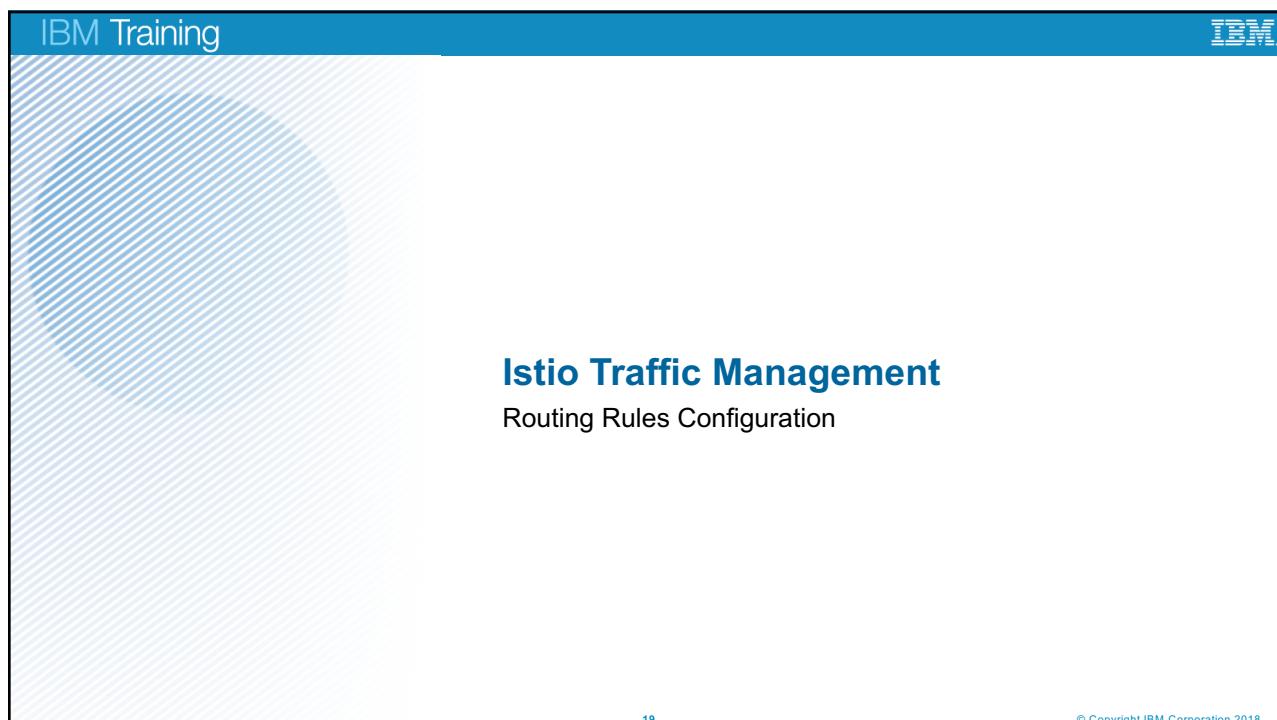
- The `istioctl` CLI
- Istio Grafana dashboards
- Istio Jaeger tracing

Istio Mesh Dashboard: Shows global request volume (1.1 ops), success rate (94.8%), and HTTP/GRPC workloads for various services like reviews-v2 and ratings-v1.

Service	Workload	Requests	P50 Latency	P90 Latency	P99 Latency	Success Rate
reviews.default.svc.cluster.local	reviews-v2.default	0.02 ops	375 ms	475 ms	498 ms	100.0%
reviews.default.svc.cluster.local	reviews-v2.default	0.02 ops	361 ms	4.48 s	4.95 s	99.01%
reviews.default.svc.cluster.local	reviews-v1.default	0.04 ops	105 ms	221 ms	247 ms	100.0%
ratings.default.svc.cluster.local	ratings-v1.default	0.08 ops	3 ms	5 ms	815 ms	100.0%
productpage.default.svc.cluster.local	productpage-v1.default	0.09 ops	303 ms	6.95 s	9.70 s	71.43%

Jaeger UI: Shows a trace for the `istio-ingressgateway` from `productpage.default.svc.cluster.local:9080/productpage`. The trace spans from July 29, 2018, 4:31 PM, with a duration of 131.08ms. It details the flow through the `productpage`, `details`, `reviews`, and `ratings` services.

Microservices traffic management and telemetry using Istio 18 © Copyright IBM Corporation 2018



IBM Training **IBM**

Traffic Management

- Pilot**
 - Define the rules you want traffic to follow, rather than which service instances should receive traffic
 - Manages and configures all the Envoy proxy instances deployed in a particular Istio service mesh
 - Maintains a canonical model of all the services in the mesh
 - Uses the model to let Envoy instances know about the other Envoy instances in the mesh via its discovery service
 - Manages the lifecycle of each Envoy instance
- Envoy**
 - Each instance maintains load balancing information (from Pilot and periodic health-checks)
 - Intelligently distributes traffic between destination instances while following its specified routing rules

The diagram illustrates the Istio architecture for traffic management. At the top, there's a stack of components: 'Platform Adapter', 'Abstract Model', and 'Envoy API'. Below this stack is a large box labeled 'Pilot'. To the left of the Pilot is a user icon connected to the 'Rules API' and 'Envoy API'. Below the Pilot, a box labeled 'Service discovery and traffic rules' has arrows pointing to four 'Proxy' boxes. Above the Pilot, there are three platform adapter boxes: 'Kubernetes', 'Mesos', and 'CloudFoundry', followed by an ellipsis (...).

IBM Training **IBM**

Traffic Routing Terms

Applications address only the destination service without knowledge of individual service subsets or other mappings. The actual choice of the version is determined by Envoy, enabling the application code to decouple itself from the evolution of dependent services.

- Service**
 - A unit of application behavior bound to a unique name in a service registry (a.k.a. downstream application)
 - Services consist of multiple network endpoints implemented by workload instances running on pods, containers, VMs, etc.
- Service versions (a.k.a. subsets)**
 - Different implementation iterations of the same logical application behavior
 - Instances implement different versions (v1, v2) or are deployed in different environments (dev, stage, prod)
- Source**
 - The client calling a service
- Host**
 - The address used by a client (source) when attempting to connect to a service

IBM Training **IBM**

Routing: Traffic Splitting and Traffic Steering

Traffic splitting decoupled from infrastructure scaling - proportion of traffic routed to a version is independent of number of instances supporting the version

Content-based traffic steering - The content of a request can be used to determine the destination of a request

- **Traffic splitting**
 - Routing not based on the request content
 - Staged rollouts with %-based traffic splits
- **Traffic steering**
 - Routing based on the contents of the request
- Either can be used for:
 - Canary rollouts
 - A/B testing
 - Red/black deployment

Microservices traffic management and telemetry using Istio 23 © Copyright IBM Corporation 2018

IBM Training **IBM**

Rule Configuration

Traffic routing—between a sources and services—is defined as rules

Traffic routing rules are configured using four types of definitions

- **VirtualService**
 - Defines the rules that control how requests for a service are routed to service subsets
- **DestinationRule**
 - Configures the policies to be applied to a request based on the selected service subset
- **ServiceEntry**
 - Commonly used to enable requests to services outside of an Istio service mesh
- **Gateway**
 - Configures a load balancer for HTTP/TCP traffic
 - Most commonly operates at the edge of the mesh to enable ingress or egress traffic for an application

Microservices traffic management and telemetry using Istio 24 © Copyright IBM Corporation 2018

IBM Training

VirtualService definition

- Rules for routing a service to its service subset
 - Configures how requests should be split across or steered to service subsets
 - Acts as a load-balanced destination from a gateway
- Used to configure
 - Traffic splitting
 - Traffic steering
- For HTTP connections
 - Timeouts and retries
 - Fault injection
 - URL rewriting

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
    - reviews
  http:
    - match:
        - headers:
            end-user:
              exact: jason
    route:
      - destination:
          host: reviews
          subset: v2
      - route:
          - destination:
              host: reviews
              subset: v3
```

Microservices traffic management and telemetry using Istio

25

© Copyright IBM Corporation 2018

IBM Training

VirtualService illustration

- Request destinations
 - Hosts that sources can invoke
 - Attributes: **hosts** and **gateways**
- Route destinations
 - Subset of the destination
 - Attribute: **route** and **destination**
- Protocol selection
 - How to connect to the destination subset
 - Attribute: **http, tcp, tls**
- Routing rules
 - Additional routing attributes to be applied for the route destinations
 - Attributes: **weight** and **match**
- HTTP traffic policy
 - Protocol-specific connection quality of service
 - Attributes: **timeout, retries, fault, rewrite, and redirect**

The diagram shows two input boxes labeled "Request Destination reviews" and "Request Destination reviews.bookinfo.com". These boxes point to a central "Protocol Selection" box containing "http, tcp, tls". From this selection, three paths lead to three separate "Route Destination" boxes: "reviews v1" (containing "Routing Rules match, weight" and "HTTP Traffic Policy timeout, retries, fault, rewrite, redirect"), "reviews-test v2" (containing "Routing Rules"), and "reviews canary" (containing "Routing Rules").

Microservices traffic management and telemetry using Istio

26

© Copyright IBM Corporation 2018

IBM Training

VIRTUALSERVICE – MODIFYING A SERVICE

How a Kubernetes service becomes an Istio virtual service

- Proxy matches request to virtual service
 - Proxy has a table of virtual services
 - Proxy selects the one whose host matches the request
- Virtual service selects the service
 - The virtual service specifies a route to a destination host
 - Proxy finds the service matching the destination host
- Proxy uses service to invoke destination
 - Proxy uses the service to find host pods
 - Proxy load balances across the host pods

Pod: productpage
xhttp.open("GET", "http://reviews/") → Service reviews → Pod: reviews

Pod: Istio proxy → productpage
xhttp.open("GET", "http://reviews/") → Host: reviews | Virtual service: reviews-vs → Service reviews → Pod: Istio proxy → reviews

Host	Virtual service
details	details-vs
productPage	productPage-vs
ratings	ratings-vs
reviews	reviews-vs

lookup (green arrow) → invocation (blue arrow)

```

apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews-vs
spec:
  hosts:
  - reviews
  http:
  - route:
    - destination:
        host: reviews
  
```

© Copyright IBM Corporation 2018

IBM Training

DESTINATIONRULE DEFINITION

- Defines subsets for a virtual service
- Policies for how to identify and invoke a service subset
 - All for the same host
 - TrafficPolicy can be specified for:
 - all subsets
 - specifically for a subset
 - certain ports for all subsets
 - certain ports within a subset
- Used to configure
 - Load balancer settings
 - Connection pool settings
 - Outlier detection
 - TLS settings

```

apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: httpbin
spec:
  host: httpbin
  trafficPolicy:
    connectionPool:
      tcp:
        maxConnections: 10
      http:
        http1MaxPendingRequests: 10
        maxRequestsPerConnection: 5
    outlierDetection:
      consecutiveErrors: 1
      interval: 1s
      baseEjectionTime: 3m
      maxEjectionPercent: 100
  subsets:
  - name: main
    labels:
      version: v23
  - name: new
    labels:
      version: v24
  - name: beta
    labels:
      usage: test
  
```

© Copyright IBM Corporation 2018

IBM Training **IBM**

DestinationRule illustration

- Destination host
 - Host that route destinations can select
 - Attribute: **host**
- Host subset
 - Identify a subset of service endpoints
 - Attribute: **labels**
- Traffic policy
 - Influence expected quality of service for destinations
 - Attributes: **trafficPolicy**
 - **loadBalancer**, **connectionPool**, **outlierDetection**, and **TLS**
 - Traffic policy can be applied to a port
 - Attribute: **portLevelSettings**

Microservices traffic management and telemetry using Istio 29 © Copyright IBM Corporation 2018

IBM Training **IBM**

Gateway definition

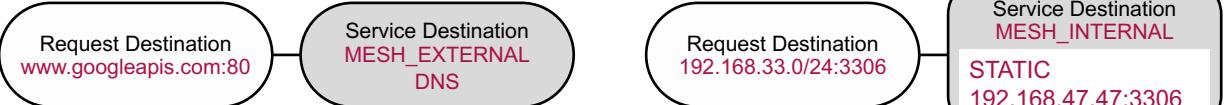
- Load balancer operating at the edge of the mesh receiving incoming HTTP/TCP connections
 - Configures ports to expose externally
 - Maps each exposed port to a request destination
 - Each gateway can have one or more Virtual Services that defines these request destinations
- Gateway is attached to Istio Ingress Controller (which can be the Kubernetes Ingress Controller)
- Hosting pod
 - Label of pod that will host the gateway
 - Attribute: **selector**
- Request destinations
 - Ports for the gateway to expose and hosts for the corresponding services
 - Attributes: **servers**
 - **port/hosts/tls** triplet

Microservices traffic management and telemetry using Istio 30 © Copyright IBM Corporation 2018

IBM Training IBM

ServiceEntry definition

- Enables adding additional entries into Istio's internal service registry
 - Makes them auto-discoverable within the mesh
 - Route to these services with traffic policies
- Used to configure additional connections
 - Connections to off-premises service endpoints
 - E.g. web APIs
 - Connections to on-premises service endpoints that are not part of the platform's service registry
 - E.g. VMs (instead of containers) hosting services



Request Destination
`www.googleapis.com:80`

Service Destination
`MESH_EXTERNAL_DNS`

Request Destination
`192.168.33.0/24:3306`

Service Destination
`MESH_INTERNAL`
`STATIC`
`192.168.47.47:3306`

Microservices traffic management and telemetry using Istio

31

© Copyright IBM Corporation 2018

IBM Training IBM

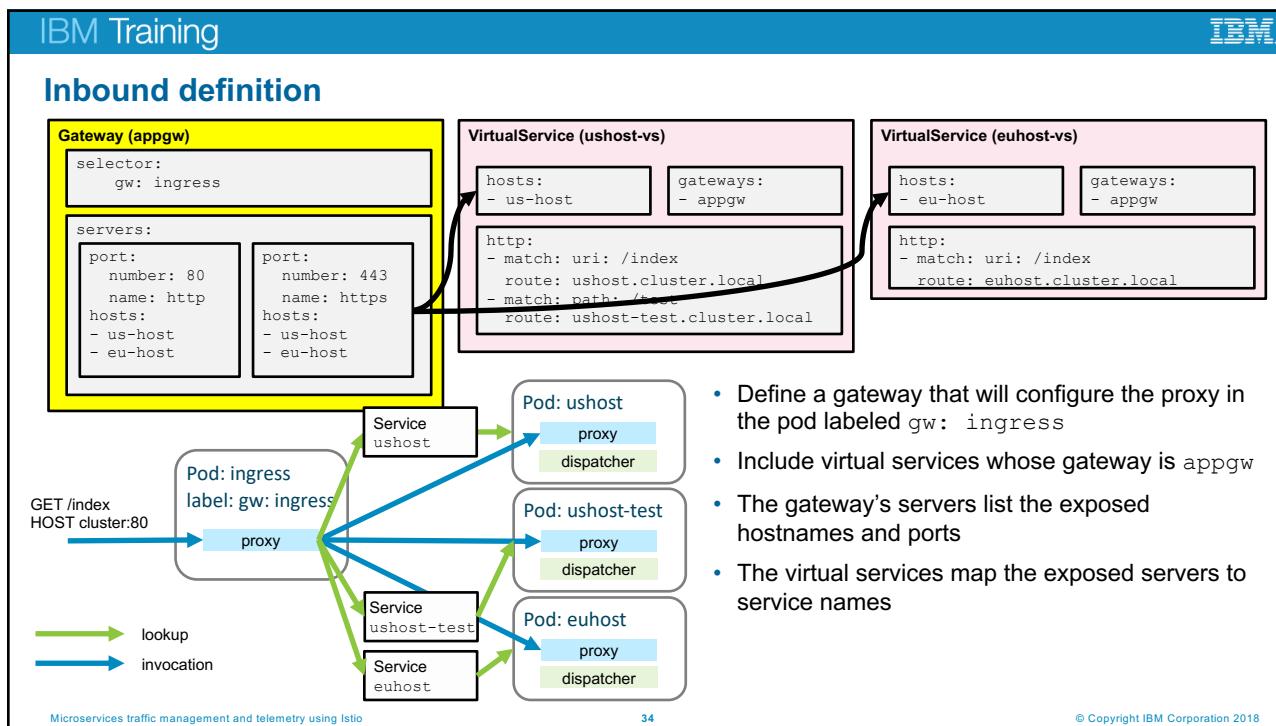
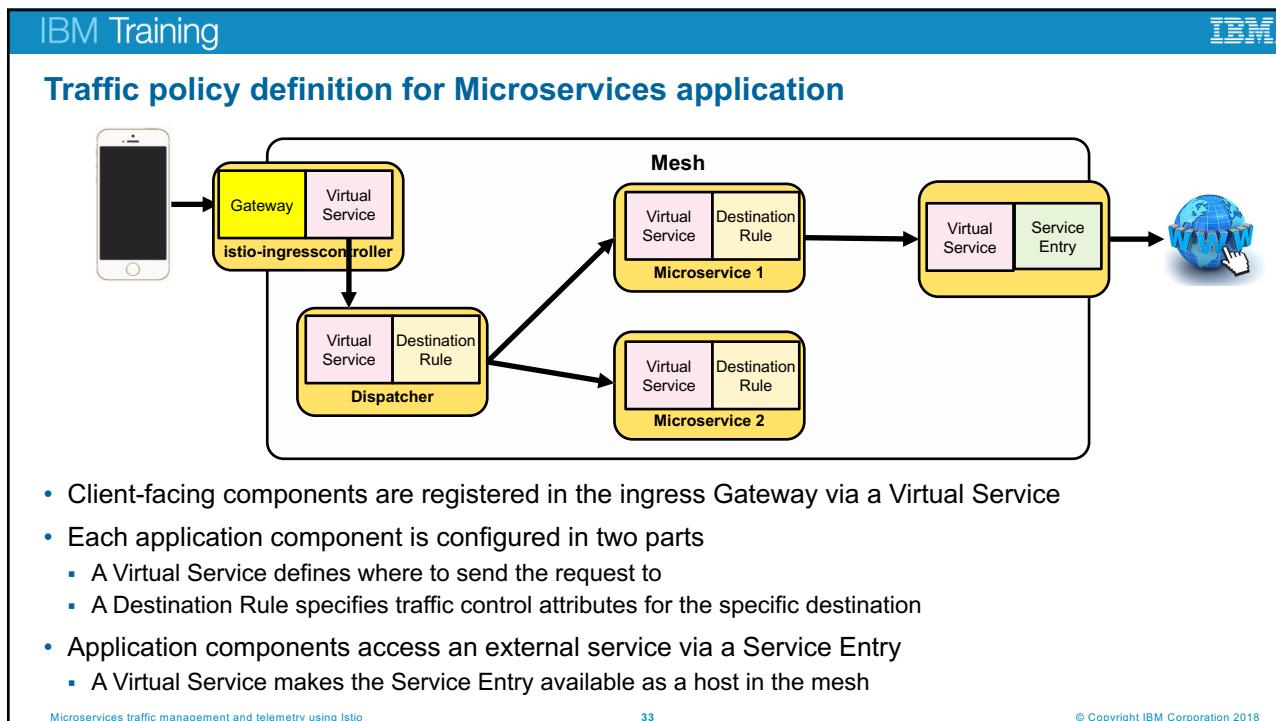
Istio Traffic Management

Traffic Routing



32

© Copyright IBM Corporation 2018



IBM Training

Traffic routing definition for in-mesh destination

The diagram illustrates the traffic flow and configuration for an in-mesh destination. It shows two main components: a **VirtualService (appone-vs)** and a **DestinationRule (appone-dr)**.

VirtualService (appone-vs) Configuration:

```

hosts:
- appone
gateways:
- mesh
http:
- match: header(source): bff
  route: appone
  subset: test
- route: appone
  subset: ver1

```

DestinationRule (appone-dr) Configuration:

```

host: appone
subsets:
- name: ver1
  labels: ver: v1
- name: test
  labels: ver: v3
trafficPolicy:
loadBalancer
connectionPool
outlierDetection
tls
portLevelSettings

```

Pod Structure:

- Pod: bff**: Contains a **proxy** and a **dispatcher**. A request **GET /app** is sent to the proxy.
- Service appone**: Contains a **proxy**.
- Pod: appone-v1** and **Pod: appone-v3**: Both contain a **proxy** and are labeled **ver: v1** and **ver: v3** respectively.

Flow and Annotations:

- A dashed blue arrow points from the **proxy** in the **bff** pod to the **proxy** in the **appone** service.
- A solid green arrow labeled **lookup** points from the **appone** service to the **proxy** in the **appone-v1** pod.
- A solid blue arrow labeled **invocation** points from the **proxy** in the **appone-v1** pod to its own **proxy**.

Annotations:

- The gateway setting in the virtual service, **mesh**, specifies that it is used for destination routing.
- Destination rule (if exists) maps the subsets to pod labels.
- The source's proxy performs load balancing and outlier detection.
- Both the source's proxy and the host's proxy performs connection pooling.

Microservices traffic management and telemetry using Istio

35 © Copyright IBM Corporation 2018

IBM Training

Outbound routing

The diagram illustrates the configuration and flow for outbound routing. It shows a **VirtualService (google-vs)** and a **ServiceEntry (google-se)**.

VirtualService (google-vs) Configuration:

```

hosts:
- google.com
gateways:
- mesh
tls:
  match:
    sni_hosts: google.com
  route: google.com
  port: 443
  weight:100

```

ServiceEntry (google-se) Configuration:

```

hosts:
- google.com
ports:
- [80, 443]
location: MESH_EXTERNAL
resolution: DNS

```

Pod Structure:

- pod: apptwo**: Contains a **proxy**.

Flow and Annotations:

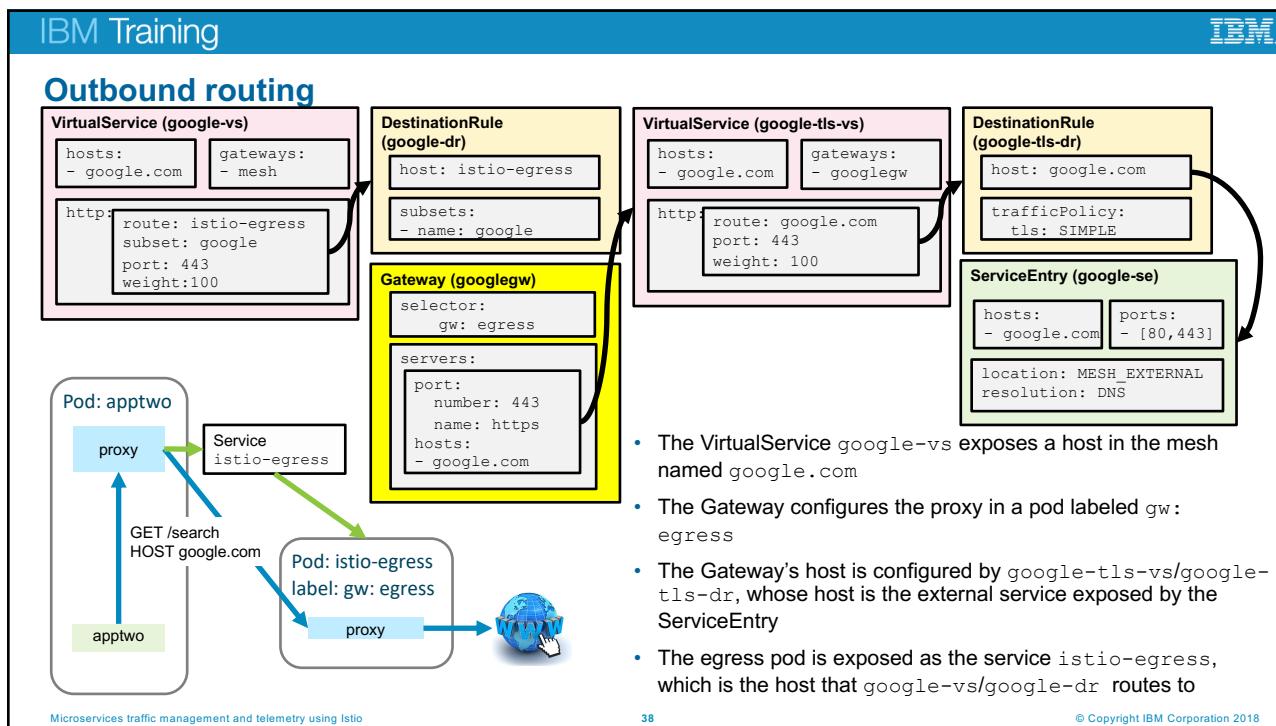
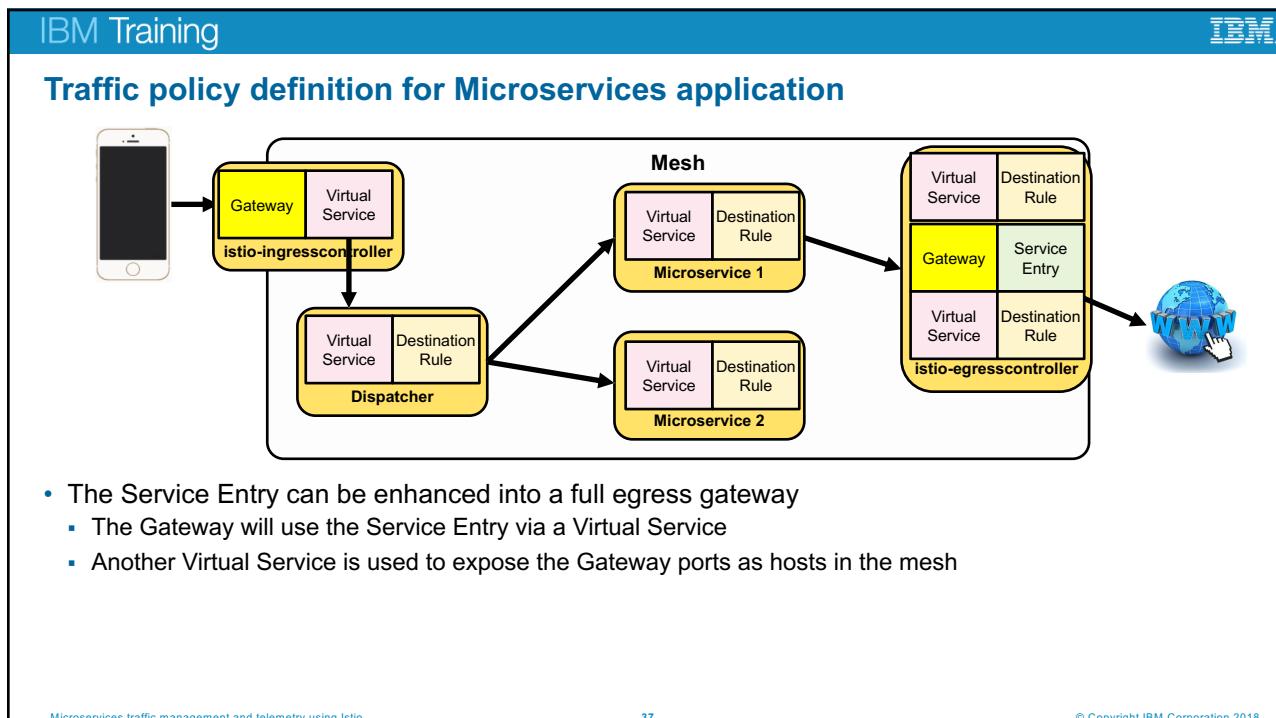
- A solid blue arrow points from the **proxy** in the **apptwo** pod to a globe icon representing the internet.
- A solid blue arrow labeled **GET /search** points from the **proxy** to the globe.

Annotations:

- The virtual service's host is **google.com**, which is what the client is invoking.
- The virtual service facilitates TLS termination using the TLS route rule.
- There is no service in the mesh named **google.com**.
- The service entry makes **google.com** available as a known host that's available externally via DNS.
- The route destination **google.com** therefore routes not to a service but to the external destination.
- Optionally (not shown): The virtual service can route to a destination rule that configures connectionPool, outlierDetection, etc.
- Optionally (not shown): An egress gateway can be added using a gateway and virtual service, which isolates outbound traffic and enables it to be measured.

Microservices traffic management and telemetry using Istio

36 © Copyright IBM Corporation 2018



IBM Training 

Traffic policy

- Each DestinationRule can configure one or more traffic policies
 - Influence expected quality of service for destinations
 - Better fault tolerance for using the service
 - Attribute: **trafficPolicy**
- Three traffic policy configurations affect resiliency
 - Load balancer (**loadBalancer**)
 - Connection pool (**connectionPool**)
 - Outlier detection (**outlierDetection**)
- While traffic policy configures resiliency, the service proxy implements it
 - Envoy implements load balancing, configured by DestinationRule's load balancing settings
 - Envoy implements circuit breaking, configured by DestinationRule's connection pool settings
 - Envoy implements outlier detection, configured by DestinationRule's outlier detection settings

Microservices traffic management and telemetry using Istio 39 © Copyright IBM Corporation 2018

IBM Training 

Traffic policy: Load balancing

- Istio configures load balancing logic
 - Load balancing is performed against the pool of pods that are targeted by the DestinationRule subset
 - The default is simple ROUND_ROBIN load balancing
- Simple load balancing is implemented by Envoy
 - Envoy has more load balancer settings than Istio supports
- Istio simple load balancer settings
 - ROUND_ROBIN
 - LEAST_CONN
 - RANDOM
 - PASSTHROUGH
- Load balancer can be simple or use session affinity
- Session affinity load balancer
 - Host is determined by a hash (**consistentHash**)
- Elements in the session affinity hash
 - httpHeaderName
 - httpCookie
 - useSourceIp
 - minimumRingSize (hash ring, with each node assigned to a number in a ring)

Microservices traffic management and telemetry using Istio 40 © Copyright IBM Corporation 2018

IBM Training 

Envoy health checking

- Envoy has passive and active health checking
 - Ideally, active health checking detects and ejects a problematic host before a service request fails
- Active health checking
 - Envoy sends tests to the host, independent of service requests
- Passive health checking
 - Envoy notes when service requests fail
- Passive health checking is configured by outlier detection

Microservices traffic management and telemetry using Istio 41 © Copyright IBM Corporation 2018

IBM Training 

Traffic policy: Connection pool and Outlier detection

Connection Pools <ul style="list-style-type: none">• Connection pool<ul style="list-style-type: none">▪ Limit the number of requests to a destination▪ Avoids overloading the destination▪ Implemented by Envoy circuit breaker• Connection pool TCP settings<ul style="list-style-type: none">▪ maxConnections▪ connectTimeout• Connection pool HTTP settings<ul style="list-style-type: none">▪ http1MaxPendingRequests▪ http2MaxRequests▪ maxRequestsPerConnection▪ maxRetries	Outlier Detection <ul style="list-style-type: none">• Outlier detection<ul style="list-style-type: none">▪ Removes hosts from the load balancer set that are not performing properly▪ Implemented by Envoy outlier detection• Outlier detection settings<ul style="list-style-type: none">▪ consecutiveErrors▪ interval▪ baseEjectionTime▪ maxEjectionPercent
--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

Microservices traffic management and telemetry using Istio 42 © Copyright IBM Corporation 2018

IBM Training

Istio Traffic Management

Bookinfo Example

43 © Copyright IBM Corporation 2018

IBM Training

The Bookinfo application

The diagram illustrates the Bookinfo application architecture. A Python 'Product page' service receives 'Requests' and sends them to three Java 'Reviews' microservices (v1, v2, v3) and a Node.js 'Ratings' microservice. The reviews are displayed with star ratings: v1 has 5 stars, v2 has 4 stars, and v3 has 3 stars. The Ratings microservice also receives requests from the Product page. To the right, a screenshot of a web browser shows a sample book page for 'The Comedy of Errors' with 'Book Details' and 'Book Reviews' sections, including reviews with 5 and 3 stars.

- Polyglot application
 - Three microservices
 - Single front end entry point
- Three versions of Reviews app
 - Allows various tests to be conducted for routing and reliability features

BookInfo Sample

The Comedy of Errors

Summary: Wikipedia Summary: The Comedy of Errors is one of William Shakespeare's early plays. It is his shortest and one of his most farcical comedies, with a major part of the humour coming from slapstick and mistaken identity, in addition to puns and word play.

Book Details

Type: paperback
Pages: 200
Publisher: Publisher A
Language: English
ISBN-10: 1234567890
ISBN-13: 123-1234567890

Book Reviews

An extremely entertaining play by Shakespeare. The slapstick humour is refreshing!

Review 1: ★★★★☆

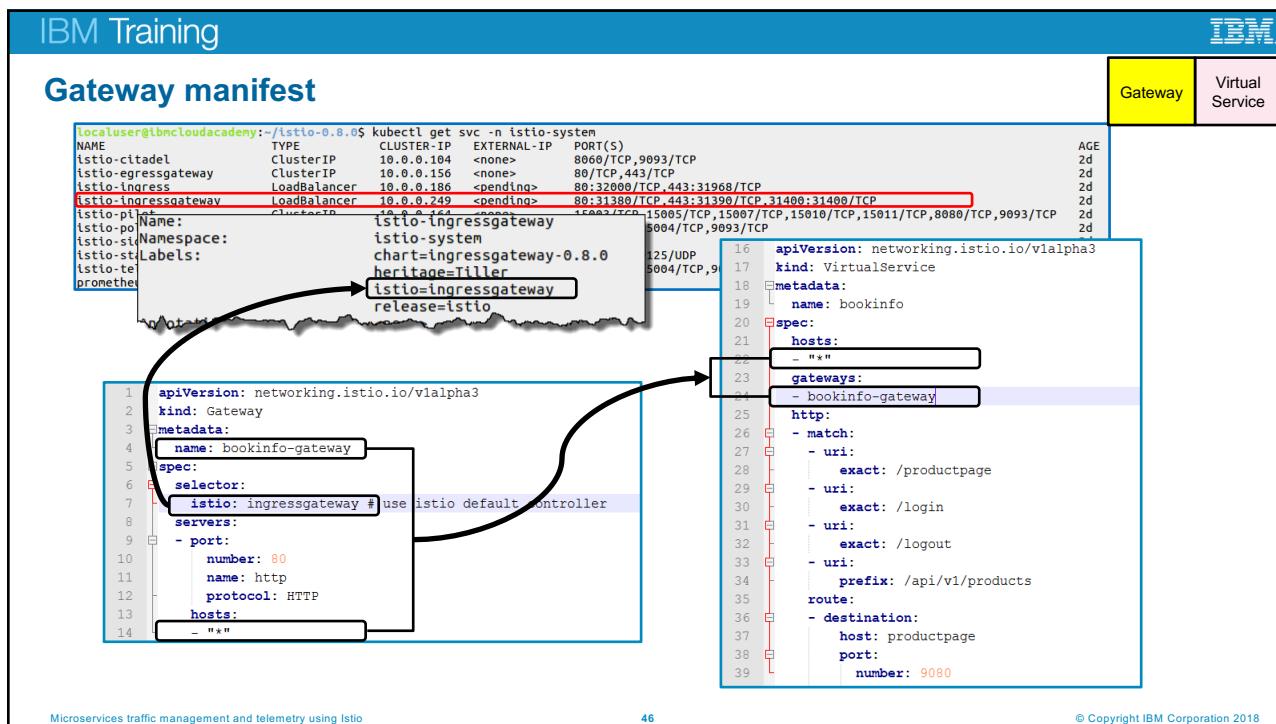
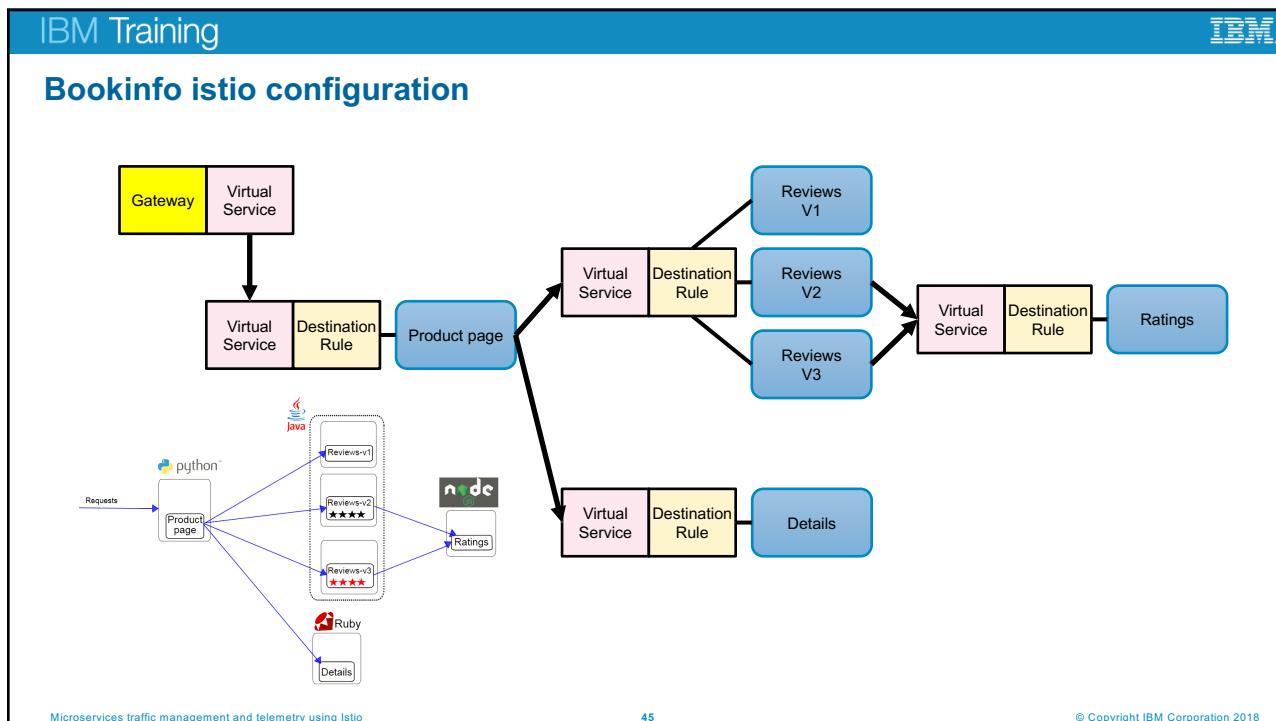
Absolutely fun and entertaining. The play lacks thematic depth when compared to other plays by Shakespeare.

Review 2: ★★★★☆

© Copyright IBM Corporation 2018

Microservices traffic management and telemetry using Istio

44



IBM Training

Product page virtual service

- VirtualService is selected based on the source destination
- The route destination links the VirtualService to the DestinationRule
- This simple DestinationRule directly refers to the pod (by finding the appropriate service name)

The diagram illustrates the flow of traffic from a Kubernetes service to an Istio component. At the top, a terminal window shows the command 'kubectl get svc' with its output. A red box highlights the 'productpage' service. Below it, two code snippets show the 'DestinationRule' and 'VirtualService' configurations. Arrows point from the highlighted 'productpage' entry in the service list to both the 'host' field in the 'DestinationRule' and the 'host' field in the 'VirtualService'. At the bottom, a box labeled 'Virtual Service' is connected by a line to a box labeled 'Destination Rule', which is then connected to a box labeled 'Product page'.

```

localuser@ibmcloudacademy:~$ kubectl get svc
NAME      TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE
details   ClusterIP   172.21.178.171 <none>       9080/TCP   3h
kubernetes   ClusterIP   172.21.0.1    <none>       443/TCP    13d
productpage ClusterIP   172.21.49.178 <none>       9080/TCP   3h
ratings   ClusterIP   172.21.74.214 <none>       9080/TCP   3h
reviews   ClusterIP   172.21.16.141 <none>       9080/TCP   3h

```

```

1 apiVersion: networking.istio.io/v1alpha3
2 kind: DestinationRule
3 metadata:
4   name: productpage
5 spec:
6   host: productpage

```

```

1 apiVersion: networking.istio.io/v1alpha3
2 kind: VirtualService
3 metadata:
4   name: productpage
5 spec:
6   hosts:
7     - productpage
8     http:
9       - route:
10         - destination:
11           host: productpage

```

Virtual Service Destination Rule Product page

Microservices traffic management and telemetry using Istio 47 © Copyright IBM Corporation 2018

IBM Training

Defining subsets using DestinationRule

- With multiple reviews application deployed, a DestinationRule defines the subsets
- Initially the VirtualService would refer to only the original version (version: v1)

The diagram shows the relationship between a 'Virtual Service' and a 'Destination Rule'. On the right, a box labeled 'Virtual Service' is connected by a line to a box labeled 'Destination Rule', which is then connected to two boxes labeled 'Reviews V1' and 'Reviews V2', representing different subsets. Below this, two code snippets are shown. An arrow points from the 'subset: v1' entry in the 'VirtualService' configuration to the 'name: v1' entry in the 'DestinationRule' configuration. Another arrow points from the 'host: reviews' entry in the 'VirtualService' to the 'host: reviews' entry in the 'DestinationRule'.

```

13 apiVersion: networking.istio.io/v1alpha3
14 kind: VirtualService
15 metadata:
16   name: reviews
17 spec:
18   hosts:
19     - reviews
20     http:
21       - route:
22         - destination:
23           host: reviews
24           subset: v1

```

```

8 apiVersion: networking.istio.io/v1alpha3
9 kind: DestinationRule
10 metadata:
11   name: reviews
12 spec:
13   host: reviews
14   subsets:
15     - name: v1
16       labels:
17         version: v1
18     - name: v2
19       labels:
20         version: v2

```

Virtual Service Destination Rule Reviews V1 Reviews V2

Microservices traffic management and telemetry using Istio 48 © Copyright IBM Corporation 2018

IBM Training

Traffic splitting across multiple destinations

- For Canary testing, a temporary change of the VirtualService would allow some traffic to be passed to v2
- For Blue-green deployment VirtualService rule can be seamlessly updated in intervals to shift workload from v1 to v2

```

graph LR
    VS[Virtual Service] --- DR[Destination Rule]
    DR --- v1[Reviews V1]
    DR --- v2[Reviews V2]
  
```

VirtualService YAML:

```

13 apiVersion: networking.istio.io/v1alpha3
14 kind: VirtualService
15 metadata:
16   name: reviews
17 spec:
18   hosts:
19     - reviews
20   http:
21     - route:
22       - destination:
23         host: reviews
24         subset: v1
25         weight: 90
26       - destination:
27         host: reviews
28         subset: v2
29         weight: 10
  
```

DestinationRule YAML:

```

8  apiVersion: networking.istio.io/v1alpha3
9   kind: DestinationRule
10  metadata:
11    name: reviews
12  spec:
13    host: reviews
14    subsets:
15      - name: v1
16        labels:
17          version: v1
18      - name: v2
19        labels:
20          version: v2
  
```

Microservices traffic management and telemetry using Istio 49 © Copyright IBM Corporation 2018

IBM Training

Content-based Traffic steering examples

- With traffic steering, depending on the content of the HTTP request, some of the traffic can be directed to a certain set of destination
- Matching can be performed against source labels, URI, headers, port, and other HTTP attributes

```

graph LR
    VS[Virtual Service] --- DR[Destination Rule]
    DR --- v1[Reviews V1]
    DR --- v2[Reviews V2]
  
```

VirtualService YAML:

```

1  apiVersion: networking.istio.io/v1alpha3
2   kind: VirtualService
3   metadata:
4     name: reviews
5   spec:
6     hosts:
7       - reviews
8     http:
9       - match:
10      - headers:
11        - end-user:
12          exact: jason
13      route:
14        - destination:
15          host: reviews
16          subset: v2
17        - destination:
18          host: reviews
19          subset: v1
  
```

DestinationRule YAML:

```

8  apiVersion: networking.istio.io/v1alpha3
9   kind: DestinationRule
10  metadata:
11    name: reviews
12  spec:
13    host: reviews
14    subsets:
15      - name: v1
16        labels:
17          version: v1
18      - name: v2
19        labels:
20          version: v2
  
```

Microservices traffic management and telemetry using Istio 50 © Copyright IBM Corporation 2018

IBM Training

HTTP fault injection and testing

- Fault injection can be used for testing
 - Faults are caused on a percentage of requests
 - Faults can cause a request delay or failure
- In this example, ratings is being invoked
 - All of the requests from bar have a 2 second timeout
 - 40% of the requests from bar also have a 5 second delay
 - 20% of the requests from foo get an HTTP 500 error
 - All other requests (not foo or bar) have a 4 second timeout
- Timeout is measured after the host is invoked, so timeout is calculated after delay period has passed
 - The 40% of requests from bar that have the delay will time out after 7 seconds (5 sec delay + 2 sec timeout)

```

apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
    - ratings
  http:
    - match:
        - headers:
            end-user: exact: bar
      fault:
        delay:
          percent: 40
          fixedDelay: 5s
      timeout: 2s
      route:
        - destination:
            host: ratings
    - match:
        - headers:
            end-user: exact: foo
      fault:
        abort:
          percent: 20
          httpStatus: 500
      route:
        - destination:
            host: ratings
    - route:
        - destination:
            host: ratings
      timeout: 4s
  
```

Microservices traffic management and telemetry using Istio 51 © Copyright IBM Corporation 2018

IBM Training

Testing circuit breaker

- Connection pool
 - Limits connections for reviews to invoke ratings
 - Limited to 1 concurrent connection, 1 request per connection
 - One concurrent request total
 - While requests are using all of the connections in a pool, any new requests are rejected
- Outlier detection
 - If there are 3 requests in 2 seconds, ratings will be ejected for 3 minutes
 - Request 1 will take more than 2 seconds, blocking the connection during that time
 - Request 2 won't get a connection, which will generate the first error
 - Request 3 won't get a connection, which will generate the second error, causing ejection
- Fixed delay
 - ratings is delayed 2 seconds before the service runs
 - Guarantees that the service will run for more than 2 seconds

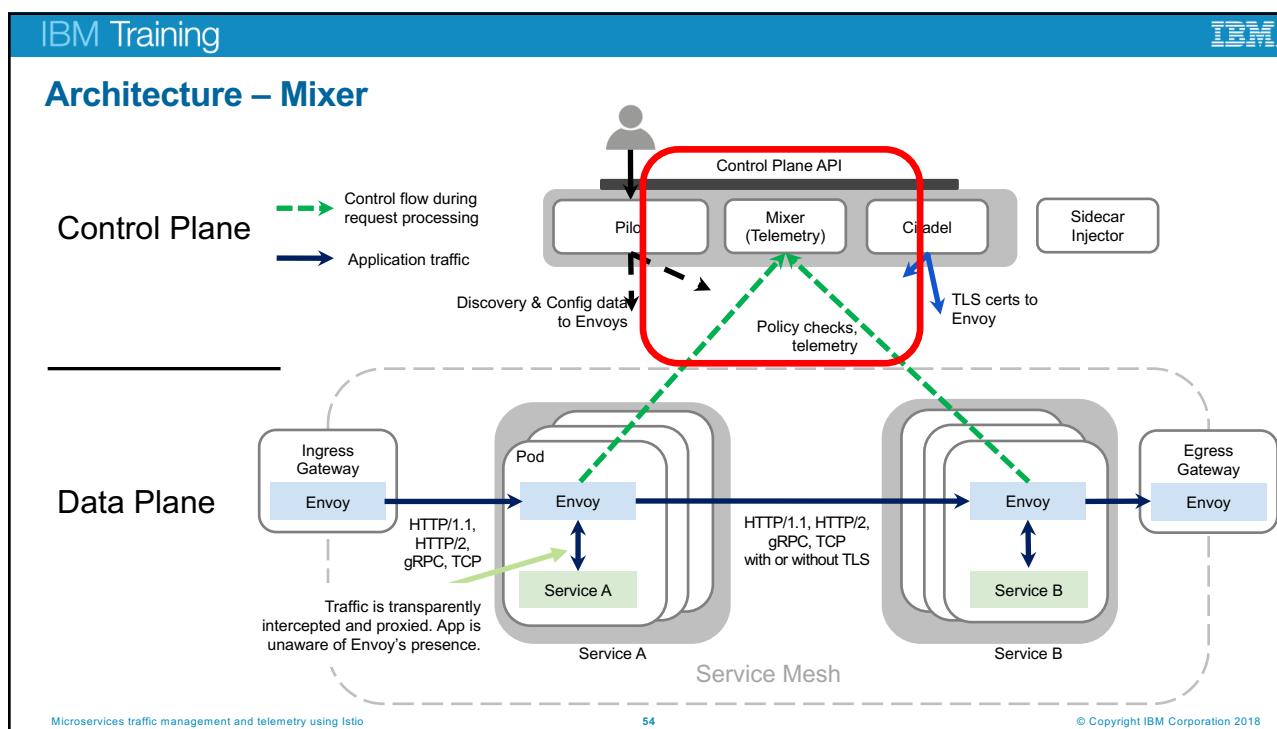
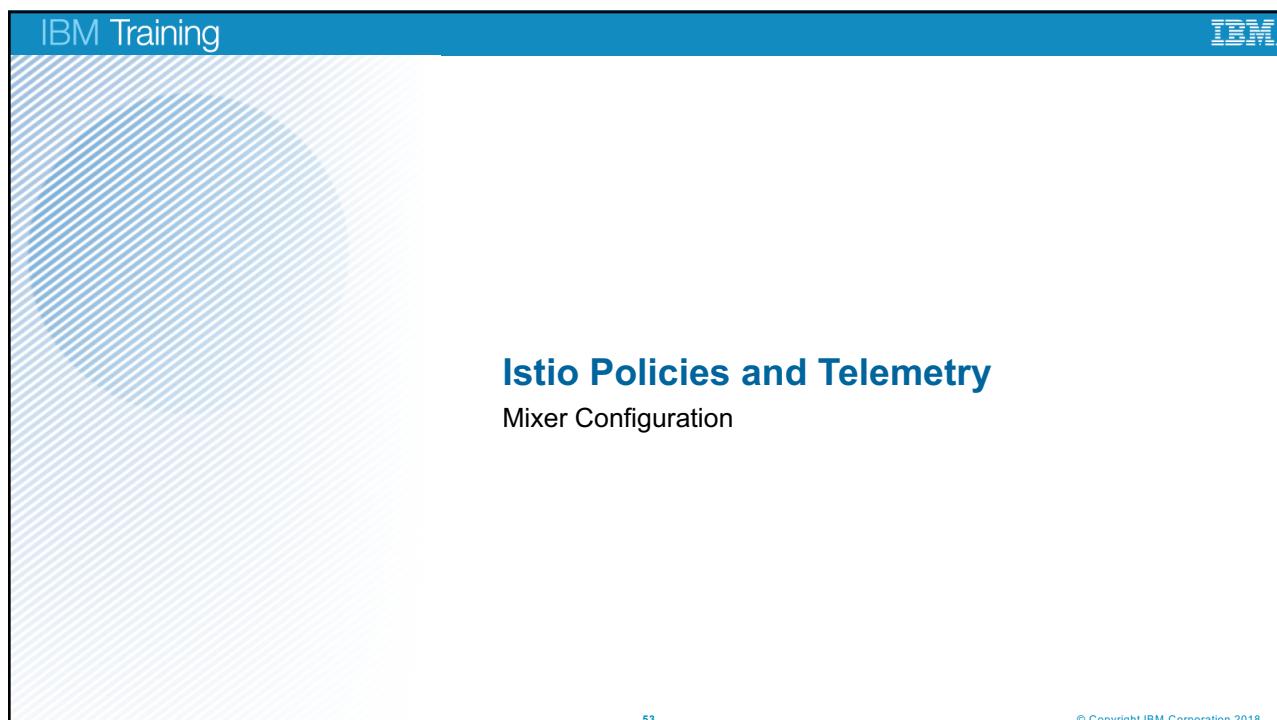
```

apiVersion: ...
kind: DestinationRule
metadata:
  name: reviews-dr
spec:
  host: reviews
  trafficPolicy:
    connectionPool:
      tcp:
        maxConnections: 1
      http:
        http1MaxPendingRequests: 1
        maxRequestsPerConnection: 1
    outlierDetection:
      consecutiveErrors: 2
      interval: 2s
      baseEjectionTime: 3m
      maxEjectionPercent: 100
  
```

```

apiVersion: ...
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
    - ratings
  http:
    - fault:
        delay:
          percent: 100
          fixedDelay: 2s
    - route:
        - destination:
            host: ratings
subset: v1
  
```

Microservices traffic management and telemetry using Istio 52 © Copyright IBM Corporation 2018



IBM Training Policies and Telemetry

- Mixer**
 - The Istio component responsible for providing policy controls and telemetry collection
 - Enforces authorization policies and collects telemetry for the services in a mesh
 - Adapter API enables it to interface with an open-ended set of infrastructure backends
 - Optional: All policy enforcement and telemetry collection can be disabled
- Envoy**
 - Performs precondition checks before each request
 - Reports telemetry after each request
 - Calls to Mixer are infrequent
 - Caching and buffering improves performance

Microservices traffic management and telemetry using Istio

55

© Copyright IBM Corporation 2018

IBM Training Attributes

- Each piece of telemetry data is an attribute
- Each attribute has a name and a type
- Mixer is in essence an attribute processing machine
- Each infrastructure backend collects the attributes it's interested in

Attribute examples	
request.path	xyz/abc
request.size	234
request.time	12:34:56.789 04/17/2017
source.ip	192.168.0.1
destination.service	example

Name	Type	Kubernetes
source.uid	string	kubernetes://...
source.labels	map[string, string]	version => v1
destination.workload.name	string	istio-telemetry
request.headers	map[string, string]	
response.duration	duration	

Microservices traffic management and telemetry using Istio

56

© Copyright IBM Corporation 2018

IBM Training

Mixer configuration definitions

- Istio Mixer offers a pluggable interface for the mesh to interact with infrastructure backends to provide:
 - Logging
 - Quota processing
 - Authorization backend
 - Metric metric
- Mixer uses rules to match adapter with template

Microservices traffic management and telemetry using Istio 57 © Copyright IBM Corporation 2018

IBM Training

Mixer CRD – custom resource definition (templates and adapters)

		Templates											
		API Key	Analytics	Authorization	Check Nothing	Kubernetes	List Entry	Log Entry	Metric	Quota	Report Nothing	Service Control	Trace Span
Adapters	Apigee	Y	Y										
	Circonus												
	CloudWatch (AWS)												Y
	Datadog												Y
	Denier				Y	Y							Y
	Fluentd								Y				
	KubernetesEnv						Y						
	List							Y					
	MemoryQuota												Y
	OPA (Open Policy Agent)				Y								
	Prometheus								Y				
	RBAC			Y									
	Redis Quota									Y			
	Service Control	Y									Y		Y
	SignalFx								Y				
	SolarWinds								Y	Y			
	Stackdriver								Y	Y			Y
	StatsD								Y				
Stdio								Y	Y				

Adapter
Interface with infrastructure backend

Template
Formats data to be sent to the adapters

Rule
Match adapter with the template it uses

Y = Rules that can be written
Y = Rules included in IBM Istio

<https://istio.io/docs/reference/config/policy-and-telemetry/>

Microservices traffic management and telemetry using Istio 58 © Copyright IBM Corporation 2018

IBM Training

OOTB definitions

- Output logging entry to telemetry's STDOUT
- Save services metrics to Prometheus
- Collect Kubernetes attributes for other adapters (note the arrow direction)
- Additional definition can be created as needed

```
localUser@ibmcloudacademy:~$ kubectl get rule --all-namespaces
NAMESPACE      NAME          AGE
istio-system   kubeattrgenrulerule   1d
istio-system   promhttp        1d
istio-system   promtcp         1d
istio-system   stdio          1d
istio-system   stdiotcp        1d
istio-system   tcpkubeattrgenrulerule 1d
```

Microservices traffic management and telemetry using Istio 59 © Copyright IBM Corporation 2018

IBM Training

Loading metric data to Prometheus

Metric data collected by Mixer, enriched using Kubernetes information and then sent to Prometheus

```
apiVersion: config.istio.io/v1alpha2
kind: metric
metadata:
  name: requestcount.metric
  namespace: istio-system
spec:
  dimensions:
    connection_security_policy: "inbound" | "outbound" | "unknown" | "mutual_tls" | "none"
    destination_app: destination.labels["app"] | "unknown"
    destination_principal: destination.principal | "unknown"
    reporter: "inbound" | "outbound" | "source" | "destination"
    request_protocol: api.protocol | context.protocol | "unknown"
    response_code: response.code | 200
    source_app: source.labels["app"] | "unknown"
    source_principal: source.principal | "unknown"
  monitored_resource_type: "UNSPECIFIED"
  value: "1"
```

```
apiVersion: config.istio.io/v1alpha2
kind: prometheus
metadata:
  name: handler
  namespace: istio-system
spec:
  metrics:
    - instance_name: requestcount.metric
      kind: COUNTER
      label_names:
        - reporter
        - source_app
        - source_principal
        - destination_app
        - destination_principal
        - request_protocol
        - response_code
        - connection_security_policy
      name: requests_total
```

```
apiVersion: config.istio.io/v1alpha2
kind: kubernetes
metadata:
  name: attributes
  namespace: istio-system
spec:
  attribute_bindings:
    destination.container.name: $out.destination_container_name | "unknown"
    destination.name: $out.destination_pod_name | "unknown"
    destination.serviceAccount: $out.destination_service_account_name | "unknown"
    source.name: $out.source_pod_name | "unknown"
    source.serviceAccount: $out.source_service_account_name | "unknown"
```

Microservices traffic management and telemetry using Istio 60 © Copyright IBM Corporation 2018

IBM Training IBM

Mixer configuration – Block a service

```
1 apiVersion: "config.istio.io/v1alpha2"
2 kind: denier
3 metadata:
4   name: handler
5   namespace: istio-system
6 spec:
7   status:
8     code: 7
9   message: Not allowed
```

```
11 apiVersion: "config.istio.io/v1alpha2"
12 kind: checknothing
13 metadata:
14   name: denyrequest
15   namespace: istio-system
16 spec:
```

```
19 apiVersion: "config.istio.io/v1alpha2"
20 kind: rule
21 metadata:
22   name: denyingress
23   namespace: istio-system
24 spec:
25   match: source.labels["istio"] == "ingressgateway" && request.headers["x-user"] == "john"
26   actions:
27     - handler: handler.denier.istio-system
28   instances: [ denyrequest.checknothing.istio-system ]
```

Microservices traffic management and telemetry using Istio © Copyright IBM Corporation 2018

IBM Training IBM

Istio Debugging

62 © Copyright IBM Corporation 2018

IBM Training

Debugging Istio

- When you have istio running, you may want to:
 - Make sure istio is running well
 - Understand how istio policy affected your application
- Debugging is performed on the control plane (Pilot or Mixer) or on the data plane (Envoy)
 - Control Plane components have an integrated controlZ interface
 - Envoy proxy has opened port 15000 for agent management

Microservices traffic management and telemetry using Istio 63 © Copyright IBM Corporation 2018

IBM Training

Debugging Istio Pilot/Policy/Telemetry

- Configuring ControlZ
 - Port forwarding controlZ port
 - Accessing controlZ configuration page
- Getting pod log


```
kubectl port-forward istio-pilot-68c9b4bdcf-x6qcd 9876:9876 -n istio-system
```

```
localuser@ibmcoudacademy:~$ kubectl get pod -n istio-system
NAME                               READY   STATUS    RESTARTS   AGE
grafana-64b7b844cc-vlmkh          1/1    Running   0          9h
istio-citadel-c8fb4f667-q48lr     1/1    Running   0          9h
istio-egressgateway-6bd595895b-57frv 0/1    Running   0          40s
istio-egressgateway-6bd595895b-8j46f 1/1    Running   0          9h
istio-egressgateway-6bd595895b-msm57 1/1    Running   0          52s
istio-egressgateway-6bd595895b-w78s9 1/1    Running   0          5m
istio-galley-57b749c55d-dkz2x2     1/1    Running   0          9h
istio-ingressgateway-57d5b78cc9-4wz8h 0/1    Running   0          38s
istio-ingressgateway-57d5b78cc9-jhnph 1/1    Running   0          9h
```

Istio ControlZ /usr/local/bin/pilot-discovery - 10.1.223.84

Logging Scopes

Scope	Description	Output Level	Stack Trace Level	Log Callers?
rbac	rbac debugging	info	none	■
ads	ads debugging	info	none	■
default	Unscoped logging messages	info	none	■
model	model	info	none	■

none
error
warn
info
debug

Microservices traffic management and telemetry using Istio 64 © Copyright IBM Corporation 2018

IBM Training

Envoy configuration

- Getting envoy configuration (JSON)


```
kubectl exec <podname> -c istio-proxy -- curl localhost:15000/config_dump -s
```

```
localuser@ibmcloudacademy:~$ kubectl get pod
NAME                               READY   STATUS    RESTARTS   AGE
details-v1-558d5fc956-7vj9g        2/2     Running   0          8h
productpage-v1-576c5fdb8-mscj      2/2     Running   0          8h
ratings-v1-556c44f648-2w2h6       2/2     Running   0          8h
reviews-v1-5ff97b656b-tdrccb     2/2     Running   0          8h
reviews-v2-948df8f54-hn2qk       2/2     Running   0          8h
reviews-v3-f9b6c94f8-clrn         2/2     Running   0          8h

[{"configs": {
  "bootstrap": {
    "@type": "type.googleapis.com/envoy.admin.v2alpha.BootstrapConfigDump",
    "bootstrap": {
      "node": {
        "id": "stdecar-10.1.235.131-details-v1-558d5fc956-7vj9g.default-default.svc.cluster.local",
        "cluster": "details",
        "metadata": {
          "ISTIO_PROXY_VERSION": "1.0.0",
          "ISTIO_PROXY_SHA": "istio-proxy:6166ae7ebac7f630206b2fe4e6767516bf198313",
          "ISTIO_VERSION": "1.0.0",
          "POD_NAME": "details-v1-558d5fc956-7vj9g",
          "istio": "sidecar",
          "INTERCEPTION_MODE": "REDIRECT"
        },
        "build_version": "0/1.8.0-dev//RELEASE"
      },
      "static_resources": {
        "clusters": [
          {
            "name": "xds-grpc",
            "type": "STRICT_DNS",
            "connect_timeout": "10s",
            "hosts": [
              {
                "host": "127.0.0.1"
              }
            ]
          }
        ]
      }
    }
  }
}]

localuser@ibmcloudacademy:~$ kubectl exec details-v1-558d5fc956-7vj9g -c istio-proxy -- curl localhost:15000/config_dump -s
```

Microservices traffic management and telemetry using Istio 65 © Copyright IBM Corporation 2018

IBM Training

Envoy logging and other agent capabilities

- Changing Envoy logging level


```
kubectl exec <podname> -c istio-proxy -- curl localhost:15000/logging?<name>=<level> -s
```

```
localuser@ibmcloudacademy:~$ kubectl exec details-v1-558d5fc956-7vj9g -c istio-proxy -- curl localhost:15000/logging -s
usage: /logging?<name>=<level> (change single level)
usage: /logging?level=<level> (change all levels)
levels: trace debug info warning error critical off
active loggers:
  admin: info
  assert: info
  backtrace: info
  client: info
  config: info
  connection: info
  file: info
  filter: info
  grpc: info
  hc: info
  health-checker: info
  http: info
  http2: info
  hystrix: info
  lua: info
  main: info
  misc: info
  mongo: info
  pool: info
  rbac: info
  redis: info
  router: info
  runtime: info
  stats: info
  testing: info
  thrift: info
  tracing: info
  upstream: info
  upstream: info
```

Command	Description
/	Admin home page
/certs	print certs on machine
/clusters	upstream cluster status
/config_dump	dump current Envoy configs (experimental)
/cpuProfiler	enable/disable the CPU profiler
/healthcheck/fail	cause the server to fail health checks
/healthcheck/ok	cause the server to pass health checks
/help	print out list of admin commands
/hot_restart_version	print the hot restart compatibility version
/listeners	print listener addresses
/logging	query/change logging levels
/quitquitquit	exit the server
/reset_counters	reset all counters to zero
/runtime	print runtime values
/runtime_modify	modify runtime values
/server_info	print server version/status information
/stats	print server stats
/stats/prometheus	print server stats in prometheus format

Microservices traffic management and telemetry using Istio 66 © Copyright IBM Corporation 2018

IBM Training 

Conclusion

- Concepts
- Traffic management
 - Rule configuration
 - Traffic routing
 - Bookinfo example
- Policies and telemetry
- Debugging Istio

Microservices traffic management and telemetry using Istio

67

© Copyright IBM Corporation 2018

IBM Training

IBM

Microservices Network Security

Using Kubernetes and Istio to secure the network in your microservice application

© Copyright IBM Corporation 2018
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

IBM Training

IBM

Objectives

- Understand network security options
- Understand Kubernetes network policy basics
- Understand network security with Istio

Prerequisites: Kubernetes, Microservices

Microservices Network Security

2

© Copyright IBM Corporation 2018

IBM Training

IBM

Network security for microservices application

3 © Copyright IBM Corporation 2018

IBM Training

IBM

Network security approaches

- Defense in depth
- Layered defense strategy, includes all aspects of IT
- Perimeter protection to stop threat at the network boundary
- Network segmentation and isolation for more granular defense

Microservices Network Security

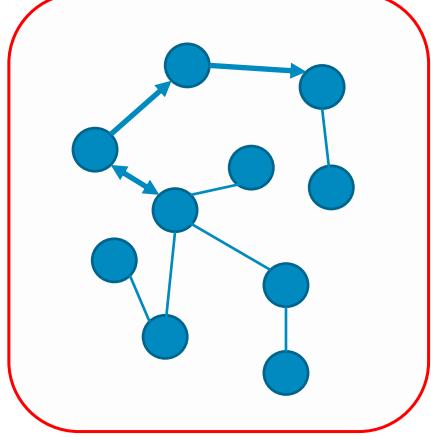
4 © Copyright IBM Corporation 2018

IBM Training IBM

Network perimeter defense

Typical perimeter controls:

- Firewall and VPN to protect system boundary
- Whitelisting: Deny all by exception, open only what is required
- Intrusion monitoring and protection at system boundary

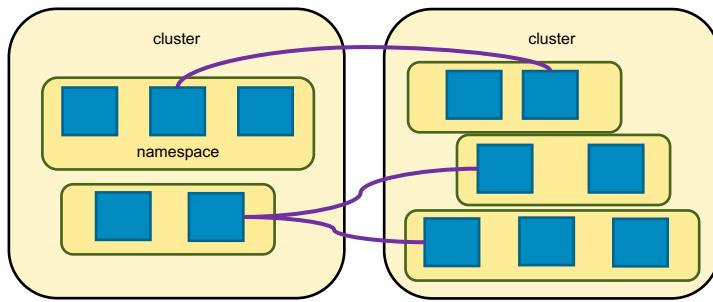


Microservices Network Security 5 © Copyright IBM Corporation 2018

IBM Training IBM

Network segmentation

- Segment the network: provides isolation
- Isolation: control what talks to what
- Use segmentation when defining network policies
- Segmentation can be done using namespaces



Microservices Network Security 6 © Copyright IBM Corporation 2018

IBM Training IBM

Communication confidentiality

- Traditional TLS issues in a microservices application
 - Applications are not tied to certain hostname or IP addresses
 - Application instances can be created or removed elastically
- Typical TLS certificate scheme are based on identifying hostname or IP address of the involved parties

The diagram shows two rectangular boxes labeled "Application A" and "Application B". A horizontal arrow connects them, with the word "mutual TLS" written above the arrow. At each end of the arrow, there is a yellow wavy shape representing a TLS certificate.

- Microservices application requires that application identity should be the same regardless of which instances the communication occurs to/from

Microservices Network Security 7 © Copyright IBM Corporation 2018

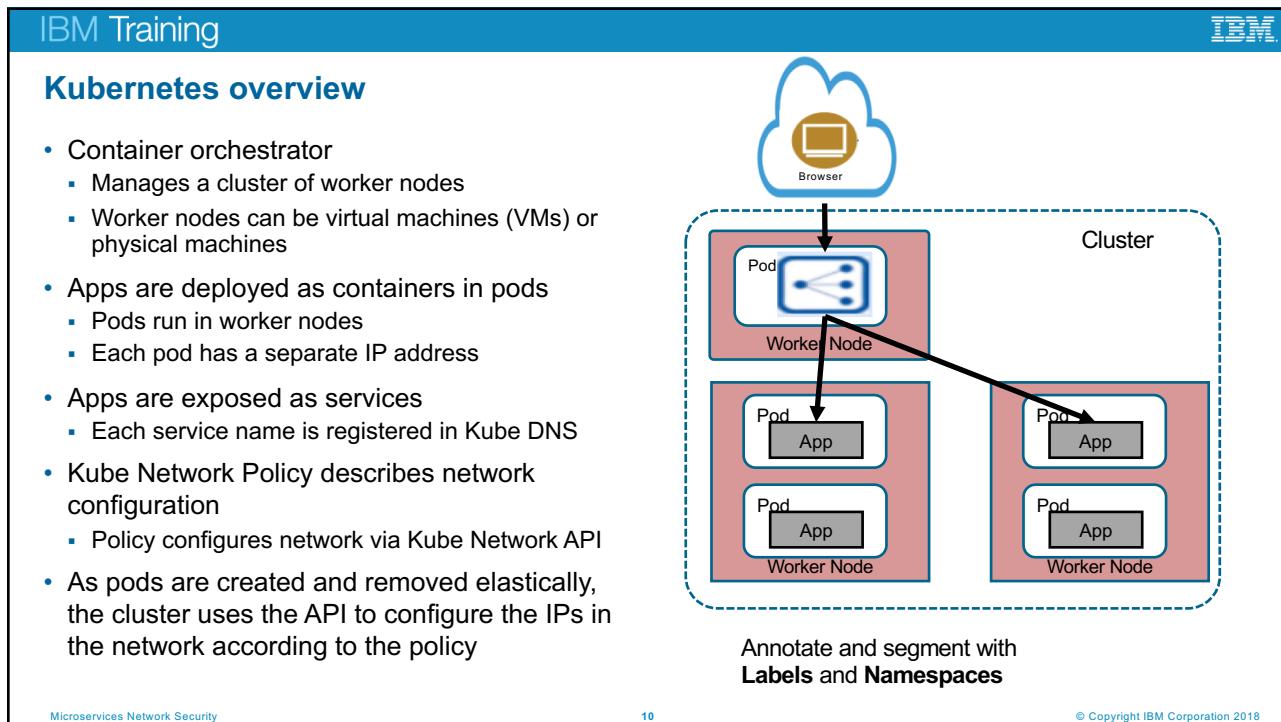
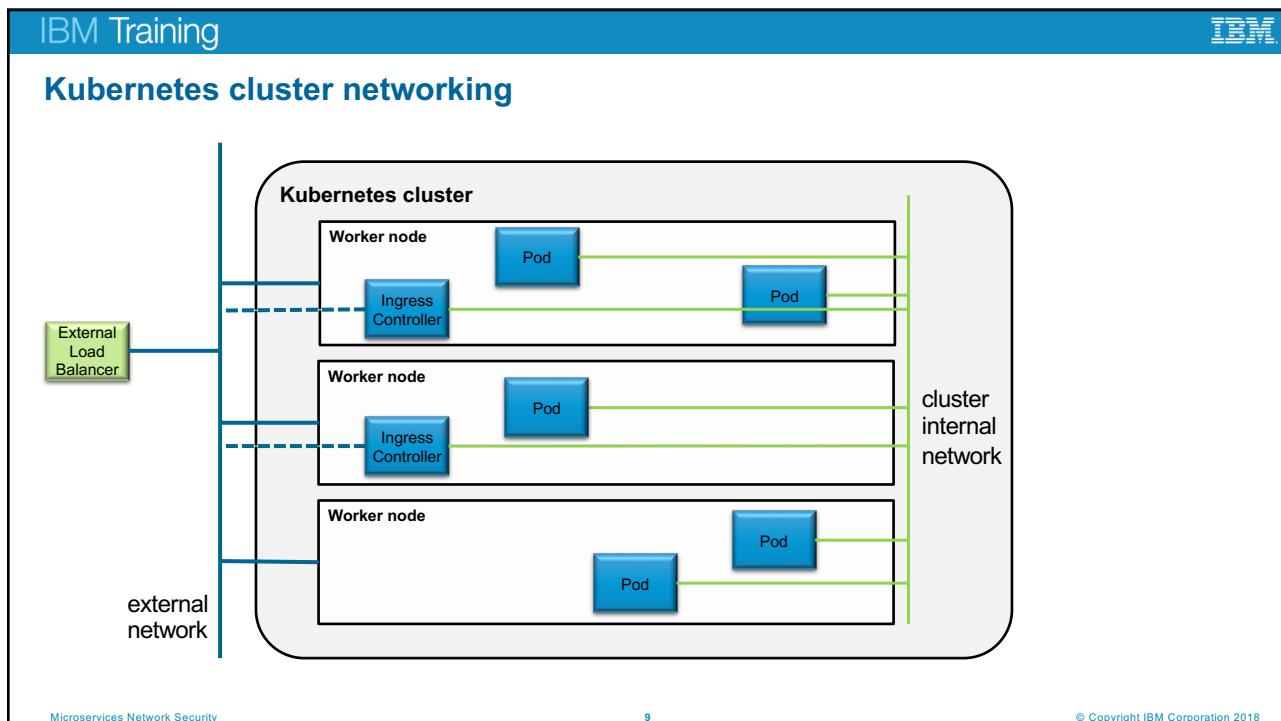
IBM Training IBM

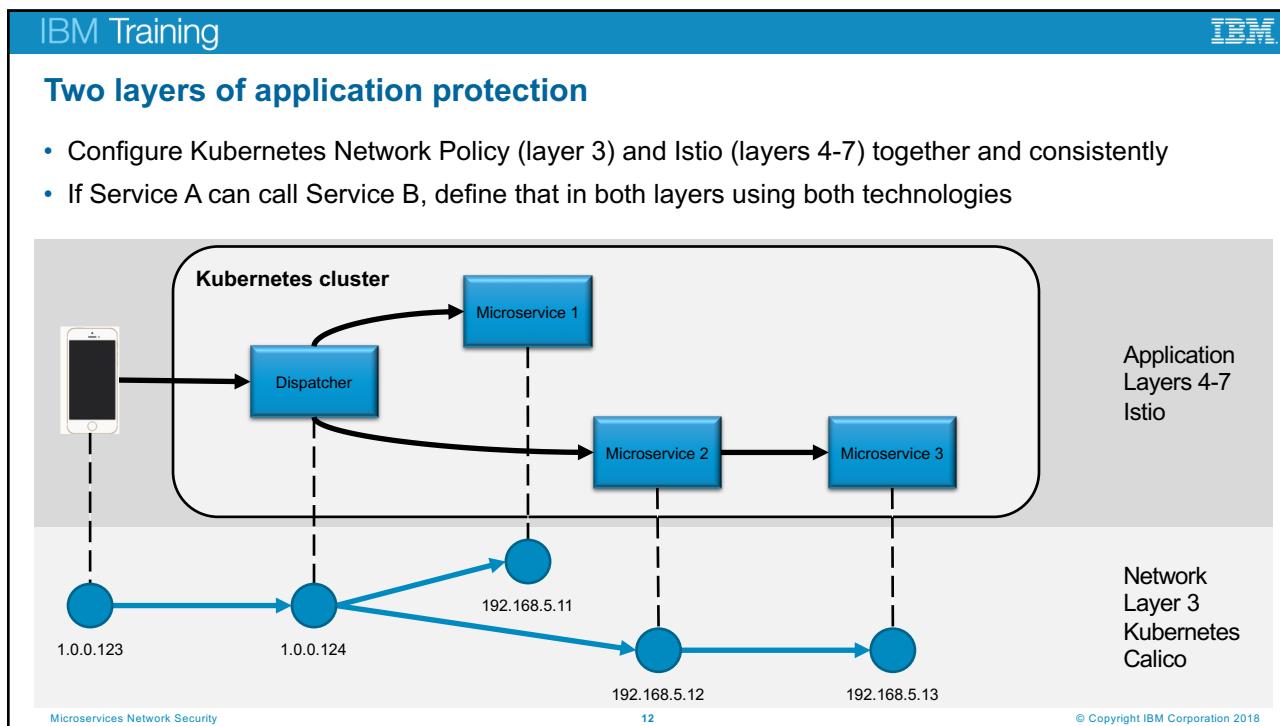
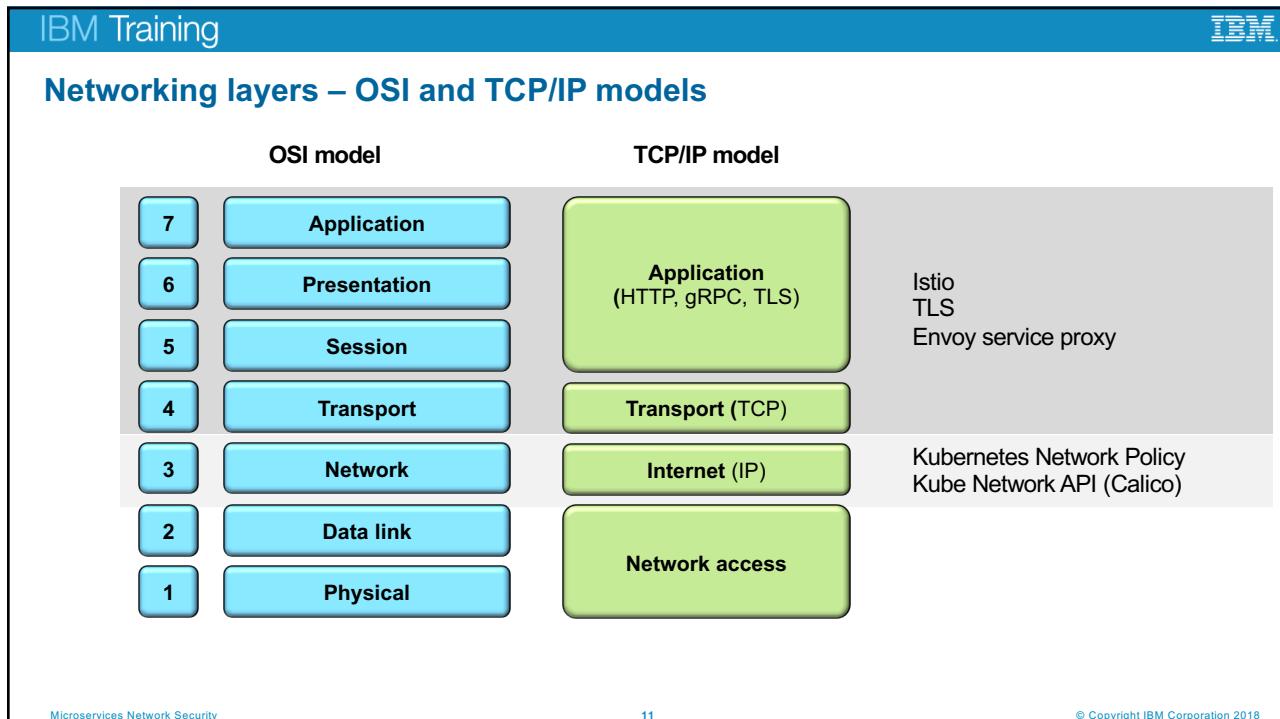
Microservices application in Kubernetes

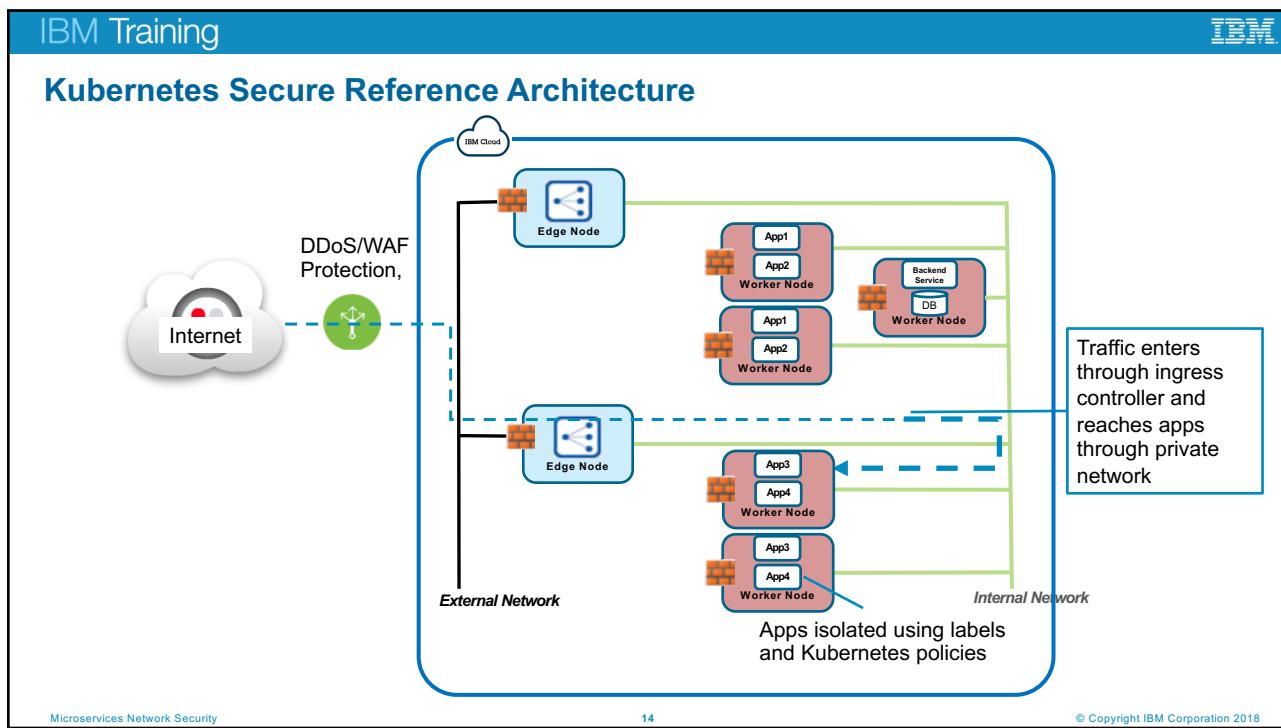
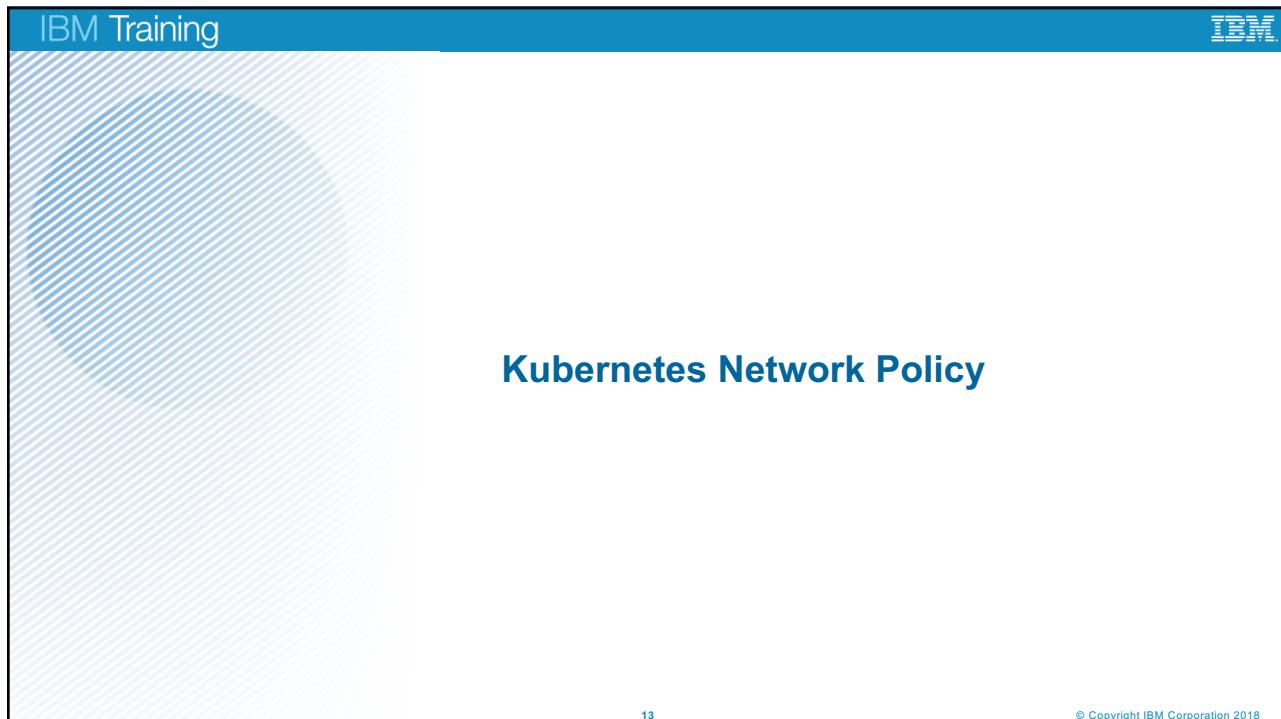
The diagram illustrates a Kubernetes cluster architecture. It features a rounded rectangle labeled "Kubernetes cluster". Inside, a blue square labeled "Dispatcher" is connected to three blue squares labeled "Microservice 1", "Microservice 2", and "Microservice 3". An arrow points from a smartphone icon outside the cluster to the "Dispatcher". Arrows also point from the "Dispatcher" to "Microservice 1" and "Microservice 2". An arrow points from "Microservice 2" to "Microservice 3".

- How the client can connect to the dispatcher?
- How to ensure that only the dispatcher can invoke Microservice 1 and 2?
- How to restrict the dispatcher from accessing Microservice 3 directly?
- How to prohibit other agents from invoking these microservices apart from the designated 'callers'?
- Must understand Kubernetes networking implementation**

Microservices Network Security 8 © Copyright IBM Corporation 2018







IBM Training

The Kubernetes Network Policy resource

- The NetworkPolicy extension is used to define private networking connectivity rules
 - Specifies which pods can communicate and how
- Cluster uses Kubernetes Network API to apply the policy to the network
- IBM uses Calico as its Kubernetes network
 - Calico is a software defined network (SDN) provider
 - Leverages Linux routing and iptables capabilities
- NetworkPolicy defines:
 - Which pod or pods the policy applies to
 - Which resources can connect to the pod (ingress)
 - Which resources the pod can connect to (egress)

```

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: test-network-policy
  namespace: default
spec:
  podSelector:
    matchLabels:
      role: db
  ingress:
  - from:
    - ipBlock:
        cidr: 172.17.0.0/16
      except:
      - 172.17.1.0/24
    - namespaceSelector:
        matchLabels:
          project: myproject
    - podSelector:
        matchLabels:
          role: frontend
  ports:
  - protocol: TCP
    port: 6379
  egress:
  - to:
    - ipBlock:
        cidr: 10.0.0.0/24
  ports:
  - protocol: TCP
    port: 5978

```

Microservices Network Security

15

© Copyright IBM Corporation 2018

IBM Training

NetworkPolicy illustration

- Object selection
 - Specifies the resources that this policy applies to
 - Attribute: **podSelector**, **namespace**
- Inbound and Outbound communication filters
 - Identify which communication is being filtered
 - Attribute: **ingress** or **egress**
- Source or target selection
 - Ingress: Access to the selected object
 - Egress: What the selected object can access
 - Attributes: **podSelector**, **namespaceSelector**, **ipBlock**
- Port selection
 - Lists the ports that are allowed access
 - Attribute: **port**

Microservices Network Security

16

© Copyright IBM Corporation 2018

IBM Training

Kubernetes network policy caveats

- Default: Allow all pods to talk to all other pods
 - Allow all traffic to and from a pod if no Kubernetes policy against it
- IBM Cloud Kubernetes Service provides a default set of Calico policies on new clusters
 - allow-node-port-dnat – Allows incoming NodePort, LoadBalancer, and Ingress service traffic to the pods that those services are exposing
 - allow-sys-mgmt – Allows incoming connections for specific IBM Cloud infrastructure (SoftLayer) systems that are used to manage the worker nodes
- Rules are OR'ed, i.e. if any ingress/egress rule matches, connection is allowed
- Connections are duplex: If connection A allowed to B, B can respond but can't create connection to A
- Network policies are scoped to the namespace they are deployed to
- Network Policy can limit both ingress and egress, but limiting ingress only may be sufficient
 - Istio Service Role limits ingress
 - Network Policy egress is still needed for accessing external resources

Microservices Network Security

17

© Copyright IBM Corporation 2018

IBM Training

Network Policy example #1

- Limit which pods can connect on the network
- Two network policies
 - Frontend can access backend
 - Backend can access sor
- Policy uses labels to find pods

```

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: backend-allow
spec:
  podSelector:
    matchLabels:
      app: App1
      tier: backend
  ingress:
  - from:
    - podSelector:
        matchLabels:
          app: App1
          tier: frontend

```

```

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: sor-allow
spec:
  podSelector:
    matchLabels:
      app: App1
      tier: sor
  ingress:
  - from:
    - podSelector:
        matchLabels:
          app: App1
          tier: backend

```

Microservices Network Security

18

© Copyright IBM Corporation 2018

IBM Training IBM

Network Policy example #2

- One team's services need to call another's
- One network policy
 - Enables all finance services to call all accounts services on specified port
- Policy uses labels to find namespaces

```

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  namespace: accounts
  name: accounts-allow
spec:
  podSelector:
    ingress:
    - from:
        - namespaceSelector:
            matchLabels:
              usage: finance
  ports:
    port: 3111
  
```

Microservices Network Security 19 © Copyright IBM Corporation 2018

IBM Training IBM

Network Policy example #3

- Limit which pods can connect to external web
- Two network policies
 - Block all egress from namespace
 - Enable backend external access on certain ports

```

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  namespace: finance
  name: finance-deny-egress
spec:
  podSelector:
    egress: []
  
```

```

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  namespace: finance
  name: finance-allow-backend
spec:
  podSelector:
    matchLabels:
      tier: backend
  egress:
    - ports:
        - port: 443
        - port: 53
      protocol: UDP
  
```

Microservices Network Security 20 © Copyright IBM Corporation 2018

IBM Training IBM

Setting up TLS from Ingress

```
apiVersion: v1
kind: Secret
metadata:
  name: tlsSecret
  namespace: default
type: Opaque
data:
  tls.crt: base64 encoded cert
  tls.key: base64 encoded key
```



```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: ingress-with-tls
  annotations:
    nginx/ssl-service: s1
spec:
  tls:
  - secretName: tlsSecret
  backend:
    serviceName: s1
    servicePort: 443
```

The diagram illustrates a network architecture for securing microservices. A 'Browser' connects to an 'Edge Node' (represented by a blue square with a white icon). The Edge Node contains a certificate authority (CA) and routes traffic through an 'External Network'. The traffic then enters a 'Worker Node' (represented by a pink rounded rectangle). Inside the Worker Node, there are two services: 'service: s2' at the top and 'service: s1' below it. The Worker Node connects to an 'Internal Network' (green line). This setup ensures end-to-end encryption from the browser to the internal services.

Encrypt all traffic end to end; TLS 1.2 mandated

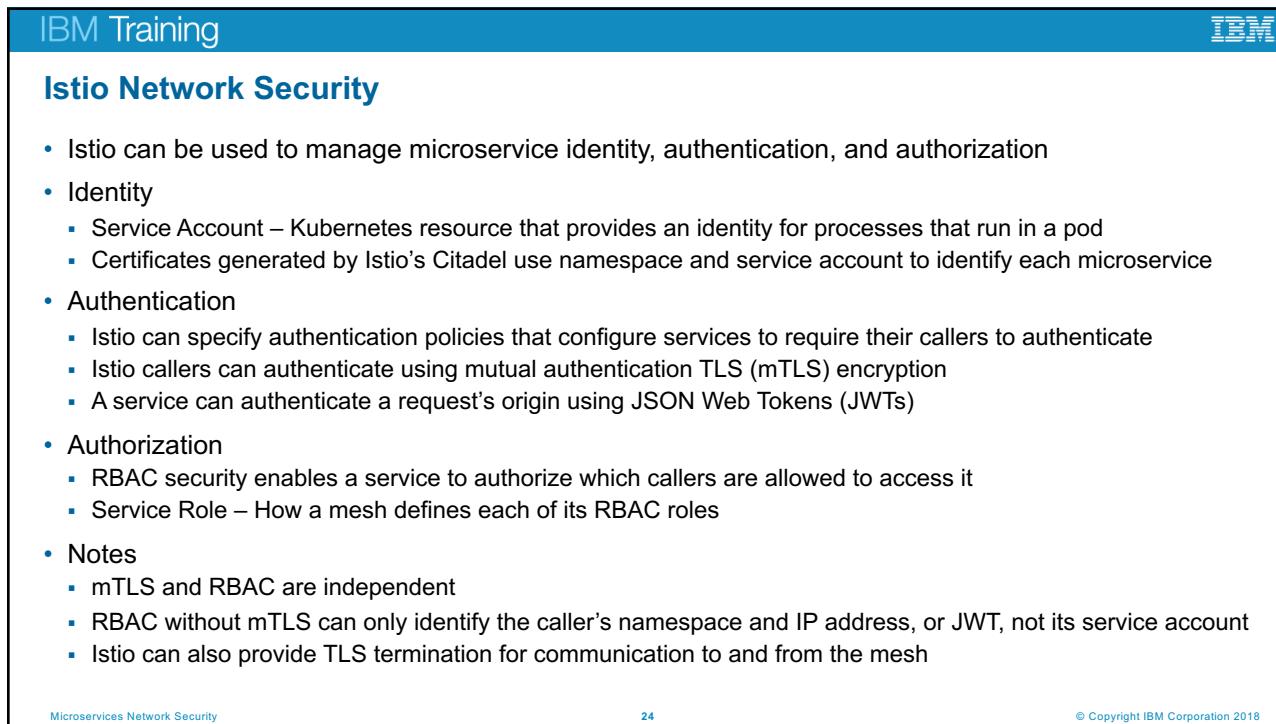
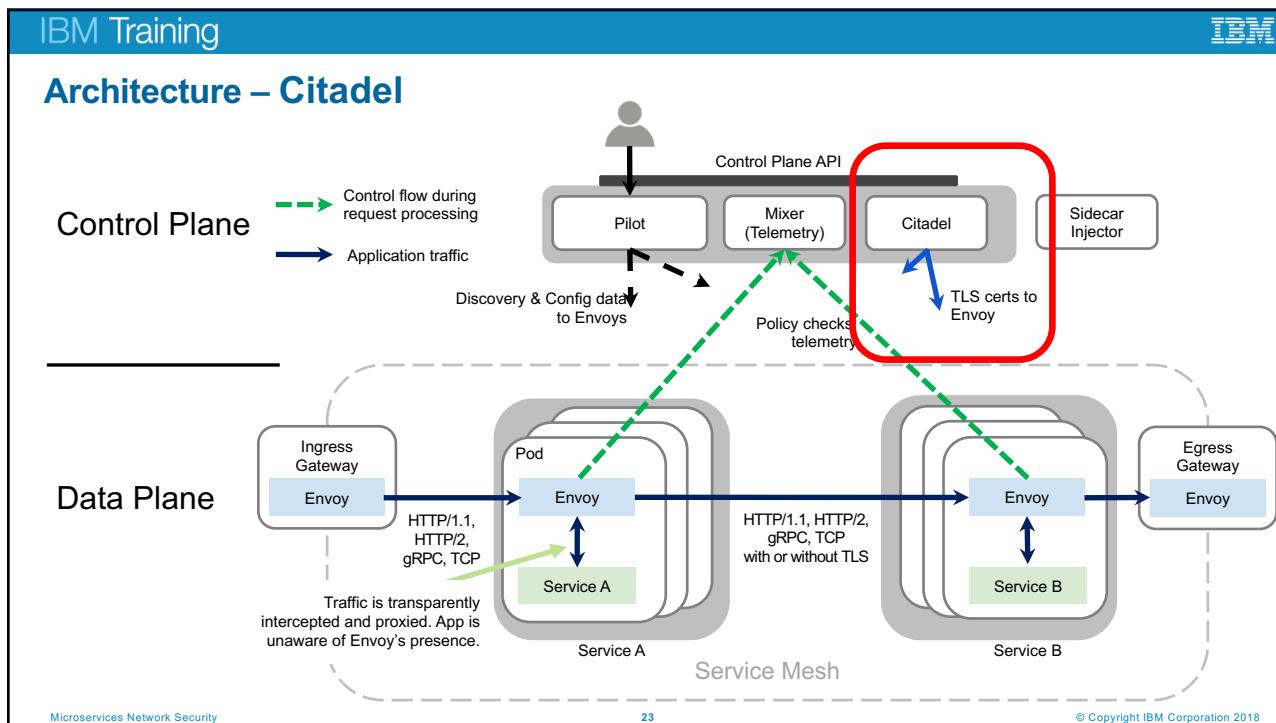
- Enforce TLS at Ingress
- Enforce TLS from Ingress to containers using Ingress 'ssl-service' annotation
- Enforce TLS from component to component
 - Self-signed CA is fine within deployment or cluster across private network
 - Certificates are stored using secrets and managed externally
- In IBM Cloud Public: CA is Digicert (LetsEncrypt in future)

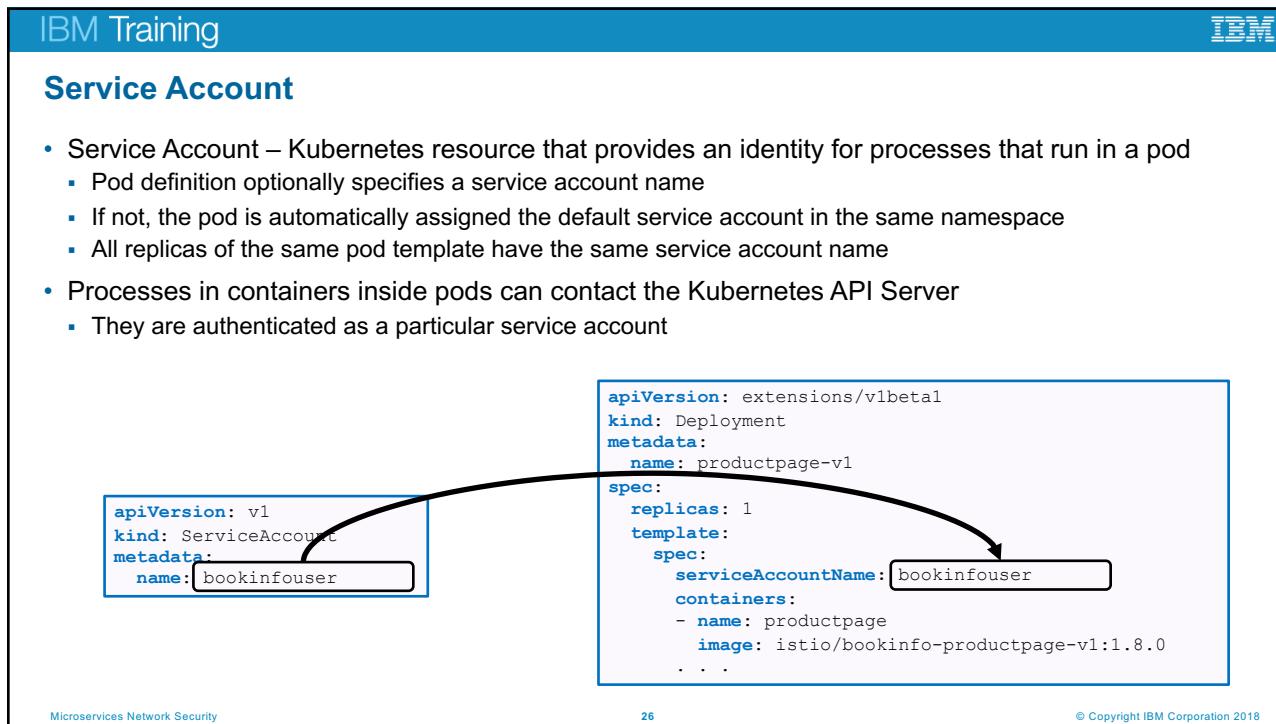
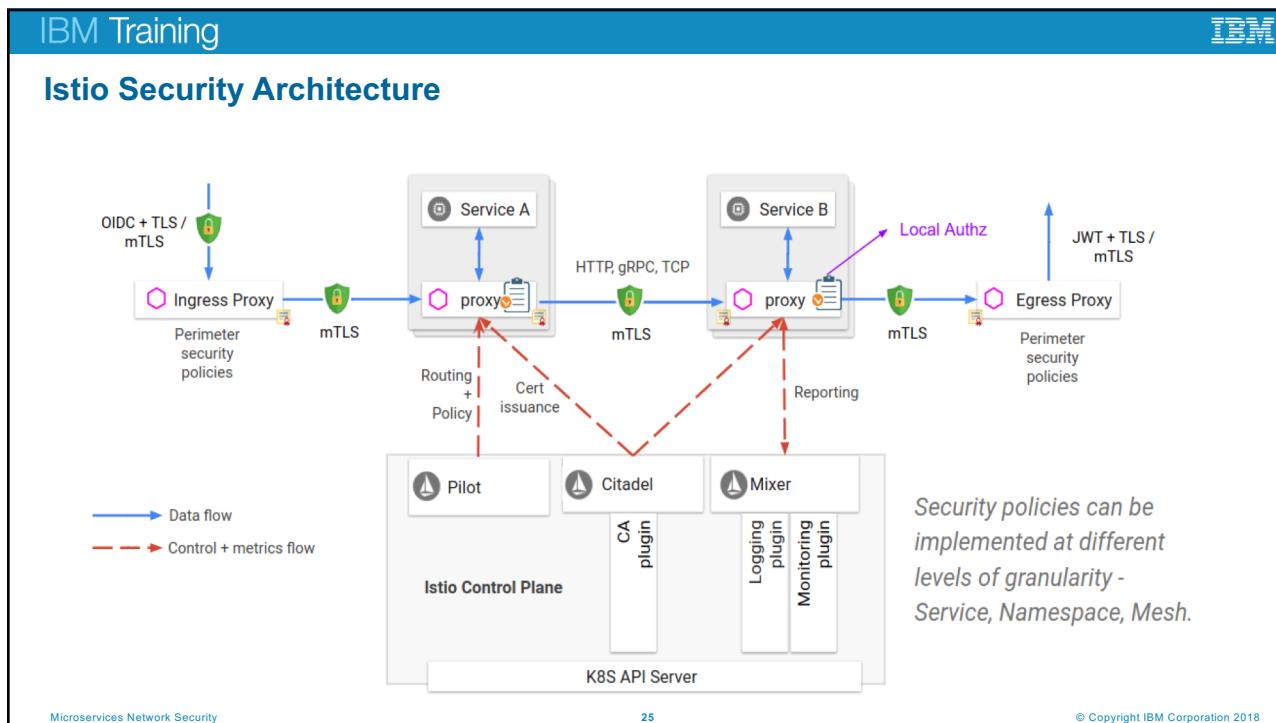
Microservices Network Security 21 © Copyright IBM Corporation 2018

IBM Training IBM

Network Security with Istio

22 © Copyright IBM Corporation 2018





IBM Training

SPIFFE workload identification

- SPIFFE ID URI format: `spiffe://trust-domain/path`
- SPIFFE – Secure Production Identity Framework For Everyone
 - Provides a secure identity, in the form of a specially crafted X.509 certificate, to every workload
 - SPIFFE IDs – Defines how services identify themselves to each other
 - SPIFFE VID – Encodes SPIFFE IDs in a cryptographically-verifiable document
 - SPIFFE Workload API – API specification for issuing and/or retrieving SVIDs
- SPIFFE Verifiable Identity Document (SVID)
 - Document that identifies a workload, can be passed to other workloads
 - Implemented as an X.509 certificate
 - SPIFFE ID is encoded in the X.509 certificate's Subject Alternative Name (SAN) extension
- SPIFFE is a Cloud Native Computing Foundation (CNCF) project
 - SPIFFE is a standard for workload identification
- SPIRE (the SPIFFE Runtime Environment) is another CNCF project, part of the SPIFFE project
 - Software that exposes the SPIFFE Workload API so workloads can retrieve their identity

Microservices Network Security

27

© Copyright IBM Corporation 2018

IBM Training

Istio Citadel (fka Istio Auth)

- Istio ID URI format (for Kubernetes)
 - `spiffe://cluster.local/ns/<namespace>/sa/<serviceaccountname>`
 - All pods with the same service account name get SPIFFE VIDs with the same identity
- Istio proxy compares SPIFFE IDs to `source.principal` and `destination.principal` fields
- Citadel is responsible for key and certificate management
 - Generation
 - Deployment
 - Rotation
 - Revocation
- Citadel certificates are SPIFFE VIDs
- Citadel is an alternative to SPIRE
 - Citadel is a more comprehensive security solution
 - Includes authentication, authorization, and auditing

Microservices Network Security

28

© Copyright IBM Corporation 2018

IBM Training IBM

Istio authentication policies

- Use authentication policies to force services to require authentication and configure it
 - Policy can be scoped to the mesh, a namespace, or selected services
- Mesh Policy – Mesh-wide
 - Scope is all pods in the mesh
 - At most one per mesh, always named `default`, no targets
- Policy – Namespace-wide
 - Scope is all pods in a namespace
 - At most one per namespace, always named `default`, no targets
- Policy – Service-specific
 - Scope is specified pods in a namespace
 - Zero, one, or many per namespace; includes target selectors
- For each service, Istio applies the narrowest matching policy
 - service-specific > namespace-wide > mesh-wide
- Only transport authentication supported is mTLS (Istio v0.7)
- Only origin authentication supported is JWT (Istio v0.7)

```
apiVersion: authentication.istio.io/v1alpha1
kind: Policy
metadata:
  name: mTLS_enable
  namespace: prod
spec:
  target:
    - name: productpage
      ports:
        - number: 9000
  peers:
    - mtls: {}
  origins:
    - jwt:
        issuer: "https://securetoken.google.com"
        audiences:
          - "productpage"
        jwksUri: "https://..../oauth2/v1/certs"
        jwt_headers:
          - "x-goog-iap-jwt-assertion"
  principalBinding: USE_ORIGIN
```

Microservices Network Security 29 © Copyright IBM Corporation 2018

IBM Training IBM

Enable mutual TLS authentication

- A (mesh) policy that specifies mTLS forces its services to require transport authorization via mTLS
 - Enabling mTLS policy causes Citadel to generate certificates for the Istio proxies
- Callers to these services must include a destination rule that enables mTLS

```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: applicationB-dr
  namespace: prod002
spec:
  host: applicationB
  trafficPolicy:
    tls:
      mode: ISTIO_MUTUAL

apiVersion: authentication.istio.io/v1alpha1
kind: Policy
metadata:
  name: prod002-policy
  namespace: prod002
spec:
  peers:
    - mtls: {}
```

Microservices Network Security 30 © Copyright IBM Corporation 2018

IBM Training 

Enabling authorization

- Istio configures authorization using Role-based Access Control (RBAC)
- RBAC Config
 - Mesh-wide singleton that enables authorization, always named default
- Mode
 - OFF – Istio authorization is disabled
 - ON – Istio authorization is enabled for all services in the mesh
 - ON_WITH_INCLUSION – Istio authorization is enabled only for specified services and namespaces
 - ON_WITH_EXCLUSION – Istio authorization is enabled for all services in the mesh except the specified services and namespaces

```
apiVersion: "rbac.istio.io/v1alpha1"
kind: RbacConfig
metadata:
  name: default
  namespace: istio-system
spec:
  mode: ON_WITH_EXCLUSION
  exclusion:
    namespaces: [ "test" ]
```

```
apiVersion: "rbac.istio.io/v1alpha1"
kind: RbacConfig
metadata:
  name: default
  namespace: istio-system
spec:
  mode: ON_WITH_INCLUSION
  inclusion:
    namespaces: [ "default" ]
```

Microservices Network Security 31 © Copyright IBM Corporation 2018

IBM Training 

Authorization policy

- Service Role
 - Represents permission to invoke these services using all or parts of their API
 - Specifies a set of services, and the methods and paths for invoking those services, and other constraints
 - Services are selected by name (including wildcards) and other constraints (typically pod labels)
- Service Role Binding
 - Grants permission to invoke the services using the specified parts of their API
 - Associates one or more subjects with a service role
 - Subjects are specified by service accounts and other identity properties
- The clients with a service role binding are allowed to call the services with a service role
 - The service blocks the call if the caller doesn't have the proper binding
 - The service blocks the call if the caller tries to use parts of the API that aren't enabled
 - The service blocks the call if the request doesn't meet the specified constraints

Microservices Network Security 32 © Copyright IBM Corporation 2018

IBM Training IBM

Role-based authorization for microservices application

- RBAC has been enabled for this mesh and these services
- Each microservice has a service role
- Each microservice client has a service role binding for each service it's allowed to invoke

Microservices Network Security 33 © Copyright IBM Corporation 2018

IBM Training IBM

Role Based Access Control (RBAC) configuration

- ServiceRole
 - rules – The set of permissions
 - services – The list of services with this role
 - methods – A list of HTTP methods
 - paths – A list of HTTP paths or gRPC methods
 - constraints – Filters based on known keys
- ServiceRoleBinding
 - roleRef – Name of the ServiceRole
 - subjects – Clients of the service
 - user – Service account name
 - properties – Filters based on known keys

Microservices Network Security 34 © Copyright IBM Corporation 2018

IBM Training IBM

Configuring ServiceRole and ServiceRoleBinding

```

graph LR
    subgraph Left_Pod [Pod]
        direction TB
        SRB[Service Role Binding] --- SRB_Signature["Service Role Binding"]
        SRB --- SSB[Service microservice1]
        SSB --- SSB_Signature["Service microservice1"]
        SSB_Signature --- SR[Service Role]
        SR --- SR_Signature["Service Role"]
        SR_Signature --- Pod_Proxy["Pod proxy"]
        Pod_Proxy --- Pod_Dispatcher["Pod dispatcher"]
    end
    subgraph Right_Pod [Pod]
        direction TB
        SR --- SR_Signature
        SR_Signature --- P[Pod]
        P --- P_Proxy["Pod proxy"]
        P_Proxy --- P_Dispatcher["Pod dispatcher"]
    end
    SSB --- SR
    SR --- P

```

ServiceRoleBinding YAML:

```

apiVersion: config.istio.io/v1alpha2
kind: ServiceRoleBinding
metadata:
  name: bind-microservice1-viewer
  namespace: appl
spec:
  subjects:
  - user: "*"
    properties:
      source.namespace: appl
  roleRef:
    kind: ServiceRole
    name: "microservice1-viewer"

```

ServiceRole YAML:

```

apiVersion: config.istio.io/v1alpha2
kind: ServiceRole
metadata:
  name: microservice1-viewer
  namespace: appl
spec:
  rules:
  - services: [ "microservice1" ]
    methods: [ "GET" ]

```

Microservices Network Security 35 © Copyright IBM Corporation 2018

IBM Training IBM

Constraint keys for service roles

Name	Description	Key Example	Values Example
destination.ip	Destination workload instance IP address, supports single IP or CIDR	destination.ip	["10.1.2.3", "10.2.0.0/16"]
destination.port	The recipient port on the server IP address, must be in the range [0, 65535]	destination.port	["80", "443"]
destination.labels	A map of key-value pairs attached to the server instance	destination.labels[version]	["v1", "v2"]
destination.name	Destination workload instance name	destination.name	["productpage*", "-test"]
destination.namespace	Destination workload instance namespace	destination.namespace	["default"]
destination.user	The identity of the destination workload	destination.user	["bookinfo-productpage"]
request.headers	HTTP request headers, The actual header name is surrounded by brackets	request.headers[X-User]	["abc123"]

Microservices Network Security 36 © Copyright IBM Corporation 2018

IBM Training IBM

Properties keys for service role bindings

Name	Description	Key Example	Value Example
source.ip	Source workload instance IP address, supports single IP or CIDR	source.ip	"10.1.2.3"
source.namespace	Source workload instance namespace	source.namespace	"default"
source.principal	The identity of the source workload	source.principal	"cluster.local/ns/default/sa/productpage"
request.headers	HTTP request headers. The actual header name is surrounded by brackets.	request.headers[User-Agent]	"Mozilla/*"
request.auth.principal	The authenticated principal of the request	request.auth.principal	"accounts.my-svc.com/104958560606"
request.auth.audiences	The intended audience(s) for this authentication information	request.auth.audiences	"my-svc.com"
request.auth.presenter	The authorized presenter of the credential	request.auth.presenter	"123456789012.my-svc.com"
request.auth.claims	Claims from the origin JWT. The actual claim name is surrounded by brackets.	request.auth.claims[iss]	"*@foo.com"

Microservices Network Security 37 © Copyright IBM Corporation 2018

IBM Training IBM

Authorization considerations

- A ServiceRole applies to Services, a ServiceRoleBinding applies to ServiceAccounts
- A single ServiceRole can be used to configure multiple services and multiple HTTP methods
- A single service can be configured with multiple ServiceRoles
 - It may be easier to configure a single service using multiple (non-conflicting) ServiceRoles
- Multiple ServiceRoles are required to configure multiple levels of access
 - Read-only: GET
 - Read/write: GET/POST
 - Administer: GET/PATCH/PUT/DELETE
- Each client needs a ServiceRoleBinding for permission to invoke a service
 - A single ServiceRoleBinding can be used to configure multiple ServiceAccounts

```

graph LR
    SRB[Service Role Binding] --> SA1[Pod sa=dispuser<br/>proxy<br/>dispatcher]
    SA1 --> S[Service bookstore]
    S --> SA2[Pod sa=storeuser<br/>proxy<br/>bookstore]
    SRB --- S
    S --- SA2
    style SRB fill:#0070C0,color:#fff
    style S fill:#FFA500,color:#fff
    style SA1 fill:#D9E1F2,color:#0070C0
    style SA2 fill:#D9E1F2,color:#0070C0
  
```

Microservices Network Security 38 © Copyright IBM Corporation 2018

IBM Training

Selecting binding subjects

- Specifying service role binding subjects
 - user
 - source.principal
- Envoy interprets wildcards differently
- user
 - “*” all users including unauthenticated (no certificate)
- source.principal
 - “*” all authenticated users

```

apiVersion: rbac.istio.io/v1alpha1
kind: ServiceRole
metadata:
  name: tester
  namespace: appl
spec:
  rules:
    - services: ["test-*"]
      methods: "*"
    - services: ["bookstore.appl.svc.cluster.local"]
      paths: ["/reviews"]
      methods: ["GET"]

```

```

apiVersion: rbac.istio.io/v1alpha1
kind: ServiceRoleBinding
metadata:
  name: bind-microservice1-viewer
  namespace: appl
spec:
  roleRef:
    kind: ServiceRole
    name: "tester"
  subjects:
    - user: "cluster.local/ns/appl/sa/dispuser"
      properties:
        source.principal: "cluster.local/ns/appl/sa/dispuser"

```

Microservices Network Security

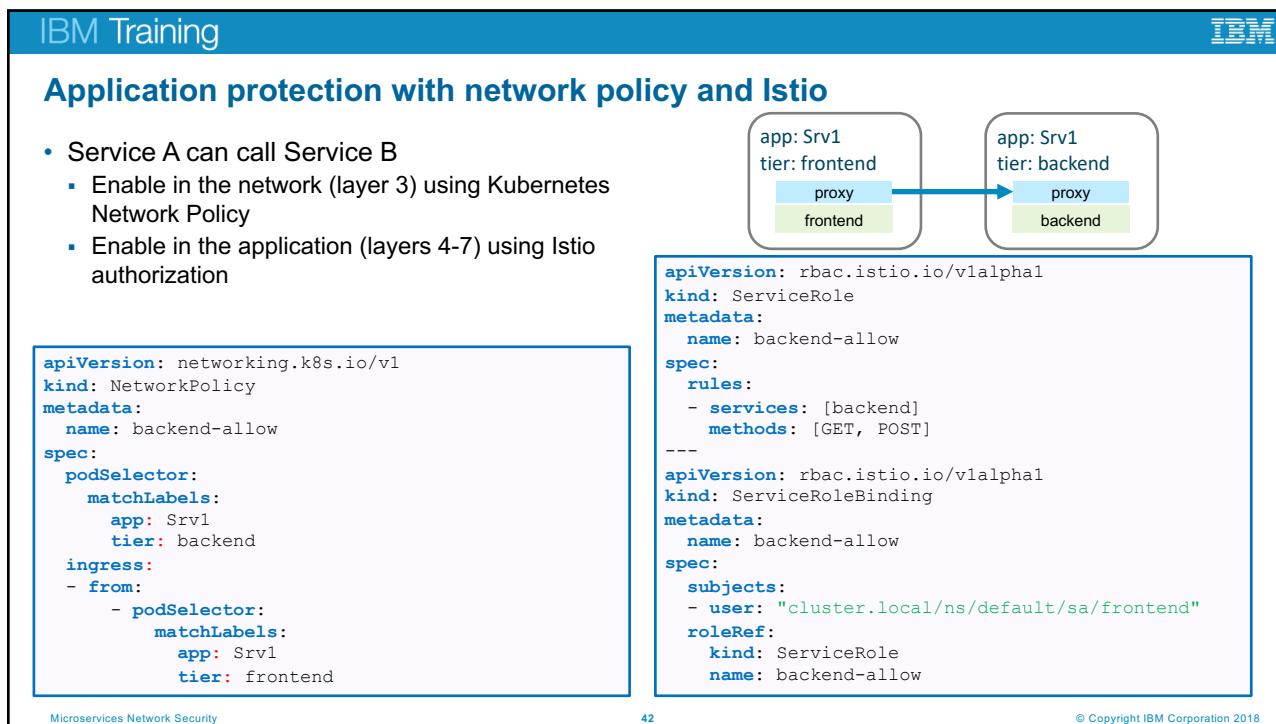
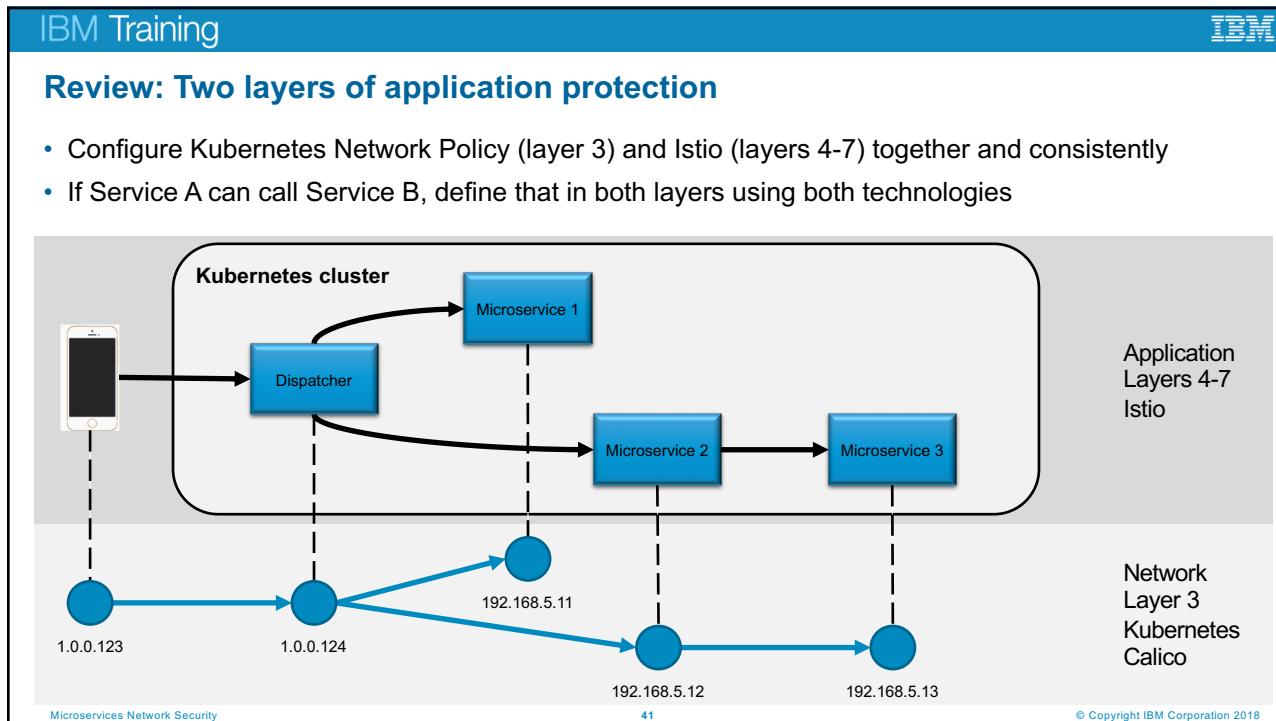
© Copyright IBM Corporation 2018

IBM Training

Bringing it all together

40

© Copyright IBM Corporation 2018



Conclusion

- Understand network security options
- Understand Kubernetes network policy basics
- Understand Network security with Istio



IBM Training



© Copyright IBM Corporation 2016. All Rights Reserved.