

Course Guide

IBM Operations Analytics Log Analysis: Scalable Collection Architecture

Course code TN631 ERC 1.0



August 2016 edition

NOTICES

This information was developed for products and services offered in the USA.

IBM may not offer the products, services, or features discussed in this document in other countries. Consult your local IBM representative for information on the products and services currently available in your area. Any reference to an IBM product, program, or service is not intended to state or imply that only that IBM product, program, or service may be used. Any functionally equivalent product, program, or service that does not infringe any IBM intellectual property right may be used instead. However, it is the user's responsibility to evaluate and verify the operation of any non-IBM product, program, or service.

IBM may have patents or pending patent applications covering subject matter described in this document. The furnishing of this document does not grant you any license to these patents. You can send license inquiries, in writing, to:

*IBM Director of Licensing
IBM Corporation
North Castle Drive, MD-NC119
Armonk, NY 10504-1785
United States of America*

The following paragraph does not apply to the United Kingdom or any other country where such provisions are inconsistent with local law: INTERNATIONAL BUSINESS MACHINES CORPORATION PROVIDES THIS PUBLICATION "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESS OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF NON-INFRINGEMENT, MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE. Some states do not allow disclaimer of express or implied warranties in certain transactions, therefore, this statement may not apply to you.

This information could include technical inaccuracies or typographical errors. Changes are periodically made to the information herein; these changes will be incorporated in new editions of the publication. IBM may make improvements and/or changes in the product(s) and/or the program(s) described in this publication at any time without notice.

Any references in this information to non-IBM websites are provided for convenience only and do not in any manner serve as an endorsement of those websites. The materials at those websites are not part of the materials for this IBM product and use of those websites is at your own risk.

IBM may use or distribute any of the information you supply in any way it believes appropriate without incurring any obligation to you.

Information concerning non-IBM products was obtained from the suppliers of those products, their published announcements or other publicly available sources. IBM has not tested those products and cannot confirm the accuracy of performance, compatibility or any other claims related to non-IBM products. Questions on the capabilities of non-IBM products should be addressed to the suppliers of those products.

This information contains examples of data and reports used in daily business operations. To illustrate them as completely as possible, the examples include the names of individuals, companies, brands, and products. All of these names are fictitious and any similarity to the names and addresses used by an actual business enterprise is entirely coincidental.

TRADEMARKS

IBM, the IBM logo, and ibm.com are trademarks or registered trademarks of International Business Machines Corp., registered in many jurisdictions worldwide. Other product and service names might be trademarks of IBM or other companies. A current list of IBM trademarks is available on the web at "Copyright and trademark information" at www.ibm.com/legal/copytrade.shtml.

Adobe, the Adobe logo, PostScript, and the PostScript logo are either registered trademarks or trademarks of Adobe Systems Incorporated in the United States, and/or other countries.

Cell Broadband Engine is a trademark of Sony Computer Entertainment, Inc. in the United States, other countries, or both and is used under license therefrom.

Intel, Intel logo, Intel Inside, Intel Inside logo, Intel Centrino, Intel Centrino logo, Celeron, Intel Xeon, Intel SpeedStep, Itanium, and Pentium are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States and other countries.

IT Infrastructure Library is a Registered Trade Mark of AXELOS Limited.

ITIL is a Registered Trade Mark of AXELOS Limited.

Java and all Java-based trademarks and logos are trademarks or registered trademarks of Oracle and/or its affiliates.

Linear Tape-Open, LTO, the LTO Logo, Ultrium, and the Ultrium logo are trademarks of HP, IBM Corp. and Quantum in the U.S. and other countries.

Linux is a registered trademark of Linus Torvalds in the United States, other countries, or both.

Microsoft, Windows, Windows NT, and the Windows logo are trademarks of Microsoft Corporation in the United States, other countries, or both.

UNIX is a registered trademark of The Open Group in the United States and other countries.

© Copyright International Business Machines Corporation 2016.

This document may not be reproduced in whole or in part without the prior written permission of IBM.

US Government Users Restricted Rights - Use, duplication or disclosure restricted by GSA ADP Schedule Contract with IBM Corp.

Contents

About this course	vii
Course objectives	ix
Audience	ix
Prerequisites	ix
Agenda	x
Course description	x
Unit 1 Overview	1-1
Unit objectives	1-2
Scalable collection architecture	1-3
Architecture components	1-4
Collection agents	1-5
Logstash	1-6
HAProxy	1-7
Kafka	1-8
Software version information	1-9
Resources included with Log Analysis	1-10
Resiliency and redundancy	1-11
About the lab exercises	1-12
Unit summary	1-13
Unit 2 Filebeat, HAProxy, and Logstash receivers.....	2-1
Unit objectives	2-2
Installing and configuring Filebeat	2-3
Other Filebeat options	2-4
Installing and configuring HAProxy	2-5
Configure HAProxy logging	2-6
HAProxy statistics page	2-7
Logstash defined	2-8
Key Logstash functions	2-9
Included Logstash plug-ins	2-10
Important plug-ins for Log Analysis	2-11
Installing Logstash	2-12
Logstash configuration	2-14
About Logstash receivers	2-15
Example of Logstash receiver configuration	2-16
Metadata for Logstash receivers	2-18
Logstash receiver redundancy	2-19
Unit summary	2-20
Student exercises	2-21
Exercise introduction: Finished lab environment	2-22
Exercise introduction: Lab environment in Unit 2 exercises	2-23

Unit 3 Kafka	3-1
Unit objectives	3-2
Kafka overview	3-3
Installing and configuring Kafka	3-5
Kafka in the scalable collection architecture	3-6
Viewing data in Kafka	3-8
Kafka Manager	3-9
Installing and using Kafka Manager	3-10
Configuring Kafka Manager	3-11
Unit summary	3-12
Student exercises	3-12
Exercise introduction: Finished lab environment	3-13
Exercise introduction: Lab environment in Unit 3 exercises	3-14
Unit 4 Logstash senders and Log Analysis core	4-1
Unit objectives	4-2
About Logstash senders	4-3
Logstash sender input section	4-4
Logstash sender filter section	4-6
Logstash sender output section	4-7
Configuring Log Analysis	4-9
Logstash sender resiliency	4-10
Unit summary	4-11
Student exercises	4-11
Exercise introduction: Finished lab environment	4-12
Exercise introduction: Lab environment in Unit 4 exercises	4-13
Unit 5 Using the Log File Agent	5-1
Unit objectives	5-2
Installing the Log File Agent	5-3
Configuring the Log File Agent	5-5
Adding a HAProxy listener for Log File Agents	5-7
Logstash receiver input section	5-8
Extra text in Log File Agent messages	5-9
Removing the extra text with the Logstash receiver	5-10
Metadata from the Log File Agent	5-12
Logstash receiver output section	5-13
Logstash sender input section	5-14
Logstash sender filter section	5-16
Logstash sender output section	5-17
Configuring Log Analysis	5-19
Unit summary	5-20
Student exercises	5-20
Exercise introduction: Finished lab environment	5-21
Exercise introduction: Lab environment in Unit 5 exercises	5-22
Unit 6 Multiline logs	6-1
Unit objectives	6-2

Multiline log example	6-3
Logstash and multiline logs	6-4
HAProxy and multiline logs	6-5
Sending data to Log Analysis with Logstash senders	6-6
Metadata and multiline messages	6-7
Unit summary	6-9
Student exercises	6-9
Exercise introduction: Finished lab environment	6-10
Exercise introduction: Lab environment in Unit 6 exercises	6-11
Unit 7 Log consolidation.....	7-1
Unit objectives	7-2
Benefits of log consolidation	7-3
Adding a custom source type	7-4
Editing the index configuration	7-5
Creating a logical data source	7-7
Configuring the collection agent: Filebeat	7-8
Configuring the collection agent: Log File Agent	7-9
Logstash sender filter section	7-10
Logstash sender output section	7-11
Searching consolidated logs	7-13
Unit summary	7-14
Student exercises	7-14
Exercise introduction: Finished lab environment	7-15
Exercise introduction: Lab environment in Unit 7 exercises	7-16
Unit 8 Troubleshooting.....	8-1
Unit objectives	8-2
Filebeat logging	8-3
Filebeat file processing	8-4
Filebeat event publishing	8-5
Filebeat connectivity	8-6
Log File Agent logging	8-7
Log File Agent log file	8-8
Log File Agent event publishing	8-9
Log File Agent connectivity	8-10
HAProxy logging	8-11
HAProxy TCP session logging	8-12
HAProxy connectivity	8-14
HAProxy statistics page	8-15
Logstash log level	8-17
Use the file output plug-in for troubleshooting	8-18
Use tags to verify filter operations	8-19
Scala output plug-in logging	8-20
Kafka troubleshooting	8-21
Zookeeper logging	8-22
Viewing data in Kafka	8-23
Kafka Manager	8-24

Using Kafka Manager to verify sender Logstash data flow	8-25
Log Analysis troubleshooting	8-26
Generic receiver accepting data	8-27
Unit summary	8-28

About this course

IBM Training

IBM



© Copyright IBM Corporation 2016
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

**IBM Operations Analytics Log Analysis:
Scalable Collection Architecture**

For enterprise-level environments, IBM Operations Analytics Log Analysis needs a reliable and scalable architecture for data collection. This course teaches you how to install and configure all of the components required for a scalable log collection architecture. In this 3-day course, you learn through lecture and hands-on exercises.

To collect log data reliably, IBM Operations Analytics Log Analysis uses several software components to ensure that log messages are not lost. During this course, you learn how to install and configure the following software:

- Filebeat, to collect log messages
- IBM Tivoli Log File Agent, to collect log messages
- HAProxy, a load balancer
- Logstash, an event pipeline
- Kafka, a message queue system
- IBM Operations Analytics Log Analysis, to index log messages and present a user search interface

The lab environment for this course uses the VMware platform.

Details	
Delivery method	Classroom or instructor-led online (ILO)
Course level	ERC 1.0 This course is a new course.
Product and version	IBM Operations Analytics Log Analysis version 1.3.3.1
Recommended duration	3 days
Skill level	Intermediate

Course objectives

- Diagram the Scalable collection architecture
- Install and configure Filebeat
- Install and configure HAProxy
- Install and configure Logstash receivers
- Install and configure Kafka
- Configure Logstash senders
- Configure Log Analysis data sources
- Install and configure the Log File Agent
- Support multi-line logs
- Consolidate physical data sources
- Troubleshoot the Scalable collection architecture

Audience

This course is designed for implementers, administrators, technical sales persons, and any others who need IBM Operations Analytics Log Analysis skills.

Prerequisites

Before taking this course, make sure that you have Linux administration skills and a working knowledge of IBM Operations Analytics Log Analysis.

Agenda

- Unit 1 Overview
- Unit 2 Filebeat, HAProxy, and Logstash receivers
- Unit 3 Kafka
- Unit 4 Logstash senders and Log Analysis core
- Unit 5 Using the Log File Agent
- Unit 6 Multi-line logs
- Unit 7 Log consolidation
- Unit 8 Troubleshooting

Agenda

Course description

The course contains the following units:

1. Overview

This unit explains the scalable collection design. You also learn about each of the software components that are used for scalable collection.

This unit has no student exercises.

2. Filebeat, HAProxy, and Logstash receivers

This unit teaches you how to install and configure Filebeat, HAProxy, and Logstash servers.

In the exercises for this unit, you install and configure a log collection agent, the HAProxy load balancer, and a Logstash receiver cluster. You configure these components to monitor a web access log from an IBM HTTP Server (IHS).

3. Kafka

This unit teaches you how to install and configure Kafka.

Apache Kafka is a high-throughput distributed messaging system. In this scalable collection architecture, Kafka acts as a durable message store with fault tolerance. Kafka requires Apache Zookeeper for coordination between brokers. In these exercises, you install Kafka and Zookeeper. You then configure the Logstash receiver cluster to send messages to Kafka and

verify the data flow. Finally, you install Kafka Manager, which allows you to view your configurations and data that Kafka is processing.

4. [Logstash senders and Log Analysis core](#)

In this unit, you learn how to install and configure Logstash senders. You also learn how to configure the IBM Operations Analytics Log Analysis core software to accept data from a Logstash sender.

Logstash senders pull messages from Kafka, then send them to the Log Analysis core software. They can also alter the messages before they send them. In these exercises, you configure two Logstash sender instances and verify that they are moving data from Kafka to Log Analysis.

5. [Using the Log File Agent](#)

You can use the IBM Tivoli Log File Agent, which is included with Log Analysis, as a collection agent. This unit teaches you how to install and configure the Log File Agent.

You can use the Log File Agent (LFA) as a collection agent to capture messages from a target log file. Earlier in this course, you installed and configured Filebeat as a collection agent. In these exercises, you install and configure a second type of log collector: the Log File Agent. You then update the other components in your lab environment to process messages from the Log File Agent.

6. [Multiline logs](#)

In this unit, you learn how to configure Logstash receivers, Logstash senders, and HAProxy load balancers to process multiline logs.

Many log files are in a multiline format. Multiline log records are single log messages that span multiple lines. To process multiline logs in the scalable collection architecture, you must ensure that all of the lines in a single message are processed together, and in the correct sequence. In the exercises for this unit, you add support for a multiline log file.

7. [Log consolidation](#)

This unit teaches you how to consolidate messages from many logs into just a few Log Analysis data sources.

In a simple configuration, IBM Log Analysis manages input from a single target log file with a single data source. Environments with many log files to monitor (hundreds or thousands) require just as many data sources. It can be problematic for administrators to create, manage, navigate, and support so many data sources. In the exercises for this unit, you alter your lab environment to consolidate several web server logs into a single Log Analysis data source.

8. [Troubleshooting](#)

In this unit, you learn how to troubleshoot each of the components in the scalable collection architecture.

This unit has no student exercises.

Unit 1 Overview

IBM Training



Overview

© Copyright IBM Corporation 2016
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

This unit explains the scalable collection design. You also learn about each of the software components that are used for scalable collection.

Unit objectives

- Describe the scalable collection architecture
- List the components in the scalable collection architecture

Scalable collection architecture

In a large environment, the method you use to collect log messages must be reliable and durable.

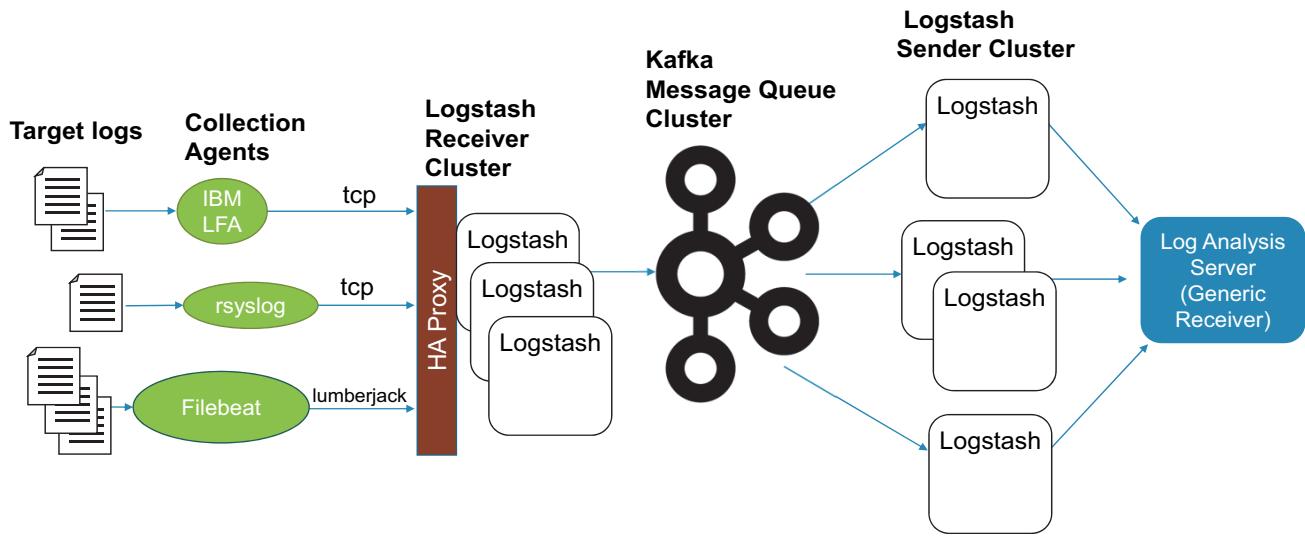
- The collection method must handle ingestion across hundreds and thousands of end points in a load-balanced and fault-tolerant manner.
- The collection method must be reliable to handle spiky log traffic, for example, several megabytes per second.
- The collection method must prevent data loss when downstream Log Analysis components are slower.
- The collection method should avoid dependence on only the IBM Log File Agent (LFA) for log shipping. Support for additional collection agents such as syslog and Filebeat should be included.

Scalable collection architecture

In large IT environments, the volume of log messages can overwhelm a log management solution. To avoid losing any log messages, the collection method that is used to collect logs must be durable, reliable, and resilient.

This course teaches you how to implement a reliable and scalable collection architecture for IBM Operations Analytics Log Analysis.

Architecture components



[Overview](#)

© Copyright IBM Corporation 2016

Architecture components

This diagram shows the flow of data through the software components that are used to reliably collect log messages.

1. Collection agents such as the Log File Agent and Filebeat collect log messages from target log files and send them to an HAProxy load balancer.
2. The HAProxy load balancer forwards log messages to a cluster of Logstash receiver instances.
3. Logstash receivers perform minimal processing on log messages and send them to the Kafka cluster.
4. The Kafka cluster acts as a durable data store of log messages. After messages arrive in Kafka, the messages are stored on disk.
5. Logstash senders retrieve messages from the Kafka cluster, process them to add or transform the data in the messages and send them to the Log Analysis core software.
6. The Log Analysis core software indexes the interesting text within each log message and presents a search interface to users.

Collection agents

- **Log File Agent:** Reads log messages from target files and sends data to HAProxy with TCP.
- **Filebeat:** Reads log messages from target files and sends data to HAProxy with lumberjack, a proprietary protocol.
- **Syslog and rsyslog:** Send data to HAProxy with TCP.

For downstream processing, collection agents should add metadata to log messages.

Collection agents

Collection agents read messages from target log files and send them to an HAProxy load balancer.

The Log File Agent and Filebeat are included with IBM Log Analysis. Collection agents must be able to add metadata to each log message that they send because the other components in the collection design use that metadata for processing and conditional logic.

Logstash

- Logstash acts as a log shipping agent or as centralized log server in a standalone or hierarchical architecture
- As a shipping agent, it generates events from input plug-ins, modifies them through filters, and sends them elsewhere based on output plug-ins
- Logstash receivers get messages from the collection agents and send them to Kafka
- Logstash senders pull messages from Kafka, alter the messages if necessary, and send the messages to the Log Analysis server



Logstash

Logstash is open source software that accepts text as input, processes text with filters, and sends text in a specific output format. In this scalable architecture, Logstash servers are used as receivers and senders:

- Receiver Logstash servers accept log messages from collection agents and store them in the Kafka cluster.
- Sender Logstash servers retrieve log messages from the Kafka message queue and send them to the Log Analysis core software.

HAProxy

- HAProxy is an open source load balancer and proxy server for TCP and HTTP-based applications that spreads requests across multiple servers.
- HAProxy has a reputation for being fast and efficient.
- HAProxy makes nonredundant services redundant.
- Health check
 - HAProxy uses health checks to determine if a back-end server is available to process requests.
 - This avoids having to manually remove a server from the back end if it becomes unavailable.
 - The default health check is to try to establish a TCP connection to the server to verify if the back-end server is listening on the configured IP address and port.
- Sticky sessions
 - Some applications require stickiness between a client and a server: it means all the requests from a client must be sent to the same server
 - Using a dedicated load balancing algorithm: A hash on the source IP

Overview

© Copyright IBM Corporation 2016

HAProxy

HAProxy is a software load balancer. In the Log Analysis scalable architecture, its role is to provide redundancy for Logstash receivers.

For Logstash receivers that process single-line messages, HAProxy sends data to each Logstash receiver in turns. For Logstash receivers that process multiline messages, HAProxy sends data to the same Logstash receiver, if the Logstash receiver is running.

Kafka

Kafka is a persistent, distributed, replicated publish–subscribe messaging system.

- Kafka is fast, scalable, and durable:
 - Fast: A single Kafka broker can handle hundreds of megabytes of reads and writes per second from thousands of clients.
 - Scalable: With Kafka, a single cluster can serve as the central data backbone for a large organization. It can be transparently expanded without downtime.
 - Durable: Messages are persisted on disk and replicated within the cluster to prevent data loss.
- Logstash receivers write data to Kafka topics for various data sources and Logstash senders read from Kafka.
- The basic unit of concurrency in Kafka is a **partition**.
A frequent practice is to create a topic for each logical data source and partitions within that topic for each physical data source.

Overview

© Copyright IBM Corporation 2016

Kafka

Kafka is an open source message queue system. In the Log Analysis scalable architecture, its role is to provide reliable storage for log messages.

Kafka is a distributed system. Kafka servers that store data are called **brokers**.

Apache Zookeeper is used to coordinate communication between different Kafka brokers. When you install Kafka, the Zookeeper software is automatically installed also.

Software version information

Component	Version
IBM Operations Analytics Log Analysis	1.3.3 Fix Pack 1
Kafka	0.8.2.2
Logstash	2.2.1
HAProxy	1.5.4

[Overview](#)

© Copyright IBM Corporation 2016

Software version information

This table lists the version of each software component that is supported with IBM Operations Analytics Log Analysis v1.3.3.1.

Resources included with Log Analysis

- These directories contain installation files:
 - <LA_HOME>/logstash-2.2.1
 - <LA_HOME>/filebeat
 - <LA_HOME>/kafka
- You can install Logstash and the Log File Agent with the remote installation tool:
 - <LA_HOME>/remote_install_tool
- You can find sample configuration files for the following components in the <LA_HOME>/kafka/test-configs directory:
 - Filebeat
 - HAProxy
 - Kafka
 - Zookeeper
 - Log File Agent
 - Logstash receivers
 - Logstash senders

Overview

© Copyright IBM Corporation 2016

Resources included with Log Analysis

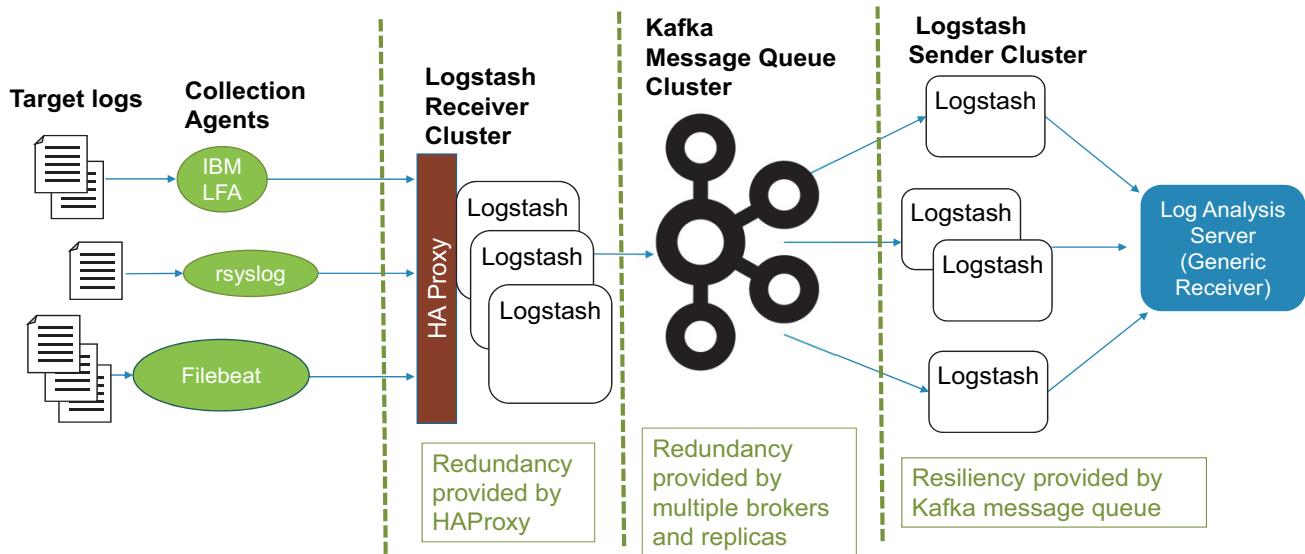
When you install IBM Operations Analytics Log Analysis, the following open source software installation files are included:

- Logstash
- Filebeat
- Kafka

Sample configuration files for the following components are included when you install IBM Operations Analytics Log Analysis:

- Filebeat
- HAProxy
- Kafka
- Zookeeper
- Log File Agent
- Logstash receivers
- Logstash senders

Resiliency and redundancy



Overview

© Copyright IBM Corporation 2016

Resiliency and redundancy

This diagram shows how the scalable collection architecture delivers redundancy and resiliency.

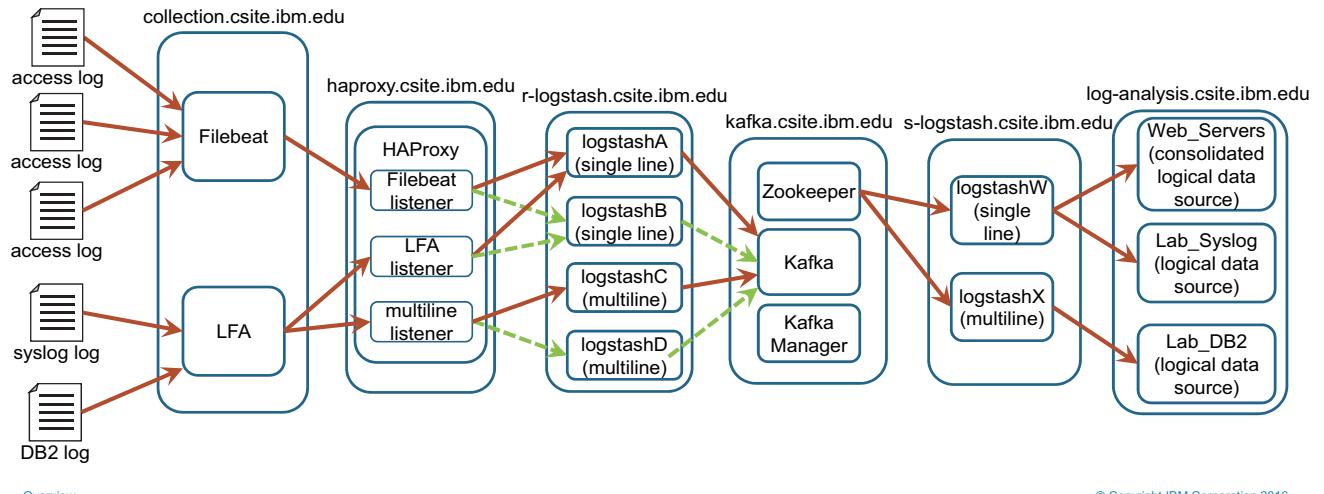
HAProxy ensures that no log messages are lost by adding redundancy to a cluster of Logstash receivers. If one Logstash receiver is down, HAProxy sends log messages to a different Logstash receiver server in the cluster.

The Kafka cluster provides reliable storage for log messages. If one broker in the Kafka cluster is down, messages can still be read from other brokers in the cluster.

Resiliency for Logstash senders is provided by the Kafka cluster. If a Logstash sender server goes down, it can read messages that are stored in the Kafka message queue when it comes back up.

About the lab exercises

- Completing the exercises is vital to understanding the scalable collection architecture
- When you are finished with the lab exercises, your environment looks like this:



[Overview](#)

© Copyright IBM Corporation 2016

About the lab exercises

In the student exercises for this course, you build an environment for scalable log message collection. You install and configure several software components and connect them together. You also verify data flow between the components.

The student exercises for this course teach you how each component works and the role of each component in the scalable collection design.

Unit summary

- Describe the scalable collection architecture
- List the components in the scalable collection architecture

Unit 2 Filebeat, HAProxy, and Logstash receivers

IBM Training



Filebeat, HAProxy, and Logstash receivers

© Copyright IBM Corporation 2016
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

This unit teaches you how to install and configure Filebeat, HAProxy, and Logstash servers.

Unit objectives

- Install and configure Filebeat
- Install and configure HAProxy
- Install and configure Logstash receivers

Installing and configuring Filebeat

- Install Filebeat by decompressing the installation file.
- Configure Filebeat with the `filebeat.yml` file.

```
filebeat:  
  prospectors:  
    -  
      paths:  
        - /opt/IBM/HTTPServer/logs/access.log  
      input_type: log  
      fields:  
        collector: filebeats-collection.csuite.ibm.edu  
        env: DEV  
        module: IBM-HTTP-Server  
        type: access-log  
        site: DALLAS  
        platform: RHEL  
      output:  
        logstash:  
          hosts: ["ha-proxy.csuite.ibm.edu:20737"]  
      shipper:  
      logging:  
        to_files: true  
        files:  
          path: /opt/filebeat-1.1.1-x86_64/  
          name: mybeat  
          rotateeverybytes: 10485760  
      level: info
```

Filebeat, HAProxy, and Logstash receivers

© Copyright IBM Corporation 2016

Installing and configuring Filebeat

To install a Filebeat server, decompress the installation files. To configure Filebeat, you edit a file named `filebeat.yml`.

In this example, Filebeat is monitoring a log file named `/opt/IBM/HTTPServer/logs/access.log`.

In this example, Filebeat adds the following metadata to each log message:

- collector
- env
- module
- type
- site
- platform

In this example, Filebeat is sending log messages to a host named **ha-proxy.csuite.ibm.edu** on port 20737. The messages are sent in a format that Logstash servers can interpret.

In this example, Filebeat is logging to a file named `/opt/filebeat-1.1.1-x86_64/mybeat`.



Hint: The indentations in the `filebeat.yml` file are very important. This file is indented with spaces, not tabs.

Other Filebeat options

- Include or exclude lines from a file
- Exclude files
- Scan frequency
- Multiline
- File as output
- Transport Layer Security

Other Filebeat options

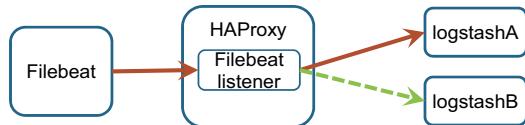
You can also configure advanced options for Filebeat servers. The following list describes other Filebeat options:

- You can configure Filebeat to include or exclude individual lines from a log file based on a regular expression match.
- You can exclude files in a directory from monitoring.
- You can increase or decrease the interval that Filebeat checks for new files.
- You can control how Filebeat deals with log messages that span multiple lines.
- You can send log messages that Filebeat collects to a file. This option is useful for troubleshooting.
- You can configure secure transport between Filebeat and Logstash.

Installing and configuring HAProxy

There are two options to install HAProxy:

- Install with a package
`yum install haproxy`
- Install from source



```
...  
defaults  
    mode          tcp  
    log           global  
    option        tcplog  
    option        dontlognull  
...  
## Listener for Filebeat Logstash cluster  
## Specify Logstash instances for processing data from Filebeat  
listen  Beats_Cluster ha-proxy.csuite.ibm.edu:20737  
    mode tcp  
        balance roundrobin  
        server receiver-logstashA logstashA.csuite.ibm.edu:18737 check fall 2 rise 3 inter 1000  
        server receiver-logstashB logstashB.csuite.ibm.edu:18738 check fall 2 rise 3 inter 1000
```

Filebeat, HAProxy, and Logstash receivers

© Copyright IBM Corporation 2016

Installing and configuring HAProxy

You can install HAProxy as an operating system package, or you can compile it from the source code.

You configure HAProxy by editing a file named `/etc/haproxy/haproxy.cfg`. In this scalable architecture, HAProxy must run in tcp mode.

You must configure a listener for HAProxy to accept data from a log collection agent and forward it to a Logstash receiver cluster. In this example, HAProxy is listening on port 20737. Any data HAProxy receives on port 20737 is forwarded to either `logstashA.csuite.ibm.edu:18737` or `logstashB.csuite.ibm.edu:18738`.

HAProxy is aware of the state of each server. In this example, HAProxy performs health checks on each Logstash server every 1000 milliseconds. HAProxy considers a Logstash server down if two consecutive health check failures occur (`fall 2`). HAProxy considers a Logstash server up again after three consecutive successful health checks (`rise 3`).

The line `balance roundrobin` configures HAProxy to send data to each Logstash server in turns.

Configure HAProxy logging

HAProxy logs to syslog or rsyslog. Follow these steps to enable logging for your HAProxy server:

1. Configure logging options in `haproxy.cfg`.
2. Configure rsyslog to receive UDP data.
3. Configure rsyslog to send haproxy messages to a log file.
4. Restart HAProxy
5. Restart rsyslog.

```
global /etc/haproxy/haproxy.cfg
...
    log      127.0.0.1 local0
...
defaults
...
    log          global
    option      tcplog
    option      dontlognull
```

```
... /etc/rsyslog.conf
# Provides UDP syslog reception
$ModLoad imudp
$UDPServerRun 514
$UDPServerAddress 127.0.0.1
...
```

```
... /etc/rsyslog.d/haproxy.conf
if ($programname == 'haproxy') then -/var/log/haproxy.log
```

Filebeat, HAProxy, and Logstash receivers

© Copyright IBM Corporation 2016

Configure HAProxy logging

HAProxy logs to syslog or rsyslog. To enable logging, you must configure HAProxy and your syslog or rsyslog server. In this example, HAProxy is logging to a local rsyslog server.

HAProxy statistics page

Statistics Report for pid 1707

> General process information

```

pid = 1707 (process #1, nbproc = 1)
User = darrin
System calls = memress = unlimited; ulimit -n = 8034
maxsock = 8034; maxconn = 4000; maxpipes = 0
current conn = 1; current pipes = 0/0; conn rate = 1/sec
Running tasks = 1/7; idle = 100 %

active UP
active UP, going down
active DOWN, going up
active or backup DOWN, going up
active or backup DOWN for maintenance (MAINT)
active or backup SOFT STOPPED for maintenance
Note: "NOLBY=DRAIN" = UP with load-balancing disabled.

Display option: [Scope: [ ]]
  • Primary site
  • Hide DOWN servers
  • Refresh now
  • CSV export
External resources:
  • Primary site
  • Updates (v1.5)
  • Online manual

```

Beats_Cluster		Sessions										Bytes										Denied										Errors										Warnings										Status										Server									
		Queue	Session rate	Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	Lb	Tot	Last	In	Out	Req	Resp	Req	Conn	Req	Retr	Redis	Status	LastChk	Wght	Act	Bck	Chk	Dwn	Dwtm	Thrte																																					
Frontend				0	1	-	0	2	3	000	21		392	615	1	314	0	0	0																																																				
receiver-logstashA	0	0	0	1	0	1	0	1	-	11	11	5m20s	379	678	1	128	0	0	0	0	0	0	0	0	0	21m18s	UP	L4OK in 0ms	1	Y	-	1	1	7m43s	-																																				
receiver-logstashB	0	0	0	1	0	1	-	10	10	6m25s	12	937	186	0	0	0	0	0	0	0	0	0	0	0	0	3m DOWN	L4CON in 0ms	1	Y	-	3	2	11m59s	-																																					
Backend	0	0	0	1	0	2	300	21	21	5m20s	392	615	1	314	0	0	0	0	0	0	0	0	0	0	0	21m18s	UP		1	1	0	1	7m43s																																						

stats		Sessions										Bytes										Denied										Errors										Warnings										Status										Server									
		Queue	Session rate	Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	Lb	Tot	Last	In	Out	Req	Resp	Req	Conn	Req	Retr	Redis	Status	LastChk	Wght	Act	Bck	Chk	Dwn	Dwtm	Thrte																																					
Frontend				1	2	-	1	1	3	000	18		24	575	1	095	239	0	0	0																																																			
Backend	0	0	0	2	0	1	300	3	0	0s	24	575	1	095	239	0	0	3	0	0	0	0	29m2s	UP		0	0	0	0																																										

```

stats socket /var/lib/haproxy/stats
...
listen stats :9000 #Listen on localhost port 9000
mode http
stats enable
stats hide-version
stats realm Haproxy\ Statistics
stats uri /haproxy_stats
stats auth netcool:object00

```

/etc/haproxy/haproxy.cfg

Filebeat, HAProxy, and Logstash receivers

© Copyright IBM Corporation 2016

HAProxy statistics page

You can configure HAProxy to display a web page that shows the status of each server that it is forwarding data to. The page also shows statistics about each of the endpoints in the cluster, such as the volume of data sent. You configure this web page in the `/etc/haproxy/haproxy.cfg` file.

This slide shows an example of the lines you add to the `/etc/haproxy/haproxy.cfg` file to enable the HAProxy statistics page.

The following table describes each line in the configuration example.

Configuration option	Description
listen stats :9000	This line configures the statistics page to use port 9000.
mode http	This line enables http mode.
stats enable	This line enables statistics.
stats hide-version	This line hides the version of HAProxy so it is not visible on the web page.
stats realm Haproxy\ Statistics	This text is what users see when they are prompted to log in to the page.
stats uri /haproxy_stats	The URI of the web page, for example, <code>http://localhost:9000/haproxy_stats</code>
stats auth netcool:object00	This line enables authentication. It also sets a user name and password. In this example, the user name is <code>netcool</code> and the password is <code>object00</code> .

Logstash defined

- Logstash acts as a log shipping agent or as centralized log server in a stand-alone or hierarchical architecture.
- As a shipping agent, it generates events from input plug-ins, modifies them through filters, and sends them elsewhere based on output plug-ins.



Logstash defined

Logstash is open source software that accepts text as input, process text with filters, and sends text in a specific output format. In this scalable architecture, Logstash servers are used as receivers and senders.

Receiver Logstash servers accept log messages from collection agents and store them in the Kafka cluster.

Logstash uses plug-ins for inputs, filters, and outputs.

Key Logstash functions

- Logstash is an event pipeline:
 - Inputs → filters/codecs → outputs
 - Inputs capture events.
 - Codecs and filters modify them.
 - Outputs send them somewhere.
- Inputs can declare a type for each stream (for example, type=WAS)
 - Inputs are used mainly for filter activation.
 - They persist with the event.
- Tags can be added throughout event processing:
 - Use for debugging purposes
 - Can be used to specify an order for event processing
- Conditional logic allows events to be processed in many ways.

Key Logstash functions

Logstash processes messages in an event pipeline. Each log message is treated as an individual event.

Logstash accepts log messages with input plug-ins. After Logstash receives a message, it can alter the text within that message with filter plug-ins. These filters can modify the message itself, or it can add and alter metadata that arrived with the log message.

Logstash sends the text in log messages to a destination with output plug-ins.

At any step in Logstash processing, you can add tags to the original log message. You can use these tags, along with metadata in the log message, to add conditional logic to the text processing pipeline. For example, you can alter the time stamp for one type of log message, while the time stamp for a different type of log message remains unchanged.

Included Logstash plug-ins

inputs	codecs	filters	outputs
<ul style="list-style-type: none"> • rabbitmq • collectd • drupal_dblog • elasticsearch • eventlog • exec • file • ganglia • gelf • gemfire • generator • graphite • heroku • imap • invalid_input • irc • jmx • log4j • lumberjack • pipe • puppet_facter 	<ul style="list-style-type: none"> • rackspace • redis • relp • s3 • snmptrap • sqlite • sqs • stdin • stomp • syslog • tcp • twitter • udp • unix • varnishlog • websocket • wmi • xmpp • zenoss • zeromq 	<ul style="list-style-type: none"> • advisor • cloudtrail • collectd • compress_spooler • dots • edn • edn_lines • fluent • graphite • json • json_lines • json_spooler • line • drop • msgpack • multiline • netflow • noop • oldlogstashjson • plain • rubydebug • spool 	<ul style="list-style-type: none"> • metaevent • metrics • multiline • mutate • noop • prune • punct • railsparrallelrequest • range • ruby • sleep • split • sumnumbers • syslog_pri • throttle • translate • unique • urldecode • useragent • uid • wms • wmts • xml • zeromq

Filebeat, HAProxy, and Logstash receivers

© Copyright IBM Corporation 2016

Included Logstash plug-ins

This slide is a partial list of Logstash plug-ins of inputs, filters, and outputs. Codecs are also listed on this slide. As open source software, plug-ins are published and maintained by a community. If necessary, you can create your own Logstash plug-ins.

Important plug-ins for Log Analysis

- Inputs
 - **file** reads text in a file.
 - **tcp** is used for accepting text over a TCP socket.
 - **kafka** is used to retrieve data from Kafka.
 - **beats** is used to accept messages from Filebeat.
- Filters
 - **drop** deletes text that is not useful.
 - **date** converts the time stamp.
 - **grok** annotates data in unstructured text.
 - **mutate** adds, modifies, or deletes fields and the data in those fields.
- Outputs
 - **file** used for testing and troubleshooting.
 - **kafka** sends data to a Kafka broker.
 - **scala** publishes data in the format that Log Analysis expects.

Important plug-ins for Log Analysis

This slide lists Logstash plug-ins that you use for Logstash receiver and Logstash sender servers. The following list provides examples of how you might use these plug-ins.

- Inputs:
 - Logstash receivers accept data from collection agents. Use the **tcp** input plug-in to accept data from a Log File Agent. Use the **beats** input plug-in to accept data from a Filebeat server.
 - Logstash senders pull data from a Kafka cluster. Use the **kafka** input plug-in to retrieve data from a Kafka cluster.
- Filters:
 - Logstash receivers must construct a topic ID for the **kafka** output. Use the **mutate** filter plug-in to build a topic ID from metadata that is attached to each log message.
 - Logstash senders must send a value for host and path to the Log Analysis server. Use the **mutate** filter plug-in to build values for host and path from metadata that is attached to each log message.
- Outputs:
 - Logstash receivers send messages to a Kafka cluster. Use the **kafka** output plug-in to send data to a list of Kafka brokers.
 - Logstash senders send messages to the Log Analysis server. Use the **scala** output plug-in to send data to Log Analysis in the correct format.
 - Use the **file** output plug-in to troubleshoot Logstash receivers and senders.

Installing Logstash

There are three options to install Logstash:

- Decompress installation files (included with Log Analysis)
- Install with Log Analysis remote installer tool (included with Log Analysis)
- Install from RPM

Installing Logstash

You can install Logstash with the methods listed on this slide.

Installing Logstash by decompressing the installation file

If you install Logstash by decompressing the installation file, you must start and stop manually. The user who decompressed the installation file owns Logstash. You can create the Logstash configuration file anywhere the owner of Logstash has read and write permissions.

Installing Logstash with the Log Analysis remote installer tool (included with Log Analysis)

Use the following steps to install Logstash with the Log Analysis remote installer tool:

1. Change to the <LA_HOME>/remote_install_tool directory
2. Edit the `ssh-config.properties` file in the config subdirectory
3. Start the remote installation tool with the command `install.sh` and follow the prompts to install Logstash.

After you install Logstash with the Log Analysis remote installer tool, your Logstash instance has the following attributes:

- You must start and stop Logstash with the logstash-util.sh utility.

`<LS_DIR>/utilities/logstash-util.sh`

- The configuration file is in the following directory:..

`<LS_DIR>/Logstash/logstash-2.2.1/logstash-scala/logstash/config/logstash-scala.conf`

- Log files are in the following directories.

`/opt/logstashA/Logstash/logs/ls-console.log`

`/opt/logstashA/Logstash/logs/scala_logstash.log`

Installing Logstash from an RPM

Use the following steps to install Logstash from an RPM Package Manager file (RPM):

1. Switch to the root user.
2. Run the following command to install Logstash.

`rpm -i logstash-2.2.1-1.noarch.rpm`

After you install Logstash from an RPM, your Logstash instance has the following attributes:

- Logstash is installed in the `/opt/logstash/` directory
- A user and a group named logstash is created
- You must start and stop Logstash with a script in the `/etc/init.d` directory
`/etc/init.d/logstash`
- Your configuration file must be in the following directory:
`/etc/logstash/conf.d/`
- Log files are in the `/var/log/logstash/` directory.

Logstash configuration

Logstash is configured with a `.conf` file.

- Logstash will not start without a valid configuration file.
- The configuration file must contain the following sections:
 - input
 - filter
 - output
- After you change the configuration, you must restart Logstash.

Logstash configuration

You configure Logstash with a text file. The configuration file name must have the extension `.conf`.

A Logstash configuration file must have an input section, a filter section, and an output section. You learn how to create a configuration file during this course.

About Logstash receivers

- Logstash receivers accept text from collection agents such as Filebeat or Log File Agent.
- Logstash receivers perform minimal processing on each message.
- Logstash receivers send messages to Kafka brokers.

About Logstash receivers

In the Log Analysis scalable architecture, the role of a Logstash receiver instance is to accept log messages from collection agents and send the messages to Kafka. Multiple Logstash receivers provide redundancy: if one Logstash receiver is down, other Logstash receivers can continue to accept data.

Logstash receivers should perform only minimal processing and alteration of each log message, so they can quickly send messages to Kafka.

Example of Logstash receiver configuration

```

input {
    beats {
        port => 18737
    } #end beats input
} #end input section

filter {
    if [fields][collector] == "filebeats-collection.csuite.ibm.edu" {
        mutate {
            add_field => [ "datasource", "%{[fields][env]}_%{[fields][module]}_%{[fields][type]}" ]
            add_field => [ "resourceID", "%{[beat][hostname]}_%{source}_1" ]
            add_tag => "mutate_filebeat"
        }# end mutate
    }# end filebeat condition
} #end filter section

output {
    if ("mutate_filebeat" in [tags]) and ! ("_grokparsefailure" in [tags]) {
        kafka {
            bootstrap_servers =>"kafka.csuite.ibm.edu:17991"
            topic_id => "%{datasource}"
            message_key => "%{resourceID}"
        } #end Kafka output
    }#end Kafka condition
} # end output section

```

Filebeat, HAProxy, and Logstash receivers

© Copyright IBM Corporation 2016

Example of Logstash receiver configuration

In this example configuration, a Logstash receiver server is accepting data from a Filebeat server with the **beats** input plug-in. The Logstash receiver adds metadata to each message with the **mutate** filter plug-in. Finally, the Logstash receiver sends log messages to a Kafka broker with the **kafka** output plug-in.

Key fields in the input section

```

input {
    beats {
        port => 18737
    } #end beats input
} #end input section

```

The diagram highlights the `input {` line with a red box and a callout pointing to it with the text "This line declares the start of the input section". It also highlights the `beats {` and `port => 18737` lines with a red box and a callout pointing to them with the text "These lines invoke the beats input plug in and accepts input from port 18737".

Figure 1 Key fields in the input section

Key fields in the filter section

```

This condition allows only messages from
the Filebeat collector into the filter

filter {
    if [fields][collector] == "filebeats-collection.csuite.ibm.edu" {
        mutate {
            add_field => [ "datasource",
                "%{[fields][env]} %{[fields][module]} %{[fields][type]}" ]
            add_field => [ "resourceID", "%{[beat][hostname]} %{source}_1" ]
            add_tag => "mutate_filebeat"
        }# end mutate
    } #end filebeat condition
} #end filter section

This tag is added to messages
that were successfully altered
by the filter

```

This condition allows only messages from the Filebeat collector into the filter

These two lines create two new fields: **datasource** and **resourceID**. These fields are used to organize data in Kafka.

Figure 2 Key fields in the filter section

Key fields in the output section

```

output {
    if ("mutate_filebeat" in [tags]) and ! ("_grokparsefailure" in [tags]) {
        kafka {
            bootstrap_servers =>"kafka.csuite.ibm.edu:17991"
            topic_id => "%{datasource}"
            message_key => "%{resourceID}"
        }# end Kafka output
    }#end Kafka condition
} # end output section

These lines use the meta data you added earlier to
create a topic in Kafka. The topic is named using the
value of %{datasource}. All messages with the same
resourceID are sent to the same partition within the topic.

```

This condition ensures that only messages that have been altered by the filebeat mutate filter and do not have any failure tags are sent as output

This line invokes the kafka output plug in.

This line sets the host and port of the Kafka broker

Figure 3 Key fields in the output section

Metadata for Logstash receivers

- Logstash receivers create topics in Kafka
- The filter section of a Logstash receiver uses data from the collection agent to set the Kafka topic ID and message key
- In this example:
 - The topic ID is **DEV_IBM-HTTP-Server_access-log**
 - The message key is **collection.csuite.ibm.edu/_opt/IBM/HTTPServer/logs/access.log_1**

```
...
filter {
    if [fields][collector] == "filebeats-collection.csuite.ibm.edu" {
        mutate {
            add_field => [ "datasource", "%{[fields][env]}_%{[fields][module]}_%{[fields][type]}" ]
            add_field => [ "resourceID", "%{[beat][hostname]}_%{source}_1" ]
            add_tag => "mutate_filebeat"
        }# end mutate
    }# end filebeat condition
}# end filter section

output {
    if ("mutate_filebeat" in [tags]) and ! ("_grokparsefailure" in [tags]) {
        kafka {
            bootstrap_servers =>"kafka.csuite.ibm.edu:17991"
            topic_id => "%{datasource}"
            message_key => "%{resourceID}"
        }# end Kafka output
    }# end Kafka condition
}# end output section
}
```

```
filebeat:
prospectors:
-
paths:
- /opt/IBM/HTTPServer/logs/access.log
input_type: log
fields:
    collector: filebeats-collection.csuite.ibm.edu
    env: DEV
    module: IBM-HTTP-Server
    type: access-log
    site: DALLAS
    platform: RHEL
...

```

filebeat.yml

logstashreceiver.conf

Filebeat, HAProxy, and Logstash receivers

© Copyright IBM Corporation 2016

Metadata for Logstash receivers

In the example on this slide, the Logstash receiver is using metadata that was added to each log message by the Filebeat collection agent.

The mutate filter adds a field named **datasource**. The value of **datasource** is constructed by using the value of the Filebeat field **env**, followed by an underscore character, followed by the value of the Filebeat field **module**, followed by an underscore character, followed by the value of the Filebeat field **type**. In this example, the value of **datasource** is **DEV_IBM-HTTP-Server_access-log**.

The mutate filter also adds a field named **resourceID**. The value of **resourceID** is constructed by using the value of the Filebeat metadata keyword **hostname**, followed by an underscore character, followed by the value of the metadata keyword **source**, followed by the string **_1**. In this example, the value of **resourceID** is

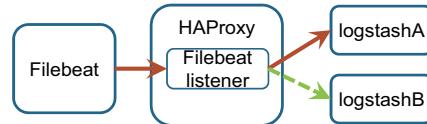
`collection.csuite.ibm.edu/_opt/IBM/HTTPServer/logs/access.log_1`.

The two fields that were created by the mutate filter, **datasource** and **resourceID**, are used in the configuration of the **kafka** output plug-in.

- The value of **datasource** is the name of the Kafka topic that this Logstash receiver writes messages to. Notice that the value of **datasource** is used as the **topic_id**.
- Any messages with the same **message_key** are written to the same partition within the topic. Notice that the value of **resourceID** is used as the **message_key**.

Logstash receiver redundancy

```
...
filebeat.yml
output:
  logstash:
    hosts: ["ha-proxy.csuite.ibm.edu:20737"]
...
```



```
listen Beats_Cluster ha-proxy.csuite.ibm.edu:20737
mode tcp
balance roundrobin
server receiver-logstashA logstashA.csuite.ibm.edu:18737 check fall 2 rise 3 inter 1000
server receiver-logstashB logstashB.csuite.ibm.edu:18738 check fall 2 rise 3 inter 1000
haproxy.cfg
```

```
input {
  beats {
    port => 18737
  } #end beats input
} #end input section
...
output {
  kafka {
    bootstrap_servers =>"kafka.csuite.ibm.edu:17991"
    topic_id => "%{datasource}"
    message_key => "%{resourceID}"
  } #end Kafka output
} # end output section
logstashReceiverA.conf
```

```
input {
  beats {
    port => 18738
  } #end beats input
} #end input section
...
output {
  kafka {
    bootstrap_servers =>"kafka.csuite.ibm.edu:17991"
    topic_id => "%{datasource}"
    message_key => "%{resourceID}"
  } #end Kafka output
} # end output section
logstashReceiverB.conf
```

Filebeat, HAProxy, and Logstash receivers

© Copyright IBM Corporation 2016

Logstash receiver redundancy

In a scalable collection environment, you should install multiple Logstash receiver servers. You use HAProxy to make your Logstash receiver servers run as a cluster for failover and redundancy.

In this example, the Filebeat collection agent sends messages from the target log file to ha-proxy.csuite.ibm.edu:20737.

HAProxy listens for traffic on port 20737. HAProxy forwards traffic from the Filebeat server on port 20737 to either logstashA.csuite.ibm.edu:18737 or logstashB.csuite.ibm.edu:18738.

Both Logstash receiver servers have identical configurations for the kafka output plug-in. Both Logstash receiver servers send data to the same list of Kafka brokers (kafka.csuite.ibm.edu:17991). This means if one Logstash receiver is down, HAProxy sends data to the other Logstash receiver, and the data is still written to the Kafka cluster.

Unit summary

- Install and configure Filebeat
- Install and configure HAProxy
- Install and configure Logstash receivers

Filebeat, HAProxy, and Logstash receivers exercises

Exercise 1 Installing and configuring Filebeat

Exercise 2 Installing and configuring HAProxy

Exercise 3 Installing and configuring the first Logstash receiver

Exercise 4 Verifying data flow

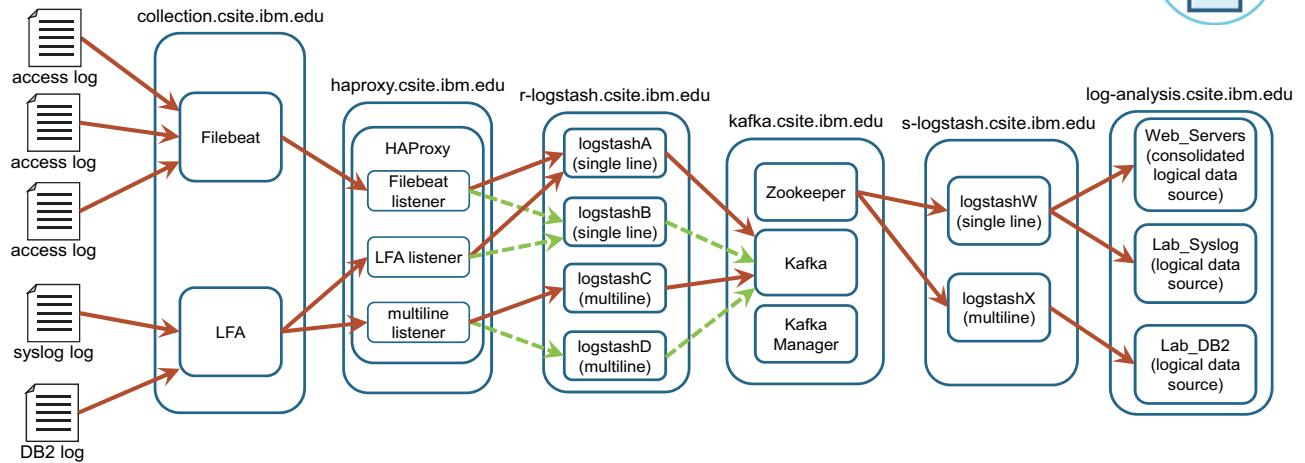
Exercise 5 Adding metadata to the log message

Exercise 6 Installing and configuring the second Logstash receiver

Perform the exercises for this unit.



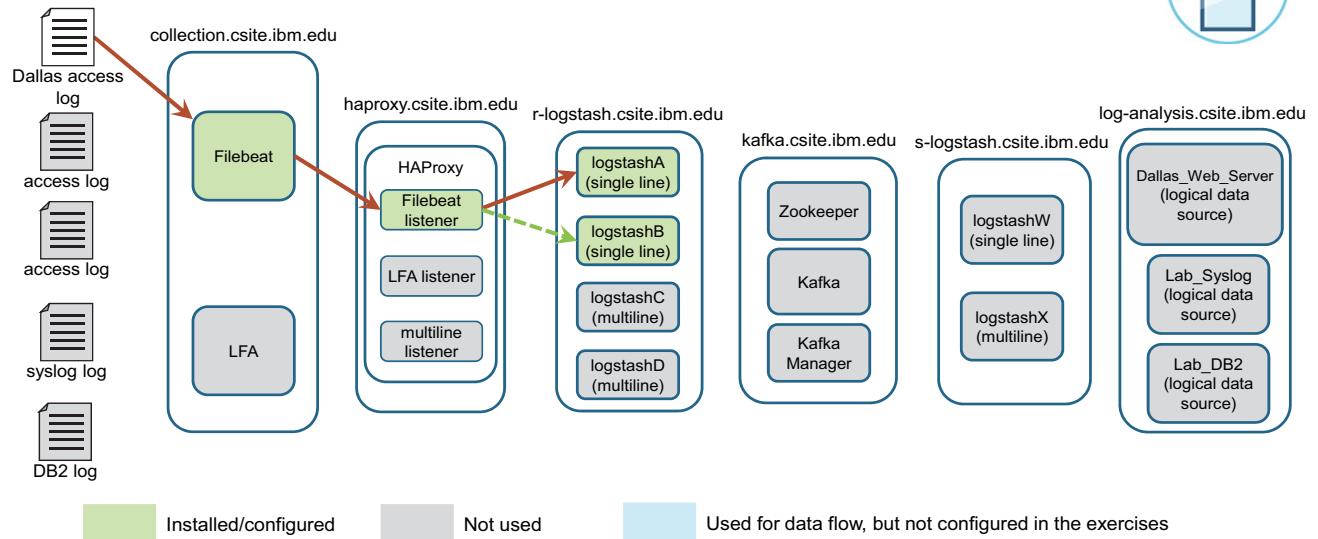
Exercise introduction: Finished lab environment



Exercise introduction: Finished lab environment

At the end of this course, your lab environment will look like this diagram. Refer to this diagram as you complete the steps in the student exercises to help you understand what you are installing and configuring.

Exercise introduction: Lab environment in Unit 2 exercises



Exercise introduction: Lab environment in Unit 2 exercises

In the exercises for this unit, you install and configure the following components:

- A Filebeat server, to monitor a web server access log
- An HAProxy load balancer, with a listener for traffic from the Filebeat server
- Two Logstash receiver servers, to accept messages from the Filebeat server

As you install and configure these components, you also verify the data flow between them.

Unit 3 Kafka

IBM Training

IBM

Kafka

© Copyright IBM Corporation 2016
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

This unit teaches you how to install and configure Kafka.

Unit objectives

- Install and configure Kafka
- Diagram data flow from a Logstash receiver server to Kafka
- Install and configure Kafka Manager

Kafka overview

- Kafka is a distributed, high-throughput message broker.
- All Kafka messages are organized into **topics**.
- To send a message to Kafka, you send it to a specific topic.
- To read a message from Kafka, you read it from a specific topic.
- **Producers** send messages to Kafka, for example, a receiver Logstash server.
- **Consumers** pull messages from Kafka, for example, a sender Logstash server.
- Topics are divided **partitions**. Partitions allow for messages within each topic to be read in parallel.
- Data is stored in transaction logs. Producers append to logs sequentially. Consumers read a range from these logs using an offset.
- Kafka, as a distributed system, runs in a cluster. Each node in the cluster is called a **broker**.
- Kafka requires Zookeeper for coordination.

Kafka overview

In the Log Analysis scalable architecture, the role of the Kafka cluster is a reliable and durable data store for log messages.

Kafka is a distributed system. Each Kafka server that stores data is called a broker. When a Logstash receiver sends output to Kafka, it sends data to a list of brokers.

Apache Zookeeper is used in the Kafka cluster to coordinate data replication and communication between brokers.

Applications that send data to Kafka are called producers. In the Log Analysis scalable architecture, the producers are Logstash receivers.

Applications that pull data out of Kafka are called consumers. In the Log Analysis scalable architecture, the consumers are Logstash senders.

The following diagram shows the relationship between brokers, topics, and partitions.

- There are two brokers in this example: Broker 0 and Broker 1.
- There are two topics in this example: ABC and XYZ. Topics are the message queues. Topics are logical collections of physical files. Partitions are the physical files where data is saved.
- The data in each topic is saved across three partitions. Partitions are the physical files where data is saved.

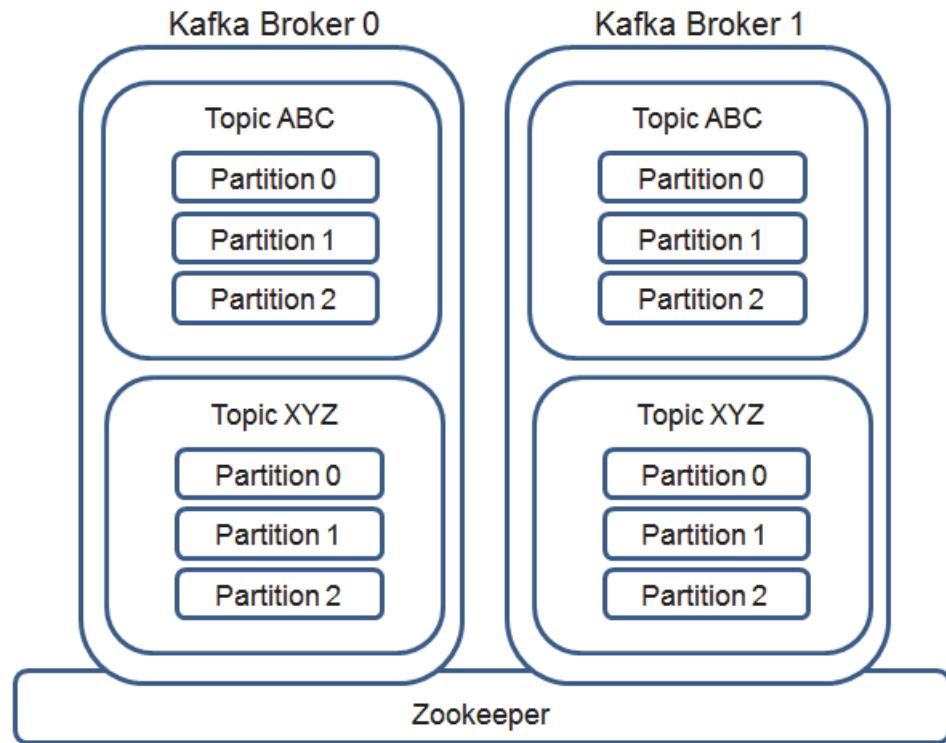


Figure 1 Brokers, topics, and partitions

Installing and configuring Kafka

- Install Kafka by decompressing the installation file.
- Zookeeper installs with Kafka.
- Configure Zookeeper with the `zookeeper.properties` file.
- Configure Kafka with the `server.properties` file.

<code>dataDir=/tmp/zookeeper</code>	<code>zookeeper.properties</code>
<code>clientPort=17981</code>	
<code>maxClientCnxns=0</code>	

<code>broker.id=0</code>	<code>server.properties</code>
<code>port=17991</code>	
<code>log.dirs=/tmp/kafka-logs-server-0</code>	
<code>num.partitions=3</code>	
<code>zookeeper.connect=kafka.csuite.ibm.edu:17981</code>	

Installing and configuring Kafka

Decompress the installation file to install a Kafka broker and Zookeeper.

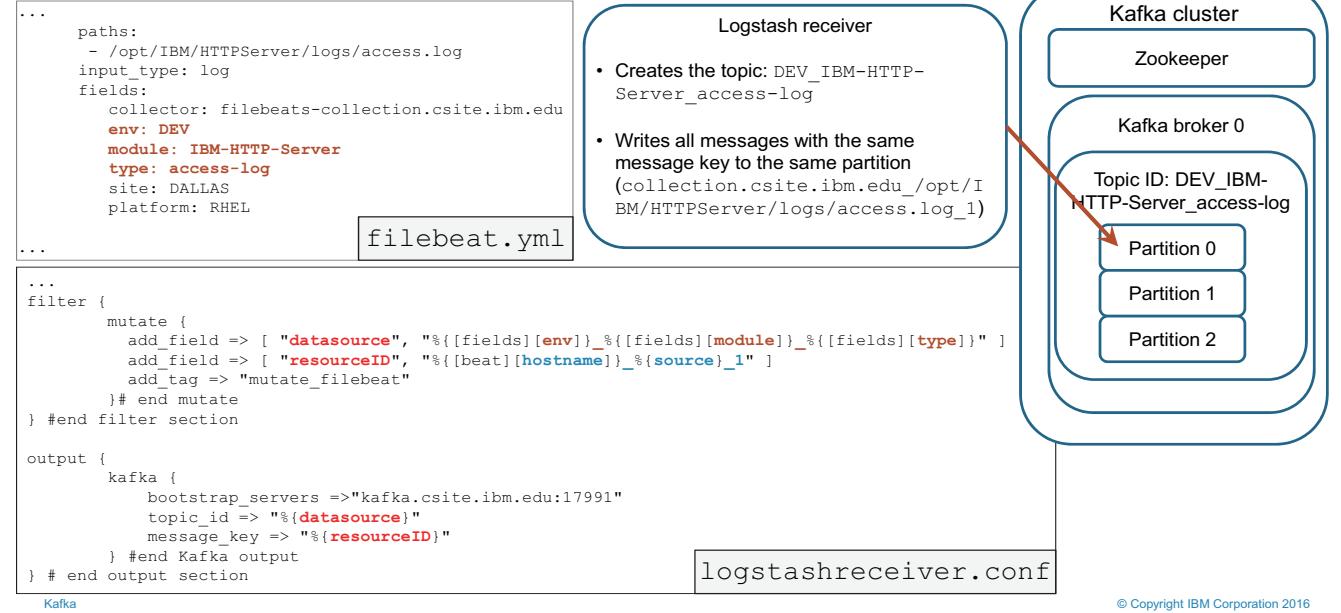
Edit the `zookeeper.properties` file to configure the Zookeeper server. The following list explains some key properties in the `zookeeper.properties` file:

- **dataDir**: The directory where Zookeeper saves state information
- **clientPort**: The port number that Kafka brokers and Logstash senders use to connect to the Zookeeper server
- **maxClientCnxns**: The number of concurrent connections that a client can make to Zookeeper

Edit the `server.properties` file to configure the Kafka broker. The following list explains some key properties in the `server.properties` file:

- **broker.id**: The unique identifier for the broker.
- **port**: The port used to connect to the broker.
- **log.dirs**: A comma-delimited list of directories where the broker saves data. Partitions are created in the directories you enter in this property.
- **num.partitions**: The number of partitions that the broker will create for each topic.
- **zookeeper.connect**: The host and port of the Zookeeper server.

Kafka in the scalable collection architecture



© Copyright IBM Corporation 2016

Kafka in the scalable collection architecture

In the example on this slide, the Logstash receiver is using metadata that was added to each log messages by the Filebeat collection agent.

The Logstash receiver mutate filter adds a new field named **datasource**. The value of **datasource** is constructed using the value of the Filebeat field **env**, followed by an underscore character, followed by the value of the Filebeat field **module**, followed by an underscore character, followed by the value of the Filebeat field **type**. In this example, the value of **datasource** is **DEV_IBM-HTTP-Server_access-log**.

The mutate filter also adds a new field named **resourceID**. The value of **resourceID** is constructed using the value of the Filebeat metadata keyword **hostname**, followed by an underscore character, followed by the value of the metadata keyword **source**, followed by the string **_1**. In this example, the value of **resourceID** is

`collection.csuite.ibm.edu/_opt/IBM/HTTPServer/logs/access.log_1`.

The two fields that were created by the mutate filter, **datasource** and **resourceID**, are used in the configuration of the kafka output plug-in.

- The value of **datasource** is the name of the Kafka topic that this Logstash receiver writes messages to. Notice that the value of **datasource** is used as the **topic_id**.
- Any messages with the same **message_key** are written to the same partition within the topic. Notice that the value of **resourceID** is used as the **message_key**.

The Logstash receiver kafka output plug-in is configured to send data to the Kafka broker at `kafka.csuite.ibm.edu:17991`. The Logstash receiver is a Kafka producer. The Logstash receiver automatically creates a topic named **DEV_IBM-HTTP-Server_access-log** in the Kafka broker.

All data is saved to the same partition in the topic, because the **message_key** option is configured in the kafka output plug-in.

Viewing data in Kafka

Kafka includes tools to list all topics and to show the data in each topic.

To list all topics, issue this command:

```
kafka-topics.sh --list --zookeeper kafka.csuite.ibm.edu:17981  
DEV_IBM-HTTP-Server_access-log  
LabDB2Instance_db2diag  
labInstance_syslog
```

To view data in a topic, issue this command:

```
kafka-console-consumer.sh --zookeeper kafka.csuite.ibm.edu:17981 --topic DEV_IBM-HTTP-  
Server_access-log --from-beginning  
  
{"message":"Apache/IHS,192.168.2.94,-,-,[10/May/2016:15:41:37 -0400],GET,\\\"GET  
/daytrader/app?action=portfolio HTTP/1.1\\\",200,3409,1308,\\\"-\\\",\\\"curl/7.21.3 (x86_64-unknown-  
linux-gnu) libcurl/7.21.3 OpenSSL/1.0.0 zlib/1.2.3\\\"","@version":"1","@timestamp":  
...}
```

Viewing data in Kafka

You can use command-line tools to show data that is saved in the Kafka cluster. These tools are in the `<KAFKA_HOME>/bin` directory.

Use the `kafka-topics.sh` tool to list all topics in the Kafka cluster.

Use the `kafka-console-consumer.sh` tool to view the data in a topic. If you use the `--from-beginning` option, the tool shows all data in a topic. If you do not use the `--from-beginning` option, the tool shows only data that is currently being written to the topic.

You can also use the `kafka-topics.sh` tool to show details about a topic. In the following example, the tool shows information about the topic named `DEV_IBM-HTTP-Server_access-log`:

```
./kafka-topics.sh --describe --zookeeper kafka.csuite.ibm.edu:17981 --topic  
DEV_IBM-HTTP-Server_access-log  
  
Topic:DEV_IBM-HTTP-Server_access-logPartitionCount:3ReplicationFactor:1Configs:  
Topic: DEV_IBM-HTTP-Server_access-logPartition: 0Leader: 0Replicas: 0Isr: 0
```

Kafka Manager

Kafka Manager is an optional web-based user interface that manages your Kafka cluster.

The screenshot shows the Kafka Manager web interface. At the top, there is a navigation bar with tabs for 'Lab-Kafka' (which is selected), 'Cluster', 'Brokers', 'Topic', 'Preferred Replica Election', 'Reassign Partitions', and 'Consumers'. Below the navigation bar, the URL path 'Clusters / Lab-Kafka / Brokers' is visible. The main content area has two sections: 'Brokers' on the left and 'Combined Metrics' on the right. The 'Brokers' section lists one broker with ID 0, host 'kafka.csite.ibm.edu', port 17991, and JMX port 8092. The 'Bytes In' and 'Bytes Out' columns both show '0.00'. The 'Combined Metrics' section displays various performance metrics with their corresponding mean values over different time intervals (Mean, 1 min, 5 min, 15 min). All metrics show values of 0.00.

Rate	Mean	1 min	5 min	15 min
Messages in /sec	0.00	0.00	0.00	0.00
Bytes in /sec	0.00	0.00	0.00	0.00
Bytes out /sec	0.00	0.00	0.00	0.00
Bytes rejected /sec	0.00	0.00	0.00	0.00
Failed fetch request /sec	0.00	0.00	0.00	0.00
Failed produce request /sec	0.00	0.00	0.00	0.00

Kafka Manager

Kafka Manager is a web-based user interface that shows the configuration of your cluster. You can also use Kafka Manager to view data and modify your cluster. Although Kafka Manager is not required, it can be useful for troubleshooting and maintenance.

Installing and using Kafka Manager

Follow these steps to install Kafka Manager:

1. Create an installation package with Simple Build Tool (sbt).
2. Decompress the installation package to install Kafka Manager.

Kafka Manager uses Java Management Extensions (JMX) to connect to Kafka brokers.

You must set the JMX port before you start your Kafka broker, for example:

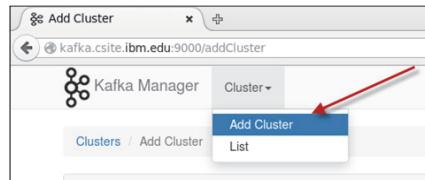
```
export JMX_PORT=8092
/opt/kafka_2.9.1-0.8.2.2/bin/kafka-server-start.sh -daemon /opt/kafka_2.9.1-0.8.2.2/config/server.properties
```

Installing and using Kafka Manager

To install Kafka Manager, you must compile it from source code with the Simple Build Tool. The Simple Build Tool creates an archive file that you use to install Kafka Manager.

Kafka Manager uses Java Management Extensions (JMX) to connect to Kafka brokers. If you want Kafka Manager to communicate with your Kafka cluster, you must set the JMX port before you start each of your Kafka brokers.

Configuring Kafka Manager



Click **Cluster > Add Cluster** to configure Kafka Manager to monitor your Kafka cluster.

Kafka

© Copyright IBM Corporation 2016

Configuring Kafka Manager

You must configure Kafka Manager before you can use it to manage your cluster. To add your cluster, open a browser and go to http://<kafka-manager_host>:9000. Click **Cluster > Add Cluster**. The following table explains how to configure the required fields on the Add Cluster page.

Field	Usage
Cluster Name	A name for your cluster. This name is used only to display the cluster in Kafka Manager.
Cluster Zookeeper Hosts	A comma-delimited list of zookeeper hosts and port numbers.
Kafka Version	Select your version of Kafka, for example 0.8.2.2.
Enable JMX Polling	Set this to true.
brokerViewThreadPoolSize	Set this to a minimum of 2.
offsetCacheThreadPoolSize	Set this to a minimum of 2.
kafkaAdminClientThreadPoolSize	Set this to a minimum of 2.

Unit summary

- Install and configure Kafka
- Diagram data flow from a Logstash receiver server to Kafka
- Install and configure Kafka Manager

Kafka

© Copyright IBM Corporation 2016

Unit summary

Kafka exercises

- Exercise 1 Installing and configuring Kafka**
Exercise 2 Configuring Logstash receiver output
Exercise 3 Verifying data flow
Exercise 4 Kafka Manager

Kafka

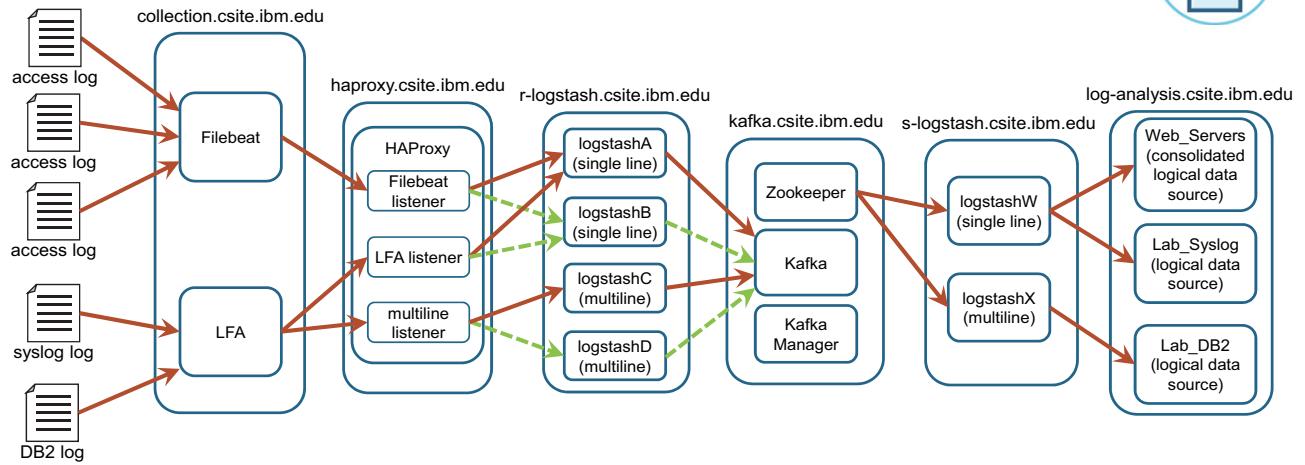
© Copyright IBM Corporation 2016

Student exercises

Perform the exercises for this unit.



Exercise introduction: Finished lab environment



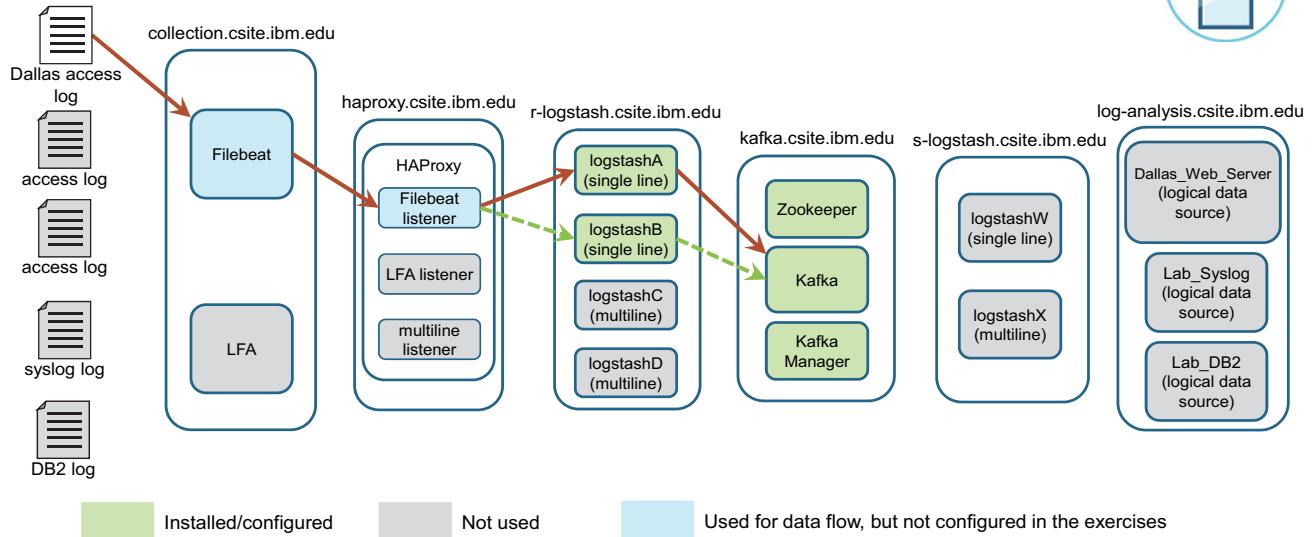
Kafka

© Copyright IBM Corporation 2016

Exercise introduction: Finished lab environment

At the end of this course, your lab environment will look like this diagram. Refer to this diagram as you complete the steps in the student exercises to help you understand what you are installing and configuring.

Exercise introduction: Lab environment in Unit 3 exercises



Kafka

© Copyright IBM Corporation 2016

Exercise introduction: Lab environment in Unit 3 exercises

In the exercises for this unit, you install and configure the following components:

- A Kafka broker, to store messages
- A Zookeeper server, to coordinate Kafka communications
- Kafka Manager, to view data in your Kafka cluster
- The output section of your logstashA and logstashB receivers, to send data to Kafka (configuration only)

As you install and configure these components, you also verify the data flow between them.

Unit 4 Logstash senders and Log Analysis core

IBM Training



Logstash senders and Log Analysis core

© Copyright IBM Corporation 2016
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

In this unit, you learn how to install and configure Logstash senders. You also learn how to configure the IBM Operations Analytics Log Analysis core software to accept data from a Logstash sender.

Unit objectives

- Configure Logstash senders
- Configure Log Analysis data sources
- Diagram data flow from Kafka to Logstash senders to Log Analysis

About Logstash senders

- Logstash senders pull messages from a Kafka topic.
- Use the Logstash sender filter section for any text transformation or annotation operations.
- Use the Logstash sender filter section to set a *host* and *path*.
- Logstash senders use a custom output plug-in named **scala** to send data to Log Analysis.

About Logstash senders

In the Log Analysis scalable architecture, the role of a Logstash sender instance is to retrieve log messages from a Kafka broker and send those messages to the Log Analysis server.

You use the input section of a Logstash sender configuration to connect to Kafka and pull messages from a specific topic.

You use the filter section of a Logstash sender configuration to alter the log message or metadata that is attached to the log message. At minimum, Logstash senders create two metadata fields for each message: **host** and **path**.

You use the output section of a Logstash sender configuration to send log messages to the Log Analysis server.

Logstash sender input section

```
input {
  kafka {
    zk_connect => "kafka.csuite.ibm.edu:17981"
    group_id => "DEV_IBM-HTTP-Server_access-log"
    topic_id => "DEV_IBM-HTTP-Server_access-log"
    consumer_threads => 1
    consumer_restart_on_error => true
    consumer_restart_sleep_ms => 100
    decorate_events => true
  }# end HTTP server log Kafka input

  kafka {
    zk_connect => "kafka.csuite.ibm.edu:17981"
    group_id => "labInstance_syslog"
    topic_id => "labInstance_syslog"
    consumer_threads => 1
    consumer_restart_on_error => true
    consumer_restart_sleep_ms => 100
    decorate_events => true
  }# end syslog server log Kafka input
}# end input section
```

[Logstash senders and Log Analysis core](#)

© Copyright IBM Corporation 2016

Logstash sender input section

Logstash senders use the `kafka` input plug-in to pull messages from a topic in the Kafka cluster. In this example, the Logstash sender server is pulling messages from two different topics:

`DEV_IBM-HTTP-Server_access-log` and `labInstance_syslog`. Notice that the `kafka` input plug-in is invoked twice: once for each topic.

The following table describes each line in the `kafka` input plug-in configuration example.

Configuration option	Description
<code>zk_connect</code>	The host and port number of the Zookeeper server.
<code>group_id</code>	A name for the consumer group that identifies the Logstash server in Kafka. Use the same value for <code>topic_id</code> and <code>group_id</code> .
<code>topic_id</code>	The name of the Kafka topic you want to pull data from.
<code>consumer_threads</code>	The number of threads to concurrently pull data from topic. This is typically set to the number of partitions for the topic.
	The <code>consumer_threads</code> parameter also specifies the number of consumers that are created in a consumer group. Each thread, or consumer, maps to a partition for the specified topic. This ensures that data is processed concurrently.
<code>consumer_restart_on_error</code>	If true, consumer threads will restart on errors.

Configuration option	Description
consumer_restart_sleep_ms	The number milliseconds to wait for consumer threads to restart after an error.
decorate_events	If true, metadata about Kafka, such as <code>topic</code> or <code>msg_size</code> is added to each log message.

Logstash sender filter section

```
filter {
  if "mutate_filebeat" in [tags] {
    mutate {
      add_field => [ "path", "%{[fields][type]}" ]
      replace => { "host" => "%{[fields][env]}_%{[fields][module]}" }
    } # end mutate
  } # end filebeat mutate condition
} # end filter section
```

Logstash senders and Log Analysis core

© Copyright IBM Corporation 2016

Logstash sender filter section

You use the filter section of a Logstash sender configuration to alter the log message or metadata that is attached to the log message. At minimum, Logstash senders create two metadata fields for each message: **host** and **path**. You use the **host** and **path** fields when you configure the data source in the Log Analysis server.

In this example, the mutate filter in a Logstash sender configuration is creating a metadata field named **path**. The value for **path** is set with metadata that was added by the Filebeat server (**type**). The mutate filter is also replacing the existing value for **host** with metadata that was added by the Filebeat server (**env_module**). The final value for **path** is **access-log**. The final value for **host** is **DEV_IBM-HTTP-Server**.

At the bottom of the slide are pictures of the Log Analysis data source wizard. Notice that you use the values of **host** and **path** that were set by the Logstash sender when you create the data source.

Logstash sender output section

```
output {  
  
    scala {  
        scala_url => "https://log-analysis.csuite.ibm.edu:9987/Unity/DataCollector"  
        scala_user => "unityadmin"  
        scala_password => "object00"  
        scala_keystore_path => ""  
        batch_size => 500000  
        idle_flush_time => 5  
        sequential_flush => true  
        num_concurrent_writers => 20  
        use_structured_api => false  
        disk_cache_path => "/opt/logstashW/training/cache/basecache"  
        date_format_string => "yyyy-MM-dd'T'HH:mm:ssX"  
        log_file => "/opt/logstashW/logstash-2.2.1/log/scala_logstashW.log"  
        log_level => "info"  
    }#end scala output  
  
}# end output section
```

Logstash sender output section

You use the output section of a Logstash sender configuration to send log messages to the Log Analysis server.

Use the scala output plug-in to send data to the Log Analysis server. The following table describes each line in the scala output plug-in configuration example.

Configuration option	Description
scala_url	The REST endpoint for the Log Analysis Generic Receiver REST API.
scala_user	The Log Analysis user name.
scala_password	The Log Analysis user password.
scala_keystore_path	The full path to the directory where the keystore file is saved. Enter a value if you want to encode the plain text scala_password.
batch_size	The maximum number of bytes that can be buffered for each data source before Logstash sends data to the Log Analysis server.
idle_flush_time	The maximum time between successive data transmissions for each data source.
sequential_flush	Defines whether batches for each data source are sent sequentially. If true, batches are sent sequentially.
num_concurrent_writers	The number of threads that concurrently transmit batches of data to the Log Analysis server.

Configuration option	Description
use_structured_api	Determines whether data is transmitted to the Log Analysis server in JSON format. If true, data is sent in JSON format.
disk_cache_path	The path on the file system that temporarily buffers data. The scala output plug-in writes data to this path before sending data to the Log Analysis server.
date_format_string	The string format that all date fields are transformed to before transmission to the Log Analysis server.
log_file	The file that is used for logging information from the scala output plug-in.
log_level	The level of logging information for the scala output plug-in. The supported levels are: fatal, error, warn, info, and debug.

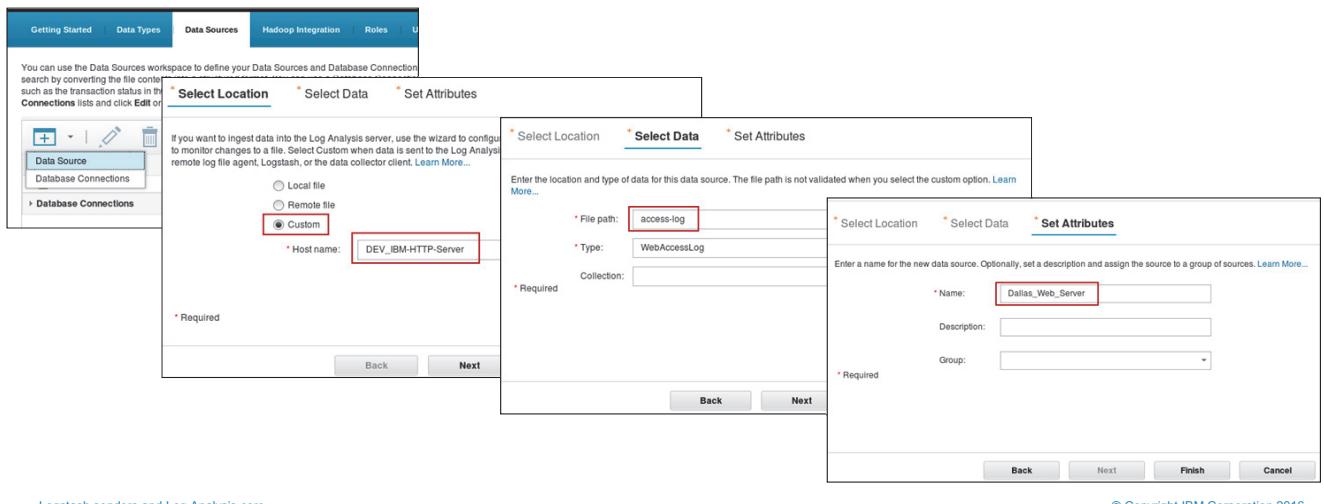


Note: If you installed Logstash by decompressing the installation file or with an RPM Package Manager file (RPM), you must manually copy the scala output plug-in files to the Logstash server. For example, copy the files to the following directory:

<LOGSTASH_HOME>/vendor/bundle/jruby/1.9/gems/logstash-core-2.2.1-java/lib/logstash/outputs/

Configuring Log Analysis

- Use the Log Analysis administrator user interface to add a data source
- Use the host and path that were set by the Logstash sender, not the actual path of the target log file



Logstash senders and Log Analysis core

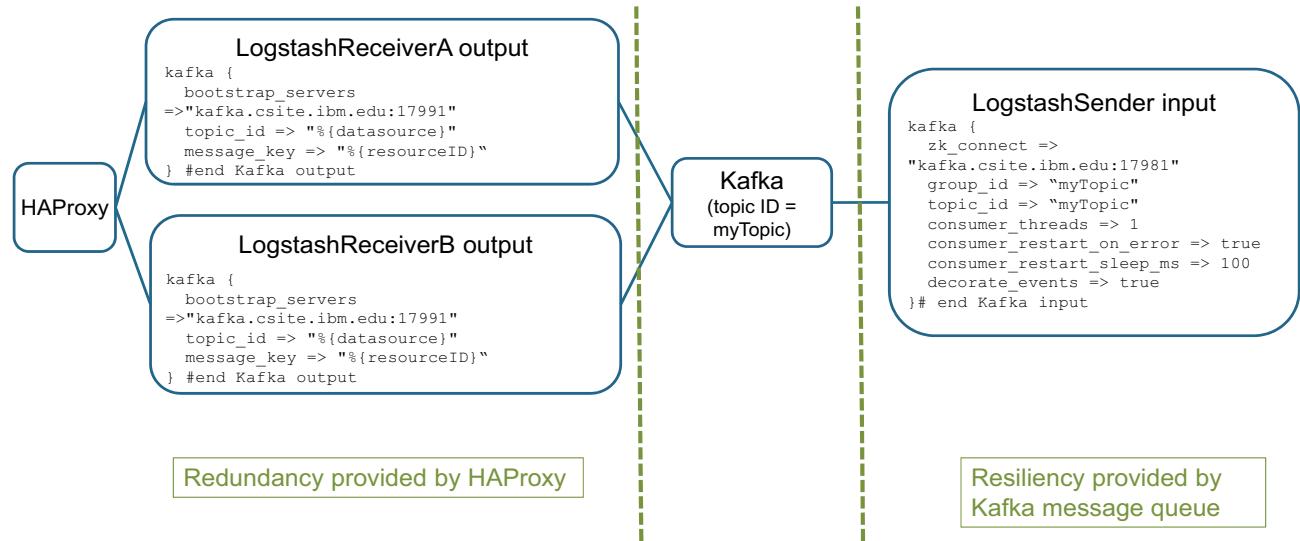
© Copyright IBM Corporation 2016

Configuring Log Analysis

To configure Log Analysis to accept data from the target log file, you must create a data source. Use the data source wizard in the Log Analysis administrator user interface to add a data source.

Use the host and path that were set by the Logstash sender filter section, not the actual path of the target log file.

Logstash sender resiliency



Resiliency for Logstash senders is provided by the Kafka cluster. If a Logstash sender server goes down, it can read messages that are stored in the Kafka message queue when it comes back up.

Unit summary

- Configure Logstash senders
- Configure Log Analysis data sources
- Diagram data flow from Kafka to Logstash senders to Log Analysis

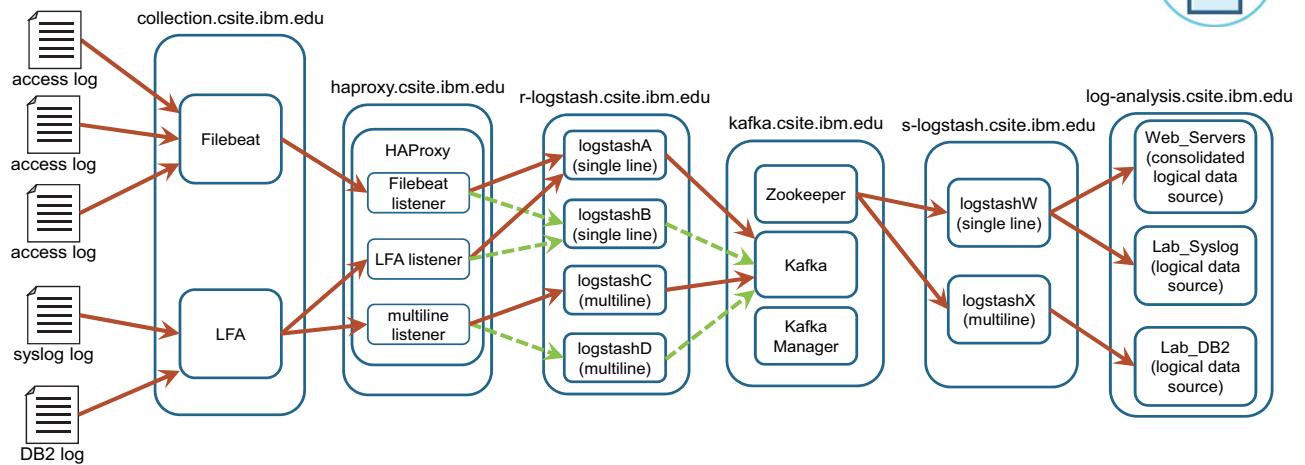
Logstash senders and Log Analysis core exercises

- Exercise 1 Installing and configuring the Logstash sender**
Exercise 2 Sending data to Log Analysis
Exercise 3 Verifying Logstash sender resiliency

Perform the exercises for this unit.



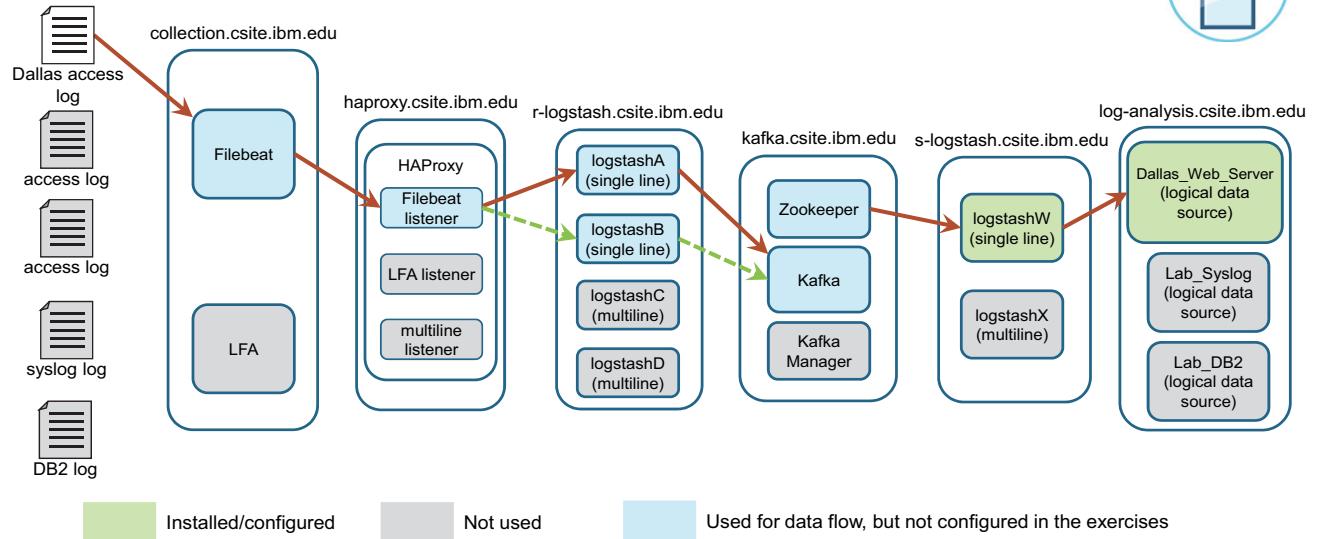
Exercise introduction: Finished lab environment



Exercise introduction: Finished lab environment

At the end of this course, your lab environment will look like this diagram. Refer to this diagram as you complete the steps in the student exercises to help you understand what you are installing and configuring.

Exercise introduction: Lab environment in Unit 4 exercises



Exercise introduction: Lab environment in Unit 4 exercises

In the exercises for this unit, you install and configure the following components:

- A Logstash sender, to pull messages from Kafka and send them to the Log Analysis server
- A data source in Log Analysis, to accept data from the target log file (configuration only)

As you install and configure these components, you also verify the data flow between them.

Unit 5 Using the Log File Agent

IBM Training

IBM

Using the Log File Agent

© Copyright IBM Corporation 2016
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

You can use the IBM Tivoli Log File Agent, which is included with Log Analysis, as a collection agent. This unit teaches you how to install and configure the Log File Agent.

Unit objectives

- Install the Log File Agent
- Configure the Log File Agent
- Configure Logstash receivers to process Log File Agent messages

Installing the Log File Agent

Use the Log Analysis remote installation tool to install the Log File Agent (LFA):

1. Edit the `ssh-config.properties` file with details about the host that you want to install to.
2. Run the `install.sh` tool.

The remote installation tool is in the `<LA_HOME>/remote_install_tool` directory.

Installing the Log File Agent

You can use the Log Analysis remote installation tool to install the Log File Agent (LFA). The following steps show you how to install the Log File Agent.

1. Change to the remote installation tool directory of the Log Analysis server.
`cd <LA_HOME>/remote_install_tool`
2. Open the **ssh-config.properties** file in a text editor. This example uses vi.
`vi config/ssh-config.properties`
3. Change the following values in bold typeface. These values are the details about how to connect to the host where you want to install the Log File Agent.

```
REMOTE_HOST=log_file_agent_target_host
PORT=ssh_port_number
TIME_OUT=60000
USER=lfa_user
#PASSWORD can be commented while using Public key based authentication
PASSWORD=lfa_user_password
```

4. Save and close the file when you are finished.
5. Run the following command to start the remote installation tool.
`<LA_HOME>/remote_install_tool/install.sh`

6. Enter the directory where you want to install the Log File Agent. In this example, the installation directory is /opt/LFA.

Enter Remote Top Level Installation Directory absolute path:
[/home/netcool/LogAnalysis]

/opt/LFA

7. Enter **n** when you are prompted to install EIF Receiver instances.

Install EIF Receiver Instances (y|Y|n|N) : [y]

n

8. Enter **y** when you are prompted to install Log File Agent (LFA) 6.3.

Install LFA 6.3 (y|Y|n|N) : [y]

y

9. Enter **n** when you are prompted to install Logstash. The installation starts.

Install logstash 2.2.1 (y|Y|n|N) : [y]

n

The installation takes several minutes to complete. The following message confirms that the installation was successful.

```
Response:=====
Response:      COMPONENT PIDSTATUS
Response:=====
Response:  Log File Agent 9784 UP
Response:=====
End Time:Thu May 26 16:06:54 UTC 2016
+++++  
Total install duration (seconds) :109
```

```
+++++ Installation Ends +++++
```

Configuring the Log File Agent

You use two files to configure the Log File Agent to monitor a target log:

- A .conf file
- A .fmt file

```
LogSources=/var/log/messages.log
BufEvtPath=/opt/LFA/IBM-LFA-6.30/logs/lab-syslog.cache
FileComparisonMode=CompareByAllMatches
ServerLocation=ha-proxy.csuite.ibm.edu
ServerPort=5530
FQDomain=yes
BufferEvents=YES
BufEvtMaxSize=102400
ConnectionMode=CO
PollInterval=3
NumEventsToCatchUp=-1
ServerSSL=NO
```

lab-syslog.conf

```
REGEX AllRecords
(.*)
hostname LABEL
-file FILENAME
RemoteHost DEFAULT
logpath PRINTF("%s",file)
type syslog
instance labInstance
cluster NONE
module syslog
env DEV
functional NONE
site NONE
text $1
END
```

lab-syslog.fmt

Configuring the Log File Agent

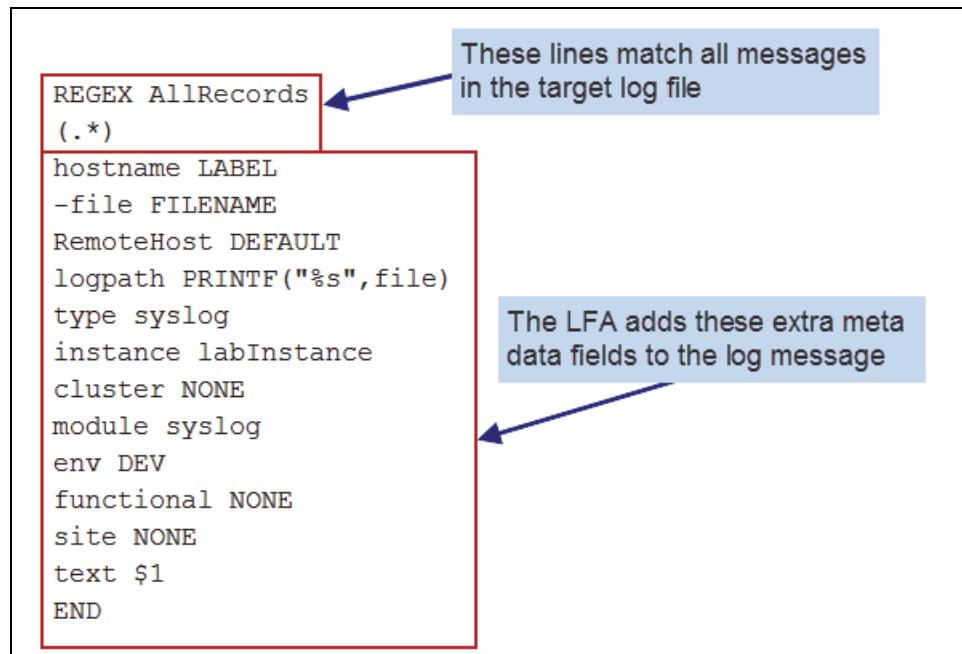
You configure the Log File Agent by creating two files: a .conf file and a .fmt file. You create these files in the <LFA_HOME>/IBM-LFA-6.30/config/lo/ directory.

The following table describes each line in the Log File Agent conf file example.

Option	Description
LogSources	Specifies the text log files to poll for messages. The complete path to each file must be specified, and file names must be separated by commas.
BufEvtPath	Specifies the full path name of the agent cache file.
FileComparisonMode	Specifies which log files are monitored when more than one matches a wildcard pattern. If set to CompareByAllMatches, all files that match the wildcard pattern that is specified in LogSources are monitored.
ServerLocation	Specifies the host name or IP address where log messages are sent.
ServerPort	Specifies the port number to use when sending log messages.
FQDomain	Specifies how and if the agent sets a domain name. <ul style="list-style-type: none"> • If set to yes, the agent determines the system domain name itself. • If set to no, the agent does not set a domain name. The host name attribute is assigned a blank string. • If set so that it does not contain a yes or no value, the domain name is accepted as the value and it is appended to the host name.

Option	Description
BufferEvents	Specifies how event buffering is enabled. The possible values are: <ul style="list-style-type: none"> • YES: Stores events in the file that is specified by the BufEvtPath option. • MEMORY_ONLY: Buffers events in memory. • NO: Does not store or buffer events.
BufEvtMaxSize	Specifies the maximum size, in KB, of the agent cache file.
EventMaxSize	Specifies in bytes, the maximum size of a generated log message event.
ConnectionMode	Specifies whether the connection is connection-oriented or connectionless.
PollInterval	Specifies the frequency, in seconds, to poll each log file that is listed in the LogSources option for new messages
NumEventsToCatchUp	Specifies the event in the log that the agent starts with. When set to -1, the agent saves its place in the file that is being monitored. When set to 0, the agent starts with the next event in the logs.
ServerSSL	Specifies whether to use Secure Sockets Layer when sending data.

The following diagram describes key fields in the .fmt file example.



Adding a HAProxy listener for Log File Agents

```
## Listener for LFA Logstash cluster
## Specify Logstash instances for processing data from LFA
listen LFA_Cluster ha-proxy.csuite.ibm.edu:5530
    mode tcp
    balance roundrobin
    server receiver-logstashA logstashA.csuite.ibm.edu:5540 check fall 2 rise 3 inter 1000
    server receiver-logstashB logstashB.csuite.ibm.edu:5541 check fall 2 rise 3 inter 1000
```

Adding a HAProxy listener for Log File Agents

You configure HAProxy by editing a file named `/etc/haproxy/haproxy.cfg`. In this scalable architecture, HAProxy must run in `tcp` mode.

You must configure a listener for HAProxy to accept data from a log collection agent and forward it to a Logstash receiver cluster. In this example, HAProxy is listening on port 5530. Any data HAProxy receives on port 5530 is forwarded to either `logstashA.csuite.ibm.edu:5540` or `logstashB.csuite.ibm.edu:5541`.

HAProxy is aware of the state of each server. In this example, HAProxy performs health checks on each Logstash server every 1000 milliseconds. HAProxy considers a Logstash server down if two consecutive health check failures occur (`fall 2`). HAProxy considers a Logstash server up again after three consecutive successful health checks (`rise 3`).

The line `balance roundrobin` configures HAProxy to send data to each Logstash server in turns.

Logstash receiver input section

```
input {  
  
    tcp {  
        port => 5540  
        type => "lfa"  
        codec => line { charset => "US-ASCII" }  
    } #end lfa input  
  
} #end input section
```

Logstash receiver input section

In the Log Analysis scalable architecture, the role of a Logstash receiver instance is to accept log messages from collection agents and send the messages to Kafka.

Use the **tcp** input plug-in to accept data from the Log File Agent. Use the **type** option to add a field named **type** to each log message. Setting a type for messages from the LFA makes it easier to alter these messages with the filter section.

Use the **codec** option to set the character encoding.

Extra text in Log File Agent messages

Extra text in Log File Agent messages

The Log File Agent sends extra text with each log message. You use the filter section of the Logstash receiver to remove this unnecessary text.

Removing the extra text with the Logstash receiver

```
filter {

    if [type] == "lfa" {
        grok {
            patterns_dir => "/opt/logstashA/logstash-2.2.1/patterns"
            match => [ "message", "%{LFAMESSAGE}" ]
            add_tag => ["grok_lfa"]
        } #end initial LFA grok
    } #end initial LFA condition

    if "grok_lfa" in [tags] {
        mutate {
            replace => [ "message", "%{LFA_ORIG_MSG}" ]
            add_field => [ "datasource", "%{LFA_INSTANCE}_%{LFA_MODULE}" ]
            add_field => [ "resourceID", "%{LFA_HOSTNAME}_%{LFA_LOGNAME}_1" ]
        }# end mutate
    }# end grok_lfa condition

} #end filter section
```

Using the Log File Agent

© Copyright IBM Corporation 2016

Removing the extra text with the Logstash receiver

Use the grok filter to remove the unnecessary text from the Log File Agent. The grok filter uses regular expression objects called patterns. The grok filter uses these patterns to match meaningful strings of text and identify them as fields within a log message. You can use a file that contains a grok pattern to match the unnecessary text from the Log File Agent. This patterns file is included with Log Analysis. You can find it in the following location:

<LA_HOME>/kafka/test-configs/logstash-configs/SCALAPATTERNS

Copy the SCALAPATTERNS file to the Logstash receiver server.

Grok filter configuration example

In the following example, line numbers have been added to the grok filter configuration for clarity.

```
1 if [type] == "lfa" {
2     grok {
3         patterns_dir => "/opt/logstashA/logstash-2.2.1/patterns"
4         match => [ "message", "%{LFAMESSAGE}" ]
5         add_tag => ["grok_lfa"]
6     } #end initial LFA grok
7 } #end initial LFA condition
```

- Line 1 is a condition. Only messages where the type is **lfa** are processed by the grok filter.
- Line 2 invokes the grok filter plug-in.
- Line 3 points to the directory where you copied the SCALAPATTERNS file.

- Line 4 uses a pattern named LFAMESSAGE to match meaningful strings of text in the data that was received from the Log File Agent.
- Line 5 adds a tag to messages that were successfully matched by the grok filter.
- Line 6 closes the grok filter configuration.
- Line 7 closes the condition.

Mutate filter configuration example

In the following example, line numbers have been added to the mutate filter configuration for clarity.

```
1 if "grok_lfa" in [tags] {  
2     mutate {  
3         replace => ["message", "%{LFA_ORIG_MSG}"]  
4         add_field => [ "datasource", "%{LFA_INSTANCE}_%{LFA_MODULE}" ]  
5         add_field => [ "resourceID", "%{LFA_HOSTNAME}_%{LFA_LOGNAME}_1" ]  
6     }# end mutate  
7 }# end grok_lfa condition
```

- Line 1 is a condition. Only messages that were successfully matched by the preceding grok filter are processed by the mutate filter.
- Line 2 invokes the mutate filter plug-in.
- Line 3 replaces the text in the **message** field with text that was identified as LFA_ORIG_MSG by the preceding grok filter. The text that was identified as LFA_ORIG_MSG is the actual log message from the target log file. This line removes the unnecessary text that was added by the Log File Agent.
- Line 4 adds a field named **datasource** and sets the value of **datasource** with metadata that was added by the Log File Agent. The **datasource** field is used in the Logstash receiver output section.
- Line 5 adds a field named **resourceID** and sets the value of **resourceID** with metadata that was added by the Log File Agent. The **resourceID** field is used in the Logstash receiver output section.
- Line 6 closes the mutate filter configuration.
- Line 7 closes the condition.

Metadata from the Log File Agent

```
"message" =>
"\u0001<START>>\u0000\u0000\u0000\u0000\u0000\u0000
0\u0000\u0000\u0000\u0000\u0000\u0000\u0000\u0000\u0000
\u0000\u0000\u0000\u0000\u0000\u0000\u0000\u0000\u0000\u0000
\u0000\u0000\u0000\u0001\xBD\u0000\u0000\u0000\u0000
000\u0004\u0003allRecords;type='syslog';text='20
16-05-12T23:39:50.764527+00:00
host=analysishost,relayHost=analysishost,
tag=kernel:, programName=kernel,procid=-,
facility=kern, sev=info,appName=kernel,
msg=Reserving 129MB of memory at 48MB for
crashkernel (System RAM:
9216MB)';RemoteHost='';site='NONE';instance='labIn
stance';hostname='collection';cluster='NONE';modul
e='syslog';env='DEV';logpath='/var/log/messages.lo
g';functional='NONE';END",
"@version" => "1",
"@timestamp" => "2016-05-26T18:20:08.149Z",
"host" => "192.168.100.176",
"port" => 45437,
"type" => "lfa"
```

Message from the LFA

```
REGEX AllRecords
(.*)
hostname LABEL
-file FILENAME
RemoteHost DEFAULT
logpath PRINTF("%s",file)
type syslog
instance labInstance
cluster NONE
module syslog
env DEV
functional NONE
site NONE
text $1
END
lab-syslog(fmt

filter {
  if [type] == "lfa" {
    grok {
      patterns_dir => "/opt/logstashA/logstash-2.2.1/patterns"
      match => [ "message", "%{LFAMESSAGE}" ]
      add_tag => ["grok_lfa"]
    } #end initial LFA grok
  } #end initial LFA condition

  if "grok_lfa" in [tags] {
    mutate {
      replace => [ "message", "%{LFA_ORIG_MSG}" ]
      add_field => [ "datasource", "%{LFA_INSTANCE}_%{LFA_MODULE}" ]
      add_field => [ "resourceID", "%{LFA_HOSTNAME}_%{LFA_LOGNAME}_1" ]
    }# end mutate
  }# end grok_lfa condition
} #end filter section
```

logstashreceiver.conf

Using the Log File Agent

© Copyright IBM Corporation 2016

Metadata from the Log File Agent

You configure the Log File Agent to add metadata to each log message with the. fmt file. In this example, the Log File Agent is adding metadata such as **hostname**, **instance**, and **module**.

The grok filter in the Logstash receiver filter section uses a pattern named LFAMESSAGE to match the metadata in the text from the Log File Agent. Logstash stores the metadata in fields named **LFA_HOSTNAME**, **LFA_INSTANCE**, **LFA_MODULE**, and so on.

The mutate filter in the Logstash receiver filter section adds a field named **datasource**. The value of **datasource** is constructed by using the fields **LFA_INSTANCE** and **LFA_MODULE**. These two fields were matched by the grok filter. In this example, the value of **datasource** is **labInstance_syslog**.

The mutate filter adds another field named **resourceID**. The value of **resourceID** is constructed by using the fields **LFA_HOSTNAME** and **LFA_LOGNAME**. These two fields were matched by the grok filter. In this example, the value of **resourceID** is

`collection_/software/log_samples/messages.log_1.`

Logstash receiver output section

```
output {  
  
    if ("grok_lfa" in [tags]) and ! ("_grokparsefailure" in [tags]) {  
        kafka {  
            bootstrap_servers =>"kafka.csuite.ibm.edu:17991"  
            topic_id => "%{datasource}"  
            message_key => "%{resourceID}"  
        } #end Kafka output  
    } #end Kafka condition  
  
} # end output section
```

Logstash receiver output section

The two fields that were created by the mutate filter, **datasource** and **resourceID**, are used in the configuration of the kafka output plug-in.

The value of **datasource** is the name of the Kafka topic that this Logstash receiver writes messages to. Notice that the value of **datasource** is used as the **topic_id**.

Any messages with the same **message_key** are written to the same partition within the topic. Notice that the value of **resourceID** is used as the **message_key**.

Logstash sender input section

```
input {
  kafka {
    zk_connect => "kafka.csuite.ibm.edu:17981"
    group_id => "DEV_IBM-HTTP-Server_access-log"
    topic_id => "DEV_IBM-HTTP-Server_access-log"
    consumer_threads => 1
    consumer_restart_on_error => true
    consumer_restart_sleep_ms => 100
    decorate_events => true
  }# end Kafka input

  kafka {
    zk_connect => "kafka.csuite.ibm.edu:17981"
    group_id => "labInstance_syslog"
    topic_id => "labInstance_syslog"
    consumer_threads => 1
    consumer_restart_on_error => true
    consumer_restart_sleep_ms => 100
    decorate_events => true
  }# end syslog server log Kafka input
}# end input section
```

Using the Log File Agent

© Copyright IBM Corporation 2016

Logstash sender input section

In the Log Analysis scalable architecture, the role of a Logstash sender instance is to retrieve log messages from a Kafka broker and send those messages to the Log Analysis server.

Logstash senders use the kafka input plug-in to pull messages from a topic in the Kafka cluster. In this example, the Logstash sender server is pulling messages from two different topics: DEV_IBM-HTTP-Server_access-log and labInstance_syslog. Notice that the kafka input plug-in is invoked twice, once for each topic.

The following table describes each line in the kafka input plug-in configuration example.

Configuration option	Description
zk_connect	The host and port number of the Zookeeper server.
group_id	A name for the consumer group that identifies the Logstash server in Kafka. Use the same value for topic_id and group_id.
topic_id	The name of the Kafka topic you want to pull data from.
consumer_threads	The number of threads to concurrently pull data from topic. This is typically set to the number of partitions for the topic.
	The consumer_threads parameter also specifies the number of consumers that are created in a consumer group. Each thread, or consumer, maps to a partition for the specified topic. This ensures that data is processed concurrently.
consumer_restart_on_error	If true, consumer threads will restart on errors.

Configuration option	Description
consumer_restart_sleep_ms	The number milliseconds to wait for consumer threads to restart after an error.
decorate_events	If true, metadata about Kafka, such as topic or msg_size is added to each log message.

Logstash sender filter section

```
filter {
  if "grok_lfa" in [tags] {
    mutate {
      add_field => [ "path", "%{LFA_INSTANCE}_%{LFA_MODULE}" ]
      replace => [ "host", "%{LFA_ENVIRONMENTNAME}_%{LFA_MODULE}" ]
    } # end mutate
  } # end lfa mutate condition
} # end filter section
```

Using the Log File Agent

© Copyright IBM Corporation 2016

Logstash sender filter section

You use the filter section of a Logstash sender configuration to alter the log message or metadata that is attached to the log message. At minimum, Logstash senders create two metadata fields for each message: **host** and **path**. You use the **host** and **path** fields when you configure the data source in the Log Analysis server.

In this example, the mutate filter in a Logstash sender configuration is creating a metadata field named **path**. The value for **path** is set with metadata that was added by the Log File Agent (**LFA_INSTANCE** and **LFA_MODULE**). The mutate filter is also replacing the existing value for **host** with metadata that was added by the Log File Agent (**LFA_ENVIRONMENTNAME** and **LFA_MODULE**). The final value for **path** is **labInstance_syslog**. The final value for **host** is **DEV_syslog**.

At the bottom of the slide are pictures of the Log Analysis data source wizard. Notice that you use the values of **host** and **path** that were set by the Logstash sender when you create the data source.

Logstash sender output section

```
output {
    scala {
        scala_url => "https://log-analysis.csuite.ibm.edu:9987/Unity/DataCollector"
        scala_user => "unityadmin"
        scala_password => "object00"
        scala_keystore_path => ""
        batch_size => 500000
        idle_flush_time => 5
        sequential_flush => true
        num_concurrent_writers => 20
        use_structured_api => false
        disk_cache_path => "/opt/logstashW/training/cache/basecache"
        date_format_string => "yyyy-MM-dd'T'HH:mm:ssX"
        log_file => "/opt/logstashW/logstash-2.2.1/log/scala_logstashW.log"
        log_level => "info"
    }#end scala output
}# end output section
```

Using the Log File Agent

© Copyright IBM Corporation 2016

Logstash sender output section

You use the output section of a Logstash sender configuration to send log messages to the Log Analysis server.

Use the scala output plug-in to send data to the Log Analysis server. The following table describes each line in the scala output plug-in configuration example.

Configuration option	Description
scala_url	The REST endpoint for the Log Analysis Generic Receiver REST API.
scala_user	The Log Analysis user name.
scala_password	The Log Analysis user password.
scala_keystore_path	The full path to the directory where the keystore file is saved. Enter a value if you want to encode the plain text scala_password.
batch_size	The maximum number of bytes that can be buffered for each data source before Logstash sends data to the Log Analysis server.
idle_flush_time	The maximum time between successive data transmissions for each data source.
sequential_flush	Defines whether batches for each data source are sent sequentially. If true, batches are sent sequentially.
num_concurrent_writers	The number of threads that concurrently transmit batches of data to the Log Analysis server.

Configuration option	Description
use_structured_api	Determines whether data is transmitted to the Log Analysis server in JSON format. If true, data is sent in JSON format.
disk_cache_path	The path on the file system that temporarily buffers data. The scala output plug-in writes data to this path before sending data to the Log Analysis server.
date_format_string	The string format that all date fields are transformed to before transmission to the Log Analysis server.
log_file	The file that is used for logging information from the scala output plug-in.
log_level	The level of logging information for the scala output plug-in. The supported levels are: fatal, error, warn, info, and debug.

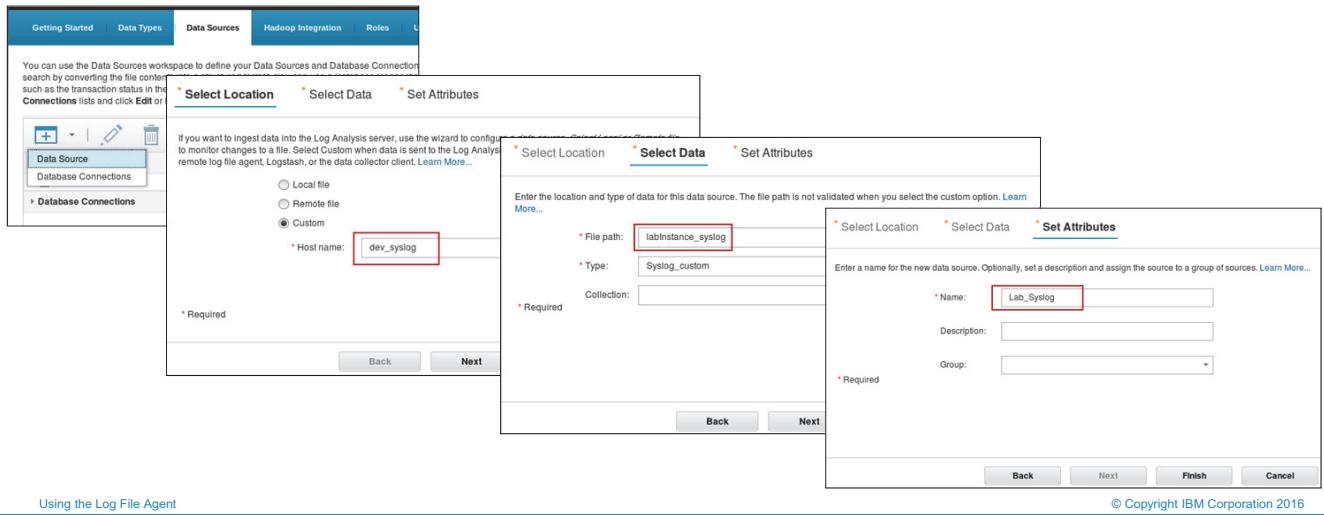


Note: If you installed Logstash by decompressing the installation file or with an RPM Package Manager file (RPM), you must manually copy the scala output plug-in files to the Logstash server. For example, copy the files to the following directory:

<LOGSTASH_HOME>/vendor/bundle/jruby/1.9/gems/logstash-core-2.2.1-java/lib/logstash/outputs/

Configuring Log Analysis

- Use the Log Analysis administrator user interface to add a data source.
- Use the host and path that were set by the Logstash sender, not the actual path of the target log file.



Using the Log File Agent

© Copyright IBM Corporation 2016

Configuring Log Analysis

To configure Log Analysis to accept data from the target log file, you must create a data source. Use the data source wizard in the Log Analysis administrator user interface to add a data source.

Use the host and path that were set by the Logstash sender filter section, not the actual path of the target log file.

Unit summary

- Install the Log File Agent
- Configure the Log File Agent
- Configure Logstash receivers to process Log File Agent messages

[Using the Log File Agent](#)

© Copyright IBM Corporation 2016

Unit summary

Using the Log File Agent exercises

- Exercise 1 Installing and configuring the Log File Agent**
Exercise 2 Configuring HAProxy
Exercise 3 Configuring the logstashA receiver
Exercise 4 Configuring the logstashB receiver
Configuring the Logstash sender and the Log Analysis data source

[Using the Log File Agent](#)

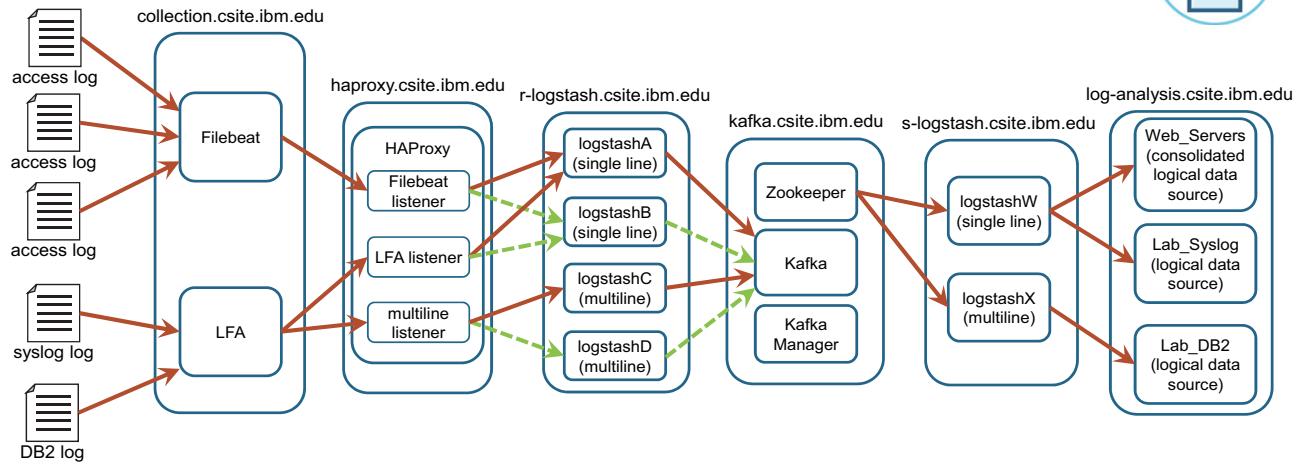
© Copyright IBM Corporation 2016

Student exercises

Perform the exercises for this unit.



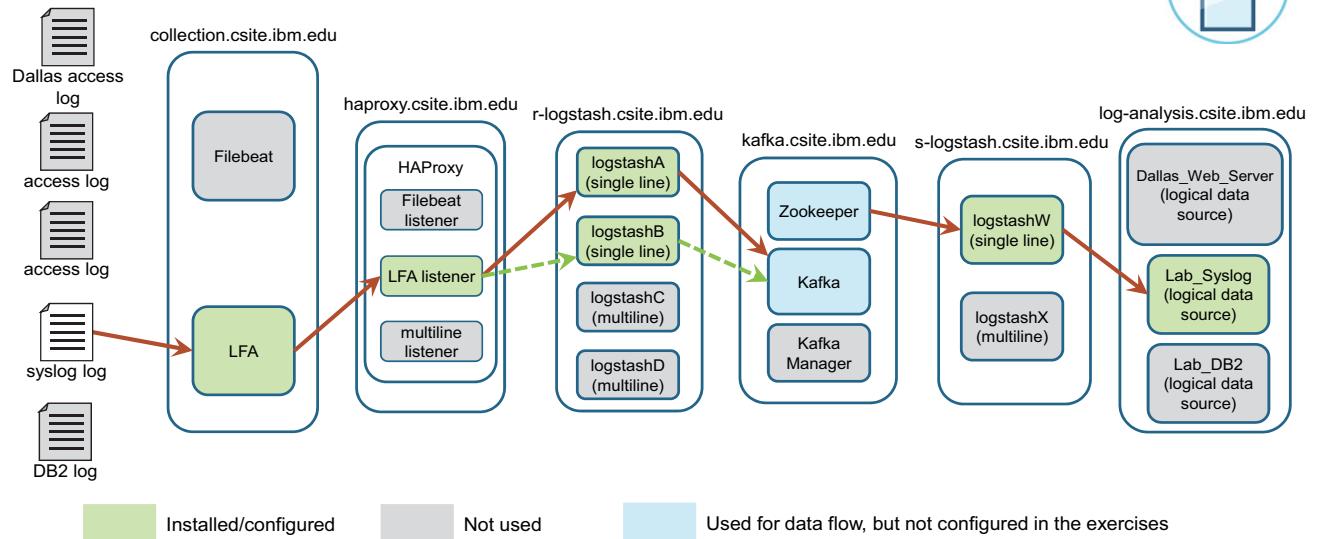
Exercise introduction: Finished lab environment



Exercise introduction: Finished lab environment

At the end of this course, your lab environment will look like this diagram. Refer to this diagram as you complete the steps in the student exercises to help you understand what you are installing and configuring.

Exercise introduction: Lab environment in Unit 5 exercises



Using the Log File Agent

© Copyright IBM Corporation 2016

Exercise introduction: Lab environment in Unit 5 exercises

In the exercises for this unit, you install and configure the following components:

- A Log File Agent, to monitor a syslog file
- An HAProxy listener for traffic from the Log File Agent (configuration only)
- Your logstashA and logstashB receivers, to accept messages from the Log File Agent (configuration only)
- Your logstashW sender, to pull messages from Kafka and send them to the Log Analysis server (configuration only)
- A data source in Log Analysis, to accept data from the target log file (configuration only)

As you install and configure these components, you also verify the data flow between them.

Unit 6 Multiline logs

IBM Training

IBM

Multiline logs

© Copyright IBM Corporation 2016
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

In this unit, you learn how to configure Logstash receivers, Logstash senders, and HAProxy load balancers to process multiline logs.

Unit objectives

- Configure Logstash for multiline logs
- Configure HAProxy for multiline logs
- Configure Logstash senders to send messages from the same log file together
- Diagram data flow for multiline logs

Multiline log example

```
2016-05-16-13.23.45.628104+000 E166427E271           LEVEL: Severe
PID      : 5731             TID   : 139822525732608PROC : db2star2
INSTANCE: db2inst1          NODE  : 000
FUNCTION: DB2 UDB, oper system services, sqloKAAnalyze, probe:10
DATA #1 : Codepath, 8 bytes
3

2016-05-16-13.23.45.858552+000 I166699E491           LEVEL: Warning
PID      : 5731             TID   : 139822525732608PROC : db2star2
INSTANCE: db2inst1          NODE  : 000
FUNCTION: DB2 UDB, SQO Memory Management, SqloMemController::calculateAutomaticSiz,
probe:130
MESSAGE  : AUTOMATIC instance_memory setting decreased due to licensing
restrictions.
DATA #1 : Size in bytes, PD_TYPE_MEM_SIZE, 8 bytes
6900779035
DATA #2 : Size in bytes, PD_TYPE_MEM_SIZE, 8 bytes
4294967296
```

Multiline log example

Many log files are in a multiline format. Multiline log records are single log messages that span multiple lines. To process multiline logs in the scalable collection architecture, you must ensure that all of the lines in a single message are processed together and in the correct sequence.

Logstash and multiline logs

- Logstash instances that processes multiline logs must run in single-threaded mode, which forces Logstash to process each line of a single message in the correct sequence.
- A good practice is to use different Logstash instances for single-line and multiline logs. This suggestion applies to Logstash receivers and senders.
- Use the `-w 1` option when you start Logstash to run it in single-threaded mode, for example:

```
/opt/logstash-2.2.1/bin/logstash -w 1 -f /opt/logstash-2.2.1/conf/logstash-multi.conf -l /opt/logstash-2.2.1/log/logstash-multi-debug.log &
```

Logstash and multiline logs

Multiline Logstash receivers must run in single-threaded mode so that they process each line of a single message in the correct sequence. This configuration helps to ensure that the lines in the log file are processed sequentially, maintaining order between the records.

So be sure to install Logstash receivers and Logstash senders that are dedicated to multiline messages.

HAProxy and multiline logs

Configure HAProxy with affinity for collection agents based on their IP information.

- The `balance source` and `hash-type consistent` lines ensure that all traffic from a collection agent is sent to the same Logstash instance, as long as that Logstash instance is up.
- This is important for multiline logs because you want all lines from the same log message to go to the same Logstash instance.

```
## Listener for LFA Multiline cluster
## Specify Logstash instances for processing multiline messages from LFA
listen Multiline_LFA_Cluster ha-proxy.csuite.ibm.edu:5980
    mode tcp
    balance source
    hash-type consistent
    server receiver-logstashC logstashC.csuite.ibm.edu:5990 check fall 2 rise 3 inter 1000
    server receiver-logstashD logstashD.csuite.ibm.edu:5991 check fall 2 rise 3 inter 1000
```

haproxy.cfg

For multiline log files, configure HAProxy with affinity for log collection agents based on their IP information.

Sending data to Log Analysis with Logstash senders

Use the `metadata_fields` option in the `scala` output plug-in to use a unique identifier for each target log file.

Use a value for `field_names` and `field_paths` that uniquely identifies the physical log data source within your entire environment.

This sample configuration ensures that messages with the same `resourceID` are sent from the Logstash sender to the Log Analysis server together in the same batches.

The `host@path` line specifies which logical data source within Log Analysis messages with the same `resourceID` are mapped to.

```

output {
  scala {
    scala_url =>
      "https://log-analysis.csuite.ibm.edu:9987/Unity/DataCollector"
    scala_user => "unityadmin"
    scala_password => "object00"
    scala_keystore_path => ""
    batch_size => 500000
    idle_flush_time => 5
    sequential_flush => true
    num_concurrent_writers => 20
    use_structured_api => false
    disk_cache_path => "/opt/logstashX/training/cache/basecache"
    date_format_string => "yyyy-MM-dd'T'HH:mm:ssX"
    log_file => "/opt/logstashX/logstash-2.2.1/log/scala_logstashX.log"
    log_level => "info"
    metadata_fields => {
      "TEST_db2diag@LabDB2Instance_db2diag" => {
        "field_names" => "resourceID"
        "field_paths" => "resourceID"
      } # end db2diag meta data
    } # end meta data fields
  }#end scala output
}# end output section

```

Multiline logs

© Copyright IBM Corporation 2016

Sending data to Log Analysis with Logstash senders

Logstash senders pull messages from Kafka, process the log messages, and send them to the Log Analysis core software.

Logstash senders that process multiline messages must use the `metadata_fields` option in the `scala` output plug-in.

Notice the following lines in this example `scala` output plug-in configuration.

```

metadata_fields => {
  "TEST_db2diag@LabDB2Instance_db2diag" => {
    "field_names" => "resourceID"
    "field_paths" => "resourceID"
}

```

The `metadata_fields` option in the Logstash sender ensures that messages with the same `resourceID` are sent from the Logstash sender to the Log Analysis server together in the same batches. This is important for multiline logs, because multiple lines from the same log message must arrive in sequence and only with other lines from the same actual log file.

Configure the value of `resourceID` to be a string that uniquely identifies the physical log source within your entire environment.

Metadata and multiline messages

In this example, the value of **resourceID** is `collection_`/`/opt/db2inst1/sql1lib/db2dump/db2diag.log_1`

From a multiline Logstash receiver filter section:

```
mutate {
    replace => [ "message", "%{LFA_ORIG_MSG}" ]
    add_field => [ "datasource", "%{LFA_INSTANCE}_%{LFA_MODULE}" ]
    add_field => [ "resourceID", "%{LFA_HOSTNAME}_%{LFA_LOGNAME}_1" ]
```

From a multiline Logstash sender filter section:

```
if "grok_lfa" in [tags] {
    mutate {
        add_field => [ "path", "%{LFA_INSTANCE}_%{LFA_MODULE}" ]
        replace => [ "host", "%{LFA_ENVIRONMENTNAME}_%{LFA_MODULE}" ]
    } # end mutate
} # end lfa mutate condition
```

From a multiline Logstash sender output section:

```
metadata_fields => {
    "TEST_db2diag@LabDB2Instance_db2diag" => {
        "field_names" => "resourceID"
        "field_paths" => "resourceID"
    } # end db2diag meta data
} # end meta data fields
```

Metadata and multiline messages

Look at the following lines in the Logstash sender output section. Notice that **field_names** and **field_paths** use the value of **resourceID**.

```
metadata_fields => {
    "TEST_db2diag@LabDB2Instance_db2diag" => {
        "field_names" => "resourceID"
        "field_paths" => "resourceID"
```

The values for host (`TEST_db2diag`) and path (`LabDB2Instance_db2diag`) are set by the Logstash sender filter section.

The value of **resourceID** is set with a filter in the multiline Logstash receiver configuration.

Example multiline Logstash receiver filter configuration:

```
mutate {
    replace => [ "message", "%{LFA_ORIG_MSG}" ]
    add_field => [ "datasource", "%{LFA_INSTANCE}_%{LFA_MODULE}" ]
    add_field => [ "resourceID", "%{LFA_HOSTNAME}_%{LFA_LOGNAME}_1" ]
```

In this example, the value of **resourceID** is:

`collection_`/`/opt/db2inst1/sql1lib/db2dump/db2diag.log_1`.

You also use **resourceID** as the Kafka message key in your Logstash receiver output configuration.

Example of a multiline Logstash receiver output configuration:

```
kafka {  
    bootstrap_servers =>"kafka.csuite.ibm.edu:17991"  
    topic_id => "%{datasource}"  
    message_key => "%{resourceID}"  
} #end Kafka output
```

The message key ensures that the Logstash sender can identify messages from Kafka for this specific physical data source, even if messages from other physical data sources are in the same Kafka topic and partition.

Configure the value of `resourceID` in the Logstash receiver filter section to be a string that uniquely identifies the physical log source within your entire environment.

Unit summary

- Configure Logstash for multiline logs
- Configure HAProxy for multiline logs
- Configure Logstash senders to send messages from the same log file together
- Diagram data flow for multiline logs

Multiline logs

© Copyright IBM Corporation 2016

Unit summary

Multiline logs exercises

Exercise 1 Configuring the Log File Agent

Exercise 2 Configuring HAProxy

Exercise 3 Installing and configuring the first multiline Logstash receiver

Exercise 4 Installing and configuring the second multiline Logstash receiver

Exercise 5 Installing and configuring the multiline Logstash sender

Exercise 6 Sending data to Log Analysis

Multiline logs

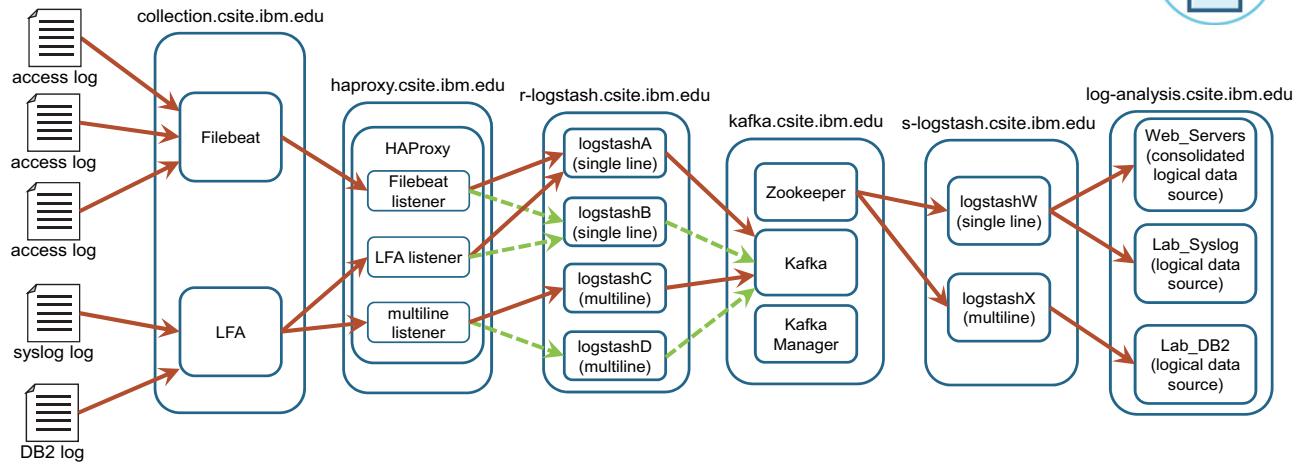
© Copyright IBM Corporation 2016

Student exercises

Perform the exercises for this unit.



Exercise introduction: Finished lab environment



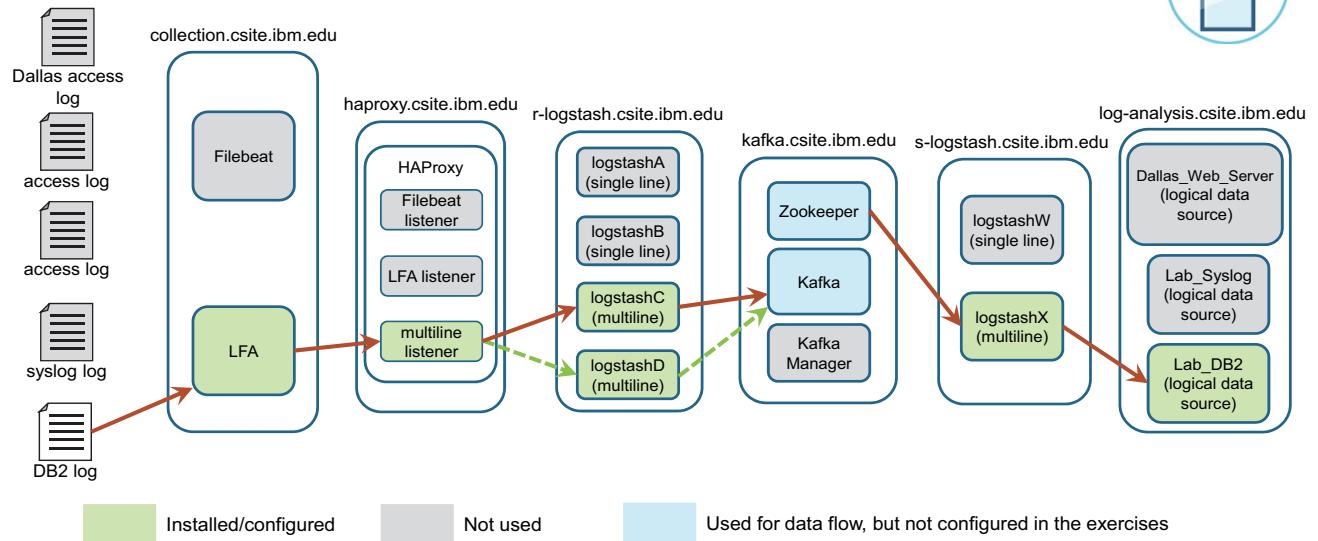
Multiline logs

© Copyright IBM Corporation 2016

Exercise introduction: Finished lab environment

At the end of this course, your lab environment will look like this diagram. Refer to this diagram as you complete the steps in the student exercises to help you understand what you are installing and configuring.

Exercise introduction: Lab environment in Unit 6 exercises



Multiline logs

© Copyright IBM Corporation 2016

Exercise introduction: Lab environment in Unit 6 exercises

In the exercises for this unit, you install and configure the following components:

- A Log File Agent, to monitor a DB2 log (configuration only)
- An HAProxy listener for traffic from the Log File Agent (configuration only)
- Two multiline Logstash receivers, to accept messages from the Log File Agent
- A multiline Logstash sender, to pull messages from Kafka and send them to the Log Analysis server
- A data source in Log Analysis, to accept data from the target log file (configuration only)

As you install and configure these components, you also verify the data flow between them.

Unit 7 Log consolidation

IBM Training

IBM

Log consolidation

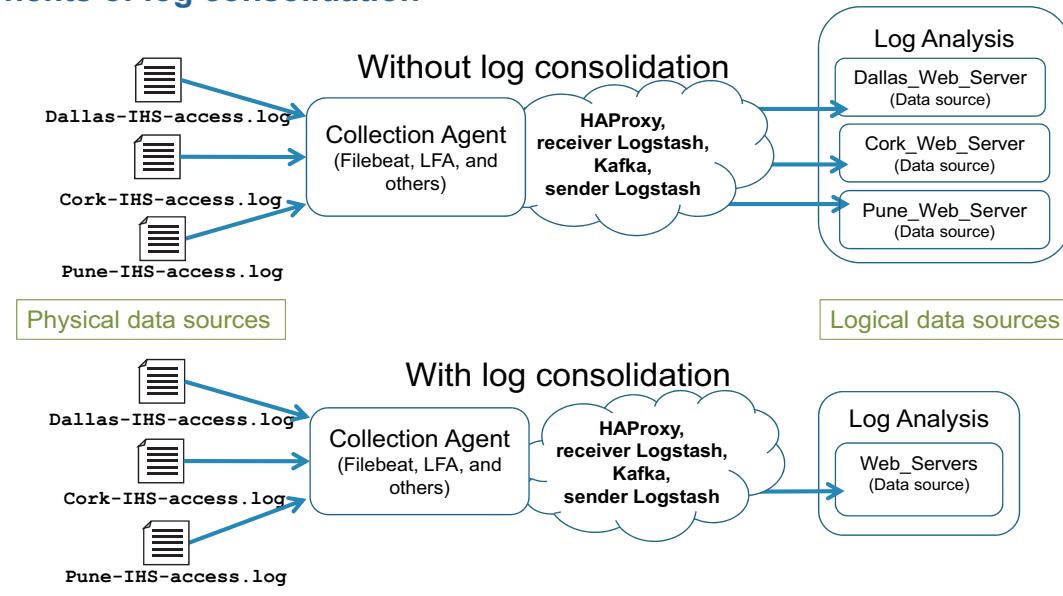
© Copyright IBM Corporation 2016
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

This unit teaches you how to consolidate messages from many logs into just a few Log Analysis data sources.

Unit objectives

- Diagram the data flow for consolidated logs
- Configure a custom source type and index configuration
- Configure log collection agents to add metadata
- Configure Logstash senders to add metadata

Benefits of log consolidation



Log consolidation

© Copyright IBM Corporation 2016

Benefits of log consolidation

In a simple configuration, IBM Log Analysis manages input from a single target log file with a single data source. Environments with many log files to monitor (hundreds or thousands) require just as many data sources. It can be problematic for administrators to create, manage, navigate, and support so many data sources.

With log collection agents and Logstash, you can consolidate logs of the same type into a single Log Analysis data source. Consolidation makes many tasks easier for several roles:

- Administrators can more easily add and manage data sources.
- Users can search through multiple log files simultaneously.

This diagram shows two different configurations: an environment without log consolidation and an environment where log consolidation is implemented.

In the environment without log consolidation, three different Log Analysis data sources are necessary to accept messages from three different web server access log files.

In the environment where log consolidation is implemented, a single Log Analysis data source is accepting messages from the three different web server access log files.

This diagram also illustrates the difference between physical data sources and logical data sources:

- Physical data sources are the actual target log files, such as Dallas-IHS-access.log, Pune-IHS-access.log, or Cork-IHS-access.log.
- Logical data sources are administrative objects that you create with the Log Analysis administrator interface, such as Web_Servers.

Adding a custom source type

Use the Log Analysis administrator interface to clone an existing source type.

The screenshot shows two windows from the IBM Log Analysis administrator interface. On the left, the 'Data Types' workspace is displayed, featuring a toolbar with 'Add', 'Edit', 'Delete', and 'New' buttons. Below the toolbar is a list of data types, with 'WebAccessLog' highlighted and selected. A red arrow points from the 'Data Types' tab in the top navigation bar to the 'WebAccessLog' entry in the list. On the right, a 'Clone WebAccessLog' dialog box is open, titled 'Clone Source Type'. It contains fields for 'Name' (set to 'Consolidated_WebAccessLog'), 'Input type' (set to 'Logs'), and checkboxes for 'Enable splitter' and 'Rule set'. The 'Rule set' checkbox is selected. A red box highlights the 'Name' field. A red arrow points from the 'Name' field in the dialog to the 'WebAccessLog' entry in the list on the left.

Adding a custom source type

To consolidate logs, you must create a custom index configuration to accept extra metadata about each log file. An index configuration is an administrative object within a Log Analysis source type. To create a custom index configuration, you must create a custom source type. Use the Log Analysis administrator user interface to add a custom source type.

To create a custom source type, create a copy of the source type for the log format you want to consolidate. In this example, you want to consolidate web access log files, so you create a copy of the existing WebAccessLog source type.

Editing the index configuration

Within the new source type, edit the index configuration and add fields that identify the physical data source. In this example, the new fields are **location**, **platform**, and **logfile**.

The screenshot shows the 'Edit Index Configuration' dialog and the 'Indexing Properties' table.

Edit Index Configuration Dialog:

- File set selected.
- Deliver data on annotator execution failure checkbox.
- Edit Index Configuration** button highlighted with a red arrow.
- * Required indicator.
- OK and Cancel buttons.

Indexing Properties Table:

- Add Field** button highlighted with a red arrow.
- Show Advanced Columns** button highlighted with a red arrow.
- Table columns: Data Type, Retrievable, Retrieve by default, Sortable, Filterable, Searchable, Combine, Paths.
- Table rows:
 - referrer: TEXT, Retrievable: checked, Retrieve by default: checked, Sortable: checked, Filterable: unchecked, Searchable: checked, Combine: ALL, Paths: annotations.WebAcc.
 - location: TEXT, Retrievable: checked, Retrieve by default: checked, Sortable: checked, Filterable: checked, Searchable: checked, Combine: ALL, Paths: metadata.Location
 - platform: TEXT, Retrievable: checked, Retrieve by default: checked, Sortable: checked, Filterable: checked, Searchable: checked, Combine: ALL, Paths: metadata.Platform
 - logfile: TEXT, Retrievable: checked, Retrieve by default: checked, Sortable: checked, Filterable: checked, Searchable: checked, Combine: ALL, Paths: metadata.Logfile

Log consolidation

© Copyright IBM Corporation 2016

Editing the index configuration

An index configuration determines which fields in each log message is indexed, and how that indexed data can be used in search queries. The Log Analysis administrator user interface includes a tool that you can use to create a custom index configuration.

In this example, you want to index three new fields in your consolidated log messages:

- **location**
- **platform**
- **logfile**

The following table explains the properties that are required for each field.

Property	Description
Field name	The name of the field that users see in their search results.
Data type	The type of data in the field. Possible values are TEXT, LONG, DOUBLE, and DATE.
Retrievable	Determines whether the contents of this field are stored for retrieval. When set to false, the content is not stored in the index. When set to True, the content is stored and available for retrieval.

Property	Description
RetrieveByDefault	When set to True, the contents of the field is always returned as part of any search response. When set to False, the field is not part of the default response. However, when required, the user can explicitly query the content of the field. The retrievable option must be set to True for this attribute to work.
Sortable	Enable or disable the field for sorting and range queries.
Filterable	Enable or disable facet counting and filtering on this field.
Searchable	Controls whether the field is enabled for searching and matching.
Combine	Determines how the values that are returned by the paths and dateFormats attributes are used. Possible values are ALL and FIRST.
Paths	JSON path expression. For indexing metadata in the field, use the string <code>metadata.</code> followed by a name that describes the metadata.

Creating a logical data source

- Use the Log Analysis administrator user interface to add a data source.
- Use the host and path that were set by the Logstash sender, not the actual path of the target log files.
- Use the new source type as the **Type**.

The figure consists of three vertically stacked screenshots of a software interface for creating a data source. The top two screenshots are side-by-side, and the third is positioned below them. All three have a header bar with tabs: 'Select Location' (underlined), 'Select Data', and 'Set Attributes'.

1. **Select Location:** Shows options for 'Local file', 'Remote file', and 'Custom'. The 'Custom' option is selected and highlighted with a red box. Below it, 'Host name:' is set to 'DEV_IBM-HTTP-Server'. A note at the bottom left says 'Required'. At the bottom are 'Back' and 'Next' buttons.

2. **Select Data:** Shows fields for 'File path:' containing 'access-log' and 'Type:' containing 'Consolidated_WebAccessLog', both highlighted with red boxes. A note at the bottom left says 'Required'. At the bottom are 'Back' and 'Next' buttons.

3. **Set Attributes:** Shows fields for 'Name:' containing 'Web_Servers' (highlighted with a red box), 'Description:', and 'Group:'. A note at the bottom left says 'Required'. At the bottom are 'Back', 'Next', 'Finish', and 'Cancel' buttons.

The bottom of the page has a footer bar with 'Log consolidation' on the left and '© Copyright IBM Corporation 2016' on the right.

Creating a logical data source

After you create a custom source type, create a new data source to accept messages from the consolidate log files. Use the name of your custom source type in the **Type** field of the data source wizard.

Configuring the collection agent: Filebeat

Add metadata to each log that the collection agent is monitoring.

In this example, metadata for site and platform are added in the `filebeat.yml` file.

```
paths:
  - /opt/IBM/HTTPServer/logs/Dallas-IHS-access.log
  input_type: log
  fields:
    collector: filebeats-collection.csuite.ibm.edu
    env: DEV
    module: IBM-HTTP-Server
    type: access-log
    site: DALLAS
    platform: RHEL

  paths:
  - /opt/IBM/HTTPServer/logs/Pune-IHS-access.log
  input_type: log
  fields:
    collector: filebeats-collection.csuite.ibm.edu
    env: DEV
    module: IBM-HTTP-Server
    type: access-log
    site: PUNE
    platform: SLES

  paths:
  - /opt/IBM/HTTPServer/logs/Cork-IHS-access.log
  input_type: log
  fields:
    collector: filebeats-collection.csuite.ibm.edu
    env: DEV
    module: IBM-HTTP-Server
    type: access-log
    site: CORK
    platform: RHEL
```

`filebeat.yml`

Log consolidation

© Copyright IBM Corporation 2016

Configuring the collection agent: Filebeat

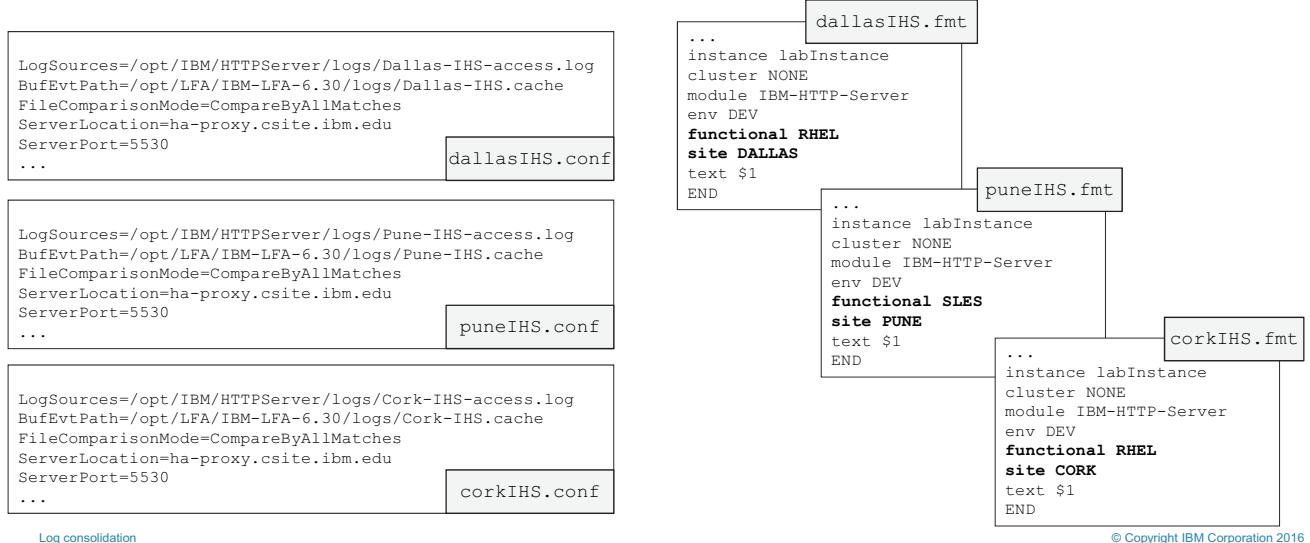
When users search a consolidated data source, the search results show messages from multiple log files. Metadata that you add to each log message can show users exactly which messages came from which log file.

When you consolidate logs, you must include metadata so that users can determine which log file a message in the search interface came from. Use the log collection agent to add metadata to log messages.

In this example configuration, Filebeat adds two metadata fields that you can include in user search results: site and platform. You configure Filebeat by editing the `filebeat.yml` file.

Configuring the collection agent: Log File Agent

In this example, metadata for site and platform are added in the .conf and .fmt file for each target log file.



Log consolidation

© Copyright IBM Corporation 2016

Configuring the collection agent: Log File Agent

In this example configuration, the Log File Agent adds to metadata fields that you can include in user search results: **site** and **platform**. You configure the Log File Agent to add metadata by editing the .fmt file.



Hint: Notice that the values for **platform** (RHEL and SLES) are in a field named **functional**.

Remember that the grok filter in a Logstash receiver configuration uses a pattern named LFAMESSAGE to match meaningful strings of text from the Log File Agent. This pattern matches the value of the **functional** field. That is why this example uses **functional** to store a value for operating system platform. If you to create a different metadata field in the .fmt file, for example, **platform**, you must change the LFAMESSAGE grok pattern to match the new field.

Logstash sender filter section

```
filter {  
  
    if "mutate_filebeat" in [tags] {  
        mutate {  
            add_field => [ "path", "%{[fields][type]}" ]  
            replace => { "host" => "%{[fields][env]}_%{[fields][module]}" }  
        } # end mutate  
    } # end filebeat mutate condition  
  
    if [fields][type] == "access-log" {  
        mutate {  
            add_field => [ "Location", "%{[fields][site]}" ]  
            add_field => [ "Platform", "%{[fields][platform]}" ]  
            add_field => [ "LogFile", "%{resourceID}" ]  
        } # end web server mutate  
    } # end web server mutate condition  
  
} # end filter section
```

Logstash sender filter section

To send metadata that was added by the log collection agents, use the Logstash sender. In this example, the Logstash sender filter section is adding three new fields to store metadata from a Filebeat collection agent:

- **Location**, to store values from the Filebeat **site** field such as DALLAS, PUNE, or CORK
- **Platform**, to store values from the Filebeat **platform** field such as RHEL or SLES.
- **LogFile**, to store a unique identifier for each target log file. This example uses the value of **resourceID**.

Logstash sender output section

```
output {
  scala {
    scala_url => "https://log-analysis.csuite.ibm.edu:9987/Unity/DataCollector"
    scala_user => "unityadmin"
    scala_password => "object00"
    scala_keystore_path => ""
    batch_size => 500000
    idle_flush_time => 5
    sequential_flush => true
    num_concurrent_writers => 20
    use_structured_api => false
    disk_cache_path => "/opt/logstashW/training/cache/basecache"
    date_format_string => "yyyy-MM-dd'T'HH:mm:ssX"
    log_file => "/opt/logstashW/logstash-2.2.1/log/scala_logstashW.log"
    log_level => "info"
    metadata_fields => {
      "DEV_IBM-HTTP-Server@access-log" => {
        "field_names" => "Location,Platform,Logfile"
        "field_paths" => "Location,Platform,Logfile"
      }
    }
  }#end scala output
}# end output section
```

Field Name	Data Type	Retrievable	Retrieve by default	Sortable	Filterable	Searchable	Combine	Paths
referrer	TEXT	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input checked="" type="checkbox"/>	ALL	annotations.WebAcc
location	TEXT	<input checked="" type="checkbox"/>	ALL	metadata.Location				
platform	TEXT	<input checked="" type="checkbox"/>	ALL	metadata.Platform				
logfile	TEXT	<input checked="" type="checkbox"/>	ALL	metadata.Logfile				

Log consolidation

© Copyright IBM Corporation 2016

Logstash sender output section

To send metadata that was added by the log collection agents, use the `metadata_fields` option in the `scala` output plug-in.

Scala output plug-in configuration example

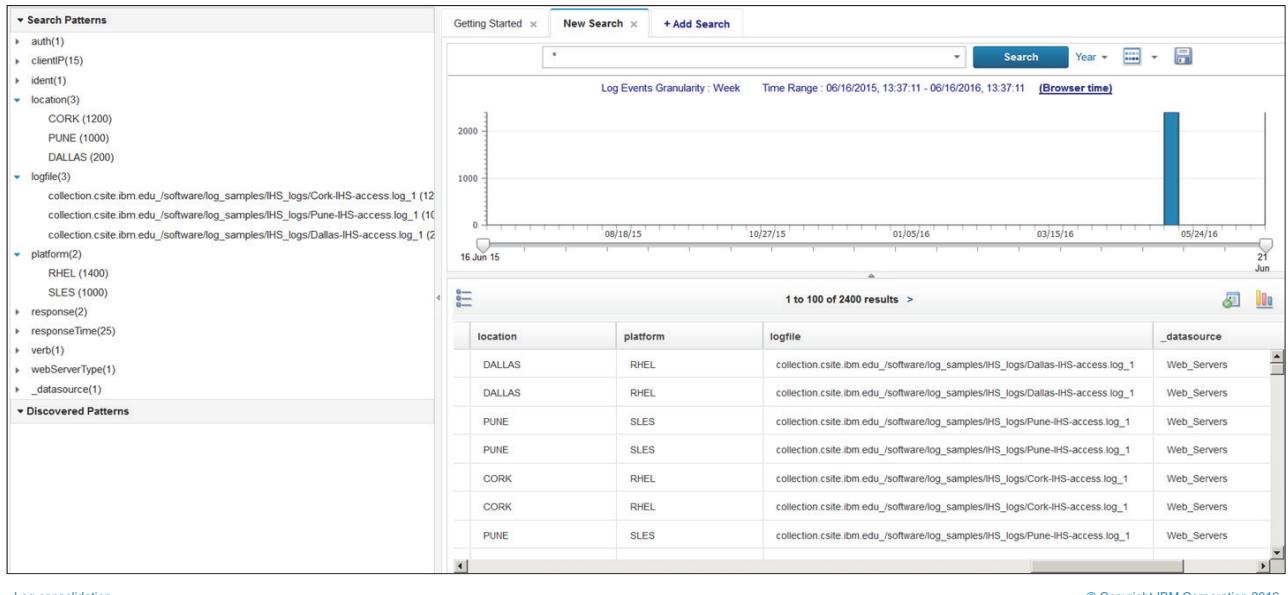
In the following example, line numbers have been added to the `scala` output plug-in configuration for clarity.

```
1 metadata_fields => {
2   "DEV_IBM-HTTP-Server@access-log" => {
3     "field_names" => "Location,Platform,Logfile"
4     "field_paths" => "Location,Platform,Logfile"
5   } # end web server meta data
6 } # end meta data fields
```

- Line 1 sends extra metadata fields to the Log Analysis server.
- Line 2 specifies the Log Analysis data source that the metadata is sent to. The data source is specified by its host and path. The format is `host@path`.
- Line 3 is a comma-delimited list of field names from the message metadata that you want to send to Log Analysis. You add these fields in the filter section of the Logstash sender.
- Line 4 is a comma-delimited list of the paths that you created in your custom index configuration. In this example, three fields were added in the index configuration: **metadata.Location**, **metadata.Platform**, and **metadata.Logfile**.

- Line 5 closes the configuration for metadata fields that are sent to the `DEV_IBM-HTTP-Server@access-log` data source.
- Line 6 closes the configuration for the `metadata_fields` option.

Searching consolidated logs



Log consolidation

© Copyright IBM Corporation 2016

Searching consolidated logs

When users search a consolidated data source, the search results show messages from multiple log files. Metadata that you add to each log message can show users exactly which messages came from which log file.

In this example, users can filter the search results for the consolidated data source by location, operating system platform, or by the actual log file the message results came from.

Unit summary

- Diagram the data flow for consolidated logs
- Configure a custom source type and index configuration
- Configure log collection agents to add metadata
- Configure Logstash senders to add metadata

Log consolidation

© Copyright IBM Corporation 2016

Unit summary

Log consolidation exercises

- Exercise 1 Deleting the current Log Analysis data source**
Exercise 2 Adding a source type and a new data source
Exercise 3 Configuring the Logstash sender
Exercise 4 Configuring Filebeat
Exercise 5 Verifying log consolidation

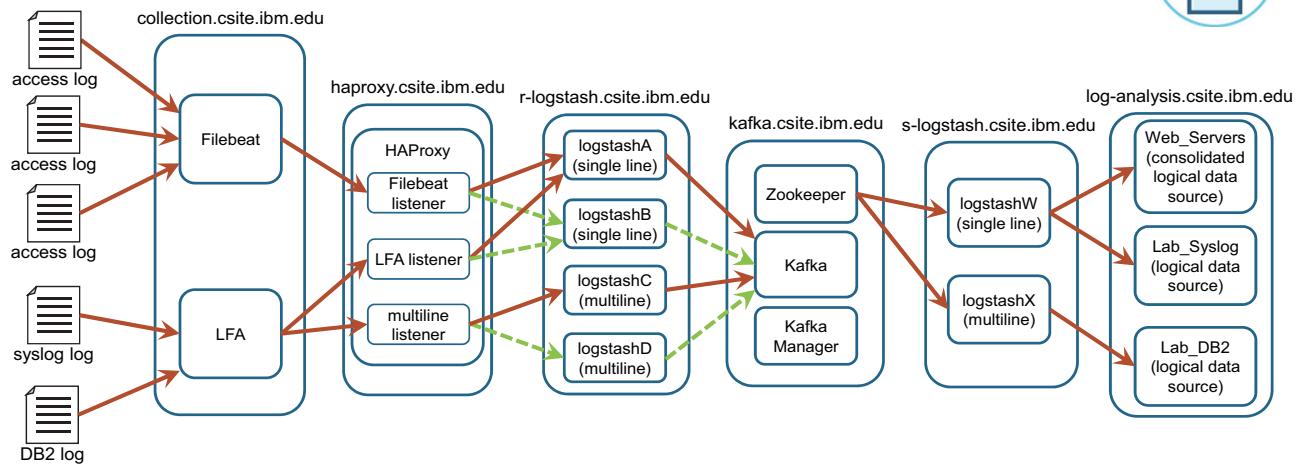
Log consolidation

© Copyright IBM Corporation 2016

Student exercises



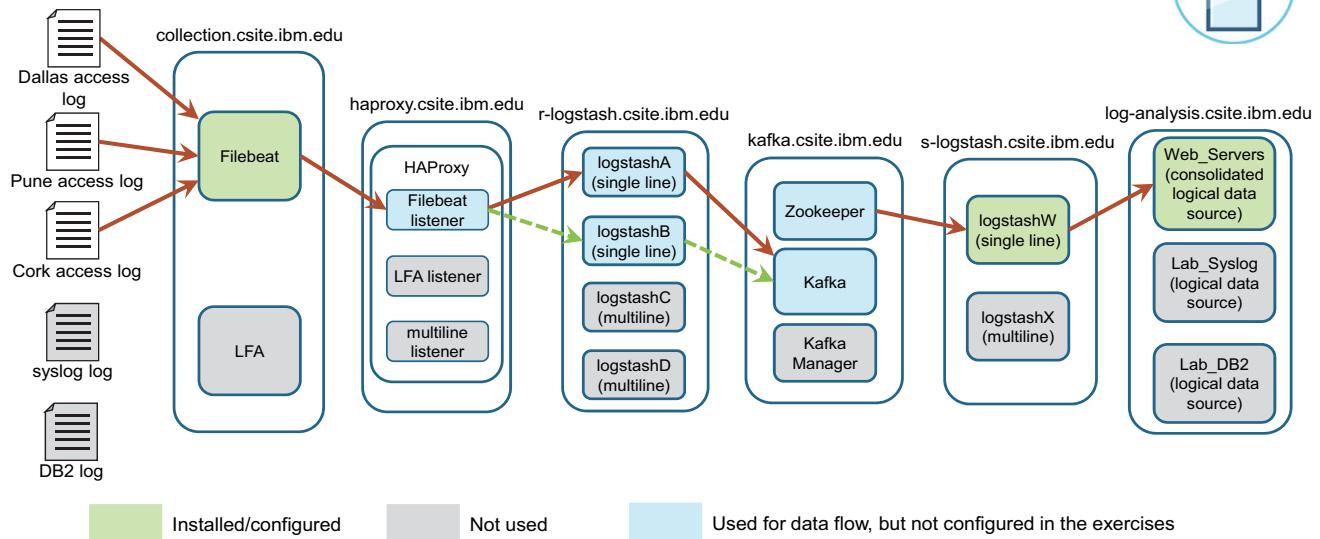
Exercise introduction: Finished lab environment



Exercise introduction: Finished lab environment

At the end of this course, your lab environment will look like this diagram. Refer to this diagram as you complete the steps in the student exercises to help you understand what you are installing and configuring.

Exercise introduction: Lab environment in Unit 7 exercises



Log consolidation

© Copyright IBM Corporation 2016

Exercise introduction: Lab environment in Unit 7 exercises

In the exercises for this unit, you configure the following components:

- Your Filebeat collection agent, to monitor two more web server log files (configuration only)
- Your logstashW sender, to send metadata to the Log Analysis server (configuration only)
- A custom source type and index configuration in Log Analysis, to index the extra metadata from Filebeat (configuration only)
- A data source in Log Analysis, to accept data from the target log files (configuration only)

As you configure these components, you also verify the data flow between them.

Unit 8 Troubleshooting

IBM Training

IBM

Troubleshooting

© Copyright IBM Corporation 2016
Course materials may not be reproduced in whole or in part without the prior written permission of IBM.

In this unit, you learn how to troubleshoot each of the components in the scalable collection architecture.

Unit objectives

- Troubleshoot Filebeat
- Troubleshoot the Log File Agent
- Troubleshoot HAProxy
- Troubleshoot Logstash
- Troubleshoot Kafka
- Troubleshoot Log Analysis

Filebeat logging

- Filebeat logging is configured in the `filebeat.yml` file.
- Filebeat log levels are: critical, error, warning, info, debug

```
...  
logging:  
  to_files: true  
  files:  
    path: /opt/filebeat-1.1.1-x86_64/  
    name: mybeat  
    rotateeverybytes: 10485760  
level: debug
```

Filebeat logging

The `filebeat.yml` file controls logging for Filebeat. Use the `level` property to set the log level. You must restart Filebeat for any changes to take effect.

Filebeat file processing

Messages like these indicate that Filebeat is reading the target log file correctly:

```
2016-06-16T17:55:54Z DBG full line read
2016-06-16T17:55:59Z DBG Publish: {
...
  "env": "DEV",
  "module": "IBM-HTTP-Server",
  "platform": "RHEL",
  "site": "CORK",
  "type": "access-log"
},
  "input_type": "log",
  "message": "Apache/IHS,10.10.255.90,-,-,[10/May/2016:14:04:42 -0400],GET,\\\"GET
//app?action=sell\u0026holdingID=5015 HTTP/1.1\\\",200,3409,1090,\\\"-\\\",\\\"curl/7.21.3 (x86_64-
linux-gnu) libcurl/7.21.3 OpenSSL/1.0.0 zlib/1.2.3\\\""",
  "offset": 280044,
  "source": "/software/log_samples/IHS_logs/Cork-IHS-access.log",
  "type": "log"
```

These messages confirm that Filebeat is reading the target log file and trying to output log messages.

Filebeat event publishing

Messages like these indicate that Filebeat is successfully sending messages:

```
2016-06-16T18:27:41Z DBG 710 events out of 952 events sent to logstash. Continue
sending ...
2016-06-16T18:27:41Z DBG Try to publish 242 events to logstash with window size 952
2016-06-16T18:27:42Z DBG 242 events out of 242 events sent to logstash. Continue
sending ...
2016-06-16T18:27:42Z DBG send completed
2016-06-16T18:27:42Z INFO Events sent: 952
2016-06-16T18:27:42Z DBG Processing 952 events
```

Filebeat event publishing

These messages confirm that Filebeat is successfully sending messages to a Logstash receiver. Notice that the number of events in the window decreases as Filebeat sends messages to the Logstash receiver.

Filebeat connectivity

Messages like these indicate that HAProxy is not responding:

```
2016-06-16T17:55:59Z DBG output worker: publish 200 events
2016-06-16T17:55:59Z DBG Try to publish 200 events to logstash with window size 120
2016-06-16T17:55:59Z DBG close connection
2016-06-16T17:55:59Z DBG 0 events out of 200 events sent to logstash. Continue sending ...
2016-06-16T17:55:59Z INFO Error publishing events (retrying): EOF
2016-06-16T17:55:59Z INFO send fail
2016-06-16T17:56:14Z INFO Connecting error publishing events (retrying): dial tcp
192.168.100.176:20737: getsockopt: connection refused
```

Messages like these indicate that the Logstash receiver cluster is not responding:

```
2016-06-16T18:09:12Z DBG Try to publish 200 events to logstash with window size 30
2016-06-16T18:09:12Z DBG close connection
2016-06-16T18:09:12Z DBG 0 events out of 200 events sent to logstash. Continue sending ...
2016-06-16T18:09:12Z INFO Error publishing events (retrying): write tcp 192.168.100.175:42323-
>192.168.100.176:20737: write: broken pipe
2016-06-16T18:09:12Z INFO send fail
```

Filebeat connectivity

The first set of messages show that Filebeat cannot connect to HAProxy. Notice the `tcp connection refused` message.

The second set of messages show that Filebeat cannot connect to any Logstash receivers on the other side of HAProxy. Notice the `broken pipe` message.

Log File Agent logging

- Log File Agent logging is configured in the <LFA_HOME>/IBM-LFA-6.30/config/lo_default_workload_instance.config file.
- Log levels are explained in the configuration file.

```
# =====
DEFAULTCFG='NO'
KBB_SIG1='-asyncoff -syncoff -dumpoff'
KBB_RAS1='ERROR (UNIT:logmonitor ALL)'
KBB_VARPREFIX='%'
KBB_RAS1_LOG='%(CTIRA_LOG_PATH)/collection_lo_default_workload_instance_%(systask)_%(sysutcstart)-.log
INVENTORY=%(CTIRA_LOG_PATH)/collection_lo_default_workload_instance_%(systask).inv COUNT=03 LIMIT=5 PRESERVE=1
MAXFILES=9'
KDC_FAMILIES='ip.pipe HTTPS_CONSOLE:N HTTPS_SERVER:N HTTPS:0'
LOGSHOME='/opt/LFA/IBM-LFA-6.30/logs'
```

Log File Agent logging

The lo_default_workload_instance.config file controls logging for the Log File Agent. Use the KBB_RAS1 property to set the log level. You must restart the Log File Agent for any changes to take effect.

Log File Agent log file

Logs are in the <LFA_HOME>/IBM-LFA-6.30/logs directory.

Look for the most recent log file name with the format:

collection_lo_default_workload_instance_kloagent_NNNNNNNN-NN.log.

```
-rw-r--r-- 1 netcool ncoadmin      0 Jun 16 19:09 collection_lo_1466104169.pid6382
-rw-r-xr-x 1 netcool ncoadmin  41866 Jun  8 19:02 collection_lo_default_workload_instance_kloagent_57584219-01.log
-rw-r--r-- 1 netcool ncoadmin  38815 Jun 16 18:58 collection_lo_default_workload_instance_kloagent_5762a250-01.log
-rw-r--r-- 1 netcool ncoadmin 5206756 Jun 16 19:00 collection_lo_default_workload_instance_kloagent_5762f6eb-01.log
-rw-r--r-- 1 netcool ncoadmin 2764151 Jun 16 19:04 collection_lo_default_workload_instance_kloagent_5762f6eb-02.log
-rw-r--r-- 1 netcool ncoadmin 5167405 Jun 16 19:05 collection_lo_default_workload_instance_kloagent_5762f837-01.log
-rw-r--r-- 1 netcool ncoadmin 4198253 Jun 16 19:09 collection_lo_default_workload_instance_kloagent_5762f837-02.log
-rw-r--r-- 1 netcool ncoadmin 5162248 Jun 16 19:07 collection_lo_default_workload_instance_kloagent_5762f837-03.log
-rw-r--r-- 1 netcool ncoadmin 36663 Jun 16 19:10 collection_lo_default_workload_instance_kloagent_5762f969-01.log
-rw-r-xr-x 1 netcool ncoadmin    736 Jun 16 19:09 collection_lo_default_workload_instance_kloagent.inv
-rw-r--r-- 1 netcool ncoadmin  2826 Jun 16 19:09 default_workload_instance:collec.LG0
-rw-r--r-- 1 netcool ncoadmin  2826 Jun 16 19:04 default_workload_instance:collec.LG1
```

Log File Agent log file

Log File Agent logs are in the <LFA_HOME>/IBM-LFA-6.30/logs directory. The name of the log files in this directory is collection_lo_default_workload_instance_kloagent, followed by a time stamp. To find the file that the Log File Agent is currently logging to, look for the file with the most recent date.

Log File Agent event publishing

Messages like these indicate that the Log File Agent is successfully sending messages to HAProxy:

```
(5762A82A.0001-C:sockeif.c,1209,"_imp_eipc_recv_data") KDE1_ReceiveOn returned 0x1DE0000B
(5762AD1F.0000-C:sockeif.c,1209,"_imp_eipc_recv_data") KDE1_ReceiveOn returned 0x1DE0000B
(5762AD1F.0001-C:sockeif.c,1209,"_imp_eipc_recv_data") KDE1_ReceiveOn returned 0x1DE0000B
(5762AD1F.0002-C:sockeif.c,1209,"_imp_eipc_recv_data") KDE1_ReceiveOn returned 0x1DE0000B
(5762AD1F.0003-C:sockeif.c,1209,"_imp_eipc_recv_data") KDE1_ReceiveOn returned 0x1DE0000B
(5762AD1F.0004-C:sockeif.c,1209,"_imp_eipc_recv_data") KDE1_ReceiveOn returned 0x1DE0000B
(5762AE00.0000-C:sockeif.c,1209,"_imp_eipc_recv_data") KDE1_ReceiveOn returned 0x1DE0000B
```

Log File Agent event publishing

This set of messages show that the Log File Agent is successfully reading messages from the target log file and sending them to HAProxy.

Log File Agent connectivity

Messages like these indicate that HAProxy is not responding:

```
(5762FCEE.0000-B:sockeif.c,390,"_imp_connect") KDE1 connection returned 0x1DE00045  
errno 107 for ha-proxy.csuite.ibm.edu port 5530
```

```
(5762FCEE.0001-B:sockeif.c,537,"_imp_eipc_create_remote_client") Cannot connect to ha-  
proxy.csuite.ibm.edu<192.168.100.176> port 5530, rc -1
```

Log File Agent connectivity

This set of messages show that the Log File Agent cannot connect to HAProxy.

HAProxy logging

- HAProxy logging is configured in the `/etc/haproxy/haproxy.cfg` file.
- HAProxy writes to syslog or rsyslog.
- HAProxy log levels are: emerg, alert, crit, err, warning, notice, info, debug.
- The `tcplog` option enables advanced logging of TCP connections with session state.

```
global                               haproxy.cfg
  log      127.0.0.1 local0  debug
  #log    127.0.0.1 local1
  #log    127.0.0.1 local2
  #log    127.0.0.1 local3
...
#-----
-
# common defaults that all the 'listen' and 'backend' sections will
# use if not designated in their block
#-----
-
defaults
  mode      tcp
  log       global
  option   tcplog
...
```

Troubleshooting

© Copyright IBM Corporation 2016

HAProxy logging

The `haproxy.cfg` file controls logging for HAProxy. Use the `global log` property to set the log level. You must restart HAProxy for any changes to take effect.

In the global defaults section of `haproxy.cfg`, the `tcplog` option configures HAProxy to record a detailed message about each TCP connection that it handles.

HAProxy TCP session logging

Messages like these are details about each TCP connection.

```
Jun 16 19:48:04 localhost haproxy[4271]: 192.168.100.175:53237 [16/Jun/2016:19:48:02.247]
LFA_Cluster LFA_Cluster/receiver-logstashA 1/0/120003 0 sD 0/0/0/0/0 0/0

Jun 16 19:48:07 localhost haproxy[4271]: 192.168.100.175:42343 [16/Jun/2016:19:48:03.777]
Beats_Cluster Beats_Cluster/receiver-logstashB 1/0/63448 30 cD 0/0/0/0/0 0/0

Jun 16 19:48:59 localhost haproxy[4271]: 192.168.100.175:42344 [16/Jun/2016:19:48:12.287]
Beats_Cluster Beats_Cluster/receiver-logstashA 1/1/61013 72 cD 0/0/0/0/0 0/0
```

HAProxy TCP session logging

HAProxy can record detailed messages about each TCP session. Each of the messages in this example describes one TCP session. The log message is not written to the log file until the TCP session has finished.

The following table explains the format of a TCP log message.

Field	Example	Description
1	Jun 16 19:48:59	Message time stamp
2	localhost	HAProxy host name
3	haproxy[4271]	Process name and process ID
4	192.168.100.175:42344	IP address and port number of the client that initiated the TCP connection
5	[16/Jun/2016:19:48:12.287]	Time stamp when HAProxy received the connection
6	Beats_Cluster	The listener that is used for the connection
7	Beats_Cluster/receiver-logstashA	The Logstash receiver that the data was forwarded to
8	1/1/61013	Total wait time / wait time to connect to the Logstash receiver / total transaction time
9	72	Bytes sent from the server to the client
10	cD	A two-letter code indicating the session termination state

Field	Example	Description
11	0/0/0/0/0	Number of concurrent HAProxy connections / number of concurrent listener connections / number of concurrent connections to the Logstash receiver / Number of active HAProxy connections at the time of the log message / session retries
12	0/0	Number of session requests before this transaction / Number of session requests before this transaction for the Logstash receiver cluster

For more information about the format of a TCP log message, refer to the HAProxy documentation at <http://www.haproxy.org/download/1.5/doc/configuration.txt>.

HAProxy connectivity

Messages like these indicate that a Logstash server is down.

```
Jun 16 20:27:25 localhost haproxy[4557]: Server Beats_Cluster/receiver-logstashB is DOWN,
reason: Layer4 connection problem, info: "Connection refused", check duration: 0ms. 1 active
and 0 backup servers left. 0 sessions active, 0 requeued, 0 remaining in queue.
Jun 16 20:27:25 localhost haproxy[4557]: Server LFA_Cluster/receiver-logstashB is DOWN,
reason: Layer4 connection problem, info: "Connection refused", check duration: 0ms. 1 active
and 0 backup servers left. 0 sessions active, 0 requeued, 0 remaining in queue.
```

HAProxy connectivity

HAProxy performs periodic health checks for the Logstash receivers in the back end cluster. If a Logstash receiver stops responding, HAProxy records that event in its log.

HAProxy statistics page

HAProxy		Statistics Report for pid 4762																																			
> General process information																																					
<small>pid = 4762 (process #1, nbproc = 1) uptime = 0d 16h39m55s system_uptime = 0d 16h39m55s maxsock = 8000 maxconn = 4000 maxpipes = 0 Current conn > 1; current pipes = 0/0; conn rate = 1/sec Running tasks: 1/13; idle = 100 %</small>																																					
<small>Note: "NOLB" "DRAIN" = UP with load-balancing disabled.</small>																																					
<small>Display option: <input type="checkbox"/> External resources: • Scope <input type="checkbox"/> Primary site • Hide DOWN servers • Refresh now • CSV export</small>																																					
Beats_Cluster		Queue			Session rate			Sessions			Bytes			Denied			Errors			Warnings																	
		Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	Last	In	Out	Req	Resp	Req	Conn	Resp	Retr	Redis															
Frontend		0	1	-	0	1	-	3	000	1				11 151	60	0	0	0																			
receiver-logstashA		0	0	-	0	1	-	0	1	-	1	1		16h32m	11 151	60	0	0	0	0	0	0															
receiver-logstashB		0	0	-	0	0	-	0	0	-	0	0		?	0	0	0	0	0	0	0	0															
Backend		0	0	-	0	1	-	300	1	1	16h32m	11 151	60	0	0	0	0	0	0	0	0	0															
LFA_Cluster		Queue			Session rate			Sessions			Bytes			Denied			Errors			Warnings																	
		Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	Last	In	Out	Req	Resp	Req	Conn	Resp	Retr	Redis															
Frontend		0	0	-	0	0	-	0	0	3 000	0			0	0	0	0	0	0	0	0	0															
receiver-logstashA		0	0	-	0	0	-	0	0	-	0	0		?	0	0	0	0	0	0	0	0															
receiver-logstashB		0	0	-	0	0	-	0	0	-	0	0		?	0	0	0	0	0	0	0	0															
Backend		0	0	-	0	0	-	300	0	-	0	0		?	0	0	0	0	0	0	0	0															
Multiline_LFA_Cluster		Queue			Session rate			Sessions			Bytes			Denied			Errors			Warnings																	
		Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	Last	In	Out	Req	Resp	Req	Conn	Resp	Retr	Redis															
Frontend		0	0	-	0	0	-	0	0	3 000	0			0	0	0	0	0	0	0	0	0															
receiver-logstashC		0	0	-	0	0	-	0	0	-	0	0		?	0	0	0	0	0	0	0	0															
receiver-logstashD		0	0	-	0	0	-	0	0	-	0	0		?	0	0	0	0	0	0	0	0															
Backend		0	0	-	0	0	-	300	0	-	0	0		?	0	0	0	0	0	0	0	0															
stats		Queue			Session rate			Sessions			Bytes			Denied			Errors			Warnings																	
		Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	Last	In	Out	Req	Resp	Req	Conn	Resp	Retr	Redis															
Frontend		1	2	-	1	1	-	3 000	5	-	1 919	100 981		0	0	0	0	0	0	0	0	0															
Backend		0	0	-	0	1	-	300	2	-	0	0s		1 919	100 981	0	0	0	2	0	0	0	16h39m UP														

Troubleshooting

© Copyright IBM Corporation 2016

HAProxy statistics page

You can configure HAProxy to display a web page that shows the status of each server that it is forwarding data to. The page also shows statistics about each of the endpoints in the cluster, such as the volume of data sent. This page is useful for troubleshooting. You configure this web page in the `haproxy.cfg` file.

The following example shows how to enable the HAProxy statistics page in the `haproxy.cfg` file.

```
...
stats socket /var/lib/haproxy/stats
...
listen stats :9000 #Listen on localhost port 9000
    mode http
    stats enable
    stats hide-version
    stats realm Haproxy\ Statistics
    stats uri /haproxy_stats
    stats auth netcool:object00
```

The following table describes each line in the configuration example.

Configuration option	Description
listen stats :9000	This line configures the statistics page to use port 9000.
mode http	This line enables http mode.
stats enable	This line enables statistics.

Configuration option	Description
stats hide-version	This line hides the version of HAProxy so it is not visible on the web page.
stats realm Haproxy\ Statistics	This text is what users see when they are prompted to log in to the page.
stats uri /haproxy_stats	The URI of the web page, for example, <code>http://localhost:9000/haproxy_stats</code>
stats auth netcool:object00	This line enables authentication. It also sets a user name and password. In this example, the user name is <code>netcool</code> and the password is <code>object00</code> .

Logstash log level

Start Logstash with the `--verbose` option to increase the log level.

```
/opt/logstashA/logstash-2.2.1/bin/logstash --verbose -f /opt/logstashA/logstash-2.2.1/conf/logstashA.conf -l /opt/logstashA/logstash-2.2.1/log/logstashA-debug.log &
```

Logstash log level

To increase the log level of a Logstash instance, start it with the `--verbose` or the `--debug` option.

With the `--debug` option, you can see Logstash accepting input, adding tags, using filters, and sending output. In this example, a Logstash sender uses the `mutate` filter to add and replace fields in a message.

```
{ :timestamp=>"2016-07-08T19:24:29.341000+0000", :message=>"config
Logstash::Filters::Mutate/@replace =
{\\"host\\":>\"%{ [fields] [env] }_%{ [fields] [module] }\\\"", :level=>:debug,
:file=>"logstash/config/mixin.rb", :line=>"122", :method=>"config_init"}
{:timestamp=>"2016-07-08T19:25:13.351000+0000",
:message=>"filters/Logstash::Filters::Mutate: adding value to field",
:field=>"path", :value=>["%{ [fields] [type] }"], :level=>:debug,
:file=>"logstash/util/decorators.rb", :line=>"31", :method=>"add_fields"}
{:timestamp=>"2016-07-08T19:25:13.366000+0000",
:message=>"filters/Logstash::Filters::Mutate: adding value to field",
:field=>"Location", :value=>["%{ [fields] [site] }"], :level=>:debug,
:file=>"logstash/util/decorators.rb", :line=>"31", :method=>"add_fields"}
{:timestamp=>"2016-07-08T19:25:13.380000+0000",
:message=>"filters/Logstash::Filters::Mutate: adding value to field",
:field=>"Platform", :value=>["%{ [fields] [platform] }"], :level=>:debug,
:file=>"logstash/util/decorators.rb", :line=>"31", :method=>"add_fields"}
```

Use the file output plug-in for troubleshooting

```
output {  
    file {  
        path => "/opt/logstashA/logstash-2.2.1/log/logstashA-debug.log"  
        codec => rubydebug  
    } #end file  
  
    if ("mutate_filebeat" or in [tags]) and ! ("_grokparsefailure" in [tags]) {  
        kafka {  
            bootstrap_servers =>"kafka.csuite.ibm.edu:17991"  
            topic_id => "%{datasource}"  
            message_key => "%{resourceID}"  
        } #end Kafka output  
    } #end Kafka condition  
}  
} # end output section
```

Use the file output plug-in for troubleshooting

Use the file output plug-in to configure Logstash to send log messages it has processed to its log file. You can look at the messages to see the result of Logstash processing, for example:

- Whether filters were successful
- The value for each field in the log message
- Tag values
- The message after Logstash is finished processing it

In this example, a Logstash sender is sending messages to a Kafka broker and to a file named /opt/logstashA/logstash-2.2.1/log/logstashA-debug.log.



Important: If you send all messages to a file, every log message in the target log file is also saved on the Logstash server. You might want to use the file output plug-in for testing only and disable file output after you are finished troubleshooting to conserve disk space.

Use tags to verify filter operations

```
filter {  
    if [fields][collector] == "filebeats-collection.csuite.ibm.edu" {  
        mutate {  
            add_field => [ "datasource", "%{[fields][env]}_%{[fields][module]}_%{[fields][type]}" ]  
            add_field => [ "resourceID", "%{[beat][hostname]}_%{source}_1" ]  
            add_tag => "mutate_filebeat"  
        } # end mutate  
    } #end filebeat condition  
  
    if [type] == "lfa" {  
        grok {  
            patterns_dir => "/opt/logstashA/logstash-2.2.1/patterns"  
            match => [ "message", "%{LFAMESSAGE}" ]  
            add_tag => ["grok_lfa"]  
        } #end initial LFA grok  
    } #end initial LFA condition  
  
    if "grok_lfa" in [tags] {  
        mutate {  
            replace => [ "message", "%{LFA_ORIG_MSG}" ]  
            add_field => [ "datasource", "%{LFA_INSTANCE}_%{LFA_MODULE}" ]  
            add_field => [ "resourceID", "%{LFA_HOSTNAME}_%{LFA_LOGNAME}_1" ]  
        } # end mutate  
    } # end grok_lfa condition  
} #end filter section
```

Troubleshooting

© Copyright IBM Corporation 2016

Use tags to verify filter operations

A good practice is to add a tag to each message after you alter it with a filter. These tags can confirm if a filter was successful.

In this example, you can confirm if the mutate and grok filters were successful by the presence or absence of the `mutate_filebeat` and `grok_lfa` tags.

Scala output plug-in logging

```
output {  
  
    scala {  
        scala_url => "https://log-analysis.csuite.ibm.edu:9987/Unity/DataCollector"  
        scala_user => "unityadmin"  
        scala_password => "object00"  
        scala_keystore_path => ""  
        batch_size => 500000  
        idle_flush_time => 5  
        sequential_flush => true  
        num_concurrent_writers => 20  
        use_structured_api => false  
        disk_cache_path => "/opt/logstashW/training/cache/basecache"  
        date_format_string => "yyyy-MM-dd'T'HH:mm:ssX"  
        log_file => "/opt/logstashW/logstash-2.2.1/log/scala_logstashW.log"  
        log_level => "info"  
    }#end scala output  
  
}# end output section
```

Scala output plug-in logging

The scala output plug-in can write messages to its own log file. This log file shows communication between the Logstash sender and the Log Analysis server.

Use the following options in the scala output plug-in configuration to control the log file:

- **log_file**: The file that is used for logging information from the scala output plug-in.
- **log_level**: The level of logging information for the scala output plug-in. The supported levels are: fatal, error, warn, info, and debug.

Kafka troubleshooting

- Kafka broker logging is configured in the `<KAFKA_HOME>/config/log4j.properties` file.
- Kafka broker log levels are: TRACE, DEBUG, INFO, WARN, ERROR, and FATAL.
- Use the following log files in the `<KAFKA_HOME>/log/server.log` directory for troubleshooting:
 - controller.log
 - kafka-request.log
 - kafkaServer-gc.log
 - kafkaServer.out
 - log-cleaner.log
 - **server.log**
 - state-change.log

```
# Turn on all our debugging info
#log4j.logger.kafka.producer.async.DefaultEventHandler=DEBUG, kafkaAppender
#log4j.logger.kafka.client.ClientUtils=DEBUG, kafkaAppender
#log4j.logger.kafka.perf=DEBUG, kafkaAppender
#log4j.logger.kafka.perf.ProducerPerformance$ProducerThread=DEBUG, kafkaAppender
#log4j.logger.org.I0Itec.zkclient.ZkClient=DEBUG
log4j.logger.kafka=DEBUG, kafkaAppender
```

log4j.properties

Kafka troubleshooting

The `log4j.properties` file controls logging for Kafka brokers. Use the `log4j.logger.kafka` property to set the log level. You must restart the Kafka broker for any changes to take effect.

Zookeeper logging

- Zookeeper logging is configured in the <KAFKA_HOME>/config/log4j.properties file.
- Zookeeper log levels are: TRACE, DEBUG, INFO, WARN, ERROR, and FATAL.
- Use the <KAFKA_HOME>/log/server.log, zookeeper-gc.log, and zookeeper.out files for troubleshooting.

```
# Turn on all our debugging info
#log4j.logger.kafka.producer.async.DefaultEventHandler=DEBUG, kafkaAppender
#log4j.logger.kafka.client.ClientUtils=DEBUG, kafkaAppender
#log4j.logger.kafka.perf=DEBUG, kafkaAppender
#log4j.logger.kafka.perf.ProducerPerformance$ProducerThread=DEBUG, kafkaAppender
log4j.logger.org.I0Itec.zkclient.ZkClient=DEBUG, kafkaAppender
log4j.logger.kafka=INFO, kafkaAppender
```

log4j.properties

Zookeeper logging

The log4j.properties file controls logging for the Zookeeper server. Use the log4j.logger.org.I0Itec.zkclient.ZkClient property to set the log level. You must restart the Zookeeper server for any changes to take effect.

Viewing data in Kafka

Use the following command to list all topics:

```
kafka-topics.sh --list --zookeeper kafka.csuite.ibm.edu:17981
```

```
DEV_IBM-HTTP-Server_access-log
LabDB2Instance_db2diag
labInstance_syslog
```

Use the following command to show all data in a topic:

```
kafka-console-consumer.sh --zookeeper kafka.csuite.ibm.edu:17981 --topic
labInstance_syslog --from-beginning
```

Viewing data in Kafka

You can use command line tools to confirm that a Kafka cluster is receiving data. These tools are in the `<KAFKA_HOME>/bin` directory.

Use the `kafka-topics.sh` tool to list all topics in the Kafka cluster.

Use the `kafka-console-consumer.sh` tool to view the data in a topic. If you use the `--from-beginning` option, the tool shows all data in a topic. If you do not use the `--from-beginning` option, the tool shows only data that is currently being written to the topic.

You can also use the `kafka-topics.sh` tool to show details about a topic. In the following example, the tool shows information about the topic named **DEV_IBM-HTTP-Server_access-log**:

```
./kafka-topics.sh --describe --zookeeper kafka.csuite.ibm.edu:17981 --topic
DEV_IBM-HTTP-Server_access-log
```

```
Topic:DEV_IBM-HTTP-Server_access-logPartitionCount:3ReplicationFactor:1Configs:
Topic: DEV_IBM-HTTP-Server_access-logPartition: 0Leader: 0Replicas: 0Isr: 0
```

Kafka Manager

The screenshot shows the Kafka Manager web interface with the following sections:

- Brokers:** Displays a table with one row for broker ID 0, host kafka.cssite.ibm.edu, port 17991, JMX port 8092, bytes in 39k, and bytes out 40k.
- Combined Metrics:** Shows rates for messages, bytes in, and bytes out per second over 1, 5, and 15 minutes.
- Topics:** A table listing topics: __consumer_offsets, DEV_JBM-HTTP-Server_access-log, LabDB2Instance_db2diag, and labInstance_syslog. Each topic has 1 partition, 1 broker, and 100 brokers spread.
- Broker Id 0:** A detailed view for broker ID 0. It includes a summary table with metrics like # of Topics (4), # of Partitions (53), and % of Messages (100.000). Below is a chart titled "Messages count" showing a steady increase from approximately 10,000 to 15,000 messages over time.

Troubleshooting

© Copyright IBM Corporation 2016

Kafka Manager

Kafka Manager is a web-based user interface that shows the configuration of your cluster. You can also use Kafka Manager to view data and modify your cluster. Although Kafka Manager is not required, it can be useful for troubleshooting and maintenance.

Using Kafka Manager to verify sender Logstash data flow

The screenshot shows the Kafka Manager interface. On the left, the 'Update Cluster' configuration page is displayed, featuring fields for Cluster Zookeeper Hosts (kafka.csite.ibm.edu:17981), Kafka Version (0.8.2.2), and various optional checkboxes related to consumer information polling and offset caching. On the right, two main sections are shown: 'Partitions by Broker' and 'Topic Summary'. The 'Partitions by Broker' section displays a table with one broker (id 0) having 1 partition, which is skewed. The 'Topic Summary' section shows a single topic (G-DEV_IBM-HTTP-Server_access-log) with 100% of partitions assigned to a single consumer instance.

Broker	# of Partitions	Partitions	Skewed?
0	1	(0)	false

Consumers consuming from this topic	
G-DEV_IBM-HTTP-Server_access-log	ZK

Topic Summary				
Total Lag	0			
% of Partitions assigned to a consumer instance	100			
DEV_IBM-HTTP-Server_access-log				
Partition	LogSize	Consumer Offset	Lag	Consumer Instance Owner
0	17000	17000	0	G-DEV_IBM-HTTP-Server_access-log_s-logstash.csite.ibm.edu-1466086619431-3abdceb5-0

Troubleshooting

© Copyright IBM Corporation 2016

Using Kafka Manager to verify sender Logstash data flow

If you enable the **Poll consumer information** option in Kafka Manager, you can see which Logstash senders are pulling data from Kafka. This information is useful for troubleshooting the data flow between Kafka and the Logstash senders.

Log Analysis troubleshooting

- The `GenericReceiver.log` file shows data coming in to the generic receiver interface.
- To increase the log level logs, edit the `<LA_HOME>/wlp/usr/servers/Unity/apps/Unity.war/WEB-INF/classes/log4j.properties` file.
- Log levels are: TRACE, DEBUG, INFO, WARN, ERROR, and FATAL

```
# Unity Generic Receiver Log configurations
log4j.logger.UnityGenericReceiver=DEBUG,UNITY_GR_FILE
#define the appender named FILE
log4j.appenders.UNITY_GR_FILE=org.apache.log4j.RollingFileAppender
log4j.appenders.UNITY_GR_FILE.MaxFileSize=50000KB
```

log4j.properties

Log Analysis troubleshooting

The generic receiver is a REST interface for loading data into IBM Operations Analytics Log Analysis.

The `log4j.properties` file controls logging for the Log Analysis generic receiver. Use the `log4j.logger.UnityGenericReceiver` property to set the log level. You must restart Log Analysis for any changes to take effect.

Generic receiver accepting data

Messages like these indicate that the generic receiver is successfully processing data.

```
06/17/16 19:07:42:146 UTC [Default Executor-thread-17915] INFO - UnityFlowController :  
Batch Status for -> Lab_DB2 , Size: 273 , Num successful: 273 , Num failures: 0 ,  
Indexed Source volume: 136484  
  
06/17/16 19:07:42:149 UTC [Default Executor-thread-17915] INFO -  
DataCollectorRestServlet : Batch of Size 273 processed and encountered 0 failures  
  
06/17/16 19:07:42:223 UTC [Thread-93] INFO - IndexStatusChecker : Updating statistics  
for data source [Lab_DB2], stream [_unity_default_stream], ingested bytes [137097],  
write date [Fri Jun 17 19:07:42 UTC 2016].
```

Generic receiver accepting data

The GenericReceiver.log file shows any attempt to load data into Log Analysis. This log file is useful for troubleshooting communication between a Logstash sender and the Log Analysis server.

This set of messages confirms that a Logstash sender is successfully sending messages to the Log Analysis server and that the data in the messages is successfully indexed.

Unit summary

- Troubleshoot Filebeat
- Troubleshoot the Log File Agent
- Troubleshoot HAProxy
- Troubleshoot Logstash
- Troubleshoot Kafka
- Troubleshoot Log Analysis



IBM Training



© Copyright IBM Corporation 2016 All Rights Reserved.